# **Readthedocs Documentation**

**Matthew Welch** 

# **PRACTICE GUIDE**

1	Introduction  1.1 Building a knowledge base	3 4 4 5
2	Glossary and list of acronyms  2.1 List of Acronyms	<b>7</b> 7 8
3	3.1 Need for SDC 3.2 The risk-utility trade-off in the SDC process 3.3 Release Types 3.4 Conditions for PUFs 3.5 Conditions for SUFs	11 12 13 14 15
4	4.1 Types of disclosure .  4.2 Classification of variables .  4.3 Disclosure scenarios .  4.4 Levels of risk .  4.5 Individual risk .  4.6 Special Uniques Detection Algorithm (SUDA) .  4.7 Risk measures for continuous variables .  4.8 Global risk .	17 17 17 19 20 21 27 32 33
5	5.1 Classification of SDC methods	37 37 38 52 66 66 67
6	6.1 General utility measures for continuous and categorical variables	<b>69</b> 70 76 79 85

7	<b>SDC</b>	with sdcMicro in R: Setting Up Your Data and more	87
	7.1	Installing <i>R</i> , <i>sdcMicro</i> and other packages	87
	7.2	Read functions in $R$	88
	7.3	Missing values	89
	7.4	Classes in <i>R</i>	89
	7.5	Objects of class sdcMicroObj	90
	7.6	Household structure	94
	7.7	Randomizing order and numbering of individuals or households	96
	7.8	Computation time	97
	7.9	Common errors	100
8	The S	SDC Process	101
	8.1	Step 1: Need for confidentiality protection	101
	8.2	Step 2: Data preparation and exploring data characteristics	
	8.3	Step 3: Type of release	
	8.4	Step 4: Intruder scenarios and choice of key variables	
	8.5	Step 5: Data key uses and selection of utility measures	
	8.6	Step 6: Assessing disclosure risk	
	8.7	Step 7: Assessing utility measures	
	8.8	Step 8: Choice and application of SDC methods	106
	8.9	Step 9: Re-measure risk	107
	8.10	Step 10: Re-measure utility	107
	8.11	Step 11: Audit and Reporting	107
	8.12	Step 12: Data release	107
9	Appe	ndices	111
	9.1	Appendix A: Overview of Case Study Variables	111
	9.2	Appendix B: Example of Blanket Agreement for SUF	112
	9.3	Appendix C: Internal and External Reports for Case Studies	
10	Biblio	ography	117
11	Lists		119
11		List of Tables	
		List of Figures	
		List of Examples	
	11.5	List of Examples	121
<b>12</b>		Study (Illustrating the SDC Process)	125
		Case study 1- SUF	
	12.2	Case study 2 - PUF	153
Bil	oliogra	aphy	179

#### Thijs Benschop

#### **Cathrine Machingauta**

#### **Matthew Welch**

Affiliations: Thijs Benschop, Consultant The World Bank, Cathrine Machingauta, The World Bank, Matthew Welch, The World Bank

Correspondence: Matthew Welch, mwelch@worldbank.org

Acknowledgments: The authors thank Olivier Dupriez (The World Bank) for his technical comments and inputs throughout the process.

The production of this guide was made possible through a World Bank Knowledge for Change II Grant: KCP II - A microdata dissemination challenge: Balancing data protection and data utility. Grant number: TF 015043, Project Number P094376. As well as from United Kingdom - DFID funding to the World Bank Multi-Donor Trust Fund - International Household Survey and Accelerated Data Program – TF071804/TF011722.

PRACTICE GUIDE 1

2 PRACTICE GUIDE

**CHAPTER** 

ONE

#### INTRODUCTION

National statistics agencies are mandated to collect microdata<sup>1</sup> from surveys and censuses to inform and measure policy effectiveness. In almost all countries, statistics acts and privacy laws govern these activities. These laws require that agencies protect the identity of respondents, but may also require that agencies disseminate the results and, in appropriate cases, the microdata. Data producers who are not part of national statistics agencies are also often subject to restrictions, through privacy laws or strict codes of conduct and ethics that require a similar commitment to privacy protection. This has to be balanced against the increasing requirement from funders that data produced using donor funds be made publically available.

This tension between complying with confidentiality requirements while at the same time requiring that microdata be released means that a demand exists for practical solutions for applying Statistical Disclosure Control (SDC), also known as microdata anonymization. The provision of adequate solutions and technical support has the potential to "unlock" a large number of datasets.

The International Household Survey Network (IHSN) and the World Bank have contributed to successful programs that have generated tools, resources and guidelines for the curation, preservation and dissemination of microdata and resulted in the documentation of thousands of surveys by countries and agencies across the world. While these programs have ensured substantial improvements in the preservation of data and dissemination of good quality metadata, many agencies are still reluctant to allow access to the microdata. The reasons are technical, legal, ethical and political, and sometimes involve a fear of being criticized for not providing perfect data. When combined with the tools and guidelines already developed by the IHSN/World Bank for the curation, preservation and dissemination of microdata, tools and guidelines for the anonymization of microdata should further reduce or remove some of these obstacles.

Working with the IHSN, PARIS21 (OECD), Statistics Austria and the Vienna University of Technology, the World Bank has contributed to the development of an open source software package for SDC, called *sdcMicro*. The package was developed for use with the open source *R* statistical software, available from the Comprehensive R Archive Network (CRAN) at http://cran.us.r-project.org. The package includes numerous methods for the assessment and reduction of disclosure risk in microdata.

Ensuring that a free open source solution is available to agencies was an important step forward, but not a sufficient one. There is still limited consolidated and reported knowledge on the impact of disclosure risk reduction methods on data utility. This limited access to knowledge combined with a lack of experience in using the tools and methods makes it difficult for many agencies to implement optimal solutions, i.e., solutions that meet their obligations towards both privacy protection and the release of data useful for policy monitoring and evaluation. This practice guide attempts to fill this critical gap by:

- 1. consolidating knowledge gained at the World Bank through experiments conducted during a large-scale evaluation of anonymization techniques
- 2. translating the experience and key results into practical guidelines

<sup>&</sup>lt;sup>1</sup> Microdata are unit-level data obtained from sample surveys, censuses and administrative systems. They provide information about characteristics of individual people or entities such as households, business enterprises, facilities, farms or even geographical areas such as villages or towns. They allow in-depth understanding of socio-economic issues by studying relationships and interactions among phenomena. Microdata are thus key to designing projects and formulating policies, targeting interventions and monitoring and measuring the impact and results of projects, interventions and policies.

It should be stressed that SDC is only one part of the data release process, and its application must be considered within the complete data release framework. The level and methods of SDC depend on the laws of the country, the sensitivity of the data and the access policy (i.e., who will gain access) considered for release. Agencies that are currently releasing data are already using many of the methods described in this guide and applying appropriate access polices to their data before release. The primary objective of this guide is to provide a primer to those new to the process who are looking for guidance on both theory and practical implementation. This guide is not intended to prescribe or advocate for changes in methods that specific data producers are already using and which they have designed to fit and comply with their existing data release policies.

The guide seeks to provide practical steps to those agencies that want to unlock access to their data in a safe way and ensure that the data remain fit for purpose.

## 1.1 Building a knowledge base

The release of data is important, as it allows researchers and policymakers to replicate officially published results, generate new insights into issues, avoid duplication of surveys and provide greater returns to the investment in the survey process.

Both the production of reports, with aggregate tables of indicators and statistics, and the release of microdata result in privacy challenges to the producer. In the past, for many agencies, the only requirement was to release a report and some key indicators. The recent movement around Open Data, Open Government and transparency means that agencies are under greater pressure to release their microdata to allow broader use of data collected through public and donor funds. This guide focuses on the methods and processes for the release of microdata.

Releasing data in a safe way is required to protect the integrity of the statistical system, by ensuring agencies honor their commitment to respondents to protect their identity. Agencies do not widely share, in substantial detail, their knowledge and experience using SDC and the processes for creating safe data with other agencies. This makes it difficult for agencies new to the process to implement solutions. To fill this experience and knowledge gap, we evaluated the use of a broad suite of SDC methods on a range of survey microdata covering important development topics related to health, labor, education, poverty and inequality. The data we used were all previously treated to make them safe for release. Given that their producers had already treated these data, it was not possible, nor was it our goal, to pass any judgment on the safety of these data, many of which are in the public domain The focus was rather on measuring the effects that various methods would have on the risk-utility trade-off for microdata produced to measure common development indicators. We used the experience from this large-scale experimentation to inform our discussion of the processes and methods in this guide.

**Important:** At no point was any attempt made to re-identify, through matching or any other method, any respondents in the surveys we used in building our knowledge base. All risk assessments were based on frequencies and probabilities.

## 1.2 Using this guide

The methods discussed in this guide originate from a large body of literature on SDC. The processes underlying many of the methods are the subject of extensive academic research and many, if not all, of them are used extensively by agencies experienced in preparing microdata for release.

Where possible, for each method and topic, we provide elaborate examples, references to the original or seminal work describing the methods and algorithms in detail and recommended readings. This, when combined with the discussion of the method and practical considerations in this guide, should allow the reader to understand the methods and their strengths and weaknesses. It should also provide enough detail for readers to use an existing software solution to implement the methods or program the methods in statistical software of their choice.

For the examples in this guide, we use the open source and free package for SDC called *sdcMicro* as well as the statistical software *R. sdcMicro* is an add-on package to the statistical software *R*. The package was developed and is

maintained by Matthias Templ, Alexander Kowarik and Bernhard Meindl. 'See and the GitHub of the developers. The GitHub repository can also be used to submit bugs found in the package.<sup>2</sup> The statistical software *R* and the *sdcMicro* package, as well as any other packages needed for the SDC process, are freely available from the Comprehensive R Archive Network (CRAN) mirrors (http://cran.r-project.org/). The software is available for Linux, Windows and Macintosh operating systems. We chose to use *R* and *sdcMicro* because it is freely available, accommodates all main data formats and is easy to adapt by the user. The World Bank, through the IHSN, has also provided funding towards the development of the *sdcMicro* package to ensure it meets the requirements of the agencies we support.

This guide does not provide a review of all other available packages for implementing the SDC process. Our concern is more with providing practical insight into the application of the methods. We would, however, like to highlight one particular other software package that is commonly used by agencies:  $\mu$ -ARGUS<sup>3</sup>.  $\mu$ -ARGUS is developed by Statistics Netherlands. sdcMicro and  $\mu$ -ARGUS are both widely used in statistics offices in the European Union and implement many of the same methods.

The user needs some knowledge of R to use sdcMicro. It is beyond the scope of this guide to teach the use of R, but we do provide throughout the guide code examples on how to implement the necessary routines in R. We also present a number of case studies that include the code for the anonymization of a number of demo datasets using R. Through these case studies, we demonstrate a number of approaches to the anonymization process in R. [#foot17]\_

# 1.3 Outline of this guide

This guide is divided into the following main sections:

- 1. Section 2 is a primer on SDC.
- 2. Section 3 gives an introduction to different release types for microdata.
- 3. Sections 4 to 6 cover SDC methods, risk and utility measurement. Here the goal is to provide knowledge that allows the reader to independently apply and execute the SDC process. This section is enriched with real examples as well as code snippets from the *sdcMicro* package. The interested reader can also find more information in the references and recommended readings at the end of each section.
- 4. Section 7 gives an overview of issues encountered when carrying out anonymization with the *sdcMicro* package in *R*, which exceed basic *R* knowledge. This section also includes tips and solutions to some of the common issues and problems that might be encountered when applying SDC methods in *R* with *sdcMicro*.
- 5. Section 8 provides a step-by-step guide to disclosure control, which draws upon the knowledge presented in the previous sections.
- 6. Section 9 presents a number of detailed case studies that demonstrate the use of the methods, their implementation in *sdcMicro* and the process that should be followed to reach the optimal risk-utility solution.

<sup>&</sup>lt;sup>2</sup> http://cran.r-project.org/web/packages/sdcMicro/index.htmlhttps://github.com/alexkowa/sdcMicro

 $<sup>^3</sup>$   $\mu$ -ARGUS is available at: http://neon.vb.cbs.nl/casc/mu.htm. The software was recently ported to open source.

<sup>&</sup>lt;sup>4</sup> There are many free resources for learning *R* available on the web. One place to start would be the CRAN *R* Project page: http://cran.r-project.org/other-docs.html

# **GLOSSARY AND LIST OF ACRONYMS**

# 2.1 List of Acronyms

COICOP	Sub Saharan Africa
COICOI	Classification of Individual Consumption by Purpose
CRAN	Comprehensive R Archive Network
CTBIL	Contingency Table-Based Information Loss
DHS	Demographic and Health Surveys
DIS	Data Intrusion Simulation
EAP	East Asia and the Pacific
ECA	Europe and Central Asia
EU	European Union
GIS	Geographical Information System
GPS	Global Positioning System
GUI	Graphical User Interface
HIV/AIDS	Human Immunodeficiency Virus/Acquired Immune Deficiency Syndrome
I2D2	International Income Distribution Database
IHSN	International Household Survey Network
LAC	Latin America and the Caribbean
LSMS	Living Standards Measurement Survey
MDAV	Maximum Distance Average Vector
MDG	Millennium Development Goal
MENA	Middle East and North America
MICS	Multiple Indicator Cluster Survey
MME	Mean Monthly Expenditures
MMI	Mean Monthly Income
MSU	Minimal Sample Uniques
NSI	National Statistical Institute
NSO	National Statistical Office
OECD	Organization for Economic Cooperation and Development
PARIS21	Partnership in Statistics for Development in the 21st century
PRAM	Post Randomization Method
PC	Principal Component
PUF	Public Use File
SA	South Asia
SDC	Statistical Disclosure Control
SSE	Sum of Squared Errors
SHIP	Survey-based Harmonized Indicators Program

Continued on next page

Table 2.1 – continued from previous page

SUDA	Special Uniques Detection Algorithm
SUF	Scientific Use File
UNICEF	United Nations Children's Fund

# 2.2 Glossary

Administrative data	Data collected for administrative purposes by government agencies. Typically, administrative
Anonymization	Use of techniques that convert confidential data into anonymized data/ removal or masking of
Attribute disclosure	Attribute disclosure occurs if an intruder is able to determine new characteristics of an indivi
Categorical variable	A variable that takes values over a finite set, e.g., gender. Also called factor in <i>R</i> .
Confidentiality	Data confidentiality is a property of data, usually resulting from legislative measures, which
Confidential data	Data that will allow identification of an individual or organization, either directly or indirectly
Continuous variable	A variable with which numerical and arithmetic operations can be performed, e.g., income.
Data protection	Data protection refers to the set of privacy-motivated laws, policies and procedures that aim
Deterministic methods	Anonymization methods that follow a certain algorithm and produce the same results if appl
Direct identifier	A variable that reveals directly and unambiguously the identity of a respondent, e.g., names,
Disclosure	Disclosure occurs when a person or an organization recognizes or learns something that they
Disclosure risk	A disclosure risk occurs if an unacceptably narrow estimation of a respondent's confidential
End user	The user of the released microdata file after anonymization. Who is the end user depends on
Factor variable	Factor variables are one way to classify categorical variables in <i>R</i> .
Hierarchical structure	Data is made up of collections of records that are interconnected through links, e.g., individu
Identifier	An identifier is a variable/ information that can be used to establish identity of an individual
Identity disclosure	Identity disclosure occurs if an intruder associates a known individual or organization with a
Indirect identification	Indirect identification occurs when the identity of an individual or organization is disclosed,
Inferential disclosure	Inferential disclosure occurs if an intruder is able to determine the value of some characterist
Information loss	Information loss refers to the reduction of the information content in the released data relativ
Interval	A set of numbers between two designated endpoints that may or may not be included. Brack
Intruder	A user who misuses released data by trying to disclose information about an individual or or
k-anonymity	The risk measure $k$ -anonymity is based on the principle that the number of individuals in a s
Key	A combination or pattern of key variables/quasi-identifiers.
Key variables	A set of variables that, in combination, can be linked to external information to re-identify re
Microaggregation	Anonymization method that is based on replacing values for a certain variable with a commo
Microdata	A set of records containing information on individual respondents or on economic entities. S
Noise addition	Anonymization method based on adding or multiplying a stochastic or randomized number t
Non-perturbative methods	Anonymization methods that reduce the detail in the data or suppress certain values (masking
Observation	A set of data derived from an object/unit of experiment, e.g., an individual (in individual-leve
Original data	The data before SDC/anonymization methods were applied. Also called "raw data" or "untre
Outlier	An unusual value that is correctly reported but is not typical of the rest of the population. Ou
Perturbative methods	Anonymization methods that alter values slightly to limit disclosure risk by creating uncertainty
Population unique	The only record in the population with a particular set of characteristics, such that the individual
Post Randomization Method (PRAM)	Anonymization method for microdata in which the scores of a categorical variable are altered
Probabilistic methods	Anonymization methods that depend on a probability mechanism or a random number-gener
Privacy	Privacy is a concept that applies to data subjects while confidentiality applies to data. The co
Public Use File (PUF)	Type of release of microdata file, which is freely available to any user, for example on the int
Quasi-identifiers Quasi-identifiers	A set of variables that, in combination, can be linked to external information to re-identify re
Raw data	The data before SDC/anonymization methods were applied. Also called "original data" or "t
Recoding	Anonymization method for microdata in which groups of existing categories/values are repla
Record	A set of data derived from an object/unit of experiment, e.g., an individual (in individual-leve
Regression	A statistical process of measuring the relation between the mean value of one variable and co
- 6	The state of the s

Re-identification risk	See Disclosure risk
Release	Dissemination – the release to users of information obtained through a statistical activity. <sup>5</sup>
Respondents	Individuals or units of observation whose information/responses to a survey make up the data
Sample unique	The only record in the sample with a particular set of characteristics, such that the individual
Scientific Use File (SUF)	Type of release of microdata file, which is only available to selected researchers under contra
sdcMicro	An R based package authored by Templ, M., Kowarik, A. and Meindl, B. with tools for the a
sdcMicroGUI	A GUI for the R based sdcMicro package, which allows users to use the sdcMicro tools with
Sensitive variables	Sensitive or confidential variables are those whose values must not be discovered for any respectively.
Statistical Disclosure Control (SDC)	Statistical Disclosure Control techniques can be defined as the set of methods to reduce the r
Suppression	Data suppression involves not releasing information that is considered unsafe because it fails
Threshold	An established level, value, margin or point at which values that fall above or below it will de
Utility	Data utility describes the value of data as an analytical resource, comprising analytical comp
Untreated data	The data before SDC/anonymization methods were applied. Also called "raw data" or "original content of the data before SDC/anonymization methods were applied.
Variable	Any characteristic, number or quantity that can be measured or counted for each unit of obse

2.2. Glossary 9

http://www.nss.gov.au/nss/home.nsf/pages/Confidentiality+-+Glossary
 http://stats.oecd.org/glossary
 Australian Bureau of Statistics, http://www.nss.gov.au/nss/home.nsf/pages/Confidentiality+-+Glossary
 Australian Bureau of Statistics, http://www.nss.gov.au/nss/home.nsf/pages/Confidentiality+-+Glossary
 Australian Bureau of Statistics, http://www.nss.gov.au/nss/home.nsf/pages/Confidentiality+-+Glossary

http://stats.oecd.org/glossary
http://stats.oecd.org/glossary
OECD, http://stats.oecd.org/glossary

# STATISTICAL DISCLOSURE CONTROL (SDC): AN INTRODUCTION

#### 3.1 Need for SDC

A large part of the data collected by statistical agencies cannot be published directly due to privacy and confidentiality concerns. These concerns are both of legal and ethical nature. SDC seeks to treat and alter the data so that the data can be published or released without revealing the confidential information it contains, while, at the same time, limit information loss due to the anonymization of the data. In this guide, we discuss only disclosure control for microdata. Microdata are datasets that provide information on a set of variables for each individual respondent. Respondents can be natural persons, but also legal entities such as companies.

The aim of anonymizing microdata is to transform the datasets to achieve an "acceptable level" of disclosure risk. The level of acceptability of disclosure risk and the need for anonymization are usually at the discretion of the data producer and guided by legislation. These are formulated in the dissemination policies and programs of the data providers and based on considerations including "[...] the costs and expertise involved; questions of data quality, potential misuse and misunderstanding of data by users; legal and ethical matters; and maintaining the trust and support of respondents" (Dupriez and Boyko, 2010). There is a moral, ethical and legal obligation for the data producers to ensure that data provided by the respondents are used only for statistical purposes.

In some cases, the dissemination of microdata is a legal obligation, but, in most cases, the legislation will formulate restrictions. Thus, a country's legislative framework will shape its microdata dissemination policy. It is crucial for data producers to "ensure there is a sound legal and ethical base (as well as the technical and methodological tools) for protecting confidentiality. This legal and ethical base requires a balanced assessment between the public good of confidentiality protection on the one hand, and the public benefits of research on the other. A decision on whether or not to provide access might depend on the merits of specific research proposals and the credibility of the researcher, and there should be some allowance for this in the legal arrangements." (Dupriez and Boyko, 2010).

"Data access arrangements should respect the legal rights and legitimate interests of all stakeholders in the public research enterprise. Access to, and use of, certain research data will necessarily be limited by various types of legal requirements, which may include restrictions for reasons of:

- National security: data pertaining to intelligence, military activities, or political decision making may be classified and therefore subject to restricted access.
- Privacy and confidentiality: data on human subjects and other personal data are subject to restricted access under
  national laws and policies to protect confidentiality and privacy. However, anonymization or confidentiality
  procedures that ensure a satisfactory level of confidentiality should be considered by custodians of such data to
  preserve as much data utility as possible for researchers.
- Trade secrets and intellectual property rights: data on, or from, businesses or other parties that contain confidential information may not be accessible for research. (...)" (Dupriez and Boyko, 2010).

Box 1, extracted from Dupriez and Boyko (2010), provides several examples of statistical legislation on microdata release.

<sup>&</sup>lt;sup>1</sup> There is another strand of literature on the anonymization of tabular data, see e.g., Hundepool et al. (2012).

#### Box 1: Examples of statistical legislation on microdata release

Besides the legal and ethical concerns and codes of conducts of agencies producing statistics, SDC is important because it guarantees data quality and response rates in future surveys. If respondents feel that data producers are not protecting their privacy, they might not be willing to participate in future surveys. "[...] one incident, particularly if it receives strong media attention, could have a significant impact on respondent cooperation and therefore on the quality of official statistics" (ibid.). At the same time, if data users are unable to gain enough utility from the data due to excessive or inappropriate SDC protection, or are unable to access the data, then the large investment in producing the data will be lost.

## 3.2 The risk-utility trade-off in the SDC process

SDC is characterized by the trade-off between risk of disclosure and utility of the data for end users. The risk—utility scale extends between two extremes; (i) no data is released (zero risk of disclosure) and thus users gain no utility from the data, to (ii) data is released without any treatment, and thus with maximum risk of disclosure, but also maximum utility to the user (i.e., no information loss). The goal of a well-implemented SDC process is to find the optimal point where utility for end users is maximized at an acceptable level of risk. Figure 2.1 illustrates this trade-off. The triangle corresponds to the raw data. The raw data have no information loss, but generally have a disclosure risk higher than the acceptable level. The other extreme is the square, which corresponds to no data release. In that case there is no disclosure risk, but also no utility from the data for the users. The points in-between correspond to different choices of SDC methods and/or parameters for these methods applied to different variables. The SDC process looks for the SDC methods and the parameters for those methods and applies these in a way that reduces the risk sufficiently, while minimizing the information loss.

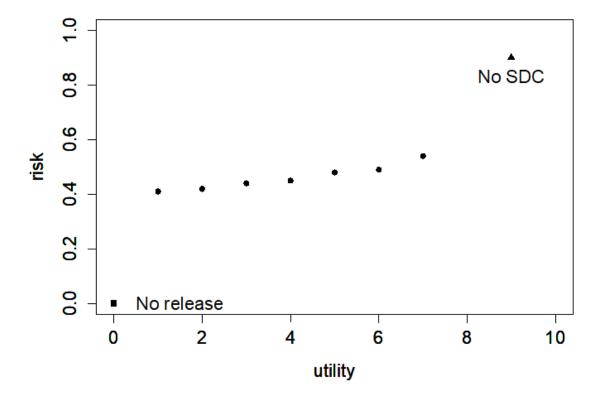


Figure 2.1: Risk-utility trade-off SDC cannot achieve total risk elimination, but can reduce the risk to an acceptable level. Any application of SDC

methods will suppress or alter values in the data and as such decrease the utility (i.e., result in information loss) when compared to the original data. A common thread that will be emphasized throughout this guide will be that the process of SDC should prioritize the goal of protecting respondents, while at the same time keeping the data users in mind to limit information loss. In general, the lower the disclosure risk, the higher the information loss and the lower the data utility for end-users.

In practice, choosing SDC methods is partially trial and error: after applying methods, disclosure risk and data utility are re-measured and compared to the results of other choices of methods and parameters. If the result is satisfactory, the data can be released. We will see that often the first attempt will not be the optimal one. The risk may not be sufficiently reduced or the information loss may be too high and the process has to be repeated with different methods or parameters until a satisfactory solution is found. Disclosure risk, data utility and information loss in the SDC context and how to measure them are discussed in subsequent chapters of this guide.

Again, it must be stressed that the level of SDC and methods applied depend to a large extent on the entire data release framework. For example, a key consideration is to whom and under what conditions the data are to be released (see also Chapter 3). If data are to be released as public use data, then the level of SDC applied will necessarily need to be higher than in the cases where data are released under license conditions to trusted users after careful vetting. With careful preparation, data may be released under both public and licensed versions. We discuss how this might be achieved later in the guide.

## 3.3 Release Types

This section discusses data release. Rather than rewriting work that has already been conducted through the World Bank and its partners at the IHSN, this section extracts from an excellent guide published by Dupriez and Boyko (2010).

The trade-off between risk and utility in the anonymization process depends greatly on who the users are<sup>2</sup> and under what conditions a microdata file is released. Generally, three types of data release methods are practiced and apply to different target groups.

- Public Use File (PUF): the data "are available to anyone agreeing to respect a core set of easy-to-meet conditions. Such conditions relate to what cannot be done with the data (e.g. the data cannot be sold), upon gaining access to the data. In some cases PUFs are disseminated with no conditions; often being made available on-line, [e.g. on the website of the statistical agency]. These data are made easily accessible because the risk of identifying individual respondents is considered minimal. Minimising the risk of disclosure involves eliminating all content that can identify respondents directly—for instance, names, addresses and telephone numbers. In addition this requires purging relevant indirect identifiers from the microdata file. These vary across survey designs, but commonly-suppressed indirect identifiers include geographical information below the sub-national level at which the sample is representative. Occasionally, certain records may be suppressed also from PUFs, as might variables characterised by extremely skewed distribution or outliers. However, in lieu of deleting entire records or variables from microdata files, alternative SDC methods can minimise the risk of disclosure while maximizing information content. Such methods include top-and-bottom coding, local suppression or using data perturbation techniques [(see Chapter 5 for an overview of anonymization methods)]. PUFs are typically generated from census data files using a sub-set [or sample] of records rather than the entire file and [from sample surveys, such as] household surveys." (Dupriez and Boyko, 2010).
- Scientific Use File (SUF) (also known as a licensed file, microdata under contract or research file): the "dissemination is restricted to users who have received authorization to access them after submitting a documented application and signing an agreement governing the data's use. While typically licensed files are also anonymised to ensure the risk of identifying individuals is minimised when used in isolation, they may still [potentially] contain identifiable data if linked with other data files. Direct identifiers such as respondents' names must be removed from a licensed dataset. The data files may, however, still contain indirect variables that could identify

3.3. Release Types

<sup>&</sup>lt;sup>2</sup> See Section 5 in Dupriez and Boyko (2010) as to who the users of microdata are and to whom microdata should be made available.

respondents by matching them to other data files such as voter lists, land registers or school records. When disseminating licensed files, the recommendation is to establish and sign an agreement between the data producer and external bona fide users – trustworthy users with legitimate need to access the data. Such an agreement should govern access and use of such microdata files<sup>3</sup>. Sometimes, licensing agreements are only entered into with users affiliated to an appropriate sponsoring institution. i.e., research centers, universities or development partners. It is further recommended that, before entering into a data access and use agreement, the data producer asks potential users to complete an application form to demonstrate the need to use a licensed file (instead of the PUF version, if available) for a stated statistical or research purpose" (Dupriez and Boyko, 2010). This also allows the data producer to learn which characteristics of the data are important for the users, which is valuable information for optimizing future anonymization processes.

• Microdata available in a controlled research data center (also known as data enclave): "Some files may be offered to users under strict conditions in a data enclave. This is a facility [(often on the premises of the data provider)] equipped with computers not linked to the internet or an external network and from which no information can be downloaded via USB ports, CD-DVD or other drives. Data enclaves contain data that are particularly sensitive or allow direct or easy identification of respondents. Examples include complete population census datasets, enterprise surveys and certain health related datasets containing highly-confidential information. Users interested in accessing a data enclave will not necessarily have access to the full dataset – only to the particular data subset they require. They will be asked to complete an application form demonstrating a legitimate need to access these data to fulfill a stated statistical or research purpose [...] The outputs generated must be scrutinised by way of a full disclosure review before release. Operating a data enclave may be expensive – it requires special premises and computer equipment. It also demands staff with the skills and time to review outputs before their removal from the data enclave in order to ensure there is no risk of disclosure. Such staff must be familiar with data analysis and be able to review the request process and manage file servers. Because of the substantial operating costs and technical skills required, some statistical agencies or other official data producers opt to collaborate with academic institutions or research centres to establish and manage data enclaves."

There are other data access possibilities besides these, such as teaching files, files for other specific purposes, remote execution or remote access. Obviously, the required level of protection depends on the type of release; a PUF file must be protected to a much larger extent than a SUF file, which in turn has to be protected more than a file which is only available in an on-site facility. Section 8.3 gives more guidance on the choice of the release type and its implications for the anonymization process. The same microdata set can be released in different ways for different users, e.g., as SUF and teaching file. Section 8.3 discusses the particular issues of multiple releases of one dataset.

The first step for any agency that wants to release data would be formulation of clear data dissemination policies for the release of microdata. We will see later that deciding on the level of anonymization needed will depend partly on knowing under what conditions the data will be released. Access policies and conditions provide the framework for the whole release process.

The following sections further specify the conditions under which microdata should be provided under different release types.

#### 3.4 Conditions for PUFs

"Generally, data regarded as public are open to anyone with access to an [National Statistical Office] (NSO) website. It is, however, normally good practice to include statements defining suitable uses for and precautions to be adopted in using the data. While these may not be legally binding, they serve to sensitise the user. Prohibitions such as attempts to link the data to other sources can be part of the 'use statement' to which the user must agree, on-line, before the data can be downloaded. [...] Dissemination of microdata files necessarily involves the application of rules or principles. [Box 2] below [taken from Dupriez and Boyko (2010)] shows basic principles normally applying to PUFs." (Dupriez and Boyko, 2010).

Box 2: Conditions for accessing and using PUFs

<sup>&</sup>lt;sup>3</sup> Appendix B provides an example of a blanket agreement.

#### 3.5 Conditions for SUFs

"For [SUFs], terms and conditions must include the basic common principles plus some additional ones applying to the researcher's organisation. There are two options: firstly, data are provided to a researcher or a team for a specific purpose; secondly, data are provided to an organization under a blanket agreement for internal use, e.g., to an international body or research agency. In both cases, the researcher's organisation must be identified, as must suitable representatives to sign the licence" (Dupriez and Boyko, 2010).

Access to a researcher or research team for a specific purpose

"If data are provided for an individual research project, the research team must be identified. This is covered by requiring interested users to complete a formal request to access the data (a model of such a request form is provided in Appendix 1 [in Dupriez and Boyko (2010)]). The conditions to obtain the data (see example in Box 3) will specify that the files will not be shared outside the organisation and that data will be stored securely. To the possible extent, the intended use of the data – including a list of expected outputs and the organisation's dissemination policy – must be identified. Access to licensed datasets is only granted when there is a legally-registered sponsoring agency, e.g., government ministry, university, research centre or national or international organization" (Dupriez and Boyko, 2010).

Box 3: Conditions for accessing and using SUFs

Blanket agreement to an organization

"In the case of a blanket agreement, where it is agreed the data can be used widely but securely within the receiving organisation, the licence should ensure compliance, with a named individual formally assuming responsibility for this. Each additional user must be made aware of the terms and conditions that apply to data files: this can be achieved by having to sign an affidavit. Where such an agreement exists, with security in place, it is not necessary for users to destroy the data after use" (Dupriez and Boyko, 2010). Appendix B provides an example of the formulation of such an agreement.

# 3.6 Conditions for microdata available in a controlled research data center

Access to microdata in research data centers is "used for particularly sensitive data or for more detailed data for which sufficient anonymisation to release them outside the NSO premises is not possible. These can be referred to also as data laboratories or research data centres. A [research data centre] may be located at the NSO headquarters or in major centres such as universities close to the research community. They are used to give researchers access to complete data files but without the risk of releasing confidential data. In a typical [research data centre], NSO staff supervise access and use of the data; the computers must not be able to communicate outside the [research data centre]; and the results obtained by the researchers must be screened for confidentiality by an NSO analyst before taken outside. A model of a data enclave access policy is provided in Appendix 2 [in Dupriez and Boyko (2010)], and a model of a data enclave access request form is in Appendix 3 [in Dupriez and Boyko (2010)]" (Dupriez and Boyko, 2010).

Research data centers "have the advantage of providing access to detailed microdata but the disadvantage of requiring researchers to work at a different location. And they are expensive to set up and operate. It is, however, quite likely that many countries have used on-site researchers as a way of providing access to microdata. These researchers are sworn in under the statistics' acts in the same way as regular NSO employees. This approach tends to favour researchers who live near NSO headquarters." (Dupriez and Boyko, 2010)

#### **Recommended Reading Material on Release Types**

Dupriez, O., & Boyko, E. (2010). *Dissemination of Microdata Files; Principles, Procedures and Practices*. International Household Survey Network (IHSN).

Readthedocs Documentation	

**CHAPTER** 

**FOUR** 

#### **MEASURING RISK**

## 4.1 Types of disclosure

Measuring disclosure risk is an important part of the SDC process: risk measures are used to judge whether a data file is safe enough for release. Before measuring disclosure risk, we have to define what type of disclosure is relevant for the data at hand. The literature commonly defines three types of disclosure; we take these directly from Lamb93 (see also HDFG12).

- Identity disclosure, which occurs if the intruder associates a known individual with a released data record.
  For example, the intruder links a released data record with external information, or identifies a respondent with extreme data values. In this case, an intruder can exploit a small subset of variables to make the linkage, and once the linkage is successful, the intruder has access to all other information in the released data related to the specific respondent.
- Attribute disclosure, which occurs if the intruder is able to determine some new characteristics of an individual
  based on the information available in the released data. Attribute disclosure occurs if a respondent is correctly reidentified and the dataset contains variables containing information that was previously unknown to the intruder.
  Attribute disclosure can also occur without identity disclosure. For example, if a hospital publishes data showing
  that all female patients aged 56 to 60 have cancer, an intruder then knows the medical condition of any female
  patient aged 56 to 60 in the dataset without having to identify the specific individual.
- **Inferential disclosure**, which occurs if the intruder is able to determine the value of some characteristic of an individual more accurately with the released data than would otherwise have been possible. For example, with a highly predictive regression model, an intruder may be able to infer a respondent's sensitive income information using attributes recorded in the data, leading to inferential disclosure.

SDC methods for microdata are intended to prevent identity and attribute disclosure. Inferential disclosure is generally not addressed in SDC in the microdata setting, since microdata is distributed precisely so that researchers can make statistical inference and understand relationships between variables. In that sense, inference cannot be likened to disclosure. Also, inferences are designed to predict aggregate, not individual, behavior, and are therefore usually poor predictors of individual data values.

#### 4.2 Classification of variables

For the purpose of the SDC process, we use the classifications of variables described in the following paragraphs (see Fig. 4.1 for an overview). The initial classification of variables into identifying and non-identifying variables depends on the way the variables can be used by intruders for re-identification (HDFG12; TeMK14):

• **Identifying variables:** these contain information that can lead to the identification of respondents and can be further categorized as:

- Direct identifiers reveal directly and unambiguously the identity of the respondent. Examples are names, passport numbers, social identity numbers and addresses. Direct identifiers should be removed from the dataset prior to release. Removal of direct identifiers is a straightforward process and always the first step in producing a safe microdata set for release. Removal of direct identifiers, however, is often not sufficient.
- Quasi-identifiers (or key variables) contain information that, when combined with other quasi-identifiers in the dataset, can lead to re-identification of respondents. This is especially the case when they can be used to match the information with other external information or data. Examples of quasi-identifiers are race, birth date, sex and ZIP/postal codes, which might be easily combined or linked to publically available external information and make identification possible. The combinations of values of several quasi-identifiers are called keys (see also *Levels of Risk*). The values of quasi-identifiers themselves often do not lead to identification (e.g. male/female), but a combination of several values of quasi-identifier can render a record unique (e.g. male, 14 years, married) and hence identifiable. It is not generally advisable to simply remove quasi-identifiers from the data to solve the problem. In many cases, they will be important variables for any sensible analysis. In practice, any variable in the dataset could potentially be used as a quasi-identifier. SDC addresses this by identifying variables as quasi-identifiers and anonymizing them while still maintaining the information in the dataset for release.
- Non-identifying variables are variables that cannot be used for re-identification of respondents. This could be
  because these variables are not contained in any other data files or other external sources and are not observable
  to an intruder. Non-identifying variables are nevertheless important in the SDC process, since they may contain
  confidential/sensitive information, which may prove damaging should disclosure occur as a result of identity
  disclosure based on identifying variables.

These classifications of variables depend partially on the availability of external datasets that might contain information that, when combined with the current data, could lead to disclosure. The identification and classification of variables as quasi-identifiers depends, amongst others, on the availability of information in external datasets. An important step in the SDC process is to define a list of possible disclosure scenarios based on how the quasi-identifiers might be combined with each other and information in external datasets and then treating the data to prevent disclosure. We discuss disclosure scenarios in more detail in *Disclosure scenarios*.

For the SDC process, it is also useful to further classify the quasi-identifiers into **categorical**, **continuous** and **semi-continuous** variables. This classification is important for determining the appropriate SDC methods for that variable, as well as the validity of risk measures.

- Categorical variables take values over a finite set, and any arithmetic operations using them are generally not meaningful or not allowed. Examples of categorical variables are gender, region and education level.
- **Continuous** variables can take on an infinite number of values in a given set. Examples are income, body height and size of land plot. Continuous variables can be transformed into categorical variables by constructing intervals (such as income bands).<sup>1</sup>
- Semi-continuous variables are continuous variables that take on values that are limited to a finite set. An example is age measured in years, which could take on values in the set {0, 1, ..., 100}. The finite nature of the values for these variables means that they can be treated as categorical variables for the purpose of SDC.<sup>2</sup>

Apart from these classifications of variables, the SDC process further classifies variables according to their sensitivity or confidentiality. Both quasi-identifiers and non-identifying variables can be classified as **sensitive** (or confidential) or **non-sensitive** (or non-confidential). This distinction is not important for direct identifiers, since direct identifiers are removed from the released data.

• Sensitive variables contain confidential information that should not be disclosed without suitable treatment using SDC methods to reduce disclosure risk. Examples are income, religion, political affiliation and variables

<sup>&</sup>lt;sup>1</sup> Recoding a continuous variable is sometimes useful in cases where the data contains only a few continuous variables. We will see in Section 3 that many methods used for risk calculation depend on whether the variables are categorical. We will also see that it is easier for the measurement of risk if the data contains only categorical or only continuous variables.

<sup>&</sup>lt;sup>2</sup> This is discussed in greater detail in the following sections. In cases where the number of possible values is large, recoding the variable, or parts of the set it takes values on, to obtain fewer distinct values is recommended.

concerning health. Whether a variable is sensitive depends on the context and country: a certain variable can be considered sensitive in one country and non-sensitive in another.

• Non-sensitive variables contain non-confidential information on the respondent, such as place of residence or rural/urban residence. The classification of a variable as non-sensitive, however, does not mean that it does not need to be considered in the SDC process. Non-sensitive variables may still serve as quasi-identifiers when combined with other variables or other external data.

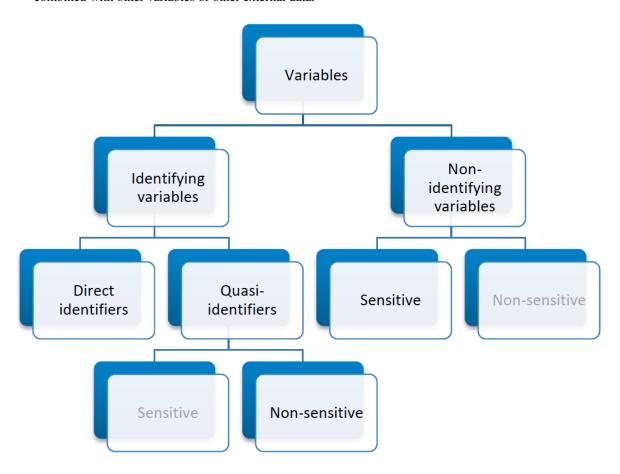


Fig. 4.1: Classification of variables

#### 4.3 Disclosure scenarios

Evaluation of disclosure risk is carried out with reference to the available data sources in the environment where the dataset is to be released. In this setting, disclosure risk is the possibility of correctly re-identifying an individual in the released microdata file by matching their data to an external file based on a set of quasi-identifiers. The risk assessment is done by identifying so-called disclosure or intrusion scenarios. A disclosure scenario describes the information potentially available to the intruder (e.g., census data, electoral rolls, population registers or data collected by private firms) to identify respondents and the ways such information can be combined with the microdata set to be released and used for re-identification of records in the dataset. Typically, these external datasets include direct identifiers. In that case, the re-identification of records in the released dataset leads to identity and, possibly, attribute disclosure. The main outcome of the evaluation of disclosure scenarios is the identification of a set of quasi-identifiers (i.e., key variables) that need to be treated during the SDC process (see Elliot et al., 2010).

An example of a disclosure scenario could be the spontaneous recognition of a respondent by a researcher. For

instance, while going through the data, the researcher recognizes a person with an unusual combination of the variables age and marital status. Of course, this can only happen if the person is well-known or is known to the researcher. Another example of a disclosure scenario for a publicly available file would be if variables in the data could be linked to a publically available electoral register. An intruder might try matching the entire dataset with individuals in the register. However, this might be difficult and take specialized expertise, or software, and other conditions have to be fulfilled. Examples are that the point in time the datasets were collected should approximately match and the content of the variables should be (nearly) identical. If these conditions are not fulfilled, exact matching is much less likely.

**Note:** Not all external data is necessarily in the public domain. Also privately owned datasets or datasets which are not released should be taken into consideration for determining the suitable disclosure scenario.

#### Info-box - Disclosure scenarios and different release types

A dataset can have more than one disclosure scenario. Disclosure scenarios also differ depending on the data access type that the data will be released under; for example, Public Use Files (PUF) or Scientific Use Files (SUF, also known as licensed) or in a data enclave. The required level of protection, the potential avenues of disclosure as well as the availability of other external data sources differ according to the access type under which the data will be released. For example, the user of a Scientific Use File (SUF) might be contractually restricted by an agreement as to what they are allowed to do with the data, whereas a Public Use File (PUF) might be freely available on the internet under a much looser set of conditions. PUFs will in general require more protection than SUFs and SUFs will require more protection than those files only released in an data enclave. Disclosure scenarios should be developed with all of this in mind.

The evaluation of disclosure risk is based on the quasi-identifiers, which are identified in the analysis of disclosure risk scenarios. The disclosure risk directly depends on the inclusion or exclusion of variables in the set of quasi-identifiers chosen. This step in the SDC process (making the choice of quasi-identifiers) should therefore be approached with great thought and care. We will see later, as we discuss the steps in the SDC process in more detail, that the first step for any agency is to undertake an exercise in which an inventory is compiled of all datasets available in the country. Both datasets released by the national statistical office and from other sources are considered and their availability to intruders as well as the variables included in these datasets is analyzed. It is this information that will serve as a key metric when deciding which variables to choose as potential identifiers, as well as dictate the level of SDC and methods needed.

#### 4.4 Levels of risk

With microdata from surveys and censuses, we often have to be concerned about disclosure at the individual or unit level, i.e., identifying individual respondents. Individual respondents are generally natural persons, but can also be units, such as companies, schools, health facilities, etc. Microdata files often have a hierarchical structure where individual units belong to groups, e.g., people belong to households. The most common hierarchical structure in microdata is the household structure in household survey data. Therefore, in this guide, we sometimes call disclosure risk for data with a hierarchical structure "household risk". The concepts, however, apply equally to establishment data and other data with hierarchical structures, such as school data with pupils and teachers or company data with employees.

We will see that this hierarchical structure is important to take into consideration when measuring disclosure risk. For hierarchical data, information collected at the higher hierarchical level (e.g., household level) would be the same for all individuals in the group belonging to that higher hierarchical level (e.g., household).<sup>3</sup> Some typical examples of variables that would have the same values for all members of the same higher hierarchical unit are, in the case of

<sup>&</sup>lt;sup>3</sup> Besides variables collected at the higher hierarchical level, also variables collected at the lower level but with no (or little) variation within the groups formed by the hierarchical structure should be treated as higher level variables. An example could be mother tongue, where most households are monolingual, but the variable is collected at the individual level.

households, those relating to housing and household income. These variables differ from survey to survey and from country to country.<sup>4</sup> This hierarchical structure creates a further level of disclosure risk for two reasons:

- 1. if one individual in the household is re-identified, the household structure allows for re-identification of the other household members in the same household,
- 2. values of variables for other household members that are common for all household members can be used for re-identification of another individual of the same household. This is discussed in more detail in *Household Risk*.

In the following pages, we first discuss risk measures used to evaluate disclosure risk in the absence of a hierarchical structure. This includes risk measures that seek to aggregate the individual risk for all individuals in the microdata file; the objective is to quantify a global disclosure risk measure for the file. We then discuss how risk measures change when taking the hierarchical structure of the data into account.

We will also discuss how risk measures differ for categorical and continuous key variables. For categorical variables, we will use the concept of uniqueness of combinations of values of quasi-identifiers (so-called "keys") used to identify individuals at risk. The concept of uniqueness, however, is not useful for continuous variables, since it is likely that all or many individuals will have unique values for that variable, by definition of a continuous variable. Risk measures for categorical variables are generally a priori measures, i.e., they can be evaluated before applying anonymization methods since they are based on the principle of uniqueness. Risk measures for continuous variables are a posteriori measures; they are based on comparing the microdata before and after anonymization and are, for example, based on the proximity of observations between the original and the treated (anonymized) datasets.

Files that are limited to only categorical or only continuous key variables are easiest for risk measurement. We will see in later sections that, in cases where both types of variables are present, recoding of continuous variables into categories is one approach to use to simplify the SDC process, but we will also see that from a utility perspective this may not be desirable. An example might be the use of income quintiles instead of the actual income variables. We will see that measuring the risk of disclosure based on the categorical and continuous variables separately is generally not a valid approach.

The risk measures discussed in the next section are based on several assumptions. In general, these measures rely on quite restrictive assumptions and will often lead to conservative risk estimates. These conservative risk measures may overstate the risk as they assume a worst-case scenario. Two assumptions should, however, be fulfilled for the risk measures to be valid and meaningful; the microdata should be a sample of a larger population (no census) and the sampling weights should be available. Special case: census data briefly discusses how to deal with census data.

#### 4.5 Individual risk

#### 4.5.1 Categorical key variables and frequency counts

The main focus of risk measurement for categorical quasi-identifiers is on identity disclosure. Measuring disclosure risk is based on the evaluation of the probability of correct re-identification of individuals in the released data. We use measures based on the actual microdata to be released. In general, the rarer a combination of values of the quasi-identifiers (i.e., key) of an observation in the sample, the higher the risk of identity disclosure. An intruder that tries to match an individual who has a relatively rare key within the sample data with an external dataset in which the same key exists will have a higher probability of finding a correct match than when a larger number of individuals share the same key. This can be illustrated with the following example that is illustrated in Table 4.1.

Table 4.1 shows values for 10 respondents for the quasi-identifiers "residence", "gender", "education level" and "labor status". In the data, we find seven unique combinations of values of quasi-identifiers (i.e., patterns or keys) of the four quasi-identifiers. Examples of keys are { 'urban', 'female', 'secondary incomplete', 'employed'} and { 'urban', 'female', 'primary incomplete', 'non-LF'}. Let  $f_k$  be the sample frequency of the  $k^{th}$  key, i.e., the number of individuals

4.5. Individual risk 21

<sup>&</sup>lt;sup>4</sup> Religion, for example, can be shared by all household members in some countries, whereas in other countries this variable is measured at the individual level and mixed-religion households exist.

in the sample with values of the quasi-identifiers that coincide with the  $k^{th}$  key. This would be 2 for the key {urban, female, secondary incomplete, employed}, since this key is shared by individuals 1 and 2 and 1 for the key ('urban', 'female', 'primary incomplete', 'non-LF'), which is unique to individual 3. By definition,  $f_k$  is the same for each record sharing a particular key.

The fewer the individuals with whom an individual shares his or her combination of quasi-identifiers, the more likely the individual is to be correctly matched in another dataset that contains these quasi-identifiers. Even when direct identifiers are removed from the dataset, that individual has a higher disclosure risk than others, assuming that their sample weights are the same. Table 4.1 reports the sample frequencies  $f_k$  of the keys for all individuals. Individuals with the same keys have the same sample frequency. If  $f_k = 1$ , this individual has a unique combination of values of quasi-identifiers and is called "sample unique". The dataset in Table 4.1 contains four sample uniques. Risk measures are based on this sample frequency.

No	Residence	Gender	Education level	Labor status	Weight	f_k	F_k	risk
1	Urban	Female	Secondary incomplete	Employed	180	2	360	0.0054
2	Urban	Female	Secondary incomplete	Employed	180	2	360	0.0054
3	Urban	Female	Primary incomplete	Non-LF	215	1	215	0.0251
4	Urban	Male	Secondary complete	Employed	76	2	152	0.0126
5	Rural	Female	Secondary complete	Unemployed	186	1	186	0.0282
6	Urban	Male	Secondary complete	Employed	76	2	152	0.0126
7	Urban	Female	Primary complete	Non-LF	180	1	180	0.0290
8	Urban	Male	Post-secondary	Unemployed	215	1	215	0.0251
9	Urban	Female	Secondary incomplete	Non-LF	186	2	262	0.0074
10	Urban	Female	Secondary incomplete	Non-LF	76	2	262	0.0074

Table 4.1: Example dataset showing sample frequencies, population frequencies and individual disclosure risk

In Listing 4.1, we show how to use the sdcMicro package to create a list of sample frequencies  $f_k$  for each record in a dataset. This is done by using the sdcMicro function freq(). A value of 2 for an observation means that in the sample, there is one more individual with exactly the same combination of values for the selected key variables. In Listing 4.1, the function freq() is applied to "sdcInitial", which is an sdcMicro object. Footnote<sup>5</sup> shows how to initialize the sdcMicro object for this example. For a complete discussion of sdcMicro objects as well as instructions on how to create sdcMicro objects, we refer to

Objects of class Objects of class \*sdcMicroObj\*. *sdcMicro* objects are used when doing SDC with *sdcMicro*. The

function freq() displays the sample frequency for the keys constructed on a defined set of quasi-identifiers. Listing 4.1 corresponds to the data in Table 4.1.

Listing 4.1: Calculating  $f_k$  using *sdcMicro* 

```
# Frequency of the particular combination of key variables (keys) for each record in the sample freq(sdcInitial, type = 'fk')
2 2 1 2 1 2 1 1 2 2
```

For sample data, it is more interesting to look at  $F_k$ , the population frequency of a combination of quasi-identifiers (key) k, which is the number of individuals in the population with the key that corresponds to key k. The population frequency is unknown if the microdata is a sample and not a census. Under certain assumptions, the expected value of the population frequencies can be computed using the sample design weight  $w_i$  (in a simple sample, this is the inverse

<sup>&</sup>lt;sup>5</sup> The code examples in this guide are based on *sdcMicro* objects. An *sdcMicro* object contains, amongst others, the data and identifies all the specified key variables. The code below creates a data.frame with the data from Table 4.1 and the *sdcMicro* objects "sdcInitial" used in most examples in this section.

of the inclusion probability) for each individual i

$$F_k = \sum_{i|key \ of \ individual \ i \ corresponds \ to \ key \ k} w_i$$

 $F_k$  is the sum of the sample weights of all records with the same key k. Hence, like f is,  $F_k$  is the same for each record with key k. The risk of correct re-identification is the probability that the key is matched to the correct individual in the population. Since every individual in the sample with key k corresponds to  $F_k$  individuals in the population, the probability of correct re-identification is  $1/F_k$ . This is the probability of re-identification in the worst-case scenario and can be interpreted as disclosure risk. Individuals with the same key have the same frequencies, i.e., the frequency of the key.

If  $F_k = 1$ , the key k is both a sample and a population unique and the disclosure risk would be 1. Population uniques are an important factor to consider when evaluating risk, and deserve special attention. Table 4.1 also shows  $F_k$  for the example dataset. This is further discussed in the case studies in Chapter 9.

Besides  $f_k$ , the sample frequency of key k (i.e., the number of individuals in the sample with the combination of quasiidentifiers corresponding to the combination specified in key k) and  $F_k$ , the estimated population frequency of key k, can be displayed in sdcMicro. Listing 4.2 illustrates how to return lists of length n of frequencies for all individuals. The frequencies are displayed for each individual and not for each key.

Listing 4.2: Calculating the sample and population frequencies using sdcMicro

```
# Sample frequency of individual's key
freq(sdcInitial, type = 'fk')
2 2 1 2 1 2 1 1 2 2

# Population frequency of individual's key
freq(sdcInitial, type = 'Fk')
3 360 360 215 152 186 152 180 215 262 262
```

In practice, this approach leads to conservative risk estimates, as it does not adequately take the sampling methods into account. In this case, the estimates of re-identification risk may be estimated too high. If this overestimated risk is used, the data may be overprotected (i.e., information loss will be higher than was necessary) when applying SDC

```
library (sdcMicro)
# Set up dataset
data <- as.data.frame(cbind(as.factor(c('Urban',</pre>
'Urban', 'Urban', 'Rural', 'Urban', 'Urban', 'Urban',
'Urban', 'Urban')),
as.factor(c('Female', 'Female', 'Female', 'Male',
'Female', 'Male', 'Female', 'Male', 'Female', 'Female')),
as.factor(c('Sec in', 'Sec in', 'Prim in', 'Sec com', 'Sec com', 'Sec com', 'Prim com', 'Post-sec
\hookrightarrow', 'Sec in', 'Sec in')),
as.factor(c('Emp', 'Emp', 'Non-LF', 'Emp', 'Unemp', 'Emp', 'Non-LF', 'Unemp', 'Non-LF', 'Non-LF'
→ ' ) ) ,
as.factor(c('yes', 'yes', 'yes', 'yes', 'no', 'no', 'yes', 'no', 'yes')),
c(180, 180, 215, 76, 186, 76, 180, 215, 186, 76)))
# Specify variable names
names(data) <- c('Residence', 'Gender', 'Educ', 'Lstat', 'Health', 'Weights')</pre>
# Set up sdcMicro object with specified quasi-identifiers and weight variable
sdcInitial <- createSdcObj(dat = data,</pre>
keyVars = c('Residence', 'Gender', 'Educ', 'Lstat'), weightVar = 'Weights')
```

4.5. Individual risk 23

measures. Instead, a Bayesian approach to risk measurement is recommended, where the posterior distribution of  $F_k$  is used (see e.g., Hundepool et al., 2012) to estimate an individual risk measure  $r_k$  for each key k.

The risk measure  $r_k$  is, as  $f_k$  and  $F_k$ , the same for all individuals sharing the same pattern of values of key variables and is referred to as individual risk. The values  $r_k$  can also be interpreted as the probability of disclosure for the individuals or as the probability for a successful match with individuals chosen at random from an external data file with the same values of the key variables. This risk measure is based on certain assumptions '6, which are strict and may lead to a relatively conservative risk measure. In sdcMicro, the risk measure  $r_k$  is automatically computed when creating an sdcMicro object and saved in the "risk" slot<sup>7</sup>. Listing 4.3 shows how to retrieve the risk measures using sdcMicro for our example. The risk measures are also presented in Table 4.1.

Listing 4.3: The individual risk slot in the sdcMicro object

```
sdcInitial@risk$individual risk fk Fk

[1,] 0.005424520 2 360
[2,] 0.005424520 2 360
[3,] 0.025096439 1 215
[4,] 0.012563425 2 152
[5,] 0.028247279 1 186
[6,] 0.012563425 2 152
[7,] 0.029010932 1 180
[8,] 0.025096439 1 215
[9,] 0.007403834 2 262
[10,] 0.007403834 2 262
```

The main factors influencing the individual risk are the sample frequencies  $f_k$  and the sampling design weights  $w_i$ . If an individual is at relatively high risk of disclosure, in our example this would be individuals 3, 5, 7 and 8 in Table 4.1 and Listing 4.3, the probability that a potential intruder correctly matches these individuals with an external data file is high **relative to the other individuals in the released data.** In our example, the reason for the high risk is the fact that these individuals are sample uniques ( $f_k = 1$ ). This risk is the worst-case scenario risk and does not imply that the person will be re-identified with certainty with this probability. For instance, if an individual included in the microdata is not included in the external data file, the probability for a correct match is zero. Nevertheless, the risk measure computed based on the frequencies will be positive.

#### 4.5.2 k-anonymity

The risk measure k-anonymity is based on the principle that, in a safe dataset, the number of individuals sharing the same combination of values (keys) of categorical quasi-identifiers should be higher than a specified threshold k. k-anonymity is a risk measure based on the microdata to be released, since it only takes the sample into account. An individual violates k-anonymity if the sample frequency count  $f_k$  for the key k is smaller than the specified threshold k. For example, if an individual has the same combination of quasi-identifiers as two other individuals in the sample, these individuals satisfy 3-anonymity but violate 4-anonymity. In the dataset in Table 4.1, six individuals satisfy 2-anonymity and four violate 2-anonymity. The individuals that violate 2-anonymity are sample uniques. The risk measure is the number of observations that violates k-anonymity for a certain value of k, which is

$$\sum_{i} I(f_k < k),$$

<sup>&</sup>lt;sup>6</sup> The assumptions for this risk measure are strict and the risk is estimated in many cases higher than the actual risk. Among other assumptions, it is assumed that all individuals in the sample are also included in the external file used by the intruder to match against. If this is not the case, the risk is much lower; if the individual in the released file is not contained in the external file, the probability of a correct match is zero. Other assumptions are that the files contain no errors and that both sets of data were collected simultaneously, i.e. they contain the same information. These assumptions will often not hold generally, but are necessary for computation of a measure. An example of a violation of the last assumptions is could occur if datasets are collected at different points in time and records have changed. This could happen when people move or change jobs and makes correct matching impossible. The assumptions assume a worst-case scenario.

<sup>&</sup>lt;sup>7</sup> See Section 7.5 for more information on slots and the *sdcMicro* object structure.

where I is the indicator function and i refers to the i<sup>th</sup> record. This is simply a count of the number of individuals with a sample frequency of their key lower than k. The count is higher for larger k, since if a record satisfies k-anonimity, it also satisfies (k+1)-anonimity. The risk measure k-anonymity does not consider the sample weights, but it is important to consider the sample weights when determining the required level of k-anonymity. If the sample weights are large, one individual in the dataset represents more individuals in the target population, the probability of a correct match is smaller, and hence the required threshold can be lower. Large sample weights go together with smaller datasets. In a smaller dataset, the probability to find another record with the same key is smaller than in a larger dataset. This probability is related to the number of records in the population with a particular key through the sample weights.

In *sdcMicro* we can display the number of observations violating a given k-anonymity threshold. In Listing 4.4, we use *sdcMicro* to calculate the number of violators for the thresholds k=2 and k=3. Both the absolute number of violators and the relative number as percentage of the number of individuals in the sample are given. In the example, four observations violate 2-anonimity and all 10 observations violate 3-anonymity.

Listing 4.4: Using the print() function to display observations violating k-anonymity

```
print(sdcInitial, 'kAnon')

Number of observations violating

- 2-anonymity: 4

- 3-anonymity: 10

------

Percentage of observations violating

- 2-anonymity: 40 %

- 3-anonymity: 100 %
```

For other levels of k-anonymity, it is possible to compute the number of violating individuals by using the sample frequency counts in the sdcMicro object. The number of violators is the number of individuals with sample frequency counts smaller than the specified threshold k. In Listing 4.5, we show an example of how to calculate any threshold for k using the already-stored risk measures available after setting up an sdcMicro object in k. k can be replaced with any required threshold. The choice of the required threshold that all individuals in the microdata file should satisfy depends on many factors and is discussed further in Section 4.3 on local suppression. In many institutions, typically required thresholds for k-anonymity are 3 and 5.

Listing 4.5: Computing k-anonymity violations for other values of k

```
sum(sdcInitial@risk$individual[,2] < k)</pre>
```

It is important to note that missing values ('NA's in  $R^8$ ) are treated as if they were any other value. Two individuals with keys {'Male', NA, 'Employed'} and {'Male', 'Secondary complete', 'Employed'} share the same key, and similarly, {'Male', NA, 'Employed'} and {'Male', 'Secondary incomplete', 'Employed'} also share the same key. Therefore, the missing value in the first key is first interpreted as 'Secondary complete', and then as 'Secondary incomplete'. This is illustrated in Table 4.2.

**Note:** The sample frequency of the third record is 3, since it is regarded to share its key both with the first and second record.

This principle is used when applying local suppression to achieve a certain level of k-anonymity (see Section 5.2.2) and is based on the fact that the value NA could replace any value.

4.5. Individual risk 25

<sup>&</sup>lt;sup>8</sup> In *sdcMicro* it is important to use the standard missing value code NA instead of other codes, such as 9999 or strings. In Chapter 6, we further discuss how to set other missing value codes to NA in *R*. This is necessary to ensure that the methods in *sdcMicro* function properly. When missing values have codes other than NA, the missing value codes are interpreted as a distinct factor level in the case of categorical variables.

Table 4.2: Example dataset to illustrate the effect of missing values on k-anonymity

No	Gender	Education level	Labor status	f_k
1	Male	Secondary complete	Employed	2
2	Male	Secondary incomplete	Employed	2
3	Male	NA	Employed	3

If a dataset satisfies k-anonymity, an intruder will always find at least k individuals with the same combination of quasi-identifiers. k-anonymity is often a necessary requirement for anonymization for a dataset before release, but is not necessarily a sufficient requirement. The k-anonymity measure is only based on frequency counts and does not take (differences in) sample weights into account. Often k-anonymity is achieved by first applying recoding and subsequently applying local suppression, and in some cases by microaggregation, before using other risk measures and disclosure methods to further reduce disclosure risk. These methods are discussed in Chapter 5.

#### 4.5.3 *l*-diversity

k-anonymity has been criticized for not being restrictive enough. Sensitive information might be disclosed even if the data satisfies k-anonymity. This might occur in cases where the data contains sensitive (non-identifying) categorical variables that have the same value for all individuals that share the same key. Examples of such sensitive variables are those containing information on an individual's health status. Table 4.3 illustrates this problem by using the same data as previously used, but adding a sensitive variable, "health". The first two individuals satisfy 2-anonymity for the key variables "residence", "gender", "education level" and "labor status". This means that an intruder will find at least two individuals when matching to the released microdata set based on those four quasi-identifiers. Nevertheless, if the intruder knows that someone belongs to the sample and has the key {'Urban', 'Female', 'Secondary incomplete' and 'Employed'}, with certainty the health status is disclosed ('yes'), because both observations with this key have the same value. This information is thus disclosed without the necessity to match exactly to the individual. This is not the case for the individuals with the key {'Urban', 'Male', 'Secondary complete', 'Employed'}. Individuals 4 and 6 have different values ('yes' and 'no') for "health", and thus the intruder would not gain information about the health status from this dataset by matching an individual to one of these individuals.

No Residence Gender Education level Labor status Health f k Fk **I-diversity** Urban Secondary incomplete 1 Female **Employed** yes 2 360 1 2 Urban Female Secondary incomplete Employed 2 360 1 yes 3 Urban Female Primary incomplete Non-LF 215 1 1 yes 4 2 2 Urban Male Secondary complete **Employed** 152 yes 5 Female Secondary complete Unemployed 1 1 Rural yes 186 6 Urban Male Secondary complete **Employed** 2 152 2 no 7 Non-LF 1 Urban Female Primary complete 1 180 8 Urban Male Post-secondary Unemployed 1 215 1 yes 9 Urban Female Secondary incomplete Non-LF no 2 262 2 10 Urban Female Secondary incomplete Non-LF 2 262 2 yes

Table 4.3: 1-diversity illustration

The concept of (distinct) l-diversity addresses this shortcoming of k-anonymity (see Machanavajjhala et al., 2007). A dataset satisfies l-diversity if for every key k there are at least l different values for each of the sensitive variables. In the example, the first two individuals satisfy only 1-diversity, individuals 4 and 6 satisfy 2-diversity. The required level of l-diversity depends on the number of possible values the sensitive variable can take. If the sensitive variable is a binary variable, the highest level if l-diversity that can be achieved is 2. A sample unique will always only satisfy 1-diversity.

To compute l-diversity for sensitive variables in sdcMicro, the function ldiversity() can be used. This is illustrated in

Listing 4.6. As arguments, we specify the names of the sensitive variables<sup>9</sup> in the file as well as a constant for recursive l-diversity, l0 and the code for missing values in the data. The output is saved in the "risk" slot of the sdcMicro object. The result shows the minimum, maximum, mean and quantiles of the l-diversity scores for all individuals in the sample. The output in Listing 4.6 reproduces the results based on the data in Table 4.3.

Listing 4.6: *l*-diversity function in *sdcMicro* 

```
# Computing 1-diversity
2
    sdcInitial <- ldiversity(obj = sdcInitial, ldiv_index = c("Health"), l_recurs_c = 2,...</pre>
    \rightarrowmissing = NA)
    # Output for 1-diversity
4
    sdcInitial@risk$ldiversitv
5
6
    L-Diversity Measures
    Min. 1st Qu. Median
10
          1.0
                                                    2.0
                    1.0
11
12
    # 1-diversity score for each record
13
    sdcInitial@risk$ldiversity[,'Health_Distinct_Ldiversity']
14
15
    [1] 1 1 1 2 1 2 1 1 2 2
```

l-diversity is useful if the data contains categorical sensitive variables that are not quasi-identifiers themselves. It is not possible to select quasi-identifiers to calculate the l-diversity. l-diversity has to be calculated for each sensitive variable separately.

# 4.6 Special Uniques Detection Algorithm (SUDA)

The previously discussed risk measures depend on identifying key variables for which there may be information available from other sources or other datasets, and which, when combined with the current data, may lead to reidentification. In practice, however, it might not always be possible to conduct an inventory of all available datasets and their variables and thus assess all known external linkages and risks.

To overcome this, an alternative heuristic measure based on special uniques has been developed to determine the riskiness of a record, which leads to a SUDA metric or score (see Elliot et al., 2002). These measures are based on the search for special uniques. To find these special uniques, algorithms, called SUDA (Special Uniqueness Detection Algorithm), have been developed. SUDA algorithms are based on the concept of special uniqueness, which is introduced in the next subsection. Since this is a heuristic approach, its performance is only tested in actual datasets, which is done in Elliot et al. (2002) for UK census data. These tests have shown that the performance of the algorithm leads to good risk estimates for these test datasets.

#### 4.6.1 Sample unique vs. special unique

The previous measures of risk focused on the uniqueness of the key of a record in the dataset. Table 4.4 reproduces the data from Table 4.1. The sample dataset has 10 records and four pre-determined quasi-identifiers {"Residence", "Gender", "Education level" and "Labor status"}. Given the four quasi-identifiers, we have seven distinct patterns in those key variables, or keys (e.g., {'Urban', 'Female', 'Secondary incomplete', 'Employed'}). The sample frequency

<sup>&</sup>lt;sup>9</sup> Alternatively, the sensitive variables can be specified when creating the *sdcMicro* object using the function createSdcObj() in the *sensibleVar* argument. This is further explained in Section 7.5. In that case, the argument *ldiv\_index* does not have to be specified in the ldiversity() function. and the variables in the *sensibleVar* argument will automatically be used to compute *l*-diversity.

<sup>&</sup>lt;sup>10</sup> Besides distinct l-diversity, there are other l-diversity methods: entropy and recursive. Distinct l-diversity is most commonly used.

counts of the first and second records equal 2, because the two records share the same pattern (i.e., {'Urban', 'Female', 'Secondary incomplete', 'Employed'}). Record 3 is a sample unique because it is the only individual in the sample who is a female living in an urban area who is employed without completing primary school. Similarly, records 5, 7 and 8 are sample uniques, because they possess distinct patterns with respect to the four key variables.

No	Residence	Gender	Education level	Labor status	Weight	f_k	F_k	risk
1	Urban	Female	Secondary incomplete	Employed	180	2	360	0.0054
2	Urban	Female	Secondary incomplete	Employed	180	2	360	0.0054
3	Urban	Female	Primary incomplete	Non-LF	215	1	215	0.0251
4	Urban	Male	Secondary complete	Employed	76	2	152	0.0126
5	Rural	Female	Secondary complete	Unemployed	186	1	186	0.0282
6	Urban	Male	Secondary complete	Employed	76	2	152	0.0126
7	Urban	Female	Primary complete	Non-LF	180	1	180	0.0290
8	Urban	Male	Post-secondary	Unemployed	215	1	215	0.0251
9	Urban	Female	Secondary incomplete	Non-LF	186	2	262	0.0074
10	Urban	Female	Secondary incomplete	Non-LF	76	2	262	0.0074

Table 4.4: Sample uniques and special uniques

In addition to the records 3, 5, 7 and 8 in Table 4.4 being sample uniques with respect to the key variable set {"Residence", "Gender", "Education level", "Labor status"}, we can find unique patterns in these records without even having to consider the complete set of key variables. For instance, a unique pattern can be found in record 5 when we look only at the variables "Education level" and "Labor status" ({'Secondary complete', 'Unemployed'}). While the values {'Secondary complete'} and {'Unemployed'} are not unique in the sample, the combination of them, {'Secondary complete', 'Unemployed'} makes record 5 unique. This variable subset is referred to as the Minimal Sample Unique (MSU) as any smaller subset of this set of variables is not unique (in this case {'Secondary complete'} and {'Unemployed'}). It is an MSU of size 2.[#foot33]\_ This holds as well for three other combinations in record 5, i.e., {'Female', 'Unemployed'} and {'Female', 'Secondary Complete'}, which are also MSUs of size 2 and {'Rural'} of size 1. In total, record 5 has four MSUs<sup>11</sup>. To determine if a set is an MSU of size k, we check whether it fulfills the minimal requirement. It suffices to check whether all subsets of size k-1 of the MSU are unique. If any of these subsets are also unique in the sample, the set found may be a sample unique, but violates the minimal requirement and is hence not an MSU. The unique subset of size k-1 could be a MSU. In our example, to determine if the MSU {'Secondary complete', 'Unemployed'} is a MSU, we checked as to whether its subsets {'Secondary complete'} and {'Unemployed'} were not unique in the sample. By definition, only sample uniques can be special uniques.

The SUDA algorithm identifies all the MSUs in the sample, which in turn are used to assign a SUDA score to each record. This score indicates how "risky" a record is. The potential risk of the records is determined based on two observations:

- The smaller the size of the MSU within a record (i.e., the fewer variables are needed to reach uniqueness), the greater the risk of the record
- The larger the number of MSUs possessed by a record, the greater the risk of the record

A record is defined as a special unique if it is a sample unique both on the complete set of quasi-identifiers (e.g., in the data in Table 4.4, the variables "Residence", "Gender", "Education level" and "Labor status") and simultaneously has at least one MSU (Elliot et al., 1998). Special uniques can be classified according to the number and size of subsets that are MSUs. Research has shown that special uniques are more likely to be population uniques than random uniques (Elliot et al., 2002) and are thus relevant for risk assessment.

<sup>11</sup> There are more combinations of quasi-identifiers that make record 5 unique (e.g., {'Rural', 'Female'} and {'Female', 'Secondary Complete', 'Unemployed'}. These combinations, however, are not considered MSUs because they do not fulfill the **minimal** subset requirement. They contain subsets that are MSUs.

#### 4.6.2 Calculating SUDA scores

The SUDA algorithm is used to search for MSUs in the data among the sample uniques to determine which sample uniques are also special uniques i.e., have subsets that are also unique (see Elliot et al., 2005). First the SUDA algorithm is used to identify the MSUs for each sample unique. To simplify the search and because smaller subsets are more important for disclosure risk, the search is limited to a maximum subset size. Subsequently, a score is assigned to each individual, which ranks the individuals according to their level of risk.

For each MSU of size k contained in a given record, a score is computed by  $\prod_{i=k}^{M} (ATT-i)$ , where M is the user-specified maximum size of MSUs<sup>12</sup>, and ATT is the total number of attributes or variables in the dataset. By definition, the smaller the size k of the MSU, the larger the score for the MSU, which reflects greater risk (see Elliot et al., 2005). The final SUDA score for each record is computed by adding the scores for each MSU in the record. In this way, records with more MSUs are assigned a higher SUDA score, which also reflects the higher risk. The SUDA score ranks the individuals according to their level of risk. The higher the SUDA score, the riskier the sample unique.

Calculating SUDA scores – a simplified example

To illustrate how SUDA scores are calculated, we compute the SUDA scores for the sample uniques in the data in Table 4.5, which replicates the data from Table 4.5. Record 5 contains four MSUs: {Rural} of size 1, and {'Secondary Complete', 'Unemployed'}, {'Female', 'Unemployed'} and {Female, Secondary Complete} of size 2. Suppose the maximum size of MSUs we search for in the data, M, is set at 3. Knowing that, ATT, the number of selected key variables in the dataset, is 4; the score assigned to {Rural} is computed by  $\prod_{i=1}^{3} (4-i) = 3*2*1 = 6$ ; and the score assigned to {Secondary complete, Unemployed}, {Female, Unemployed} and {Female, Secondary Complete} is  $\prod_{i=2}^{3} (4-i) = 2*1 = 2$ . The SUDA score for the fifth record in Table 4.5 is then 6+2+2+2=12, which is the sum of these four scores per MSU. The SUDA scores for the other sample uniques are computed accordingly<sup>13</sup>. The values that are in the MSUs in the sample uniques are shaded in Table 4.5. Records that are not sample uniques  $(f_k > 1)$  cannot be special uniques and are assigned the score 0.

No	Resi-	Gen-	Education level	Labor sta-	Weight	f_k	SUDA	DIS-
	dence	der		tus			score	SUDA
1	Urban	Female	Secondary incom-	Employed	180	2	0	0.0000
			plete					
2	Urban	Female	Secondary incom-	Employed	180	2	0	0.0000
			plete					
3	Urban	Female	Primary incomplete	Non-LF	215	1	6	0.0051
4	Urban	Male	Secondary complete	Employed	76	2	0	0.0000
5	Rural	Female	Secondary complete	Unem-	186	1	12	0.0107
				ployed				
6	Urban	Male	Secondary complete	Employed	76	2	0	0.0000
7	Urban	Female	Primary complete	Non-LF	180	1	6	0.0051
8	Urban	Male	Post-secondary	Unem-	215	1	10	0.0088
				ployed				
9	Urban	Female	Secondary incom-	Non-LF	186	2	0	0.0000
			plete					
10	Urban	Female	Secondary incom-	Non-LF	76	2	0	0.0000
			plete					

Table 4.5: Illustrating the calculation of SUDA and DIS-SUDA scores

To estimate record-level disclosure risks, SUDA scores can be used in combination with the Data Intrusion Simulation (DIS) metric (Elliot and Manning, 2003), a method for assessing disclosure risks for the entire dataset (i.e., file-level disclosure risks). Roughly speaking, the DIS-SUDA method distributes the file-level risk measure generated by the

<sup>12</sup> OECD, http://stats.oecd.org/glossary

<sup>&</sup>lt;sup>13</sup> The third record has one MSU, {'Primary incomplete'}; the seventh record has one MSU, {'Primary complete'}; and the eighth record has three MSUs, {'Urban, Unemployed'}, {'Male, Unemployed'} and {'Post-secondary'}.

DIS metric between records according to the SUDA scores of each record. This way, SUDA scores are calibrated against a consistent measure to produce the DIS-SUDA scores, which provide the record-level disclosure risk. These scores are used to compute the conditional probability that a unique match found by an intruder between the sample unique in the released microdata and an external data source is also a correct match, and hence a successful disclosure. The DIS-SUDA measure can be computed in *sdcMicro*. Since the DIS score is a probability, its values are in the interval [0, 1]. A full description of the DIS-SUDA method is provided by Elliot and Manning (2003).

Note that unlike the risk methods discussed earlier, the DIS-SUDA score does not fully account for the sampling weights. Risk measures based on the previous methods (i.e., negative binomial models) will in general have lower risks for those records with greater sampling weight, given the same sample frequency count, than those measured using DIS-SUDA. Therefore, instead of replacing the risk measures introduced in the previous section, the SUDA scores and DIS-SUDA approach should be used as a complementary method. As mentioned earlier, DIS-SUDA is particularly useful in situations where taking an inventory of all already available datasets and their variables is difficult.

#### 4.6.3 Application of SUDA, DIS-SUDA using sdcMicro

Both SUDA and DIS-SUDA scores can be computed using *sdcMicro* (Templ et al., 2014). Given that the search for MSUs with the SUDA algorithm can be computationally demanding, *sdcMicro* uses an improved SUDA2 algorithm, which more effectively locates the boundaries of the search space for MSUs (Manning et al., 2008).

SUDA scores can be calculated using the suda2() function in *sdcMicro*. It is important to specify the missing argument in suda2(). This should match the code for missing values in your dataset. In *R* this is most likely the *R* standard missing value, NA. We mention this because **the default missing value code in the sdcMicro suda2() function is -999 and will most likely need to be changed to 'NA' when using most R datasets.** The scores are saved in the risk slot of the *sdcMicro* object. The syntax in Listing 4.7 shows how to retrieve the output.

Listing 4.7: Evaluating SUDA scores

```
# Evaluating SUDA scores for the specified variables
    sdcInitial <- suda2(obj = sdcInitial, missing = NA)</pre>
2
    # The results are saved in the risk slot of the sdcMicro object
    # SUDA scores
5
    sdcInitial@risk$suda2$score
6
    [1] 0.00 0.00 1.75 0.00 3.25 0.00 1.75 2.75 0.00 0.00
    # DIS-SUDA scores
10
    sdcInitial@risk$suda2$disScore
11
12
    [1] 0.000000000 0.000000000 0.005120313 0.000000000 0.010702061
13
    [6] 0.000000000 0.005120313 0.008775093 0.000000000 0.000000000
14
15
    # Summary of DIS-SUDA scores
    sdcInitial@risk$suda2
17
18
    Dis suda scores table:
19
20
    thresholds number
21
            > 0
                        6
    1
22
           > 0.1
                        4
23
24
    3
           > 0.2
                        0
            > 0.3
                        0
25
           > 0.4
                        0
26
    6
           > 0.5
                        0
```

(continues on next page)

(continued from previous page)

```
28 7 > 0.6 0
29 8 > 0.7 0
30 - - - - - - - - - -
```

To compare DIS scores before and after applying SDC methods, it may be useful to use histograms or density plots of these scores. Listing 4.8 shows how to generate histograms of the SUDA scores summarized in Listing 4.7. The histogram is shown in Fig. 4.2. All outputs relate to the data used in the example. In our case, we have not applied any SDC method to the data yet and thus have only the plots for the initial values. Typically, after applying SDC methods, one would recalculate the SUDA scores and compare them to the original values. One way to quickly see the differences would be to rerun these visualizations and compare them to the base for risk changes.

Listing 4.8: Histogram and density plots of DIS-SUDA scores

```
# Plot a histogram of disScore
hist(sdcInitial@risk$suda2$disScore, main = 'Histogram of DIS-SUDA scores')

# Density plot
density <- density(sdcInitial@risk$suda2$disScore)
plot(density, main = 'Density plot of DIS-SUDA scores')</pre>
```

# Histogram of DIS-SUDA scores

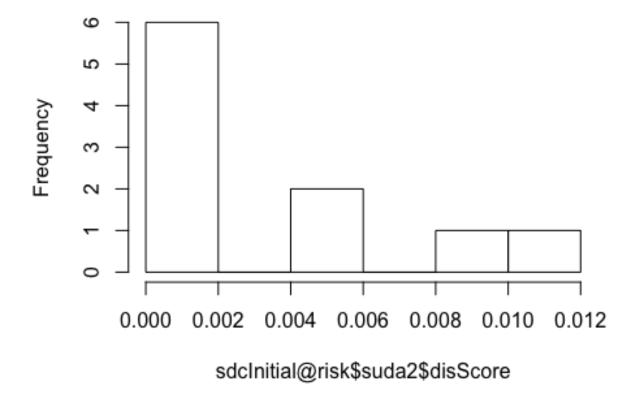


Fig. 4.2: Visualizations of DIS-SUDA scores

#### 4.7 Risk measures for continuous variables

The principle of rareness or uniqueness of combinations of quasi-identifiers (keys) is not useful for continuous variables, because it is likely that all or many individuals will have unique keys. Therefore, other approaches are exploited for measuring the disclosure risk of continuous variables. These methods are based on uniqueness of the values in the neighborhood of the original values. The uniqueness is defined in different ways: in absolute terms (interval measure) or relative terms (record linkage). Most measures are a posteriori measures: they are evaluated after anonymization of the raw data, compare the treated data with the raw data and evaluate for each individual the distance between the values in the raw and the treated data. This means that these methods are not useful for identifying individuals at risk within the raw data, but rather show the distance/difference between the dataset before and after anonymization and can therefore be interpreted as evaluation of the anonymization method. For that reason, they resemble the information loss measures discussed in Chapter 6. Finally, risk measures for continuous quasi-identifiers are also based on outlier detection. Outliers play an important role in the re-identification of these records.

#### 4.7.1 Record linkage

Record linkage is an a posteriori method that evaluates the number of correct linkages when linking the perturbed values with the original values. The linking algorithm is based on the distance between the original and the perturbed values (i.e., distance-based record linkage). The perturbed values are matched with the closest individual. It is important to note that this method does not give information on the initial risk, but is rather a measure to evaluate the perturbation algorithm (i.e., it is designed to indicate the level of uncertainty introduced into the variable by counting the number of records that could be correctly matched).

Record linkage algorithms differ with respect to which distance measure is used. When a variable has very different scaling than other continuous variables in the dataset, rescaling the variables before using record linkage is recommended. Very different scales may lead to undesired results when measuring the multivariate distance between records based on several continuous variables. Since these methods are based on both the raw data and treated data, examples of their applications require the introduction of SDC methods and are therefore postponed to the case studies in Chapter 9.

Besides distance-based record linkage, another method for linking is probabilistic record linkage (see Domingo-Ferrer and Torra, 2003). The literature shows, however, that results from distance-based record linkage are better than the results from probabilistic record linkage. Individuals in the treated data that are linked to the correct individuals in the raw data are considered at risk of disclosure.

#### 4.7.2 Interval measure

Successful application of an SDC method should result in perturbed values that are considered not too close to their initial values; if the value is relatively close, re-identification may be relatively easy. In the application of interval measures, intervals are created around each perturbed value and then a determination is made as to whether the original value of that perturbed observation is contained in this interval. Values that are within the interval around the initial value after perturbation are considered too close to the initial value and hence unsafe and need more perturbation. Values that are outside of the intervals are considered safe. The size of the intervals is based on the standard deviation of the observations and a scaling parameter. This method is implemented in the function dRisk() in sdcMicro. Listing 4.9 shows how to print or display the risk value computed by sdcMicro by comparing the income variables before and after anonymization. "sdcObj" is an sdcMicro object and "compExp" is a vector containing the names of the income variables. The size of the intervals is k times the standard deviation, where k is a parameter in the function dRisk(). The larger k, the larger the intervals are, and hence the larger the number of observations within the interval constructed around their original values and the higher the risk measure. The result 1 indicates that all (100 percent) the observations are outside the interval of 0.1 times the standard deviation around the original values.

Listing 4.9: Example with the function dRisk()

```
dRisk(obj = sdcObj@origData[,compExp], xm = sdcObj@manipNumVars[,compExp], k = 0.1)
[1] 1
```

For most values, this is a satisfactory approach. It is not a sufficient measure for outliers, however. After perturbation, outliers will stay outliers and are easily re-identifiable, even if they are sufficiently far from their initial values. Therefore, outliers should be treated with caution.

### 4.7.3 Outlier detection

Outliers are important for measuring re-identification risk in continuous microdata. Continuous data are often skewed, especially right-skewed. This means that there are a few outliers with very high values relative to the other observations of the same variable. Examples are income in household data, where only few individuals/households may have very high incomes, or turnover data for firms that are much larger than other firms in the sample are. In cases like these, even if these values are perturbed, it may still be easy to identify these outliers, since they will stay the largest values even after perturbation. (The perturbation will have created uncertainty as to the exact value, but because the value started out so much further away from other observations, it may still be easy to link to the high-income individual or very large firm.). Examples would be the only doctor in a geographical area with a high income or one single large firm in one industry type. Therefore, identifying outliers in continuous data is an important step when identifying individuals at high risk. In practice, identifying the values of a continuous variable that are larger than a predetermined p%-percentile might help identify outliers, and thus units at greater risk of identification. The value of p depends on the skewness of the data.

We can calculate the p%-percentile of a continuous variable in R and show the individuals who have income larger than this percentile. Listing 4.10 provides an illustration for the  $90^{th}$  percentile.

Listing 4.10: Computing 90 % percentile of variable income

```
# Compute the 90 % percentile for the variable income*
perc90 <- quantile(file[,'income'], 0.90, na.rm = TRUE)

# Show the ID of observations with values for income larger than the 90 % percentile
file[(file[, 'income'] >= perc90), 'ID']
```

A second approach for outlier detection is a posteriori measure comparing the treated and raw data. An interval is constructed around the perturbed values as described in the previous section. If the original values fall into the interval around the perturbed values, the perturbed values are considered unsafe since they are too close to the original values. There are different ways to construct such intervals, such as rank-based intervals and standard deviation-based intervals. Templ and Meindl (2008) propose a robust alternative for these intervals. They construct the intervals based on the squared Robust Mahalanobis Distance (RMD) of the individual values. The intervals are scaled by the RMD such that outliers obtain larger intervals and hence need to have a larger perturbation in order to be considered safe than values that are not outliers. This method is implemented in *sdcMicro* in the function dRiskRMD(), which is an extension of the dRisk() function. This method is illustrated in the case studies in Chapter 9.

# 4.8 Global risk

To construct one aggregate risk measure at the global level for the complete dataset, we can aggregate the measures for risk at the individual level in several ways. Global risk measures should be used with caution: behind an acceptable global risk can hide some very high-risk records that are compensated by many low risk records.

4.8. Global risk 33

### 4.8.1 Mean of individual risk measures

A straightforward way of aggregating the individual risk measures is taking the mean of all individuals in the sample, which is equal to summing over all keys in the sample if multiplied by the sample frequencies of these keys and dividing by the sample size n:

$$R_1 = \frac{1}{n} \sum_{i} r_k = \frac{1}{n} \sum_{k} f_k r_k$$

 $r_k$  is the individual risk of key k that the i<sup>th</sup> individual shares (see Section 4.5.1). This measure is reported as global risk in *sdcMicro*, is stored in the "risk" slot and can be displayed as shown in Listing 4.11. It indicates that the average re-identification probability is 0.01582 or 0.1582 %.

Listing 4.11: Computation of the global risk measure

```
# Global risk (average re-identification probability)
sdcInitial@risk$global$risk

[1] 0.01582
```

The global risk in the example data in Table 4.1 is 0.01582, which is the expected proportion of all individuals in the sample that could be re-identified by an intruder. Another way of expressing the global risk is the number of expected re-identifications,  $n * R_1$ , which is in the example 10 \* 0.01582. The expected number of re-identifications is also saved in the *sdcMicro* object. Listing 4.12 shows how to display this.

**Note:** This global risk measure should be used with caution. The average risk can be relatively low, but a few individuals could have a very high probability of re-identification.

An easy way to check for this is to look at the distribution of the individual risk values or the number of individuals with risk values above a certain threshold, as shown in the next section.

Listing 4.12: Computation of expected number of re-identifications

```
# Global risk (expected number of reidentifications)
sdcInitial@risk$global$risk_ER

[1] 0.1582
```

# 4.8.2 Count of individuals with risks larger than a certain threshold

All individuals belonging to the same key have the same individual risk,  $r_k$ . Another way of expressing the total risk in the sample is the total number of observations that exceed a certain threshold of individual risk. Setting the threshold can be absolute (e.g., all those individuals who have a disclosure risk higher than 0.05 or 5%) or relative (e.g., all those individuals with risks higher than the upper quartile of individual risk). Listing 4.13 shows how, using R, one would count the number of observations with an individual re-identification risk higher than 5%. In the example, no individual has a higher disclosure risk than 0.05.

Listing 4.13: Number of individuals with individual risk higher than the threshold 0.05

```
sum(sdcInitial@risk$individual[,1] > 0.05)

[1] 0
```

These calculations can then be used to treat data for individuals whose risk values are above a predetermined threshold. We will see later that there are methods in *sdcMicro*, such as localSupp(), that can be used to suppress values of certain key variables for those individuals with risk above a specified threshold. This is explained further in Section 5.2.2.

## 4.9 Household risk

In many social surveys, the data have a hierarchical structure where an individual belongs to a higher-level entity (see Section 4.4). Typical examples are households in social surveys or pupils in schools. Re-identification of one household member can lead to re-identification of the other household members, too. It is therefore easy to see that if we take the household structure into account, the re-identification risk is the risk that at least one of the household members is re-identified.

$$r^h = P(A_1 \cup A_2 \cup ... \cup A_J) = 1 - \prod_{j=1}^J 1 - P(A_j),$$

where  $A_j$  is the event that the  $j^{th}$  member of the household is re-identified and  $P(A_j) = r_k$  is the individual disclosure risk of the  $jt^h$  member. For example, if a household member has three members with individual disclosure risks based on their respective keys 0.02, 0.03 and 0.03, respectively, the household risk is

$$1 - (1 - 0.02)(1 - 0.03)(1 - 0.03)) = 0.078$$

The hierarchical or household risk cannot be lower than the individual risk, and the household risk is always the same for all household members. The household risk should be used in cases where the data contain a hierarchical structure, i.e., where a household structure is present in the data. Using *sdcMicro*, if a household identifier is specified (in the argument *hhId* in the function createSdcObj()) while creating an *sdcMicro* object, the household risk will automatically be computed. Listing 4.14 shows how to display these risk measures.

Listing 4.14: Computation of household risk and expected number of re-identifications

```
# Household risk
cdcInitial@risk$global$hier_risk

# Household risk (expected number of reidentifications
cdcInitial@risk$global$hier_risk_ER
```

**Note:** The size of a household is an important identifier itself, especially for large households. Suppression of the actual size variable (e.g., number of household members), however, does not suffice to remove this information from the dataset, as a simple count of the household members for a particular household will allow reconstructing this variable as long as a household ID is in the data, which allows assigning individuals to households. We flag this for the reader's attention as it is important. Further discussion on approaches to the SDC process that take into account the household structure where it exists can be found in Section\*\* **5.4.** 

#### **Recommended Reading Material on Risk Measurement**

Elliot, Mark J, Anna Manning, Ken Mayes, John Gurd, and Michael Bane. 2005. "SUDA: A Program for Detecting Special Uniques." *Joint UNECE/Eurostat Work Session on Statistical Data Confidentiality*. Geneva.

Hundepool, Anco, Josep Domingo-Ferrer, Luisa Franconi, Sarah Giessing, Eric Schulte Nordholt, Keith Spicer, and Peter Paul de Wolf. 2012. *Statistical Disclosure Control*. Chichester: John Wiley & Sons Ltd. doi:10.1002/9781118348239.

Lambert, Diane. 1993." Measures of Disclosure Risk and Harm." Journal of Official Statistics 9(2): 313-331.

4.9. Household risk 35

#### **Readthedocs Documentation**

Machanavajjhala, Ashwin, Daniel Kifer, Johannes Gehrke, and Muthuramakrishnan Venkitasubramaniam. 2007. "L-diversity: Privacy Beyond K-anonymity." *ACM Trans. Knowl. Discov. Data* 1 (Article 3) (1556-4681). doi:10.1145/1217299.1217302. http://www.truststc.org/pubs/465/L%20Diversity%20Privacy.pdf. Accessed October 5, 2015.

Templ, Matthias, Bernhard Meindl, Alexander Kowarik, and Shuang Chen. 2014. "Introduction to Statistical Disclosure Control (SDC)." http://www.ihsn.org/home/sites/default/files/resources/ihsn-working-paper-007-Oct27.pdf. August 1. Accessed November 13, 2014.

**CHAPTER** 

**FIVE** 

## ANONYMIZATION METHODS

This chapter describes the SDC methods most commonly used. All methods are implementable in *R* by using the *sdcMicro* package. We discuss for every method for what type of data the method is suitable, both in terms of data characteristics and type of data. Furthermore, options such as specific parameters for each method are discussed as well as their impacts. These findings are meant as guidance but should be used with caution, since every dataset has different characteristics and our findings may not always address your particular dataset. The last three sections of this chapter are on the anonymization of variables and datasets with particular characteristics that deserve special attention. Section 5.4 deals with for anonymizing geographical data, such as GPS coordinates, Section 5.5 discusses the anonymization of data with a hierarchical structure (household structure) and Section 5.6 describes the peculiarities of dealing with and releasing census microdata.

To determine which anonymization methods are suitable for specific variables and/or datasets, we begin by presenting some classifications of SDC methods.

# 5.1 Classification of SDC methods

SDC methods can be classified as **non-perturbative** and **perturbative** (see Hundepool et al., 2012).

- **Non-perturbative methods** reduce the detail in the data by generalization or suppression of certain values (i.e., masking) without distorting the data structure.
- **Perturbative methods** do not suppress values in the dataset but perturb (i.e., alter) values to limit disclosure risk by creating uncertainty around the true values.

Both non-perturbative and perturbative methods can be used for categorical and continuous variables.

We also distinguish between **probabilistic** and **deterministic** SDC methods.

- **Probabilistic methods** depend on a probability mechanism or a random number-generating mechanism. Every time a probabilistic method is used, a different outcome is generated. For these methods it is often recommended that a seed be set for the random number generator if you want to produce replicable results.
- **Deterministic methods** follow a certain algorithm and produce the same results if applied repeatedly to the same data with the same set of parameters.

SDC methods for microdata intend to prevent identity and attribute disclosure. Different SDC methods are used for each type of disclosure control. Methods such as recoding and suppression are applied to quasi-identifiers to prevent identity disclosure, whereas top coding a quasi-identifier (e.g., income) or perturbing a sensitive variable prevent attribute disclosure.

As this practice guide is written around the use of the *sdcMicro* package, we discuss only SDC methods that are implemented in the *sdcMicro* package or can be easily implemented in *R*. These are the most commonly applied methods from the literature and used in most agencies experienced in using these methods. Table 5.1 gives an overview

<sup>&</sup>lt;sup>1</sup> We also show code examples in R, which are drawn from findings we gathered by applying these methods to a large collection of datasets.

of the SDC methods discussed in this chapter, their classification, types of data to which they are applicable and their function names in the *sdcMicro* package.

Table 5.1: SDC methods and corresponding functions in sdcMicro

Method	Classification of SDC method	Data Type	Function in sdcMicro
	metriod		
Global recoding	non-perturbative, determinitic	continuous and categori-	globalRecode, groupVars
		cal	
Top and bottom cod-	non-perturbative, determinitic	continuous and categori-	topBotCoding
ing		cal	
Local suppression	non-perturbative, determinitic	categorical	localSuppression, local-
			Supp
PRAM	perturbative, probabilistic	categorical	pram
Micro aggregation	perturbative, probabilistic	continuous	microaggregation
Noise addition	perturbative, probabilistic	continuous	addNoise
Shuffling	perturbative, probabilistic	continuous	shuffle
Rank swapping	perturbative, probabilistic	continuous	rankSwap

# 5.2 Non-perturbative methods

## 5.2.1 Recoding

Recoding is a deterministic method used to decrease the number of distinct categories or values for a variable. This is done by combining or grouping categories for categorical variables or constructing intervals for continuous variables. Recoding is applied to all observations of a certain variable and not only to those at risk of disclosure. There are two general types of recoding: global recoding and top and bottom coding.

#### Global recoding

Global recoding combines several categories of a categorical variable or constructs intervals for continuous variables. This reduces the number of categories available in the data and potentially the disclosure risk, especially for categories with few observations, but also, importantly, it reduces the level of detail of information available to the analyst. To illustrate recoding, we use the following example. Assume that we have five regions in our dataset. Some regions are very small and when combined with other key variables in the dataset, produce high re-identification risk for some individuals in those regions. One way to reduce risk would be to combine some of the regions by recoding them. We could, for example, make three groups out of the five, call them 'North', 'Central' and 'South' and re-label the values accordingly. This reduces the number of categories in the variable region from five to three. **NOTE:** Any grouping should be some logical grouping and not a random joining of categories. Examples would be grouping districts into provinces, municipalities into districts, or clean water categories together. Grouping all small regions without geographical proximity together is not necessarily the best option from the utility perspective. Table 5.2 illustrates this with a very simplified example dataset. Before recoding, three individuals have distinct keys, whereas after recoding (grouping 'Region 1' and 'Region 2' into 'North', 'Region 3' into 'Central' and 'Region 4' and 'Region 5' into 'South'), the number of distinct keys is reduced to four and the frequency of every key is at least two, based on the three selected quasi-identifiers. The frequency counts of the keys  $f_k$  are shown in the last column of Table 5.2. An intruder would find at least two individuals for each key and cannot distinguish any more between individuals 1-3, individuals 4 and 6, individuals 5 and 7 and individuals 8 – 10, based on the selected key variables.

Table 5.2: Illustration of effect of recoding on frequency counts of keys

Before recoding					
Individual	Region	Gender	Religion	f_k	
1	Region 1	Female	Catholic	1	
2	Region 2	Female	Catholic	2	
3	Region 2	Female	Catholic	2	
4	Region 3	Female	Protestant	2	
5	Region 3	Male	Protestant	1	
6	Region 3	Female	Protestant	2	
7	Region 3	Male	Protestant	2	
8	Region 4	Male	Muslim	2	
9	Region 4	Male	Muslim	2	
10	Region 5	Male	Muslim	1	
After recodi	ng		,		
Individual	Region	Gender	Religion	f_k	
1	North	Female	Catholic	3	
2	North	Female	Catholic	3	
3	North	Female	Catholic	3	
4	Central	Female	Protestant	2	
5	Central	Male	Protestant	2	
6	Central	Female	Protestant	2	
7	Central	Male	Protestant	2	
8	South	Male	Muslim	3	
9	South	Male	Muslim	3	
10	South	Male	Muslim	3	

Recoding is commonly the first step in an anonymization process. It can be used to reduce the number of unique combinations of values of key variables. This generally increases the frequency counts for most keys and reduces the risk of disclosure. The reduction in the number of possible combinations is illustrated in Table 5.3 with the quasi-identifiers "region", "marital status" and "age". Table 5.3 shows the number of categories of each variable and the number of theoretically possible combinations, which is the product of the number of categories of each quasi-identifier, before and after recoding. "Age" is interpreted as a semi-continuous variable and treated as a categorical variable. The number of possible combinations and hence the risk for re-identification are reduced greatly by recoding. One should bear in mind that the number of possible combinations is a theoretical number; in practice, these may include very unlikely combinations such as age = 3 and marital status = widow and the actual number of combinations in a dataset may be lower.

Table 5.3: Illustration of the effect of recoding on the theoretically possible number of combinations an a dataset

Number of categories	Region	Marital status	Age	Possible combination s
before recoding	20	8	100	16,000
after recoding	6	6	15	540

The main parameters for global recoding are the size of the new groups, as well as defining which values are grouped together in new categories. **NOTE:** Care should be taken to choose new categories in line with the data use of the end users and to minimize information loss as a result of recoding. We illustrate this with three examples:

Age variable: The categories of age should be chosen so that they still allow data users to make calculations relevant for the subject being studied. For example, if indicators need to be calculated for children of school going ages 6 – 11 and 12 – 17, and age needs to be grouped to reduce risk, then care should be taken to create age intervals that still allow the calculations to be made. A satisfactory grouping could be, for example, 0 –

5, 6 - 11, 12 - 17, etc., whereas a grouping 0 - 10, 11 - 15, 16 - 18 would destroy the data utility for these users. While it is common practice to create intervals (groups) of equal width (size), it is also possible (if data users require this) to recode only part of the variables and leave some values as they were originally. This could be done, for example, by recoding all ages above 20, but leaving those below 20 as they are. If SDC methods other than recoding will be used later or in a next step, then care should be taken when applying recoding to only part of the distribution, as this might increase the information loss due to the other methods, since the grouping does not protect the ungrouped variables. Partial recoding followed by suppression methods such as local suppression may, for instance, leads to a higher number of suppressions than desired or necessary in case the recoding is done for the entire value range (see the next section on local suppression). In the example above, the number of suppressions of values below 20 will likely be higher than for values in the recoded range. The disproportionately high number of suppressions in this range of values that are not recoded can lead to higher utility loss for these groups.

- Geographic variables: If the original data specify administrative level information in detail, e.g., down to municipality level, then potentially those lower levels could be recoded or aggregated into higher administrative levels, e.g., province, to reduce risk. In doing so, the following should be noted: Grouping municipalities into abstract levels that intersect different provinces would make data analysis at the municipal or provincial level challenging. Care should be taken to understand what the user requires and the intention of the study. If a key component of the survey is to conduct analysis at the municipal level, then aggregating up to provincial level could damage the utility of the data for the user. Recoding should be applied if the level of detail in the data is not necessary for most data users and to avoid an extensive number of suppressions when using other SDC methods subsequently. If the users need information at a more detailed level, other methods such as perturbative methods might provide a better solution than recoding.
- Toilet facility: An example of a situation where a high level of detail might not be necessary and recoding may do very little harm to utility is the case of a detailed household toilet facility variable that lists responses for 20 types of toilets. Researchers may only need to distinguish between improved and unimproved toilet facilities and may not require the exact classification of up to 20 types. Detailed information of toilet types can be used to re-identify households, while recoding to two categories improved and unimproved facilities reduces the re-identification risk and in this context, hardly reduces data utility. This approach can be applied to any variable with many categories where data users are not interested in detail, but rather in some aggregate categories. Recoding addresses aggregation for the data users and at the same time protects the microdata. Important is to take stock of the aggregations used by data users.

Recoding should be applied only if removing the detailed information in the data will not harm most data users. If the users need information at a more detailed level, then recoding is not appropriate and other methods such as perturbative methods might work better.

In *sdcMicro* there are different options for global recoding. In the following paragraphs, we give examples of global recoding with the functions groupVars() and globalRecode(). The function groupVars() is generally used for categorical variables and the function globalRecode() for continuous variables. Finally, we discuss the use of rounding to reduce the detail in continuous variables.

Recoding a categorical variable using the sdcMicro function groupVars()

Assume that an object of class *sdcMicro* was created, which is called "sdcInitial" (see Section 7.5 how to create objects of class *sdcMicro*). In Example 5.1, the variable "sizeRes" has four different categories: 'capital, large city', 'small city', town', and 'countryside'). The first three are recoded or regrouped as 'urban' and the category 'countryside' is renamed 'rural'. In the function arguments, we specify the categories to be grouped (before) and the names of the categories after recoding (after). It is important that the vectors "before" and "after" have the same length. Therefore, we have to repeat 'urban' three times in the "after" vector to match the three different values that are recoded to 'urban'. **NOTE: the function groupVars() works only for variables of class factor.** We refer to Section 7.4 on classes in *R* and how to change the class of a variable.

Example 5.1: Using the sdcMicro function groupVars() to recode a categorical variable

<sup>&</sup>lt;sup>2</sup> Here the *sdcMicro* object "sdcIntial" contains a dataset with 2,500 individuals and 103 variables. We selected five quasi-identifiers: "sizeRes", "age", "gender", "region", and "ethnicity".

```
# Frequencies of sizeRes before recoding
table(sdcInitial@manipKeyVars$sizeRes)
## capital, large city
                        small city
                                                     town
                                                                  countryside
                                 310
                                                        146
##
                                                                         1358
# Recode urban
sdcInitial <- groupVars(obj = sdcInitial, var = c("sizeRes"), before = c("capital,...</pre>
→large city", "small city", "town"), after = c("urban", "urban", "urban"))
# Recode rural
sdcInitial <- groupVars(obj = sdcInitial, var = c("sizeRes"), before = c(</pre>
→"countryside"), after = c("rural"))
# Frequencies of sizeRes before recoding
table (sdcInitial@manipKeyVars$sizeRes)
## urban rural
## 1142 1358
```

Figure 5.1 illustrates the effect of recoding the variable "sizeRes" and show respectively the frequency counts before and after recoding. We see that the number of categories has reduced from 4 to 2 and the small categories ('small city' and 'town') have disappeared.

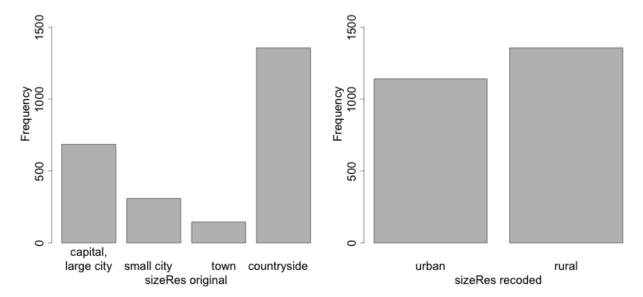


Figure 5.1 Effect of recoding – frequency counts before and after recoding

Recoding a continuous variable using the sdcMicro function: globalRecode()

Global recoding of numerical (continuous) variables can be achieved in *sdcMicro* by using the function globalRecode(), which allows specifying a vector with the break points between the intervals. Recoding a continuous variable changes it into a categorical variable. One can additionally specify a vector of labels for the new categories. By default, the labels are the intervals, e.g., "(0, 10]". Example 5.2 shows how to recode the variable age in 10-year intervals for age values between 0 and 100. **NOTE: Values that fall outside the specified intervals are assigned a missing value** (NA). Therefore, the intervals should cover the entire value range of the variable.

Example 5.2: Using the *sdcMicro* function globalRecode() to recode a continuous variable (age)

```
sdcInitial \leftarrow globalRecode(sdcInitial, column = c('age'), breaks = 10 * c(0:10))
# Frequencies of age after recoding
table (sdcInitial@manipKeyVars$age)
     (0,10] (10,20] (20,30] (30,40]
                                             (40,50]
                                                        (50,60]
                                                                  (60,70]
                                                                            (70, 80]
                                                                                      (80,90]
\hookrightarrow (90, 1001
##
        462
                   483
                             344
                                       368
                                                 294
                                                            214
                                                                      172
                                                                                 94
               3
→26
```

Figure 5.2 shows the effect of recoding the variable "age".

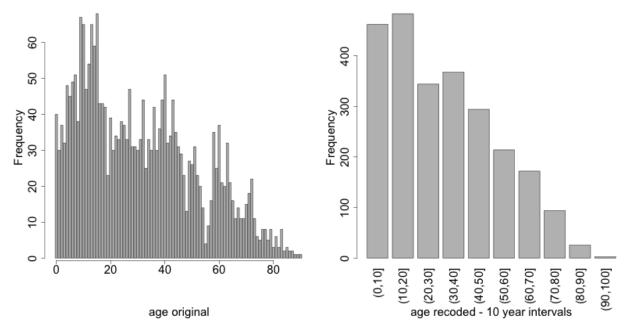


Figure 5.2 Age variable before and after recoding

Instead of creating intervals of equal width, we can also create intervals of unequal width. This is illustrated in Example 5.3, where we use the age groups 1-5, 6-11, 12-17, 18-21, 22-25, 26-49, 50-64 and 65+. In this example, this is a useful step, since even after recoding in 10-year intervals, the categories with high age values have low frequencies. We chose the intervals by respecting relevant school age and employment age values (e.g., retirement age is 65 in this example) such that the data can still be used for common research on education and employment. Figure 5.3 shows the effect of recoding the variable "age".

Example 5.3: Using globalRecode() to create intervals of unequal width

```
\leftrightarrow 25, 49, 65, 100))
# Frequencies of age after recoding
table(sdcInitial@manipKeyVars$age)
##
     (0, 51)
           (5,111)
                 (11, 17)
                         (17, 211)
                                 (21, 251)
                                        (25, 491)
                                               (49,65) (65,100)
                                                  350
##
      192
             317
                     332
                            134
                                    142
                                           808
                                                          185
```

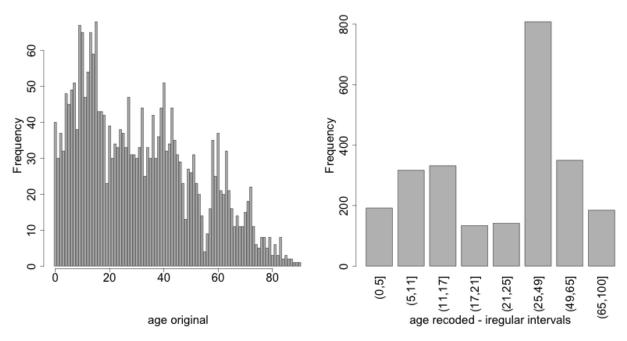


Figure 5.3 Age variable before and after recoding

Caution about using the globalRecode() function in *sdcMicro*: In the current implementation of *sdcMicro*, the intervals are defined as **left-open**. In mathematical terms, this means that, in our example, age 0 is excluded from the specified intervals. In interval notation, this is denoted as (0, 5] (as in x-axis labels in Figure 5.2 and Figure 5.3 graph labels for the recoded variable). The interval (0, 5] is interpreted as from 0 to 5 and does not include 0, but does include 5. *R* recodes values that are not contained in any of the intervals as missing (NA). This implementation would set in our example all age values 0 (children under 1 year) to missing and could potentially mean a large data loss. The globalRecode() function allows only constructing intervals, which are left-open. This may not be a desirable result and the loss of the zero ages from the data is clearly problematic for a real-world dataset.

To construct **right-open** intervals, e.g., in our example, for age intervals [0,14), [15, 65), [66, 100), we present two alternatives for global recoding:

• A work-around for semi-continuous variables<sup>3</sup> that would allow for the globalRecode() to be used would be subtracting a small number from the boundary intervals, thus allowing the desired intervals to be created. In the following example, subtracting 0.1 from each interval forces globalRecode() to include 0 in the lowest interval and allow for breaks where we want them. We set the upper interval boundary to be larger than the maximum value for the "age" variable. We can use the option *labels* to define clear labels for the new categories. This is illustrated in Example 5.4.

Example 5.4: Constructing right-open intervals for semi-continuous variables using built-in *sdcMicro* function globalRecode()

```
sdcInitial <- globalRecode(sdcInitial, column = \mathbf{c}('age'), breaks = \mathbf{c}(-0.1, 14.9, 64.9, \leftrightarrow 99.9), labels = \mathbf{c}('[0,15)', '[15,65)', '[65,100)'))
```

• It is also possible to use *R* code to manually recode the variables without using *sdcMicro* functions. When using the built-in *sdcMicro* functions, the change in risk after recoding is automatically recalculated, but if recoded manually it is not. In this case, we need to take an extra step and recalculate the risk after manually changing the variables in the *sdcMicro* object. This approach is also valid for continuous variables and is illustrated in Example 5.5.

Example 5.5: Constructing intervals for semi-continuous and continuous variables using manual recoding in R

<sup>&</sup>lt;sup>3</sup> This approach works only for semi-continuous variables, because in the case of continuous variables, there might be values that are between the lower interval boundary and the lower interval boundary minus the small number. For example, using this for income, we would have an interval (9999, 19999] and the value 9999.5 would be misclassified as belonging to the interval [10000, 19999].

```
# Group age 0-14
sdcInitial@manipKeyVars$age[sdcInitial@manipKeyVars$age >= 0 &
sdcInitial@manipKeyVars$age < 15] <- 0

# Group age 15-64
sdcInitial@manipKeyVars$age[sdcInitial@manipKeyVars$age >= 15 &
sdcInitial@manipKeyVars$age < 65] <- 1

# Group age 65-100
sdcInitial@manipKeyVars$age[sdcInitial@manipKeyVars$age >= 65 &
sdcInitial@manipKeyVars$age <= 100] <- 2

# Add labels for the new values
sdcInitial@manipKeyVars$age <-ordered(sdcInitial@manipKeyVars$age,
levels = c(0,1,2), labels = c("0-14", "15-64", "65-100"))

# Recalculate risk after manual manipulation
sdcInitial <- calcRisks(sdcInitial)</pre>
```

### Top and bottom coding

Top and bottom coding are similar to global recoding, but instead of recoding all values, only the top and/or bottom values of the distribution or categories are recoded. This can be applied only to ordinal categorical variables and (semi-)continuous variables, since the values have to be at least ordered. Top and bottom coding is especially useful if the bulk of the values lies in the center of the distribution with the peripheral categories having only few observations (outliers). Examples are age and income; for these variables, there will often be only a few observations above certain thresholds, typically at the tails of the distribution. The fewer the observations within a category, the higher the identification risk. One solution could be grouping the values at the tails of the distribution into one category. This reduces the risk for those observations, and, importantly, does so without reducing the data utility for the other observations in the distribution.

Deciding where to apply the threshold and what observations should be grouped requires:

- Reviewing the overall distribution of the variable to identify at which point the frequencies drop below the desired number of observations and identify outliers in the distribution. Figure 5.4 shows the distribution of the age variable and suggests 65 (red vertical line) for the top code age.
- Taking into account the intended use of the data and the purpose for which the survey was conducted. For example, if the data are typically used to measure labor force participation for those aged 15 to 64, then top and bottom coding should not interfere with the categories 15 to 64. Otherwise the analyst would find it impossible to create the desired measures for which the data were intended. In the example, we consider this and code all age larger than 64.

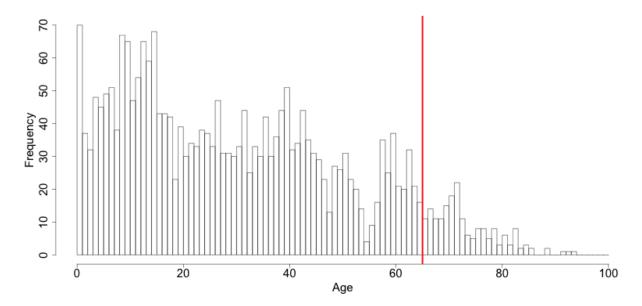


Figure 5.4: Utilizing the frequency distribution of variable age to determine threshold for top coding

Top and bottom coding can be easily done with the function topBotCoding() in *sdcMicro*. Top coding and bottom coding cannot be done simultaneously in *sdcMicro*. Example 5.6 illustrates how to recode values of age higher than 64 and values of age lower than 5; 65 and 5 replace the values respectively. To construct several top or bottom coding categories, e.g., age 65 – 80 and higher than age 80, one can use the groupVars() function in *sdcMicro* or manual recoding as described in the previous subsection.

Example 5.6: Top coding and bottom coding in sdcMicro using topBotCoding() function

#### Rounding

Rounding is similar to grouping, but used for continuous variables. Rounding is useful to prevent exact matching with external data sources. In addition, it can be used to reduce the level of detail in the data. Examples are removing decimal figures or rounding to the nearest 1,000.

The next section discusses the method local suppression. Recoding is often used before local suppression to reduce the number of necessary suppressions.

### **Recommended Reading Material on Recoding**

Hundepool, Anco, Josep Domingo-Ferrer, Luisa Franconi, Sarah Giessing, Rainer Lenz, Jane Naylor, Eric Schulte Nordholt, Giovanni Seri, and Peter Paul de Wolf. 2006. *Handbook on Statistical Disclosure Control*. ESSNet SDC. http://neon.vb.cbs.nl/casc/handbook.htm.

Hundepool, Anco, Josep Domingo-Ferrer, Luisa Franconi, Sarah Giessing, Eric Schulte Nordholt, Keith Spicer, and Peter Paul de Wolf. 2012. *Statistical Disclosure Control*. Chichester: John Wiley & Sons Ltd.

doi:10.1002/9781118348239.

Templ, Matthias, Bernhard Meindl, Alexander Kowarik, and Shuang Chen. 2014. Statistical Disclosure Control (SDCMicro). http://www.ihsn.org/home/software/disclosure-control-toolbox. (accessed November 13, 2014).

De Waal, A.G., and Willenborg, L.C.R.J. 1999. *Information loss through global recoding and local suppression*. Netherlands Official Statistics, 14:17-20, 1999. Special issue on SDC

# 5.2.2 Local suppression

It is common in surveys to encounter values for certain variables or combinations of quasi-identifiers (keys) that are shared by very few individuals. When this occurs, the risk of re-identification for those respondents is higher than the rest of the respondents (see k-anonymity in Section 4.5.2). Often local suppression is used after reducing the number of keys in the data by recoding the appropriate variables. Recoding reduces the number of necessary suppressions as well as the computation time needed for suppression. Suppression of values means that values of a variable are replaced by a missing value (NA in R). Section 4.5.2 on k-anonymity discusses how missing values influence frequency counts and k-anonymity. It is important to note that not all values for all individuals of a certain variable are suppressed, which would be the case when removing a direct identifier, such as "name"; only certain values for a particular variable and a particular respondent or set of respondents are suppressed. This is illustrated in the following example and Table 5.4.

Table 5.4 presents a dataset with seven respondents and three quasi-identifiers. The combination {'female', 'rural', 'higher'} for the variables "gender", "region" and "education" is an unsafe combination, since it is unique in the sample. By suppressing either the value 'female' or 'higher', the respondent cannot be distinguished from the other respondents anymore, since that respondent shares the same combination of key variables with at least three other respondents. Only the value in the unsafe combination of the single respondent at risk is suppressed, not the values for the same variable of the other respondents. The freedom to choose which value to suppress can be used to minimize the total number of suppressions and hence the information loss. In addition, if one variable is very important to the user, we can choose not to suppress values of this variable, unless strictly necessary. In the example, we can choose between suppressing the value 'female' or 'higher' to achieve a safe data file; we chose to suppress 'higher'. This choice should be made taking into account the needs of data users. In this example we find "gender" more important than "education".

l suppression illustration -	

Variable	Before local suppression			able Before local suppression After local suppression		
ID	Gender	Region	Educat ion*	Gender	Region	Educat ion
1	female	rural	higher	female	rural	
						NA/miss
						ing
						5
2	male	rural	higher	male	rural	higher
3	male	rural	higher	male	rural	higher
4	male	rural	higher	male	rural	higher
5	female	rural	lower	female	rural	lower
6	female	rural	lower	female	rural	lower
7	female	rural	lower	female	rural	lower

Since continuous variables have a high number of unique values (e.g., income in dollars or age in years), k-anonymity and local suppression are not suitable for continuous variables or variables with a very high number of categories. A possible solution in those cases might be to first recode to produce fewer categories (e.g., recoding age in 10-year intervals or income in quintiles). Always keep in mind, though, what effect any recoding will have on the utility of the data.

<sup>&</sup>lt;sup>5</sup> In R suppressed values are recoded NA, the standard missing value code.

The sdcMicro package includes two functions for local suppression: localSuppression() and localSupp(). The function localSuppression() is most commonly used and allows the use of suppression on specified quasi-identifiers to achieve a certain level of k-anonymity for these quasi-identifiers. The algorithm used seeks to minimize the total number of suppressions while achieving the required k-anonymity threshold. By default, the algorithm is more likely to suppress values of variables with many different categories or values, and less likely to suppress variables with fewer categories. For example, the values of a geographical variable, with 12 different areas, are more likely to be suppressed than the values of the variable "gender", which has typically only two categories. If variables with many different values are important for data utility and suppression is not desired for them, it is possible to rank variables by importance in the localSuppression() function and thus specify the order in which the algorithm will seek to suppress values within quasi-identifiers to achieve k-anonymity. The algorithm seeks to apply fewer suppressions to variables of high importance than to variables with lower importance. Nevertheless, suppressions in the variables with high importance might be inevitable to achieve the required level of k-anonymity.

In Example 5.7, local suppression is applied to achieve the k-anonymity threshold of 5 on the quasi-identifiers "gender", "region", "religion", "age" and "ethnicity" Without ranking the importance of the variables, the value of the variable "age" is more likely to be suppressed, since this is the variable with most categories. The variable "age" has 10 categories after recoding. The variable "gender" is least likely to be suppressed, since it has only two different values: 'male' and 'female'. The other variables have 4 ("sizeRes"), 2 ("region"), and 8 ("ethnicity") categories. After applying the localSuppression() function, we display the number of suppressions per variable with the built-in print() function with the option 'ls' for the local suppression output. As expected, the variable "age" has most suppressions (80). In fact, only the variable "ethnicity" of the other variables also needed suppressions (8) to achieve the k-anonymity threshold of 5. The variable "ethnicity" is the variable with the second highest number of suppressions. Subsequently, we undo and redo local suppression on the same data and reduce the number of suppressions on "age" by specifying the importance vector with high importance (little suppression) on the quasi-identifier "age". We also assign importance to the variable "gender". This is done by specifying an importance vector. The values in the importance vector can range from 1 to k, the number of quasi-identifiers. In our example k is equal to 5. Variables with lower values in the importance vectors have high importance and, when possible, receive fewer suppressions than variables with higher values.

To assign high importance to the variables "age" and "gender", we specify the importance vector as c(5, 1, 1, 5, 5), with the order according to the order of the specified variables in the *sdcMicro* object. The effect is clear: there are no suppressions in the variables "age" and "gender". For that, the other variables, especially "sizeRes" and "ethnicity", received many more suppressions. The total number of suppressed values has increased from 88 to 166. **NOTE: Fewer suppressions in one variable increase the number of necessary suppressions in other variables (cf. Example 5.7).** Generally, the total number of suppressed values needed to achieve the required level of k-anonymity increases when specifying an importance vector, since the importance vector prevents to use the optimal suppression pattern. The importance vector should be specified only in cases where the variables with many categories play an important role in data utility for the data users<sup>7</sup>.

Example 5.7: Application of local suppression with and without importance vector

```
# local suppression without importance vector
sdcInitial <- localSuppression(sdcInitial, k = 5)</pre>
print(sdcInitial, 'ls')
       KeyVar | Suppressions (#) | Suppressions (%)
##
      sizeRes |
                                 0 1
                                                 0.000
                                80 |
##
                                                  3.200
          age |
                                                  0.000
##
       gender |
                                 0 1
                                 0 1
                                                  0.000
##
       region |
                                 8 1
                                                  0.320
   ethnicity |
# Undoing the supressions
                                                                               (continues on next page)
```

<sup>&</sup>lt;sup>6</sup> Here the *sdcMicro* object "sdcIntial" contains a dataset with 2,500 individuals and 103 variables. We selected five quasi-identifiers: "sizeRes", "age", "gender", "region", and "ethnicity".

<sup>&</sup>lt;sup>7</sup> This can be assessed with utility measures.

(continued from previous page)

```
sdcInitial <- undolast(sdcInitial)</pre>
# Local suppression with importance vector to avoid suppressions in the first.
→ (gender) and fourth (age) variables
sdcInitial \leftarrow localSuppression(sdcInitial, importance = c(5, 1, 1, 5, 5), k = 5)
print(sdcInitial, 'ls')
       KeyVar | Suppressions (#) | Suppressions (%)
                                87 /
##
      sizeRes |
                                                 3.480
##
                                 0 1
                                                 0.000
          age |
##
                                 0 1
                                                 0.000
       gender |
                                17 /
##
       region |
                                                 0.680
    ethnicity |
                                62 |
                                                 2.480
```

Figure 5.5 demonstrates the effect of the required k-anonymity threshold and the importance vector on the data utility by using several labor market-related indicators from an I2D2 $^8$  dataset before and after anonymization. Figure 5.5 displays the relative changes as a percentage of the initial value after re-computing the indicators with the data to which local suppression was applied. The indicators are the proportion of active females and males, and the number of females and males of working age. The values computed from the raw data were, respectively, 68%, 12%, 8,943 and 9,702. The vertical line at 0 is the benchmark of no change. The numbers indicate the required k-anonymity threshold (3 or 5) and the colors indicate the importance vector: red (no symbol) is no importance vector, blue (with \* symbol) is high importance on the variable with the employment status information and dark green (with + symbol) is high importance on the age variable.

A higher k-anonymity threshold leads to greater information loss (i.e., larger deviations from the original values of the indicators, the 5's are further away from the benchmark of no change than the corresponding 3's) caused by local suppression. Reducing the number of suppressions on the employment status variable by specifying an importance vector does not improve the indicators. Instead, reducing the number of suppressions on age greatly reduces the information loss. Since specific age groups have a large influence on the computation of these indicators (the rare cases are in the extremes and will be suppressed), high suppression rates on age distort the indicators. It is generally useful to compare utility measures (see Chapter 6) to specify the importance vector, since the effects can be unpredictable.

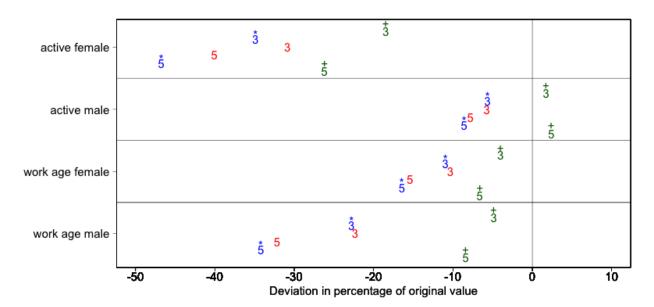


Figure 5.5: Changes in labor market indicators after anonymization of I2D2 data

<sup>&</sup>lt;sup>8</sup> I2D2 is a dataset with data related to the labor market.

The threshold of k-anonymity to be set depends on several factors, which are amongst others: 1) the legal requirements for a safe data file; 2) other methods that will be applied to the data; 3) the number of suppressions and related information loss resulting from higher thresholds; 4) the type of variable; 5) the sample weights and sample size; and 6) the release type (see Chapter 3). Commonly applied levels for the k-anonymity threshold are 3 and 5.

Table 5.5 illustrates the influence of the importance vector and k-anonymity threshold on the running time, global risk after suppression and total number of suppressions required to achieve this k-anonymity threshold. The dataset contains about 63,000 individuals. The higher the k-anonymity threshold, the more suppressions are needed and the lower the risk after local suppression (expected number of re-identifications). In this particular example, the computation time is shorter for higher thresholds. This is due the higher number of necessary suppressions, which reduces the difficulty of the search for an optimal suppression pattern.

The age variable is recoded in five-year intervals and has 20 age categories. This is the variable with the highest number of categories. Prioritizing the suppression of other variables leads to a higher total number of suppressions and a longer computation time.

Threshold Id	Importance	Total number of	Number of sup-	Global risk	Running
k-anony mity	nce vector**	suppressi ons	pressi ons age	measure	time (hours)
3	none (default)	6,676	5,387	293.0	11.8
3	employmen t	7,254	5,512	356.5	13.1
	status				
3	age variable	8,175	60	224.6	4.5
5	none (default)	9,971	7,894	164.6	8.5
5	employmen t	11,668	8,469	217.0	10.2
	status				
5	age variable	13,368	58	123.1	3.8

Table 5.5: How importance vectors and k-anonymity thresholds affect running time and total number of suppressions

In cases where there are a large number of quasi-identifiers and the variables have many categories, the number of possible combinations increases rapidly (see *k*-anonymity). If the number of variables and categories is very large, the computation time of the localSuppression() algorithm can be very long (see Section 7.7 on computation time). Also, the algorithm may not reach a solution, or may come to a solution that will not meet the specified level of *k*-anonymity. Therefore, reducing the number of quasi-identifiers and/or categories before applying local suppression is recommended. This can be done by recoding variables or selecting some variables for other (perturbative) methods, such as PRAM. This is to ensure that the number of suppressions is limited and hence the loss of data is limited to only those values that pose most risk.

In some datasets, it might prove difficult to reduce the number of quasi-identifiers and even after reducing the number of categories by recoding, the local suppression algorithm takes a long time to compute the required suppressions. A solution in such cases can be the so-called 'all-m approach' (see de Wolf, 2015). The all-m approach consists of applying the local suppression algorithm as described above to all possible subsets of size m of the total set of quasi-identifiers. The advantage of this approach is that the partial problems are easier to solve and computation time will be slower. Caution should be applied since this method does not necessarily lead to k-anonymity in the complete set of quasi-identifiers. There are two possibilities to reach the same level of protection: 1) to choose a higher threshold for k or 2) to re-apply the local suppression algorithm on the complete set of quasi-identifiers after using the all-m approach to achieve the required threshold. In the second case, the all-m approach leads to a shorter computation time at the cost of a higher total number of suppressions. **NOTE: The required level is not achieved automatically on the entire set of quasi-identifiers if the all-m approach is used. Therefore, it is important to evaluate the risk measures carefully after using the all-m approach.** 

In sdcMicro the all-m approach is implemented in the 'combs' argument in the localSuppression() function. The value for m is specified in the 'combs' argument and can also take on several values. The subsets of different sizes are then used sequentially in the local suppression algorithm. For example if 'combs' is set to c(3,9), first all subsets of size 3 are considered and subsequently all subsets of size 9. Setting the last value in the combs argument to the total number of key variables guarantees the achievement of k-anonymity for the complete dataset. It is also possible to specify

different values for k for each subset size in the 'k' argument. If we would want to achieve 5-anonimity on the subsets of size 3 and subsequently 3-anonimity on the subsets of size 9, we would set the 'k' argument to c(5,3). Example 5.8 illustrates the use of the all-m approach in sdcMicro.

### Example 5.8 The all-m approach in sdcMicro

```
# Apply k-anonymity with threshold 5 to all subsets of two key variables and \rightarrow subsequently to the complete dataset sdcInitial <- localSuppression(sdcInitial, k = 5, combs = \mathbf{c}(2, 5)) # Apply k-anonymity with threshold 5 to all subsets of three key variables and \rightarrow subsequently with threshold 2 to the complete dataset sdcInitial <- localSuppression(sdcInitial, k = \mathbf{c}(3, 5), combs = \mathbf{c}(5, 2))
```

Table 5.6 presents the results of using the all-m approach of a test dataset with 9 key variables and 4,000 records. The table shows the arguments 'k' and 'combs' of the localSuppression() function, the number of k-anonymity violators for different levels of k as well as the total number of suppressions. We observe that the different combinations do not always lead to the required level of k-anonimity. For example, when setting k=3, and combs 3 and 7, there are still 15 records in the dataset (with a total of 9 quasi-identifiers) that violate 3-anonimity after local suppression. Due to the smaller sample size, the gains in running time are not yet apparent in this example, since the rerunning algorithm several times takes up time. A larger dataset would benefit more from the all-m approach, as the algorithm would take longer in the first place.

Table 5.6 Effect of the all-*m* approach on k-anonymity

Argumen ts	**Numbe r of on complet e	violato rs for dit set	Total num- ber of suppres sio ns	Running time (sec- ond s)		
k	combs	k = 2	k = 3	k = 5		
before local suppres sion	2,464	3,324	3,877	0	0.00	
3	•	0	0	1,766	2,264	17.08
5	•	0	0	0	3,318	10.57
3	3	2,226	3,202	3,819	3,873	13.39
3	3, 7	15	108	1,831	6,164	46.84
3	3, 9	0	0	1,794	5,982	31.38
3	5, 9	0	0	1,734	6,144	62.30
5	3	2,047	3,043	3,769	3,966	12.88
5	3, 7	0	6	86	7,112	46.57
5	3, 9	0	0	0	7,049	24.13
5	5, 9	0	0	0	7,129	54.76
5, 3	3, 7	11	108	1,859	6,140	45.60
5, 3	3, 9	0	0	1,766	2,264	30.07
5, 3	5, 9	0	0	0	3,318	51.25

Often the dataset contains variables that are related to the key variables used for local suppression. Examples are rural/urban to regions in case regions are completely rural or urban or variables that are only answered for specific categories (e.g., sector for those working, schooling related variables for certain age ranges). In those cases, the variables rural/urban or sector might not be quasi-identifiers themselves, but could allow the intruder to reconstruct suppressed values in the quasi-identifiers region or employment status. For example, if region 1 is completely urban, and all other regions are only semi-urban or rural, a suppression in the variable region for a record in region 1 can be simply reconstructed by the rural/urban variable. Therefore, it is useful to suppress the values corresponding to

the suppressions in those linked variables. Example 5.9 illustrates how to suppress the values in the variable "rururb" corresponding to the suppressions in the region variable. All values of "rururb", which correspond to a suppressed value (NA) in the variable "region" are suppressed (set to NA).

### Example 5.9: Manually suppressing values in linked variables

Alternatively, the linked variables can be specified when creating the *sdcMicro* object. The linked variables are called ghost variables. Any suppression in the key variable will lead to a suppression in the variables linked to that key variable. Example 5.10 shows how to specify the linkage between "region" and "rururb" with ghost variables.

Example 5.10: Suppressing values in linked variables by specifying ghost variables

The simpler alternative for the localSuppression() function in *sdcMicro* is the localSupp() function. The localSupp() function can be used to suppress values of certain key variables of individuals with risks above a certain threshold. In this case, all values of the specified variable for respondents with a risk higher than the specified threshold will be suppressed. The risk measure used is the individual risk (see Section 4.5). This is useful if one variable has sensitive values that should not be released for individuals with high risks of re-identification. What is considered high re-identification probability depends on legal requirements. In the following example, the values of the variable "education" are suppressed for all individuals whose individual risk is higher than 0.1, which is illustrated in Example 5.11. For an overview of the individual risk values, it can be useful to look at the summary statistics of the individual risk values as well as the number of suppressions.

Example 5.11: Application of built-in *sdcMicro* function localSupp()

```
# Summary statistics
summary(sdcInitial@risk$individual[,1])
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.05882 0.10000 0.14290 0.26480 0.33330 1.00000

# Number of individuals with individual risk higher than 0.1
sum(sdcInitial@risk$individual[,1] > 0.1)
## [1] 1863

# local suppression
sdcInitial <- localSupp(sdcInitial, threshold = 0.1, keyVar = 'education')</pre>
```

# 5.3 Perturbative methods

Perturbative methods do not suppress values in the dataset, but perturb (alter) values to limit disclosure risk by creating uncertainty around the true values. An intruder is uncertain whether a match between the microdata and an external file is correct or not. Most perturbative methods are based on the principle of matrix masking, i.e., the altered dataset Z is computed as

$$Z = AXB + C$$

where X is the original data, A is a matrix used to transform the records, B is a matrix to transform the variables and C is a matrix with additive noise.

**NOTE:** Risk measures based on frequency counts of keys are no longer valid after applying perturbative methods. This can be seen in Table 5.7, which displays the same data before and after swapping some values. The swapped values are in grey. Both before and after perturbing the data, all observations violate k-anonymity at the level 3 (i.e., each key does not appear more than twice in the dataset). Nevertheless, the risk of **correct** re-identification of the records is reduced and hence information contained in other (sensitive) variables possibly not disclosed. With a certain probability, a match of the microdata with an external data file will be wrong. For example, an intruder would find one individual with the combination {'male', 'urban', 'higher'}, which is a sample unique. However, this match is not correct, since the original dataset did not contain any individual with these characteristics and hence the matched individual cannot be a correct match. The intruder cannot know with certainty whether the information disclosed from other variables for that record is correct.

Table 5.7: Sample data before and after perturbation

Variable	Original data			After perturbing the data		
ID	Gender	Region	Education	Gender	Region	Education
1	female	rural	higher	female	rural	higher
2	female	rural	higher	female	rural	lower
3	male	rural	lower	male	rural	lower
4	male	rural	lower	female	rural	lower
5	female	urban	lower	male	urban	higher
6	female	urban	lower	female	urban	lower

One advantage of perturbative methods is that the information loss is reduced, since no values will be suppressed, depending on the level of perturbation. One disadvantage is that data users might have the impression that the data was not anonymized before release and will be less willing to participate in future surveys. Therefore, there is a need for reporting both for internal and external use (see Section 8.11).

An alternative to perturbative methods is the generation of synthetic data files with the same characteristics as the original data files. Synthetic data files are not discussed in these guidelines. For more information and an overview of the use of synthetic data as SDC method, we refer to Drechsler (2011) and Section 3.8 in Hundepool et al. (2012). We discuss here five perturbative methods: Post Randomization Method (PRAM), microaggregation, noise addition, shuffling and rank swapping.

# 5.3.1 PRAM (Post RAndomization Method)

PRAM is a perturbative method for categorical data. This method reclassifies the values of one or more variables, such that intruders that attempt to re-identify individuals in the data do so, but with positive probability, the re-identification made is with the wrong individual. This means that the intruder might be able to match several individuals between external files and the released data files, but cannot be sure whether these matches are to the correct individual.

PRAM is defined by the transition matrix P, which specifies the transition probabilities, i.e., the probability that a value of a certain variable stays unchanged or is changed to any of the other k-1 values. k is the number of categories

or factor levels within the variable to be PRAMmed. For example, if the variable region has 10 different regions, k equals 10. In case of PRAM for a single variable, the transition matrix is size k \* k. We illustrate PRAM with an example of the variable "region", which has three different values: 'capital', 'rural1' and 'rural2'. The transition matrix for applying PRAM to this variable is size 3\*3:

$$P = \begin{bmatrix} 1 & 0 & 0 \\ 0.05 & 0.8 & 0.15 \\ 0.05 & 0.15 & 0.8 \end{bmatrix}$$

The values on the diagonal are the probabilities that a value in the corresponding category is not changed. The value 1 at position (1,1) in the matrix means that all values 'capital' stay 'capital'; this might be a useful decision, since most individuals live in the capital and no protection is needed. The value 0.8 at position (2,2) means that an individual with value 'rural1' will stay with probability 0.8 'rural1'. The values 0.05 and 0.15 in the second row of the matrix indicate that the value 'rural1' will be changed to 'capital' or 'rural2' with respectively probability 0.05 and 0.15. If in the initial file we had 5,000 individuals with value 'capital' and resp. 500 and 400 with values 'rural1' and 'rural2', we expect after applying PRAM to have 5,045 individuals with capital, 460 with rural1 and 395 with rural29. The recoding is done independently for each individual. We see that the tabulation of the variable "region" yields different results before and after PRAM, which are shown in Table 5.8. The deviation from the expectation is due to the fact that PRAM is a probabilistic method, i.e., the results depend on a probability-generating mechanism; consequently, the results can differ every time we apply PRAM to the same variables of a dataset. **NOTE: The number of changed** values is larger than one might think when inspecting the tabulations in Table 5.8. Not all 5,000 individuals with value captial after PRAM had this value before PRAM and the 457 individuals in rural1 after PRAM are not all included in the 500 individuals before PRAM. The number of changes is larger than the differences in the tabulation (cf. transition matrix). Given that the transition matrix is known to the end users, there are several ways to correct statistical analysis of the data for the distortions introduced by PRAM.

Table 5.8: Tabulation of variable "region" before and after PRAM

Value	Tabulation before PRAM	Tabulation after PRAM
capital	5,000	5,052
rural1	500	457
rural2	400	391

One way to guarantee consistency between the tabulations before and after PRAM is to choose the transition matrix so that, in expectation, the tabulations before and after applying PRAM are the same for all variables.[#foot43]\_ This method is called invariant PRAM and is implemented in *sdcMicro* in the function pram(). The method pram() determines the transition matrix that satisfies the requirements for invariant PRAM. **NOTE: Invariant does not guarantee** that cross-tabulations of variables (unlike univariate tabulations) stay the same.

In Example 5.12, we give an example of invariant PRAM using *sdcMicro*.<sup>11</sup> PRAM is a probabilistic method and the results can differ every time we apply PRAM to the same variables of a dataset. To overcome this and make the results reproducible, it is good practice to set a seed for the random number generator in *R*, so the same random numbers will be generated every time.<sup>12</sup> The number of changed records per variable is also shown.

Example 5.12: Producing reproducible PRAM results by using set.seed()

```
# Set seed for random number generator set.seed(123)
```

(continues on next page)

<sup>&</sup>lt;sup>9</sup> The 5,045 is the expectation computed as 5,000 \* 1 + 500 \* 0.05 + 400 \* 0.05.

<sup>&</sup>lt;sup>11</sup> In this example and the following examples in this section, the *sdcMicro* object "sdcIntial" contains a dataset with 2,000 individuals and 39 variables. We selected five categorical quasi-identifiers and 9 variables for PRAM: "ROOF", "TOILET", "WATER", "ELECTCON", "FUELCOOK", "OWNMOTORCYCLE", "CAR", "TV", and "LIVESTOCK". These PRAM variables were selected according to the requirements of this particular dataset and for illustrative purposes.

<sup>&</sup>lt;sup>12</sup> The PRAM method in *sdcMicro* sometimes produces the following error: Error in factor(xpramed, labels = lev): invalid 'labels'; length 6 should be 1 or 5. Under some circumstances, changing the seed can solve this error.

(continued from previous page)

Table 5.9 shows the tabulation of the variable after applying invariant PRAM. We can see that the deviations from the initial tabulations, which are in expectation 0, are smaller than with the transition matrix that does not fulfill the invariance property. The remaining deviations are due to the randomness.

Table 5.9: Tabulation of variable "region" before and after (invariant) PRAM

Value	Tabulation before PRAM	Tabulation after PRAM	Tabulation after invariant PRAM
capital	5,000	5,052	4,998
rural1	500	457	499
rural2	400	391	403

Table 5.10 presents the cross-tabulations with the variable gender. Before applying invariant PRAM, the share of males in the city is much higher than the share of females (about 60%). This property is not maintained after invariant PRAM (the shares of males and females in the city are roughly equal), although the univariate tabulations are maintained. One solution is to apply PRAM separately for the males and females in this example <sup>13</sup>. This can be done by specifying the strata argument in the pram() function in *sdcMicro* (see below).

Table 5.10: Cross-tabulation of variable "region" and variable "gender" before and after invariant PRAM

	Tabulation before PRAM		Tabulation after invariant PRAM	
Value	male	female	male	female
capital	3,056	1,944	2,623	2,375
rural1	157	343	225	274
rural2	113	287	187	216

The pram() function in *sdcMicro* has several options. **NOTE:** If no options are set and the PRAM method is applied to an sdcMicro object, all PRAM variables selected in the sdcMicro object are automatically used for PRAM and PRAM is applied within the selected strata (see Section 7.5 on *sdcMicro* objects for more details). Alternatively, PRAM can also be applied to variables that are not specified in the *sdcMicro* object as PRAM variables, such as key variables, which is shown in Example 5.13. In that case, however, the risk measures that are automatically computed will not be correct anymore, since the variables are perturbed. Therefore, if during the SDC process PRAM will be applied to some key variables, it is recommended to create a new *sdcMicro* object where the variables to be PRAMmed are selected as PRAM variables in the function createSdcObj().

Example 5.13: Selecting the variable "toilet" to apply PRAM

<sup>&</sup>lt;sup>13</sup> This can also be achieved with multidimensional transition matrices. In that case, the probability is not specified for 'male' -> 'female', but for 'male' + 'rural' -> 'female' + 'rural' and for 'male' + 'urban' -> 'female' + 'urban'. This is not implemented in sdcMicro but can be achieved by PRAMming the males and females separately. In the example here, this could be done by specifying gender as strata variable in the pram() function in *sdcMicro*.

The results for PRAM differ if applied simultaneously to several variables or subsequently to each variable separately. It is not possible to specify the entire transition matrix in *sdcMicro*, but we can set minimum values (between 0 and 1) for the diagonal entries. The diagonal entries specify the probability that a certain value stays the same after applying PRAM. Setting the minimum value to 1 will yield no changes to this category. By default, this value is 0.8, which applies for all categories. It is also possible to specify a vector with value for each diagonal element of the transformation matrix/category. In Example 5.14 values of the first region are less likely to change than values of the other regions. **NOTE: The invariant PRAM method requires that the transition matrix has a unit eigenvalue.** Not all sets of restrictions can therefore be used (e.g., the minimum value 1 on any of the categories).

Example 5.14: Specifying minimum values for diagonal entries in PRAM transition matrix

In the invariant PRAM method, we can also specify the amount of perturbation by specifying the parameter alpha. This choice is reflected in the transition matrix. By default, the alpha value is 0.5. The larger alpha, the larger the perturbations. Alpha equal to zero leads to no changes. The maximum value for alpha is 1.

PRAM is especially useful when a dataset contains many variables and applying other anonymization methods, such as recoding and local suppression, would lead to significant information loss. Checks on risk and utility are important after PRAM.

To do statistical inference on variables to which PRAM was applied, the researcher needs knowledge about the PRAM method as well as about the transition matrix. The transition matrix, together with the random number seed, can, however, lead to disclosure through reconstruction of the non-perturbed values. Therefore, publishing the transition matrix but not the random seed is recommended.

A disadvantage of using PRAM is that very unlikely combinations can be generated, such as a 63-year- old who goes to school. Therefore, the PRAMmed variables need to be audited to prevent such combinations from happening in the released data file. In principal, the transition matrix can be designed in such a way that certain transitions are not possible (probability 0). For instance, for those that go to school, the age must range within 6 to 18 years and only such changes are allowed. In *sdcMicro* the transition matrix cannot be exactly specified. A useful alternative is constructing strata and applying PRAM within the strata. In this way, the changes between variables will only be applied within the strata. Example 5.15 illustrates this by applying PRAM to the variable "toilet" within the strata generated by the "region" education. This prevents changes in the variable "toilet", where toilet types in a particular region are exchanged with those in other regions. For instance, in the capital region certain types of unimproved toilet types are not in use and therefore these combinations should not occur after PRAMming. Values are only changed with those that are available in the same strata. Strata can be formed by any categorical variable, e.g., gender, age groups, education level.

Example 5.15: Minimizing unlikely combinations by applying PRAM within strata

```
# Applying PRAM within the strata generated by the variable region sdcInitial <- pram(obj = sdcInitial, variables = c("TOILET"), strata_variables = c( \( \to \) "REGION")) ## Number of changed observations:
```

(continues on next page)

(continued from previous page)

```
## ----- ## TOILET != TOILET_pram : 179 (8.95%)
```

#### **Recommended Reading Material on PRAM**

Gouweleeuw, J. M, P Kooiman, L.C.R.J Willenborg, and P.P de Wolf. "Post Randomization for Statistical Disclosure Control: Theory and Implementation." *Journal of Official Statistics* 14, no. 4 (1998a): 463-478. Available at http://www.jos.nu/articles/abstract.asp?article=144463

Gouweleeuw, J. M, P Kooiman, L.C.R.J Willenborg, and Peter Paul de Wolf. "The Post Randomization Method for Protecting Microdata." *Qüestiió, Quaderns d'Estadística i Investigació Operativa* 22, no. 1 (1998b): 145-156. Available at http://www.raco.cat/index.php/Questiio/issue/view/2250

Marés, Jordi, and Vicenç Torra. 2010."PRAM Optimization Using an Evolutionary Algorithm." *In Privacy in Statistical Databases*, by Josep Domingo-Ferrer and Emmanouil Magkos, 97-106. Corfú, Greece: Springer.

Warner, S.L. "Randomized Response: A Survey Technique for Eliminating Evasive Answer Bias." *Journal of American Statistical Association* 57 (1965): 622-627.

## 5.3.2 Microaggregation

Microaggregation is most suitable for continuous variables, but can be extended in some cases to categorical variables. [#foot47]\_\_ It is most useful where confidentiality rules have been predetermined (e.g., a certain threshold for k-anonymity has been set) that permit the release of data only if combinations of variables are shared by more than a predetermined threshold number of respondents (k). The first step in microaggregation is the formation of small groups of individuals that are homogeneous with respect to the values of selected variables, such as groups with similar income or age. Subsequently, the values of the selected variables of all group members are replaced with a common value, e.g., the mean of that group. Microaggregation methods differ with respect to (i) how the homogeneity of groups is defined, (ii) the algorithms used to find homogeneous groups, and (iii) the determination of replacement values. In practice, microaggregation works best when the values of the variables in the groups are more homogeneous. When this is the case, then the information loss due to replacing values with common values for the group will be smaller than in cases where groups are less homogeneous.

In the univariate case, and also for ordinal categorical variables, formation of homogeneous groups is straightforward: groups are formed by first ordering the values of the variable and then creating g groups of size  $n_i$  for all groups i in  $1, \ldots, g$ . This maximizes the within-group homogeneity, which is measured by the within-groups sum of squares (SSE)

$$SSE = \sum_{i=1}^{g} \sum_{j=1}^{n_i} (x_{ij} - \overline{x}_i)^T (x_{ij} - \overline{x}_i)$$

The lower the SSE, the higher the within-group homogeneity. The group sizes can differ amongst groups, but often groups of equal size are used to simplify the search<sup>15</sup>.

The function microaggregation() in *sdcMicro* can be used for univariate microaggregation. The argument 'aggr' specifies the group size. Forming groups is easier if all groups – except maybe the last group of remainders – have the same size. This is the case in the implementation in *sdcMicro* as it is not possible to have groups of different sizes. Example 5.16 shows how to use the function microaggregation() in *sdcMicro*. The default group size is 3 but the user can specify any desired group size. Choice of group size depends on the homogeneity within the groups and the

<sup>15</sup> Here all groups can have different sizes (i.e., number of individuals in a group). In practice, the search for homogeneous groups is simplified by imposing equal group sizes for all groups.

<sup>&</sup>lt;sup>16</sup> In this example and the following examples in this section, the *sdcMicro* object "sdcIntial" contains a dataset with 2,000 individuals and 39 variables. We selected five categorical quasi-identifiers and three continuous quasi-identifiers: "INC", "EXP" and "WEALTH".

required level of protection. In general it holds that the larger the group, the higher the protection. A disadvantage of groups of equal sizes is that the data might be unsuitable for this. For instance, if two individuals have a low income (e.g., 832 and 966) and four individuals have a high income (e.g., 3,313,3,211,2,987,3,088), the mean of two groups of size three (e.g., (832 + 966 + 2,987) / 3 = 1,595 and (3,088 + 3,211 + 3,313) / 3 = 3,204) would represent neither the low nor the high income.

Example 5.16: Applying univariate microaggregation with sdcMicro function microaggregation()

```
sdcInitial <- **microaggregation**\ (obj = sdcInitial, variables =
'INC', aggr = 3, method = mafast, measure = "mean")</pre>
```

By default, the microaggregation function replaces values with the group mean. An alternative, more robust approach is to replace group values with the median. This can be specified in the argument *measure* of the function microaggregation(). In cases where the median is chosen, one individual in every group keeps the same value if groups have odd sizes. In cases where there is a high degree of heterogeneity within the groups (this is often the case for larger groups), the median is preferred to preserve the information in the data. An example is income, where one outlier can lead to multiple outliers being created when using microaggregation. This is illustrated in Table 5.11. If we choose the mean as replacement for all values, which are grouped with the outlier (6,045 in group 2), these records will be assigned values far from their original values. If we chose the median, the incomes of individuals 1 and 2 are not perturbed, but no value is an outlier. Of course, this might in itself present problems. **NOTE: If microaggregation alters outlying values, this can have a significant impact on the computation of some measures sensitive to outliers, such as the GINI index.** In the case where microaggregation is applied to categorical variables, the median is used to calculate the replacement value for the group.

Table 5.11: Illustrating the effect of choosing mean vs. median for microaggregation where outliers are concerned

ID	Group	Income	Microaggr egation (mean)	Microaggr egation (median)
1	1	2,300	2,245	2,300
2	2	2,434	3,608	2,434
3	1	2,123	2,245	2,300
4	1	2,312	2,245	2,300
5	2	6,045	3,608	2,434
6	2	2,345	3,608	2,434

In case of multiple variables that are candidates for microaggregation, one possibility is to apply univariate microaggregation to each of the variables separately. The advantage of univariate microaggregation is minimal information loss, since the changes in the variables are limited. The literature shows, however, that disclosure risk can be very high if univariate microaggregation is applied to several variables separately and no additional anonymization techniques are applied (Domingo-Ferrer et al., 2002). To overcome this shortcoming, an alternative to univariate microaggregation is multivariate microaggregation.

Multivariate microaggregation is widely used in official statistics. The first step in multivariate aggregation is the creation of homogeneous groups based on several variables. Groups are formed based on multivariate distances between the individuals. Subsequently, the values of all variables for all group members are replaced with the same values. Table 5.12 illustrates this with three variables. We see that the grouping by income, expenditure and wealth leads to a different grouping, as in the case in Table 5.11, where groups were formed based only on income.

Table 5.12: Illustration of multivariate microaggregation

ID	Group	Before microaggregation			After microaggregation		
		Incom e	Exp	Wealt h	Incom e	Exp	Wealt h
1	1	2,300	1,714	5.3	2,285. 7	1,846. 3	6.3
2	1	2,434	1,947	7.4	2,285. 7	1,846. 3	6.3
3	1	2,123	1,878	6.3	2,285. 7	1,846. 3	6.3
4	2	2,312	1,950	8.0	3,567. 3	2,814. 0	8.3
5	2	6,045	4,569	9.2	3,567. 3	2,814. 0	8.3
6	2	2,345	1,923	7.8	3,567. 3	2,814. 0	8.3

There are several multivariate microaggregation methods that differ with respect to the algorithm used for creating groups of individuals. There is a trade-off between speed of the algorithm and within-group homogeneity, which is directly related to information loss. For large datasets, this is especially challenging. We discuss the Maximum Distance to Average Vector (MDAV) algorithm here in more detail. The MDAV algorithm was first introduced by Domingo-Ferrer and Torra (2005) and represents a good choice with respect to the trade-off between computation time and the group homogeneity, computed by the within-group SSE. The MDAV algorithm is implemented in *sdcMicro*.

The algorithm computes an average record or centroid C, which contains the average values of all included variables. We select an individual A with the largest squared Euclidean distance from C, and build a group of k records around A. The group of k records is made up of A and the k-1 records closest to A measured by the Euclidean distance. Next, we select another individual B, with the largest squared Euclidean distance from individual A. With the remaining records, we build a group of k records around B. In the same manner, we select an individual D with the largest distance from B and, with the remaining records, build a new group of k records around D. The process is repeated until we have fewer than 2k records remaining. The MDAV algorithm creates groups of equal size with the exception of maybe one last group of remainders. The microaggregated dataset is then computed by replacing each record in the original dataset by the average values of the group to which it belongs. Equal group sizes, however, may not be ideal for data characterized by greater variability. In sdcMicro multivariate microaggregation is also implemented in the function microaggregation(). Example 5.17 shows how to choose the MDAV algorithm in sdcMicro.

Example 5.17: Multivariate microaggregation with the Maximum Distance to Average Vector (MDAV) algorithm in *sdcMicro* 

```
sdcInitial <- microaggregation(obj = sdcInitial, variables = c("INC", "EXP", "WEALTH →"), method = "mdav")
```

It is also possible to group variables only within strata. This reduces the computation time and adds an extra layer of protection to the data, because of the greater uncertainty produced<sup>17</sup>. In *sdcMicro* this can be achieved by specifying the strata variables, as shown in Example 5.18.

Example 5.18: Specifying strata variables for microaggregation

Besides the method MDAV, there are few other grouping methods implemented in *sdcMicro* (Templ, Meindl and Kowarik, 2014). Table 5.13 gives an overview of these methods. Whereas the method 'MDAV' uses the Euclidian distance, the method 'rmd' uses the Mahalanobis distance instead. An alternative to these methods is sorting the respondents based on the first principal component (PC), which is the projection of all variables into a one-dimensional space maximizing the variance of this projection. The performance of this method depends on the share of the total variance in the data that is explained by the first PC. The 'rmd' method is computationally more intensive due to the computation of Mahalanobis distances, but provides better results with respect to group homogeneity. It is recommended for smaller datasets (ibid.).

Table 5.13: Grouping methods for microaggregation that are implemented in sdcMicro

<sup>&</sup>lt;sup>17</sup> Also the homogeneity in the groups will be generally lower, leading to larger changes, higher protection, but also more information loss, unless the strata variable correlates with the microaggregation variable.

Method / option in sd-	Description
cMicro	
mdav	grouping is based on classical (Euclidean) distance measures
rmd	grouping is based on robust multivariate (Mahalanobis) distance measures
pca	grouping is based on principal component analysis whereas the data are sorted on the
	first principal component
clustpppca	grouping is based on clustering and (robust) principal component analysis for each
	cluster
influence	grouping is based on clustering and aggregation is performed within clusters

In case of several variables to be used for microaggregation, looking first at the covariance or correlation matrix of these variables is recommended. If not all variables correlate well, but two or more sets of variables show high correlation, less information loss will occur when applying microaggregation separately to these sets of variables. In general, less information loss will occur when applying multivariate microaggregation, if the variables are highly correlated. The advantage of replacing the values with the mean of the groups rather than other replacement values has the advantage that the overall means of the variables are preserved.

#### **Recommended Reading Material on Microaggregation**

Domingo-Ferrer, Josep, and Josep Maria Mateo-Sanz. 2002."Practical data-oriented microaggregation for statistical disclosure control." *IEEE Transactions on Knowledge and Data Engineering 14* (2002): 189-201.

Hansen, Stephen Lee, and Sumitra Mukherjee. 2003. "A polynomial algorithm for univariate optimal." *IEEE Transactions on Knowledge and Data Engineering* 15 (2003): 1043-1044.

Hundepool, Anco, Josep Domingo-Ferrer, Luisa Franconi, Sarah Giessing, Rainer Lenz, Jane Naylor, Eric Schulte Nordholt, Giovanni Seri, and Peter Paul de Wolf. 2006. *Handbook on Statistical Disclosure Control*. ESSNet SDC. http://neon.vb.cbs.nl/casc/handbook.htm

Hundepool, Anco, Josep Domingo-Ferrer, Luisa Franconi, Sarah Giessing, Eric Schulte Nordholt, Keith Spicer, and Peter Paul de Wolf. 2012. *Statistical Disclosure Control*. Chichester: John Wiley & Sons Ltd. doi:10.1002/9781118348239.

Templ, Matthias, Bernhard Meindl, Alexander Kowarik, and Shuang Chen. 2014, August. "International Household Survey Network (IHSN)." http://www.ihsn.org/home/software/disclosure-control-toolbox. (accessed November 13, 2014).

### 5.3.3 Noise addition

Noise addition, or noise masking, means adding or subtracting (small) values to the original values of a variable, and is most suited to protect continuous variables (see Brand (2002) for an overview). Noise addition can prevent exact matching of continuous variables. The advantages of noise addition are that the noise is typically continuous with mean zero, and exact matching with external files will not be possible. Depending on the magnitude of noise added, however, approximate interval matching might still be possible.

When using noise addition to protect data, it is important to consider the type of data, the intended use of the data and the properties of the data before and after noise addition, i.e., the distribution – particularly the mean – covariance and correlation between the perturbed and original datasets.

Depending on the data, it may also be useful to check that the perturbed values fall within a meaningful range of values. Figure 5.7 on page 68 illustrates the changes in data distribution with increasing levels of noise. For data that has outliers, it is important to note that when the perturbed data distribution is similar to the original data distribution (e.g., at low noise levels), noise addition will not protect outliers. After noise addition, these outliers can generally still be detected as outliers and hence easily be identified. An example is a single very high income in a certain region.

After perturbing this income value, the value will still be recognized as the highest income in that region and can thus be used for re-identification. This is illustrated in Figure 5.6, where 10 original observations (open circles) and the anonymized observations (red triangles) are plotted. The tenth observation is an outlier. The values of the first nine observations are sufficiently protected by adding noise: their magnitude and order has changed and exact or interval matching can be successfully prevented. The outlier is not sufficiently protected since, after noise addition, the outlier can still be easily identified. The fact that the absolute value has changed is not sufficient protection. On the other hand, at high noise levels, protection is higher even for the outliers, but the data structure is not preserved and the information loss is large, which is not an ideal situation. One way to circumvent the outlier problem is to add noise of larger magnitude to outliers than to the other values.

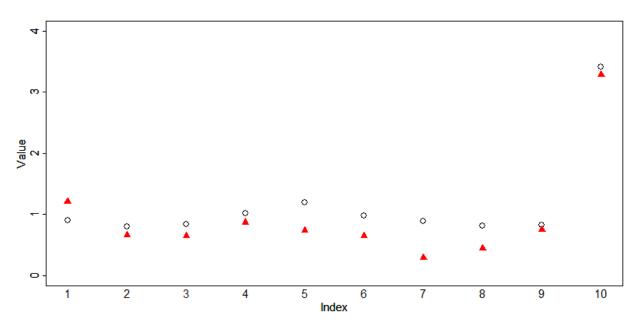


Figure 5.6: Illustration of effect of noise addition to outliers

There are several noise addition algorithms. The simplest version of noise addition is uncorrelated additive normally distributed noise, where  $x_j$ , the original values of variable j are replaced by

$$z_j = x_j + \varepsilon_j,$$

where  $\varepsilon_j \sim N(0, \ \sigma_{\varepsilon_j}^2)$  and  $\sigma_{\varepsilon_j} = \alpha \bullet \sigma_j$  with  $\sigma_j$  the standard deviation of the original data. In this way, the mean and the covariances are preserved, but not the variances and correlation coefficient. If the level of noise added,  $\alpha$ , is disclosed to the user, many statistics can be consistently estimated from the perturbed data. The added noise is proportional to the variance of the original variable. The magnitude of the noise added is specified by the parameter  $\alpha$ , which specifies this proportion. The standard deviation of the perturbed data is  $1 + \alpha$  times the standard deviation of the perturbed data. A decision on the magnitude of noise added should be informed by the legal situation regarding data privacy, data sensitivity and the acceptable levels of disclosure risk and information loss. In general, the level of noise is a function of the variance of the original variables, the level of protection needed and the desired value range after anonymization 18. An  $\alpha$  value that is too small will lead to insufficient protection, while an  $\alpha$  value that is too high will make the data useless for data users.

In *sdcMicro* noise addition is implemented in the function addNoise(). The algorithm and parameter can be specified as arguments in the function addNoise(). Simple noise addition is implemented in the function addNoise() with the value "additive" for the argument 'method'. Example 5.19 shows how to use *sdcMicro* to add uncorrelated noise to expenditure variables, where the standard deviation of the added noise equals half the standard deviation of the original variables.<sup>19</sup> Noise is added to all selected variables.

<sup>&</sup>lt;sup>18</sup> Common values for  $\alpha$  are between 0.5 and 2. The default value in the *sdcMicro* function addNoise() is 150, which is too large for most datasets; the level of noise should be set in the argument 'noise'.

<sup>&</sup>lt;sup>19</sup> In this example and the following examples in this section, the sdcMicro object "sdcIntial" contains a dataset with 2,000 individuals and 39

#### Example 5.19: Uncorrelated noise addition

```
sdcInitial <- addNoise(obj = sdcInitial, variables = c('TOTFOOD', 'TOTHLTH', 'TOTALCH

→', 'TOTCLTH', 'TOTHOUS', 'TOTFURN', 'TOTTRSP', 'TOTCMNQ', 'TOTRCRE', 'TOTEDUC',

→'TOTHOTL', 'TOTMISC'), noise = 0.5, method = "additive")
```

Figure 5.7 shows the frequency distribution of a numeric continuous variable and the distribution before and after noise addition with different levels of noise (0.1, 0.5, 1, 2 and 5). The first plot shows the distribution of the original values. The histograms clearly show that noise of large magnitudes (high values of alpha) lead to a distribution of the data far from the original values. The distribution of the data changes to a normal distribution when the magnitude of the noise grows respective to the variance of the data. The mean in the data is preserved, but, with an increased level of noise, the variance of the perturbed data grows. After adding noise of magnitude 5, the distribution of the original data is completely destroyed.

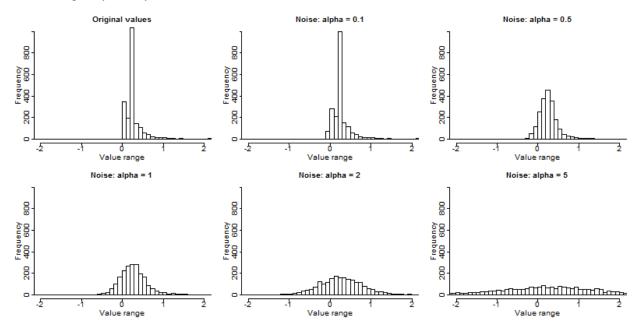


Figure 5.7: Frequency distribution of a continuous variable before and after noise addition

Figure 5.8 shows the value range of a variable before adding noise (no noise) and after adding several levels of noise ( $\alpha$  from 0.1 to 1.5 with 0.1 increments). In the figure, the minimum value, the  $20^{th}$ ,  $30^{th}$ ,  $40^{th}$  percentiles, the median, the  $60^{th}$ ,  $70^{th}$ ,  $80^{th}$  and  $90^{th}$  percentiles and the maximum value are plotted. The median (50:sup:th percentile) is indicated with the red "+" symbol. From Figure 5.7 and Figure 5.8, it is apparent that the range of values expands after noise addition, and the median stays roughly at the same level, as does the mean by construction. The larger the magnitude of noise added, the wider the value range. In cases where the variable should stay in a certain value range (e.g., only positive values, between 0 and 100), this can be a disadvantage of noise addition. For instance, expenditure variables typically have non-negative values, but adding noise to these variables can generate negative values, which are difficult to interpret. One way to get around this problem is to set any negative values to zero. This truncation of values below a certain threshold, however, will distort the distribution (mean and variance matrix) of the perturbed data. This means that the characteristics that were preserved by noise addition, such as the conservation of the mean and covariance matrix, are destroyed and the user, even with knowledge of the magnitude of the noise, can no longer use the data for consistent estimation.

Another way to avoid negative values is the application of multiplicative rather than additive noise. In that case, variables are multiplied by a random factor with expectation 1 and a positive variance. This will also lead to larger perturbations (in absolute value) of large initial values (outliers). If the variance of the noise added is small, there will

5.3. Perturbative methods

variables. We selected five categorical quasi-identifiers and 12 continuous quasi-identifiers. These are the expenditure components "TFOODEXP", "TALCHEXP", "TCLTHEXP", "THOUSEXP", "TFURNEXP", "THLTHEXP", "TTRANSEXP", "TCOMMEXP", "TRECEXP", "TEDUEXP", "TRESTHOTEXP", "TMISCEXP".

be no or few negative factors and thus fewer sign changes than in case of additive noise masking. Multiplicative noise masking is not implemented in *sdcMicro*, but can be relatively easily implemented in base *R* by generating a vector of random numbers and multiplying the data with this vector. For more information on multiplicative noise masking and the properties of the data after masking, we refer to Kim and Winkler (2003).

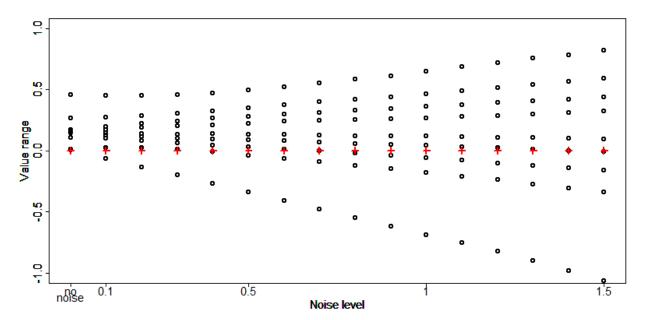


Figure 5.8: Noise levels and the impact on the value range (percentiles)

If two or more variables are selected for noise addition, correlated noise addition is preferred to preserve the correlation structure in the data. In this case, the covariance matrix of noise  $\Sigma_{\varepsilon}$  is proportional to the covariance matrix of the original data  $\Sigma_X$ :

$$\Sigma_{\varepsilon} = \alpha \Sigma_X$$

In the addNoise() function of the *sdcMicro* package, correlated noise addition can be used by specifying the methods 'correlated' or 'correlated'. The method "correlated" assumes that the variables are approximately normally distributed. The method 'correlated' is a version of the method 'correlated', which is robust against the normality assumption. Example 5.20 shows how to use the 'correlated' method. The normality of variables can be investigated in *R*, with, for instance, a Jarque-Bera or Shapiro-Wilk test<sup>20</sup>.

### Example 5.20: Correlated noise addition

```
sdcInitial <- addNoise(obj = sdcInitial, variables = c('TOTFOOD', 'TOTHLTH', 'TOTALCH
         ', 'TOTCLTH', 'TOTHOUS', 'TOTFURN', 'TOTTRSP', 'TOTCMNQ', 'TOTRCRE', 'TOTEDUC',
          'TOTHOTL', 'TOTMISC'), noise = 0.5, method = "correlated2")</pre>
```

In many cases, only the outliers have to be protected, or have to be protected more. The method 'outdect' adds noise only to the outliers, which is illustrated in Example 5.21. The outliers are identified with univariate and robust multivariate procedures based on a robust Mahalanobis distance calculated by the MCD estimator (Templ et al., 2014). Nevertheless, noise addition is not the most suitable method for outlier protection.

### Example 5.21: Noise addition for outliers using the 'outdect' method

```
sdcInitial <- addNoise(obj = sdcInitial, variables = c('TOTFOOD', 'TOTHLTH', 'TOTALCH

→', 'TOTCLTH', 'TOTHOUS', 'TOTFURN', 'TOTTRSP', 'TOTCMNQ', 'TOTRCRE', 'TOTEDUC',

→'TOTHOTL', 'TOTMISC'), noise = 0.5, method = "outdect")
```

<sup>&</sup>lt;sup>20</sup> The Shapiro-Wilk test is implemented in the function shapiro.test() from the package *stats* in *R*. The Jarque-Bera test has several implementations in *R*, for example, in the function jarque.bera.test() from the package *tseries*.

If noise addition is applied to variables that are a ratio of an aggregate, this structure can be destroyed by noise addition. Examples are income and expenditure data with many income and expenditure categories. The categories add up to total income or total expenditures. In the original data, the aggregates match with the sum of the components. After adding noise to their components (e.g., different expenditure categories), however, their new aggregates will not necessarily match the sum of the categories anymore. One way to keep this structure is to add noise only to the aggregates and release the components as ratio of the perturbed aggregates. Example 5.22 illustrates this by adding noise to the total of expenditures. Subsequently, the ratios of the initial expenditure categories are used for each individual to reconstruct the perturbed values for each expenditure category.

Example 5.22: Noise addition to aggregates and their components

#### **Recommended Reading Material on Noise Addition**

Brand, Ruth. 2002. "Microdata Protection through Noise Addition." In *Inference Control in Statistical Databases - From Theory to Practice*, edited by Josep Domingo-Ferrer. Lecture Notes in Computer Science Series 2316, 97-116. Berlin Heidelberg: Springer. http://link.springer.com/chapter/10.1007%2F3-540-47804-3 8

Kim, Jay J, and William W Winkler. 2003. "Multiplicative Noise for Masking Continuous Data." *Research Report Series* (Statistical Research Division. US Bureau of the Census). https://www.census.gov/srd/papers/pdf/rrs2003-01.pdf

Torra, Vicenç, and Isaac Cano. 2011. "Edit Constraints on Microaggregation and Additive Noise." In *Privacy and Security Issues in Data Mining and Machine Learning*, edited by C. Dimitrakakis, A. Gkoulalas-Divanis, A. Mitrokotsa, V. S. Verykios, Y. Saygin. Lecture Notes in Computer Science Volume 6549, 1-14. Berlin Heidelberg: Springer. http://link.springer.com/book/10.1007/978-3-642-19896-0

Mivule, K. 2013. "Utilizing Noise Addition for Data Privacy, An Overview." *Proceedings of the International Conference on Information and Knowledge Engineering (IKE 2012)*, (pp.65-71).Las Vegas, USA. http://arxiv.org/ftp/arxiv/papers/1309/1309.3958.pdf

### 5.3.4 Rank swapping

Data swapping is based on interchanging values of a certain variable across records. Rank swapping is one type of data swapping, which is defined for ordinal and continuous variables. For rank swapping, the values of the variable are first ordered. The possible number of values for a variable to swap with is constrained by the values in a neighborhood around the original value in the ordered values of the dataset. The size of this neighborhood can be specified, e.g., as a percentage of the total number of observations. This also means that a value can be swapped with the same or very similar values. This is especially the case if the neighborhood is small or there are only a few different values in the variable (ordinal variable). An example is the variables "education" with only few categories: ('none', 'primary', 'secondary', 'tertiary'). In these cases, rank swapping is not a suitable method.

If rank swapping is applied to several variables simultaneously, the correlation structure between the variables is preserved. Therefore, it is important to check whether the correlation structure in the data is plausible. Rank swapping is implemented in the function rankSwap() in sdcMicro. The variables, which have to be swapped, should be specified in the argument 'variables'. By default, values below the  $5^{th}$  percentile and above the  $95^{th}$  percentile are top and bottom coded and replaced by their average value (see Section 5.2.1.2 on top and bottom coding). By specifying the options 'TopPercent' and 'BottomPercent' we can choose these percentiles. The argument 'P' defines the size of the neighborhood as percentage of the sample size. If the value 'p' is 0.05, the neighborhood will be of size 0.05 \* n, where n is the sample size. Since rank swapping is a probabilistic method, i.e., the swapping depends on a random number generating mechanism, specifying a seed for the random number generator before using rank swapping is recommended to guarantee reproducibility of results. The seed can also be specified as a function argument in the function rankSwap(). Example 5.23 shows how to apply rank swapping with sdcMicro. If the variables contain missing values (NA in R), the function rankSwap() will automatically recode those to the value specified in the 'missing' argument. This value should not be in the value range of any of the variables. After using the function rankSwap(), these values should be recoded NA. This is shown in the Example 5.23.

### Example 5.23: Rank swapping using sdcMicro

```
# Check correlation structure between the variables
cor(file$TOTHOUS, file$TOTFOOD)
## [1] 0.3811335

# Set seed for random number generator
set.seed(12345)

# Apply rank swapping
rankSwap(sdcInitial, variables = c("TOTHOUS", "TOTFOOD"), missing = NA)
```

Rank swapping has been found to yield good results with respect to the trade-off between information loss and data protection (Domingo-Ferrer and Torra, 2001). Rank swapping is not useful for variables with few different values or many missing values, since the swapping in that case will not result in altered values. Also, if the intruder knows to whom the highest or lowest value of a specific variable belongs (e.g., income), the level of this variable will be disclosed after rank swapping, because the values themselves are not altered and the original values are all disclosed. This can be solved by top and bottom coding the lowest and/or highest values.

#### Recommended Reading Material on Rank Swapping

Dalenius T. and Reiss S.P. 1978. Data-swapping: a technique for disclosure control (extended abstract). In Proc. ASA Section on Survey Research Methods. American Statistical Association, Washington DC, 191–194.

Domingo-Ferrer J. and Torra V. 2001. "A Quantitative Comparison of Disclosure Control Methods for Microdata." In *Confidentiality, Disclosure and Data Access: Theory and Practical Applications for Statistical Agencies*, edited by P. Doyle, J.I. Lane, J.J.M. Theeuwes, and L. Zayatz, 111–134. Amsterdam, North-Holland.

Hundepool A., Van de Wetering A., Ramaswamy R., Franconi F., Polettini S., Capobianchi A., De Wolf P.-P., Domingo-Ferrer J., Torra V., Brand R. and Giessing S. 2007. μ-Argus User's Manual version 4.1.

# 5.3.5 Shuffling

Shuffling as introduced by Muralidhar and Sarathy (2006) is similar to swapping, but uses an underlying regression model for the variables to determine which variables are swapped. Shuffling can be used for continuous variables and is a deterministic method. Shuffling maintains the marginal distributions in the shuffled data. Shuffling, however, requires a complete ranking of the data, which can be computationally very intensive for large datasets with several variables.

The method is explained in detail in Muralidhar and Sarathy (2006). The idea is to rank the individuals based on their original variables. Then fit a regression model with the variables to be protected as regressands and a set of variables that predict this variable well (i.e., are correlated with) as regressors. This regression model is used to generate n synthetic (predicted) values for each variable that has to be protected. These generated values are also ranked and each original value is replaced with another original value with the rank that corresponds to the rank of the generated value. This means that all original values will be in the data. Table 5.14 presents a simplified example of the shuffling method. The regressands are not specified in this example.

ID	Income (orig)	Rank (orig)	Income (pred)	Rank (pred)	Shuffle d values
1	2,300	2	2,466.56	4	2,345
2	2,434	6	2,583.58	7	2,543
3	2,123	1	2,594.17	8	2,643
4	2,312	3	2,530.97	6	2,434
5	6,045	10	5,964.04	10	6,045
6	2,345	4	2,513.45	5	2,365
7	2,543	7	2,116.16	1	2,123
8	2,854	9	2,624.32	9	2,854
9	2,365	5	2,203.45	2	2,300
10	2,643	8	2,358.29	3	2,312

The method 'ds' (the default method of data shuffling in sdeMicro) is recommended for use (Templ et al., 2014)<sup>21</sup>. A regression function with regressors for the variables to be protected must be specified in the argument 'form'. At least two regressands should be specified and the regressors should have predictive power for the variables to be predicted. This can be checked with goodness-of-fit measures such as the  $R^2$  of the regression. The  $R^2$  captures only linear relations, but these are also the only relations that are captured by the linear regression model used for shuffling. Following is an example for shuffling expenditure variables, which are predicted by total household expenditures and household size.

Example 5.24: Shuffling using a specified regression equation

```
# Evaluate R-squared (goodness-of-fit) of the regression model
summary(lm(file, form = TOTFOOD + TOTALCH + TOTCLTH + TOTHOUS + TOTFURN + TOTHLTH +

TOTTRSP + TOTCMNQ + TOTRCRE + TOTEDUC + TOTHOTL + TOTMISC ~ EXP + HHSIZE))

# Shuffling using the specified regression equation
sdcInitial <- shuffle(sdcInitial, method='ds', form = TOTFOOD + TOTALCH + TOTCLTH +

TOTHOUS + TOTFURN + TOTHLTH + TOTTRSP + TOTCMNQ + TOTRCRE + TOTEDUC + TOTHOTL +

TOTMISC ~ EXP + HHSIZE)
```

#### **Recommended Reading Material on Shuffling**

K. Muralidhar and R. Sarathy. 2006."Data shuffling - A new masking approach for numerical data," *Management Science*, 52, 658-670.

## 5.3.6 Comparison of PRAM, rank swapping and shuffling

PRAM, rank swapping and shuffling are all perturbative methods, i.e., they change the values for individual records and are mainly used for continuous variables. After rank swapping and shuffling, the original values are all contained

<sup>&</sup>lt;sup>21</sup> In *sdcMicro*, there are several other methods for shuffling implemented, including 'ds', 'mvn' and 'mlm'. See the Help option for the shuffle function in *sdcMicro* for details on methods 'ds', 'mvm' and 'mlm'.

in the treated dataset but might be assigned to other records. This implies that univariate tabulations are not changed. This also holds in expectation for PRAM, if a transition matrix is chosen that has the invariant property.

Choosing a method is based on the structure to be preserved in the data. In cases where the regression model fits the data well, data shuffling would work very well, as there should be sufficient (continuous) regressors available. Rank swapping works well if there are sufficient categories in the variables. PRAM is preferred if the perturbation method should be applied to only one or few variables; the advantage is the possibility of specifying restrictions on the transition matrix and applying PRAM only within strata, which can be user defined.

# 5.4 Anonymization of geospatial variables

Recently, geospatial data has become increasingly popular with researchers and wide-spread. Georeferenced data identifies the geographical location for each record with the help of a Geographical Information System (GIS), that uses for instance GPS (Global Positioning System) coordinates or address data. The advantages of geospatial data are manifold: 1) researchers can create their own geographical areas, such as the service area of a hospital; 2) it enables researchers to measure the proximity to facilities, such as schools; 3) researchers can use the data to extract geographical patterns; and 4) it enables linking of data from different sources (see e.g., Burgert et al., 2015). However, geospatial data, due to the precise reference to a location, also pose a challenge to the privacy of the respondents.

One way to anonymize georeferenced data is removing the GIS variables and instead leaving in or creating other geographical variables, such as province, region. However, this approach also removes the benefits of geospatial data. Another option is the geographical displacement of areas and/or records. Burgert et al. (2013) describe a geographical displacement procedure for a health dataset. This paper also includes the code in Python. Hu and Drechsler (2015) propose three different strategies for generating synthetic geocodes.

#### Recommended Reading Material on Anonymization of Geospatial Data

C.R. Burgert, J. Colston, T. Roy and B. Zachary. 2013. "DHS Spatial Analysis Report No. 7 - Geographic Displacement Procedure and Georeferenced Data Release Policy for the Demographic and Health Surveys" (USAID). http://dhsprogram.com/pubs/pdf/SAR7/SAR7.pdf

J. Hu and J. Drechsler. 2015. "Generating synthetic geocoding information for public release." http://www.iab.de/389/section.aspx/Publikation/k150601301

# 5.5 Anonymization of the quasi-identifier household size

The size of a household is an important identifier, especially for large households.<sup>22</sup> Suppression of the actual size variable, if available (e.g., number of household members), however, does not suffice to remove this information from the dataset, as a simple count of the household members for a particular household will allow this variable to be reconstructed as long as a household ID is in the data. In any case, households of a very large size or with a unique or special key (i.e., combination of values of quasi-identifiers) should be checked manually. One way to treat them is to remove these households from the dataset before release. Alternatively, the households can be split, but care should be taken to suppress or change values for these households to prevent an intruder from immediately understanding that these households have been split and reconstructing them by combining the two households with the same values.

<sup>&</sup>lt;sup>22</sup> Even if the dataset does not contain an explicit variable with household size, this information can be easily extracted from the data and should be taken into account. Section 7.6 shows how to create a variable "household size" based on the household IDs.

# 5.6 Special case: census data

Census microdata are a special case because the user (and intruder) knows that all respondents are included in the dataset. Therefore, risk measures that use the sample weights and are based on uncertainty of the correctness of a match are no longer applicable. If an intruder has identified a sample unique and successfully matched, there is no doubt whether the match is correct, as it would be in the case of a sample. One approach to release census microdata is to release a stratified sample of the sample (1-5%) of the total census. NOTE: After sampling, the anonymization process has to be followed; sampling alone is not sufficient to guarantee confidentiality.

Several statistical offices release microdata based on census data. A few examples are:

• The British Office for National Statistics (ONS) released several

files based on the 2011 census: 1) A microdata teaching file for educational purposes. This file is a 1% sample of the total census with a limited set of variables. 2) Two scientific use files with 5% samples are available for registered researchers who accept the terms and conditions of their use. 3) Two 10% samples are available in controlled research data centers for approved researchers and research goals. All these files have been anonymized prior to release.<sup>23</sup>

• The U.S. Census Bureau released two samples of the 2000 census: a 5%

sample on the national level and a 1% sample on the state level. The national level file is more detailed, but the most detailed geographical area has at least 400,000 people. This, however, allows representation of all states from the dataset. The state-level file has less detailed variables but a more detailed geographical structure, which allows representation of cities and larger counties from the dataset (the minimum size of a geographical area is 100,000). Both files have been anonymized by using data swapping, top coding, perturbation and reducing detail by recoding.[#foot57]\_

<sup>23</sup> More information on census microdata at ONS is available on their website: http://www.ons.gov.uk/ons/guide-method/census/2011/census-data/census-microdata/index.html

### MEASURING UTILITY AND INFORMATION LOSS

SDC is a trade-off between risk of disclosure and loss of data utility and seeks to minimize the latter, while reducing the risk of disclosure to an acceptable level. Data utility in this context means the usefulness of the anonymized data for statistical analyses by end users as well as the validity of these analyses when performed on the anonymized data. Disclosure risk and its measurement are defined in Chapter 4 of this guide. In order to make a trade-off between minimizing disclosure risk and maximizing utility of data for end users, it is necessary to measure the utility of the data after anonymization and compare it with the utility of the original data. This chapter describes measures that can be used to compare the data utility before and after anonymization, or alternatively quantify the information loss. Information loss is the inverse of data utility: the larger the data utility after anonymization, the smaller the information loss. NOTE: If the microdata to be anonymized is based on a sample, the data will incur a sampling error. Also other errors may be present in the data, such as nonresponse error. NOTE: The methods discussed here only measure the information loss caused by the anonymization process relative to the original sample data and do not attempt to measure the error caused by other sources.

Ideally, the information loss is evaluated with respect to the needs and uses of the end users of the microdata. However, different end users of anonymized data may have very diverse uses for the released data and it might not be possible to collect an exhaustive list of these uses. Even if many uses can be identified, the characteristics in the data needed for these uses can be contradictory (e.g., one user needs a detailed geographical level whereas another is interested in a detailed age structure and does not need a detailed geographical structure). Nevertheless, as pointed out earlier, only one anonymized dataset can be released for each dataset and every type of release to avoid unintended disclosure. Releasing multiple anonymized datasets for different purposes may lead to unintended disclosure. <sup>1</sup> Therefore, it is not possible to anonymize and release a file tailored to each user's needs.

Since collecting and taking into account all data uses is often impossible, we also look at general (or generic) information loss measures besides user- and data-specific information loss measures. These measures do not take into account the specific data use, but can be used as guiding measures for information loss and evaluating whether a dataset is still analytically valid after anonymization. The main idea for such measures is to compare records between the original and treated datasets and compare statistics computed from both datasets (Hundepool et al., 2012). Examples of such measures are the number of suppressions, number of changed values, changes in contingency tables and changes in mean and covariance matrices.

Many of the SDC methods discussed earlier are parametric, in the sense that their outcome depends on parameters chosen by the user. Examples are the cluster size for microaggregation (see Section 5.3.2) or the importance vector in local suppression (see Section 5.2.2). Data utility and information loss measures are useful for choosing these parameters by comparing the impact of different parameters on the information loss. Figure 6.1 illustrates this by showing the trade-off between the disclosure risk and data utility of a hypothetical dataset. The triangle represents the original data with full utility and a certain level of disclosure risk, which is too high for disclosure. The square represents no release of microdata. Although there is no risk of disclosure, there is also no utility from the data for users since no data is released. The points in between represent the result of applying different SDC methods with

<sup>&</sup>lt;sup>1</sup> It is possible to release data files for different groups of users, e.g., PUF and SUF. All information in the less detailed file, however, must also be included in the more detailed file to prevent unintended disclosure. Datasets released in data enclaves can be customized for the user, since the risk that they will be combined with other version is zero.

different parameter specifications. We would select the SDC method corresponding to the point, which maximizes the utility, while keeping disclosure risk at an acceptable level.

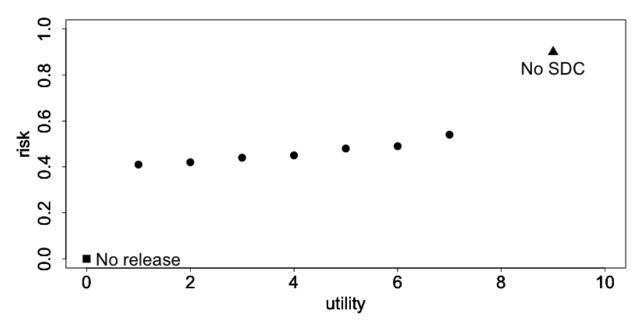


Figure 6.1: The trade-off between risk and utility in a hypothetical dataset

In the following sections, we first propose general utility measures independent of data use, and later present an example of a specific measure useful to measure information loss with respect to specific data uses. Finally, we show how to visualize changes in the data caused by anonymization and discuss the selection of utility measures for a particular dataset.

# 6.1 General utility measures for continuous and categorical variables

General or generic measures of information loss can be divided into those comparing the actual values of the raw and anonymized data, and those comparing statistics from both datasets. All measures are a posteriori, since they measure utility after anonymization and require both the data before and after the anonymization process. General utility measures are different for categorical and continuous variables.

### 6.1.1 General utility measures for categorical variables

#### Number of missing values

An informative measure is to compare the number of missing values in the data. Missing values are often introduced after suppression and more suppressions indicate a higher degree of information loss. After using the local suppression function on an *sdcMicro* object, the number of suppressions for each categorical key variable can be retrieved with the print() function, which is illustrated in Example 6.1[#foot59]. The argument 'ls' in the print() function stands for local suppression. The output shows both the absolute and relative number of suppressions.

Example 6.1: Using the print() function to retrieve the total number of suppressions for each categorical key variable

```
sdcInitial <- localSuppression(sdcInitial, k = 5, importance = NULL)
print(sdcInitial, 'ls')</pre>
```

(continues on next page)

(continued from previous page)

```
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
##
    URBRUR |
                              0 1
                                              0.000
##
    REGION /
                             81 /
                                              4.050
##
      RELIG |
                              0 1
                                              0.000
                              0 1
   MARITAL /
                                              0.000
```

More generally, it is possible to count and compare the number of missing values in the original data and the treated data. This can be useful to see the proportional increase in the number of missing values. Missing values can also have other sources, such as nonresponse. Example 6.2 shows how to display the number of missing values for each of the categorical key variables in an *sdcMicro* object. Here it is assumed that all missing values are coded 'NA'. If the missing values are not coded 'NA', but instead another value, it is possible to use the alternative missing values code. The results agree with the number of missing values introduced by local suppression in the previous example, but also shows that the variable "RELIG" has 1,000 missing values in the original data.

Example 6.2: Displaying the number of missing values for each categorical key variable in an *sdcMicro* object

```
| *# Store the names of all categorical key variables in a vector*
| namesKeyVars <- **names**\ (sdcInitial@manipKeyVars)</pre>
| *# Matrix to store the number of missing values (NA) before and after
  anonymization*
| NAcount <- **matrix**\ (NA, nrow = 2, ncol =
  **length**\ (namesKeyVars))
**colnames**\ (NAcount) <- **c**\ (**paste0**\ ('NA', namesKeyVars)) *#
column names*
**rownames**\ (NAcount) <- **C**\ ('initial', 'treated') *# row names*
| *# NA count in all key variables (NOTE: only those coded NA are
  counted) *
| for(i in 1:\ **length**\ (namesKeyVars))
| NAcount[1, i] <-
  **sum**\ (**is.na**\ (sdcInitial@origData[,namesKeyVars[i]]))
| NAcount[2, i] <- **sum**\ (**is.na**\ (sdcInitial@manipKeyVars[,i]))
*# Show results*
NAcount
 ``##
               NAURBRUR NAREGION NARELIG NAMARITAL``
 ``## initial
                      0
                               0
                                    1000
 ``## treated
                                                 51``
                      0
                              81
                                    1000
```

#### Number of records changed

Another useful statistic is the number of records changed per variable. These can be counted in a similar way as the missing values and include suppressions (i.e., changes to missing/'NA' in R). The number of records changed gives a good indication of the impact of the anonymization methods on the data. Example 6.3 illustrates how to compute the number of records changed for the PRAMmed variables.

#### Example 6.3: Computing number of records changed per variable

```
# Store the names of all pram variables in a vector
namesPramVars <- **names**\ (sdcInitial@manipPramVars)</pre>
# Dataframe to save the number of records changed
recChanged <- **rep**\ (0, **length**\ (namesPramVars))</pre>
names(recChanged) <- c(paste0('RC', namesPramVars))</pre>
# Count number of records changed*
for(j in 1:length(namesPramVars)) # for all key variables
comp <- sdcInitial@origData[namesPramVars[j]] !=_</pre>

→sdcInitial@manipPramVars[namesPramVars[j]]
temp1 <- sum(comp, na.rm = TRUE) # all changed variables without NAs
temp2 <- sum(is.na(comp)) # if NA, changed, unless NA initially
temp3 <- sum(is.na(sdcInitial@origData[namesPramVars[j]])</pre>
+ **is.na**\ (sdcInitial@manipPramVars[namesPramVars[j]])==2)
# both NA, no change, but counted in temp2
recChanged[j] <- temp1 + temp2 - temp3</pre>
# Show results
recChanged
    ## RCWATER RCROOF RCTOILET
        125
                86
```

#### Comparing contingency tables

A useful way to measure information loss in categorical variables is to compare univariate tabulations and, more interestingly, contingency tables (also cross tabulations or two-way tables) between pairs of variables. To maintain the analytical validity of a dataset, the contingency tables should stay approximately the same. The function table() produces contingency tables of one or more variables. Example 6.4 creates a contingency table of the variables "REGION" and "URBRUR". We observe small differences between the tables before and after anonymization.

Example 6.4: Comparing contingency tables of categorical variables

```
# Contingency table (cross tabulation) of the variables region and urban/rural
     table**\ (sdcInitial@origData[, **c**\ ('REGION', 'URBRUR')]) *
before anonymization*
```

```
## URBRUR
## REGION 1 2
## 1 235 89
## 2 261 73
## 3 295 76
## 4 304 71
## 5 121 139
## 6 100 236
```

table(sdcInitial@manipKeyVars[, c('REGION', 'URBRUR')]) # after anonymization

```
## URBRUR
## REGION 1 2
## 1 235 89
## 2 261 73
## 3 295 76
## 4 304 71
## 5 105 130
## 6 79 201
```

Domingo-Ferrer and Torra (2001) propose a Contingency Table-Based Information Loss (CTBIL) measure, which quantifies the distance between the contingency tables in the original and treated data. Alternatively, visualizations of the contingency table with mosaic plots can be used to compare the impact of anonymization methods on the tabulations and contingency tables (see Section 6.4.3).

### 6.1.2 General utility measures for continuous variables

#### Statistics: mean, covariance, correlation

The statistics characterizing the dataset should not change after the anonymization. Examples of such statistics are the mean, variance, and covariance and correlation structure of the most important variables in the dataset. Other statistics characterizing the data include the principal components and the loadings. Domingo-Ferrer and Torra (2001) give an overview of statistics that can be considered. In order to evaluate the information loss caused by the anonymization, one should compare the appropriate statistics for continuous variables computed from the data before and after anonymization. There are several ways to evaluate the loss of utility with respect to the changes in these statistics, for instance, by comparing means and (co-)variances in the data or comparing the (multivariate) distributions of the data. Especially changes in the correlations gives valuable information on the validity of the data for regressions. Functions from the *R* base package or any other statistical package can be used to do this. Following are a few examples in *R*.

To compute the mean of each numerical variable we use the function colMeans(). To ignore missing values, it is necessary to use the option na.rm = TRUE. "numVars" is a vector with the names of the numerical variables. Example 6.5 shows how to compute the means for all numeric variables. The untreated data is extracted from the 'origData' slot of the *sdcMicro* object and the anonymized data from the 'manipNumVars' slot, which contains the manipulated numeric variables. We observe small changes in each of the three variables.

Example 6.5: Comparing the means of continuous variables

# untreated data

**colMeans**(sdcInitial@origData[, numVars], na.rm = TRUE)

```
## INC INCRMT INCWAGE
## 479.7710 961.0295 1158.1330

# anonymized data
colMeans(sdcInitial@manipNumVars[, numVars], na.rm = TRUE)

## INC INCRMT INCWAGE
```

## 489.6030 993.8512 1168.7561

In the same way, one can compute the covariance and correlation matrices of the numerical variables in the *sdcMicro* object from the untreated and anonymized data. This is shown in Example 6.6. We observe that the variance of each variable (the diagonal elements in the covariance matrix) have increased by the anonymization. These functions also allow computing confidence intervals in the case of samples. The means and covariances of subsets in the data also should not differ. An example is the mean of income by gender, by age group or by region. These characteristics of the data are important for analysis.

Example 6.6: Comparing covariance structure and correlation matrices of numeric variables

```
# untreated data
cov(sdcInitial@origData[, numVars])
## INC INCRMT INCWAGE
## INC 1645926.1 586975.6 2378901
## INCRMT 586975.6 6984502.3 1664257
## INCWAGE 2378900.7 1664257.4 16169878
cor(sdcInitial@origData[, numVars])
## INC INCRMT INCWAGE
## INC 1.0000000 0.1731200 0.4611241
## INCRMT 0.1731200 1.0000000 0.1566028
## INCWAGE 0.4611241 0.1566028 1.0000000
# anonymized data
cov(sdcInitial@manipNumVars[, numVars])
## INC INCRMT INCWAGE
## INC 2063013.1 649937.5 2382447
## INCRMT 649937.5 8566169.1 1778985
## INCWAGE 2382447.4 1778985.1 19925870
cor(sdcInitial@manipNumVars[, numVars])
## INC INCRMT INCWAGE
## INC 1.0000000 0.1546063 0.3715897
## INCRMT 0.1546063 1.0000000 0.1361665
## INCWAGE 0.3715897 0.1361665 1.0000000
```

Domingo-Ferrer and Torra (2001) propose several measures for the discrepancy between the covariance and correlation matrices. These measures are based on the mean squared error, the mean absolute error or the mean variation of the individual cells. We refer to Domingo-Ferrer and Torra (2001) for a complete overview of these measures.

#### **IL1s information loss measure**

Alternatively, we can also compare the actual data and quantify the distance between the original dataset X and the treated dataset Z. Here X and Z contain only continuous variables. Yancey, Winkler and Creecy (2002) introduce the distance measure IL1s, which is the sum of the absolute distances between the corresponding observations in the raw and anonymized datasets, which are standardized by the standard deviation of the variables in the original data. For the continuous variables in the dataset, the IL1s measure is defined as

$$IL1s = \frac{1}{\mathrm{pn}} \sum_{j=1}^p \sum_{i=1}^n \frac{|x_{\mathrm{ij}} - z_{\mathrm{ij}}|}{\sqrt{2}S_j}$$
 ,

where p is the number of continuous variables; n is the number of records in the dataset;  $x_{ij}$  and  $z_{ij}$ , respectively, are the values before and after anonymization for variable j and individual i; and  $S_j$  is the standard deviation of variable j in the original data (Yancey, Winkler and Creecy, 2002).

When using *sdcMicro*, the IL1s data utility measure can be computed for all numerical quasi-identifiers with the function dUtility(), which is illustrated in Example 6.7. If required, the measure can also be computed on subsets of the complete set of numerical quasi-identifiers. The function is called dUtility(), but returns a measure of information loss. The result is saved in the utility slot of the *sdcMicro* object. Example 6.7 also illustrates how to call the result.

Example 6.7: Using dUtility() to compute IL1s data utility measure in sdcMicro

```
# Evaluating IL1s measure for all variables in the sdcMicro object sdcInitial
sdcInitial <- dUtility(sdcInitial)
# Calling the result of IL1s
sdcInitial@utility$il1

## [1] 0.2203791

# IL1s for a subset of the numerical quasi-identifiers
subset <- c('INCRMT', 'INCWAGE', 'INCFARMBSN')
dUtility(obj = sdcInitial@origData[,subset], xm = sdcInitial@manipNumVars[,subset], method = 'IL1')
```

```
## [1] 0.5641103
```

The measure is useful for comparing different methods. The smaller the value of the measure, the closer the values are to the original values and the higher the utility. **NOTE: This measure is related to risk measures based on distance and intervals (see Section 4.7).** The greater the distance between the original and anonymized values, the lower the data utility. Greater distance, however, also reduces the risk of re-identification.

#### **Eigenvalues**

Another way to evaluate the information loss is to compare the robust eigenvalues of the data before and after anonymization. Example 6.8 illustrates how to use this approach with *sdcMicro*. Here "contVars" is a vector with the names of the continuous variables in which we are interested. "obj" is the argument that specifies the untreated data and "xm" is the argument that specifies the anonymized data. The function's output is the difference in eigenvalues. Therefore, the minimum value is 0. Again, the main use is to compare different methods. The greater the value, the greater the changes in the data and the information loss.

Example 6.8: Using dUtility() to compute eigenvalues in sdcMicro

```
# Comparison of eigenvalues of continuous variables

dUtility(obj = sdcInitial@origData[,contVars], xm = sdcInitial@manipNumVars[,contVars], method = 'eigen')

## [1] 2.482948

# Comparison of robust eigenvalues of continuous variables

dUtility(obj = sdcInitial@origData[,contVars], xm = sdcInitial@manipNumVars[,contVars], method = 'robeigen')

## [1] -4.297621e+14
```

## 6.2 Utility measures based on the end user's needs

Not all needs and uses of a certain dataset can be inventoried. Nevertheless, some types of data have similar uses or important characteristics, which can be evaluated before and after anonymization. Examples of such "benchmarking indicators" (Templ et al., 2014) are different for each dataset. Examples include poverty measures for income datasets and school attendance ratios. Often ideas for selecting such indicators come from the reports data users publish based on previously released microdata.

The approach is to compare the indicators calculated on the untreated data and the data after anonymization with different methods. If the differences between the indicators are not too large, the anonymized dataset can be released for use by researchers. It should be taken into account that indicators calculated on samples are estimates with a certain variance and confidence interval. Therefore, for sample data, it is more informative to compare the overlap of confidence intervals and/or to evaluate whether the point estimate calculated after anonymization is contained within the confidence interval of the original estimate. Examples of benchmark indicators and their confidence intervals and how to compute these in *R* are included in the case studies in these guidelines. Here we give the example of the GINI coefficient.

The GINI coefficient is a measure of statistical dispersion, which is often used to measure inequality in income. A way to measure the information loss in income data is to compare the income distribution, which can be easily done by comparing the GINI coefficients. Several *R* packages have functions to compute the GINI coefficient. We chose the *laeken* package, which computes the GINI coefficient as the area between the 45-degree line and the Lorenz curve. To use the gini() function, we first have to install and load the *laeken* library. To calculate the GINI coefficient for the variable income, we use the sample weights in the data. This is shown in Example 6.9. The GINI coefficient of sample data is a random variable. Therefore, it is useful to construct a confidence interval around the coefficient to evaluate the significance of any change in the coefficient after anonymization. The gini() function computes a 1-alpha confidence interval for the GINI coefficient by using bootstrap.

Example 6.9: Computing the GINI coefficient from the income variable to determine income inequality

```
# Gini coefficient before anonymization

gini(inc = sdcInitial@origData[selInc,'INC'], weights = curW[selInc], na.rm = TRUE)$value # before

## [1] 34.05928

# Gini coefficient after anonymization
gini(inc = sdcInitial@manipNumVars[selInc,'INC'], weights = curW[selInc], na.rm = TRUE)$value # after
```

## [1] 67.13218

Regression

Besides comparing covariance and correlation matrices, regressions are a useful tool to evaluate whether the structure in the data is maintained after anonymization. By comparing regressions parameters, it is also possible to compare relations between non-continuous variables (e.g., by introducing dummy variables or regression with ordinal variables). If it is known for what purpose and in what field the data is used, common regressions can be used to compare the change in coefficients and confidence intervals.

An example of using regression to evaluate the data utility in income data is the Mincer equation. The Mincer equation explains earnings as a function of education and experience while controlling for other variables. The Mincer equation is often used to evaluate the gender pay gap and gender wage inequality by including a gender dummy. Here we show how to evaluate the impact of anonymization methods on the gender coefficient. We regress the log income on a constant, a gender dummy, years of education, years of experience, years of experience squared and other factors influencing wage.

$$\ln (\text{wage}) = \beta_0 + \beta_1 gender + \beta_2 education + \beta_3 experience + \beta_3 experience^2 + \beta X$$

The parameter of interest here is  $\beta_1$ , the effect of gender on the log wage. X is a matrix with several other factors influencing wage and  $\beta$  the coefficients of these factors. Example 6.10 illustrates how to run a Mincer regression in R using the function lm() and evaluate the coefficients and confidence intervals around the coefficients. We run the regression as specified for paid employees with a positive wage in the age groups 15-65 years.

Example 6.10: Estimating the Mincer equation (regression) to evaluate data utility before and after anonymization

# Mincer equation variables before anonymization

Mlwage <- log(sdcMincer@origData\$wage) # log wage

Mempstat <- sdcMincer@origData\empstat=='Paid employee' # TRUE if 'paid employee', else FALSE or NA

Mage <- sdcMincer@origData\$age # age in years

Meducy <- sdcMincer@origData\$educy # education in years

Mexp <- sdcMincer@origData\$exp # experience in years

Mexp2 <- Mexp^2 # squared experience

Mgender <- sdcMincer@origData\$gender # gender dummy

Mwgt <- sdcMincer@origData\$wgt # weight variable for regression

MfileB <- as.data.frame(cbind(Mlwage, Mempstat, Mage, Meducy, Mexp, Mexp2, Mgender, Mwgt))

# Mincer equation variables after anonymization

Mlwage <- log(sdcMincer@manipNumVars\$wage) # log wage

Mempstat <- sdcMincer@manipKeyVars\$empstat=='Paid employee'

#TRUE if 'paid employee', else FALSE or NA

Mage <- sdcMincer@manipKeyVars\$age # age in years

Meducy <- sdcMincer@manipKeyVars\$educy # education in years

Mexp <- sdcMincer@manipKeyVars\$exp # experience in years

 $Mexp2 <- Mexp^2 \# squared \ experience$ 

Mgender <- sdcMincer@manipKeyVars\$gender # gender dummy

Mwgt <- sdcMincer@origData\$wgt # weight variable for regression

MfileA <- as.data.frame(cbind(Mlwage, Mempstat, Mage, Meducy, Mexp, Mexp2, Mgender, Mwgt))

# Specify regression formula

```
Mformula <- 'Mlwage ~ Meducy + Mexp + Mexp2 + Mgender'
# Regression Mincer equation
mincer1565B <- lm(Mformula, data = subset(MfileB, MfileB$Mage >= 15 & MfileB$Mage <= 65 &
MfileB$Mempstat==TRUE & MfileB$Mlwage != -Inf), na.action = na.exclude, weights = Mwgt) # before
mincer1565A <- lm(Mformula, data = subset(MfileA, MfileA$Mage >= 15 & MfileA$Mage <= 65 &
MfileA$Mempstat==TRUE & MfileA$Mlwage != -Inf), na.action = na.exclude, weights = Mwgt) # after
# The objects mincer1565B and mincer1565A contain the results of the regressions before and after anonymization
mincer1565B$coefficients # before
## (Intercept) Meducy Mexp Mexp2 Mgender
## 3.9532064886 0.0212367075 0.0255962570 -0.0005682651 -0.4931289413
mincer1565A$coefficients # after
## (Intercept) Meducy Mexp Mexp2 Mgender
## 4.0526250282 0.0141090329 0.0326711056 -0.0007605492 -0.5393641862
# Compute the 95 percent confidence interval
confint(obj = mincer1565B, level = 0.95) # before
## 2.5 % 97.5 %
## (Intercept) 3.435759991 4.4706529860
## Meducy -0.018860497 0.0613339120
## Mexp 0.004602597 0.0465899167
## Mexp2 -0.000971303 -0.0001652273
## Mgender -0.658085143 -0.3281727396
confint(obj = mincer1565A, level = 0.95) # after
## 2.5 % 97.5 %
## (Intercept) 3.46800378 4.6372462758
## Meducy -0.03305743 0.0612754964
## Mexp 0.01024867 0.0550935366
## Mexp2 -0.00119162 -0.0003294784
```

If the new estimates fall within the original confidence interval and the new and original confidence intervals are greatly overlapping, the data can be considered valid for this type of regression after anonymization. Figure 6.2 shows the point estimates and confidence intervals for the gender coefficient in this trade-off for a sample income dataset and several SDC methods and parameters. The red dot and confidence bar (on the top) correspond to the estimates for the untreated data, whereas the other confidence bars correspond to the respective SC methods and different parameters. The anonymization reduces the number of expected re-identifications in the data (left axis) and the point estimates and confidence intervals vary greatly for the different SDC methods. We would choose a method, which reduces

## Mgender -0.71564602 -0.3630823543

the expected number of identifications, while not changing the gender coefficient and having a great overlap of the confidence interval with the confidence interval estimated from the original data.

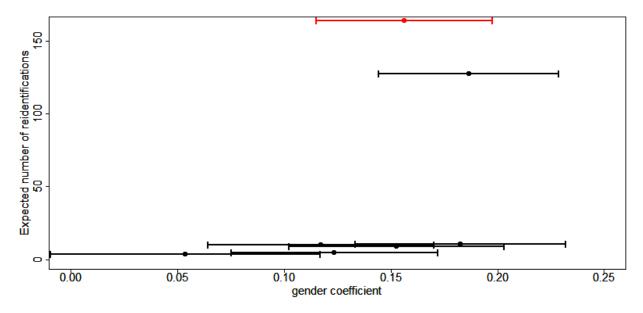


Figure 6.2: Effect of anonymization on the point estimates and confidence interval of the gender coefficient in the Mincer equation

## 6.3 Assessing data utility with the help of data visualizations (in R)

The use of graphs and other visualization techniques is a good way to assess at a glance how much the data have changed after anonymization, and can aid the selection of appropriate anonymization techniques for the data. Visualizations can be a useful tool to assess the impact on data utility of anonymization methods and helps choose among anonymization methods. The software package *R* provides several functions and packages that can help visualize the results of anonymization. This section lists a few of these functions and packages and provides code examples to illustrate how to implement them. We present the following visualizations:

- histograms and density plots
- boxplots
- · mosaic plots

To make appropriate visualizations, we need to use the raw data and the anonymized data. When using an *sdcMicro* object for the anonymization process, the raw data are stored in the "origData" slot of the object and the anonymized variables are in the slots "manipKeyVars", "manipPramVars", "manipNumVars" and "manipStrataVar" slots. See Section 7.5 for more information on *sdcMicro* objects, slots and how to access slots.

### 6.3.1 Histograms and density plots

Histograms and density plots are useful for quick comparisons of variable distribution before and after anonymization. The advantage of histograms is that the results are exact. Visualization depends on the bin widths and the start point of the first bin, however. Histograms can be used for continuous and semi-continuous variables. Density plots display the kernel density of the data; therefore, the plot depends on the kernel that is chosen and whether the data fits the kernel well. Nevertheless, density plots are a good tool to illustrate the change of values and value ranges of continuous variables.

Histograms can be plotted with function hist() and kernel densities with the functions plot() and density() in *R*. Example 6.11 provides examples of how to use these functions to illustrate the changes in the variable "INC", an income variable. The function hist() needs as argument the break points for the histogram. The results are shown in Figure 6.3 and Figure 6.4. The histograms and density plots give a clear indication how the values have changed: the variability of the data has increased and the shape of the distribution has changed. **NOTE: The vertical axes of the histograms have different scales.** 

Example 6.11: Plotting histograms and kernel densities

# Plot histograms

# Plot histogram before anonymization

hist(sdcInitial@origData\$INC, breaks = (0:180)\*1e2, main = "Histogram income - original data")

# Plot histogram after anonymization (noise addition)

hist(sdcInitial@manipNumVars\$INC, breaks = (-20:190)\*1e2, main = "Histogram income - anonymized data")

# Plot densities

# Plot original density curve

plot(density(sdcInitial@origData\$INC), xlim = c(0, 8000), ylim = c(0, 0.006), main = "Density income", xlab = "income")

par(new = TRUE)

# Plot density curve after anonymization (noise addition)

plot(density(sdcInitial@manipNumVars\$INC), xlim = c(0, 8000), ylim = c(0, 0.006), main = "Density income", xlab = "income")

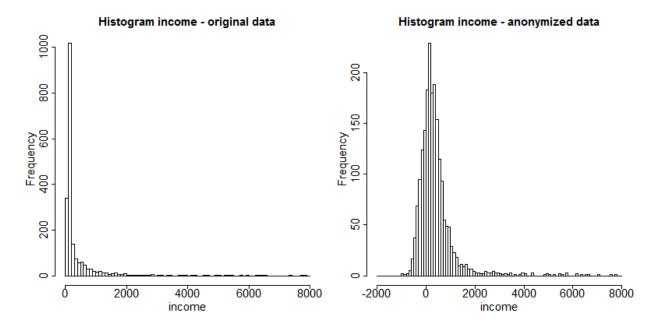


Figure 6.3: Histograms of income before and after anonymization

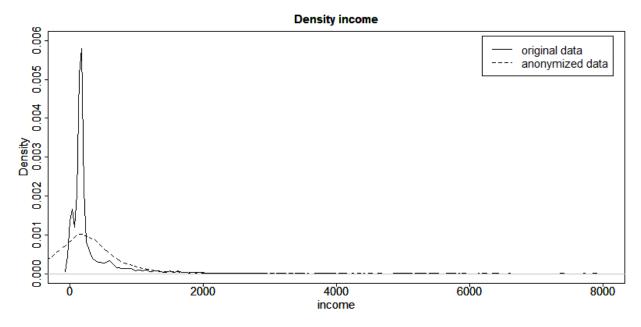


Figure 6.4: Density plots of income before and after anonymization

### 6.3.2 Box plots

Box plots give a quick overview of the changes in the spread and outliers of continuous variables before and after anonymization. Example 6.12 shows how to generate box plots in *R* with the function boxplot(). The result in Figure 6.5 shows an example for an expenditure variable after adding noise. The box plot shows clearly that the variability in the expenditure variable increased as a result of the anonymization methods applied.

Example 6.12: Creating boxplots for continuous variables

**boxplot**(sdcObj@origData\$TOTFOOD, sdcObj@manipNumVars\$TOTFOOD, xaxt = 'n', ylab = "Expenditure") **axis**(1, at = **c**(1,2), labels = **c**('before, 'after'))

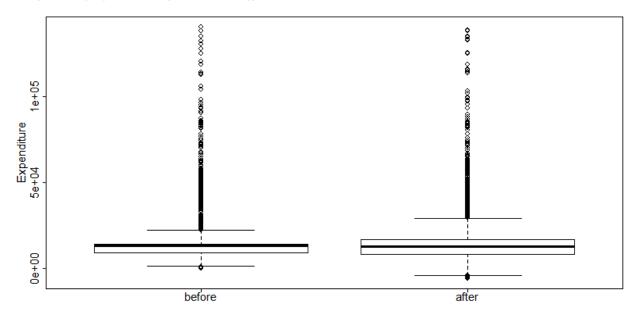


Figure 6.5: Example of box plots of an expenditure variable before and after anonymization

### 6.3.3 Mosaic plots

Univariate and multivariate mosaic plots are useful for showing changes in the tabulations of categorical variables, especially when comparing several "scenarios" next to one another. A scenario here refers to the choice of anonymization methods and their parameters. With mosaic plots we can, for instance, quickly see the effect of different levels of k-anonymity or differences in the importance vectors in the local suppression algorithm (see Section 5.2.2).

We illustrate the changes in tabulations with an example of the variable "WATER" before and after applying PRAM. We can use mosaic plots to quickly see the changes for each category. Example 6.13 shows the code in R. The function mosaic plot() is available in base R. To plot a tabulation, first the tabulation must be made with the table() function. To show the labels in the mosaic plot(), we change the class of the variables to 'factor' (see Section 7.4 on classes in R). Looking at the mosaic plot in Figure 6.6 we see invariant PRAM has virtually no influence on the univariate distribution.

Example 6.13: Creating univariate mosaic plots

```
# Collecting data of variable WATER before and after anonymization, assigning factor levels for labels in plot dataWater <- t(cbind(table(factor(sdcHH@origData$WATER, levels = c(1, 2, 3, 4, 5, 6, 7, 8, 9), labels = c("Pipe (own tap)", "Public standpipe", "Borehole", "Wells (protected)", "Wells (unprotected)", "Surface water", "Rain water", "Vendor/truck", "Other"))), table(factor(sdcHH@manipPramVars$WATER, levels = c(1, 2, 3, 4, 5, 6, 7, 8, 9), labels = c("Pipe (own tap)", "Public standpipe", "Borehole", "Wells (protected)", "Wells (unprotected)", "Surface water", "Rain water", "Vendor/truck", "Other")))))

rownames(dataWater) <- c("before", "after")

# Plotting mosaic plot
mosaicplot(dataWater, main = "", color = 2:10, las = 2)
```

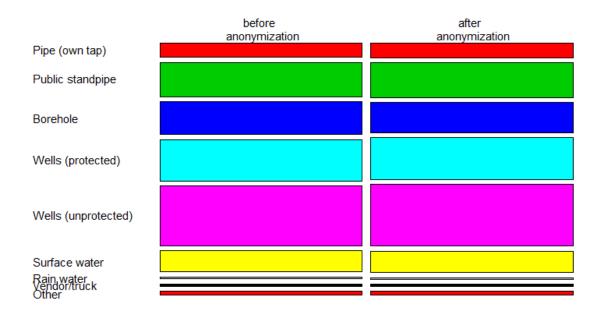


Figure 6.6: Mosaic plot to illustrate the changes in the WATER variable

We use the variables "gender" and "relationship status" to illustrate the use of mosaic plots for the illustration of changes in univariate tabulations introduced by several sets of anonymization methods. Table 6.1 provides the methods applied in each scenario. Scenario 0, the base scenario, shows the original categories of the gender and relationship status variables, while scenarios 1 to 6 show shifts in the categories after applying different anonymization techniques. Table 6.1 provides a description of the anonymization methods used in each scenario. In total we visualize the impact of six different sets of anonymization methods. We can use mosaic plots to quickly see which set of methods has what impact on the gender and relationship status variables, which can be used to select the best scenario. Looking at the mosaic plots in Figure 6.7, we see that scenarios 2, 5 and 6 give the smallest changes for the gender variable and scenarios 3 and 4 for the relationship status variable.

Table 6.1: Description of anonymization methods by scenario

Scenario	Description of anonymization methods applied
0 (base)	Original data, no treatment

0 (base) | Original data, no treatment |

- Recode age (five-year intervals), plus local suppression (required k = 3, high importance on toilet), while also recoding region, urban, education level and occupation variables
   Recode age (five-year steps), plus local suppression (required k = 5, high importance on water, toilet and
- literacy), while also recoding region, urban, education level and occupation variables
- Recode age (five-year intervals), plus local suppression (required k = 3, no importance vector), microaggregation (wealth index), while also recoding region, urban, education level and occupation variables
- Recode age (five-year intervals) plus local suppression (required k=3, no importance vector), PRAM literacy, while also recoding region, urban, education level and occupation variables

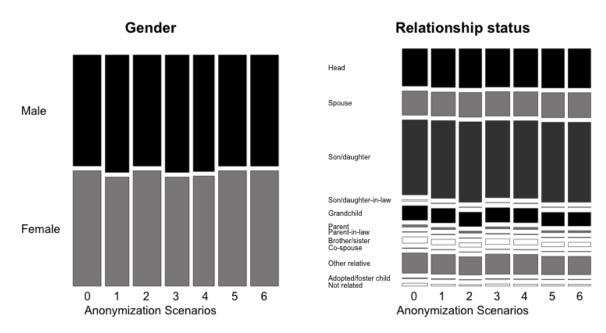


Figure 6.7: Comparison of treated vs. untreated gender and relationship status variables with mosaic plots

As we discussed in Section 5.3.1, invariant PRAM preserves the univariate distributions. Therefore, in this case it is more interesting to look at the multivariate mosaic plots. Mosaic plots are also a powerful tool to show changes in cross-tabulations/contingency tables. Example 6.14 shows how to generate mosaic plots for two variables. To compare the changes, we need to compare two different plots. Figure 6.8 and Figure 6.9 illustrate that (invariant) PRAM does not preserve the two-way tables in this case.

Example 6.14: Creating multivariate mosaic plots

```
# Before anonymization: contingency table and mosaic plot
ROOFTOILETbefore \leftarrow t(table(factor(sdcHH@origData$ROOF, levels = c(1,2,3,4,5,9),
labels = \mathbf{c}("Concrete/cement/\n brick/stone", "Wood",
"Bamboo/thatch", "Tiles/shingles",
"Tin/metal sheets", "Other")),
factor(sdcHH@origData$TOILET, levels = c(1,2,3,4,9),
labels = \mathbf{c}("Flush \n toilet", "Improved \n pit \n latrine",
"Pit \n latrine", "No \n facility", "Other"))))
mosaicplot(ROOFTOILETbefore, main = "", las = 2, color = 2:6)
# After anonymization: contingency table and mosaic plot
ROOFTOILETafter \leftarrow t(table(factor(sdcHH@manipPramVars$ROOF, levels = c(1,2,3,4,5,9),
labels = \mathbf{c}("Concrete/cement/ \n brick/stone", "Wood",
"Bamboo/thatch", "Tiles/shingles",
"Tin/metal sheets", "Other")),
factor(sdcHH@manipPramVarsTOILET, levels = c(1,2,3,4,9),
labels = \mathbf{c}("Flush \n toilet", "Improved \n pit \n latrine",
"Pit \n latrine", "No \n facility", "Other"))))
mosaicplot(ROOFTOILETafter, main = "", las = 2, color = 2:6)
```

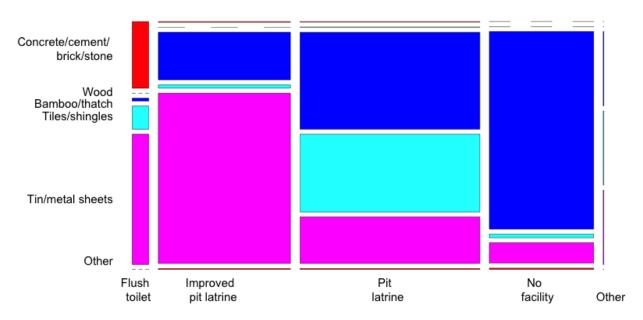


Figure 6.8: Mosaic plot of the variables ROOF and TOILET before anonymization

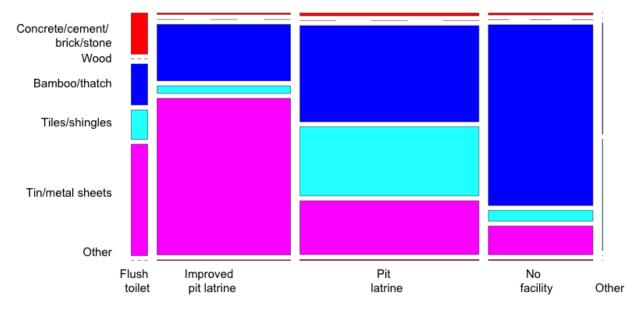


Figure 6.9: Mosaic plot of the variables ROOF and TOILET after anonymization

# 6.4 Choice of utility measure

Besides the users' requirements on the data, the utility measures should be chosen in accordance with the variable types and anonymization methods employed. The employed utility measures can be a combination of both general and user-specific measures. As discussed earlier, different utility measures should be used for continuous and categorical data. Furthermore, some utility measures are not informative after certain anonymization methods have been applied. For example, after applying perturbative methods that interchange data values, comparing values directly is not useful because they will give the impression of high levels of information loss. In such cases, it is more informative to look at

#### **Readthedocs Documentation**

means, covariances and benchmarking indicators that can be computed from the data. Furthermore, it is important not only to focus on the characteristics of variables one by one, but also on the interactions between variables. This can be done by cross-tabulations and regressions. In general, when anonymizing sampled data, it is advisable to compute confidence intervals around estimates to interpret the magnitude of changes.

#### **Recommended Reading Material on Measuring Utility and Information Loss**

- A.G. De Waal and L.C.R.J. Willenborg. 1999. "Information Loss through Global Recoding and Local Suppression" In Netherlands Official Statistics: Special Issue on SDC, 14, 17-10.
- J. Domingo-Ferrer, J.M. Mateo-Sanz and V. Torra. 2001. "Comparing SDC Methods for Microdata on the basis of Information Loss and Disclosure Risk". In Pre-proceedings of ETK-NTTS 2001 (vol. 2), 807-826. http://neon.vb.cbs.nl/casc/NTTSJosep.pdf
- J. Domingo-Ferrer and V. Torra. 2001. "Disclosure Protection Methods and Information Loss for Microdata". In P. Doyle, J.I. Lane, J.J.M. Theeuwes and L. Zayatz (eds.) *Theory and Practical Applications for Statistical Agencies*, 91-110, Amsterdam. http://crises-deim.urv.cat/webCrises/publications/bcpi/cliatpasa01Disclosure.pdf

### SDC WITH SDCMICRO IN R: SETTING UP YOUR DATA AND MORE

# 7.1 Installing R, sdcMicro and other packages

This guide is based on the software package sdcMicro, which is an add-on package for the statistical software R. Both R and sdcMicro, as well as other R packages, are freely available from the CRAN (Comprehensive R Archive Network) website for Linux, Mac and Windows (http://cran.r-project.org). This website also offers descriptions of packages. Besides the standard version of R, there is a more user-friendly user interface for R: RStudio. RStudio is also freely available for Linux, Mac and Windows (http://www.rstudio.com). The sdcMicro package is dependent on (i.e., uses) other R packages that must be installed on your computer before using sdcMicro. Those will automatically be installed when installing sdcMicro. For some functionalities, we use still other packages (such as foreign for reading data and some graphical packages). If so, this is indicated in the appropriate section in this guide. R, RStudio, the sdcMicro package and its dependencies and other packages have regular updates. It is strongly recommended to regularly check for updates: this requires installing a new version for an update of R; with the update.packages() command or using the menu options in R or RStudio one can update the installed packages.

When starting *R* or *RStudio*, it is necessary to specify each time which packages are being used by loading those. This loading of packages can be done either with the library() or the require() function. Both options are illustrated in Example 7.1.

Example 7.1: Loading required packages

library(sdcMicro) # loading the sdcMicro package

require(sdcMicro) # loading the sdcMicro package

All packages and functions are documented. The easiest way to access the documentation of a specific function is to use the built-in help, which generally gives an overview of the parameters of the functions as well as some examples. The help of a specific function can be called by a question mark followed by the function name without any arguments. Example 7.2 shows how to call the help file for the microaggregation() function of the *sdcMicro* package.[#foot60]\_ The download page of the each package on the CRAN website also provides a reference manual with a complete overview of the functions in the package.

Example 7.2: Displaying help for functions

?microaggregation # help for microaggregation function

When issues or bugs in the *sdcMicro* package are encountered, comments, remarks or suggestions can be posted for the developers of *sdcMicro* on their GitHub: https://github.com/alexkowa/sdcMicro/issues

### 7.2 Read functions in R

The first step in the SDC process when using sdcMicro is to read the data into R and create a dataframe.  $^2$  R is compatible with most statistical data formats and provides read functions for most types of data. For those read functions, it is sometimes necessary to install additional packages and their dependencies in R. An overview of data formats, functions and the packages containing these functions is provided in Table 7.1. These functions are also available as write (e.g., write.dta()) to save the anonymized data in the required format. #foot[62]\_

Table 7.1: Packages and functions for reading data in R

**Type/software **	Extension	Package	Function
SPSS	.sav	foreign	read.spss()
STATA (v. 5-12)	.dta	foreign	read.dta()
STATA (v. 13) [#foot64]	.dta	readstata13	read.dta13()
SAS	.xpt	foreign	read.xport()
EXCEL	.csv	utils (base package)	read.csv()
EXCEL	.xls	xlsReadWrite	read.xls()
EXCEL	.xlsx	xlsx	read.xlsx()

Most of these functions have options that specify how to handle missing values and variables with factor levels and value labels. Example 7.3, Example 7.4 and Example 7.5 provide example code for reading in a *STATA* (.dta) file, an *Excel* (.csv) file and a *SPSS* (.sav) file.

Example 7.3: Reading in a STATA file

setwd("/Users/World Bank") # working directory with data file

fname = "data.dta" # name of data file

library(foreign) # loads required package for read/write function for STATA files

file <- read.dta(fname, convert.factors = F)

# reads the data into the data frame called file, factor levels read as numeric codes

Example 7.4: Reading in an Excel file

setwd("/Users/World Bank") # working directory with data file

fname = "data.csv" # name of data file

file <- read.csv(fname, header = TRUE, sep = ",", dec = ".")

# reads the data into the data frame called file, the first line contains the variable names, fields are separated with commas, decimal points are indicated with '.'

Example 7.5: Reading in an SPSS file

setwd("/Users/World Bank") # working directory with data file

fname = "data.sav" # name of data file

library(foreign) # loads required package for read/write function for SPSS files

file <- read.spss(fname, use.value.labels = FALSE)

<sup>&</sup>lt;sup>2</sup> A dataframe is an object class in *R*, which is similar to a data table or matrix.

# reads the data into the data frame called file, factor levels are read as numeric codes

The maximum data size in *R* is technically restricted. The maximum size depends on the *R* build (32-bit or 64-bit) and the operating system. Some SDC methods require long computation times for large datasets (see Section 7.7 on computation times).

### 7.3 Missing values

The standard way missing values are represented in *R* is by the symbol 'NA', which is different to impossible values, such as division by zero or the log of a negative number, which are represented by the symbol 'NaN'. The value 'NA' is used for both numeric and categorical variables.[#foot65]\_ Values suppressed by the localSuppression() routine are also replaced by the 'NA' symbol. Some datasets and statistical software might use different values for missing values, such as '999' or strings. It is possible to include arguments in read functions to specify how missing values in the dataset should be treated and automatically recode missing values to 'NA'. For instance, the function read.table() has the 'na.strings' argument, which replaces the specified strings with 'NA' values.

Missing values can also be recoded after reading the data into *R*. This may be necessary if there are several different missing value codes in the data, different missing value codes for different variables or the read function for the datatype does not allow specifying the missing value codes. When preparing data, it is important to recode any missing values that are not coded as 'NA' to 'NA' in *R* before starting the anonymization process to ensure the correct measurement of risk (e.g., *k*-anonymity), as well as to ensure that many of the methods are correctly applied to the data. Example 7.6 shows how to recode the value '99' to 'NA' for the variable "toilet".

Example 7.6: Recoding missing values to NA

file[file[,'toilet'] == 99,'toilet'] <- NA # Recode missing value code 99 to NA for variable toilet

#### 7.4 Classes in R

All objects in *R* are of a specific class, such as integer, character, matrix, factor or dataframe. The class of an object is an attribute from which the object inherits. To find out the class of an object, one can use the function class(). Functions in *R* might require objects or arguments of certain classes or functions might have different functionality depending on the class of the argument. Examples are the write functions that require dataframes and most functions in the *sdcMicro* package that require either dataframes or *sdcMicro* objects. The functionality of the functions in the *sdcMicro* package differs for dataframes and *sdcMicro* objects. It is easy to change the class attribute of an object with functions that start with "as.", followed by the name of the class (e.g., as.factor(), as.matrix(), as.data.frame()). Example 7.7 shows how to check the class of an object and change the class to "data.frame". Before changing the class attribute of the object "file", it was in the class "matrix". An important class defined and used in the *sdcMicro* package is the class named *sdcMicroObj*. This class is described in the next section.

Example 7.7: Changing the class of an object in R

# Finding out the class of the object 'file' class(file)

# Changing the class to data frame

"matrix"

7.3. Missing values 89

```
file <- as.data.frame(file)

# Checking the result

class(file)

"data.frame"
```

## 7.5 Objects of class sdcMicroObj

The *sdcMicro* package is built around objects[#foot66]\_ of class *sdcMicroObj*, a class especially defined for the *sdcMicro* package. Each member of this class has a certain structure with slots that contain information regarding the anonymization process (see Table 7.2 on page 100 for a description of all slots). Before evaluating risk and utility and applying SDC methods, creating an object of class *sdcMicro* is recommended. All examples in this guide are based on these objects. The function used to create an *sdcMicro* object is createSdcObj(). Most functions in the *sdcMicro* package, such as microaggregation() or localSuppression(), automatically use the required information (e.g., quasi-identifiers, sample weights) from the *sdcMicro* object if applied to an object of class *sdcMicro*.

The arguments of the function createSdcObj() allow one to specify the original data file and categorize the variables in this data file before the start of the anonymization process. **NOTE:** For this, disclosure scenarios must already have been evaluated and quasi-identifiers selected. In addition, one must ensure there are no problems with the data, such as variables containing only missing values.

In Example 7.8, we show all arguments of the function createSdcObj(), and first define vectors with the names of the different variables. This practice gives a better overview and later allows for quick changes in the variable choices if required. We choose the categorical quasi-identifiers (keyVars); the variables linked to the categorical quasi-identifiers that need the same suppression pattern (ghostVars, see Section 5.2.2); the numerical quasi-identifiers (numVars); the variables selected for applying PRAM (pramVars); a variable with sampling weights (weightVar); the clustering ID (hhId, e.g., a household ID, see Section 4.9); a variable specifying the strata (strataVar) and the sensitive variables specified for the computation of *l*-diversity (sensibleVar, see Section 4.5.3). **NOTE: Most SDC methods in the sdcMicro package are automatically applied within the strata, if the 'strataVar' argument is specified.** Examples are local suppression and PRAM. Not all variables must be specified, e.g., if there is no hierarchical (household) structure, the argument 'hhId' can be omitted. The names of the variables correspond to the names of the variables in the dataframe containing the microdata to be anonymized. The selection of variables is important for the risk measures that are automatically calculated. Furthermore, several methods are by default applied to all variables of one sort, e.g., microaggregation to all key variables.[#foot67]\_ After selecting these variables, we can create the *sdcMicro* object. To obtain a summary of the object, it is sufficient to write the name of the object.

Example 7.8: Selecting variables and creating an object of class *sdcMicroObj* for the SDC process in R

```
# Select variables for creating sdcMicro object

# All variable names should correspond to the names in the data file

# selected categorical key variables

selectedKeyVars = c('region', 'age', 'gender', 'marital', 'empstat')

# selected linked variables (ghost variables)

selectedGhostVars = c('urbrur')

# selected categorical numerical variables

selectedNumVar = c('wage', 'savings')

# weight variable

selectedWeightVar = c('wgt')
```

```
# selected pram variables
selectedPramVars = c('roof', 'wall')
# household id variable (cluster)
selectedHouseholdID = c('idh')
# stratification variable
selectedStrataVar = c('strata')
# sensitive variables for l-diversity computation
selectedSensibleVar = c('health')
# creating the sdcMicro object with the assigned variables
sdcInitial <- createSdcObj(dat = file,
keyVars = selectedKeyVars,
ghostVars = selectedGhostVars,
numVar = selectedNumVar,
weightVar = selectedWeightVar,
pramVars = selectedPramVars,
hhId = selectedHouseholdID,
strataVar = selectedStrataVar,
sensibleVar = selectedSensibleVar)
# Summary of object
sdcInitial
## Data set with 4580 rows and 14 columns.
## --> Categorical key variables:region, age, gender, marital, empstat
## --> Numerical key variables: wage, savings
## -->Weight variable: wgt
## Information on categorical Key-Variables:
## Reported is the number, mean size and size of the smallest category for
recoded variables.
## In parenthesis, the same statistics are shown for the unmodified data.
## Note: NA (missings) are counted as seperate categories!
## Key Variable Number of categories Mean size
##region2 (2) 2290.000 (2290.000)
## age5 (5) 916.000 (916.000)
## gender3 (3) 1526.667 (1526.667)
## marital8 (8) 572.500 (572.500)
## empstat3 (3) 1526.667 (1526.667)
##
## Size of smallest
```

```
## 646 (646)
## 16 (16)
## 50 (50)
## 26 (26)
## 107 (107)
## Infos on 2/3-Anonymity:
##
## Number of observations violating
## - 2-anonymity: 157
## - 3-anonymity: 281
##
## Percentage of observations violating
## - 2-anonymity: 3.428 %
## - 3-anonymity: 6.135 %
## -----
## Numerical key variables:wage, savings
##
## Disclosure risk is currently between [0.00%; 100.00]
## Current Information Loss:
## IL1: 0.00
## Difference of Eigenvalues: 0.000%
 _____
```

Table 7.2 presents the names of the slots and their respective contents. The slot names can be listed using the function slotNames(), which is illustrated in Example 7.9. Not all slots are used in all cases. Some slots are filled only after applying certain methods, e.g., evaluating a specific risk measure. Certain slots of the objects can be accessed by accessor functions (e.g., extractManipData for extracting the anonymized data) or print functions (e.g., print()) with the appropriate arguments. The content of a slot can also be accessed directly with the '@' operator and the slot name. This is illustrated for the risk slot in Example 7.9. This functionality can be practical to save intermediate results and compare the outcomes of different methods. Also, for manual changes to the data during the SDC process, such as changing missing value codes or manual recoding, the direct accession of the data in the slots with the manipulated data (i.e., slot names starting with 'manip') is useful. Within each slot there are generally several elements. Their names can be shown with the names() function and they can be accessed with the '\$' operator. This is shown for the element with the individual risk in the risk slot.

Example 7.9: Displaying slot names and accessing slots of an S4 object

```
# List names of all slots of sdcMicro object
slotNames(sdcInitial)
## [1] "origData" "keyVars" "pramVars"
```

```
## [4] "numVars" "ghostVars" "weightVar"
## [7] "hhId" "strataVar" "sensibleVar"
## [10] "manipKeyVars" "manipPramVars" "manipNumVars"
## [13] "manipGhostVars" "manipStrataVar" "originalRisk"
## [16] "risk" "utility" "pram"
## [19] "localSuppression" "options" "additionalResults"
## [22] "set" "prev" "deletedVars"

#Accessing the risk slot
sdcInitial@risk

#List names within the risk slot
names(sdcInitial@risk)

## [1] "global"" \ ``"individual"" \ ``"numeric"
#Two ways to access the individual risk within the risk slot
sdcInitial@risk$individual
```

Table 7.2: Slot names and slot description of sdcMicro object

get.sdcMicroObj(sdcInitial, "risk")\$individual

Slotname	Content
origData	original data as specified in the dat argument of the createSdcObj() function
keyVars	indices of columns in origData with specified categorical key variables
pramVars	indices of columns in origData with specified PRAM variables
numVars	indices of columns in origData with specified numerical key variables
ghostVars	indices of columns in origData with specified ghostVars
weightVar	indices of columns in origData with specified weight variable
hhId	indices of columns in origData with specified cluster variable
strataVar	indices of columns in origData with specified strata variable
sensibleVar	indices of columns in origData with specified sensitive variables for lDiversity
manipKeyVars	manipulated categorical key variables after applying SDC methods (cf. keyVars slot)
manipPramVars	manipulated PRAM variables after applying PRAM (cf. pramVars slot)
manipNumVars	manipulated numerical key variables after applying SDC methods (cf. numVars slot)
manipGhostVars	manipulated ghost variables (cf. ghostVars slot)
manipStrataVar	manipulated strata variables (cf. strataVar slot)
originalRisk	global and individual risk measures before anonymization
risk	global and individual risk measures after applied SDC methods
utility	utility measures (il1 and eigen)
pram	details on PRAM after applying PRAM
localSuppression	number of suppression per variable after local suppression
options	options specified
additionalResults	additional results
set	list of slots currently in use (for internal use)
prev	information to undo one step with the undo() function
deletedVars	variables deleted (direct identifiers)

There are two options to save the results after applying SDC methods:

- Overwriting the existing sdcMicro object, or
- Creating a new *sdcMicro* object. The original object will not be altered and can be used for comparing results. This is especially useful for comparing several methods and selecting the best option.

In both cases, the result of any function has to be re-assigned to an object with the '<-' operator. Both methods are illustrated in Example 7.10.

Example 7.10: Saving results of applying SDC methods

# Applying local suppression and reassigning the results to the same sdcMicro object sdcInitial <- localSuppression(sdcInitial)

# Applying local suppression and assigning the results to a new sdcMicro object sdc1 <- localSuppression(sdcInitial)

If the results are reassigned to the same *sdcMicro* object, it is possible to undo the last step in the SDC process. This is useful when changing parameters. The results of the last step, however, are lost after undoing that step. **NOTE: The undolast() function can be used to go only one step back, not several.** The result must also be reassigned to the same object. This is illustrated in Example 7.11.

Example 7.11: Undo last step in SDC process

# Undo last step in SDC process sdcInitial <- undolast(sdcInitial)

#### 7.6 Household structure

If the data has a hierarchical structure and some variables are measured on the higher hierarchical level and others on the lower level, the SDC process should be adapted accordingly (see also Sections 4.9 and 5.5). A common example in social survey data is datasets with a household structure. Variables that are measured on the household level are, for example, household income, type of house and region. Variables measured on the individual level are, for example, age, education level and marital status. Some variables are measured on the individual level, but are nonetheless the same for all household members in almost all households. These variables should be treated as measured on the household level from the SDC perspective. An example is the variable religion for some countries.

The SDC process should be divided into two stages in cases where the data have a household structure. First, the variables on the higher (household) level should be anonymized; subsequently, the treated higher-level variables should be merged with the individual variables and anonymized jointly. In this section, we explain how to extract household variables from a file and merge them with the individual levels variables after treatment in *R*. We illustrate this process with an example of household and individual-level variables.

These steps are illustrated in Example 7.12. We require both an individual ID and a household ID in the dataset; if they are lacking, they must be generated. The individual ID has to be unique for every individual in the dataset and the household ID has to be unique across households. The first step is to extract the household variables and save them in a new dataframe. We specify the variables that are measured at the household level in the string vector "HHVars" and subtract only these variables from the dataset. This dataframe will have for each household the same number of entries as it has household members (e.g., if a household has four members, this household will appear four times in the file). We next apply the function unique() to select only one record per household. This argument of the unique function is the household ID, which is the same for all household members, but unique across households.

Example 7.12: Create a household level file with unique records (remove duplicates)

# Create subset of file with only variables measured at household level

```
HHVars <- c('region', 'hhincome')
fileHH <- file[,HHVars]
# Remove duplicated rows based on the household ID / only every household once in fileHH
fileHH <- unique(fileHH, by = c('HID'))
# Dimensions of fileHH (number of households)
dim(fileHH)
```

After anonymizing the household variables based on the dataframe "fileHH", we recombine the anonymized household variables with the original variables, which are measured on the individual level. We can extract the individuallevel variables from the original dataset using "INDVars" – a string vector with the individual-level variable names. For extracting the anonymized data from the sdcMicro object, we can use the extractManipData() function from the sdcMicro package. Next, we merge the data using the merge function. The 'by' argument in the merge function specifies the variable used for merging - in this case the household ID, which has the same variable name in both datasets. All other variables should have different names in both datasets. These steps are illustrated in Example 7.13.

Example 7.13 Merging anonymized household-level variables with individual-level variables

```
# Extract manipulated household level variables from the SDC object
HHmanip <- extractManipData(sdcHH)
# Create subset of file with only variables measured at individual level
fileIND <- file[,INDVars]
# Merge the file by using the household ID
fileCombined \leftarrow merge(HHmanip, fileIND, by = \mathbf{c}('HID'))
```

The file *fileCombined* is used for the SDC process with the entire dataset. How to deal with data with household structure is illustrated in the case studies in Chapter 9.

As discussed in Section 5.5, the size of a household can also be a quasi-identifier, even if the household size is not included in the dataset as variable. For the purpose of evaluating the disclosure risk, it might be necessary to create such a variable by a headcount of the members of each household. Example 7.14 shows how to generate a variable household size with values for each individual based on the household ID (HID). Two cases are shown: 1) the file sorted by household ID and 2) the file not sorted.

```
Example 7.14 Generating the variable household size
```

```
# Sorted by HID
file \$ hhsize <- \ \textbf{rep(unname(table(file \$ HID)), unname(table(file \$ HID)))}
# Unsorted
filehhsize < - rep(diff(c(1, 1 + which(diff(file HID) != 0), length(b)+1)), diff(c(1, 1 + which(diff(file HID) != 0), length(b)+1))
length(file$HID) + 1)))
```

**NOTE:** In some cases, the order of the individuals within the households can provide information that could lead to re-identification. An example is information on the relation to the household head. In many countries, the first individual in the household is the household head, the second the partner of the household head and the next few are children. Therefore, the line number within the household could correlate well with a variable that contains information on the relation to the household head. One way to avoid this unintended release of information is to change the order of the individuals within each household at random. Example 7.15 illustrates a way to do this in *R*.

Example 7.15 Changing the order of individuals within households

```
# List of household sizes by household

hhsize <- diff(c(1, 1 + which(diff(file$HID) != 0), length(file$HID) + 1))

# Line numbers randomly assigned within each household

set.seed(123)

dataAnon$INDID <- unlist(lapply(hhsize, function(n){sample(1:n, n, replace = FALSE, prob = rep(1/n, n))}))

# Order the file by HID and randomized INDID (line number)

dataAnon <- dataAnon[order(dataAnon$HID, dataAnon$INDID),]
```

## 7.7 Randomizing order and numbering of individuals or households

Often the order and numbering of individuals, households, and also geographical units contains information that could be used by an intruder to re-identify records. For example, households with IDs that are close to one another in the dataset are likely to be geographically close as well. This is often the case in a census, but also in a household survey households close to one another in the dataset likely share the same low level geographical unit if the dataset is sorted in that way. Another example is a dataset that is alphabetically sorted by name. Here, removing the direct identifier name before release is not sufficient to guarantee that the name information cannot be used (e.g. first record has a name which likely starts with 'a'). Therefore, it is often recommended to randomize the order of records in a dataset before release. Randomization can also be done within subsets of the dataset, e.g., within regions. If suppressions were made in the geographical variable used for creating the subsets, randomization within the geographical subsets implies that the geographical variable is the same for all records in the subset and the suppressed value can be easily derived (for instance, in cases where the geographical unit is included in the randomized ID). Therefore, if the variable used for the subsets has suppressed values, randomization should be done at the dataset level and not at the subset level.

Table 7.3 illustrates the need and process of randomizing the order of records in a dataset. The first three columns in Table 7.3 show the original dataset. Some suppressions were made in the variable "district", as shown in columns 4 to 6 ('NA' values). This dataset also already shows the randomized household IDs. The order of the records in the columns 1-3 and columns 4-6 is unchanged. By the order of the records, it is easy to guess the values of the two suppressed values. Both the record before and after have the same value for district as the suppressed values, respectively 3 and 5. After reordering the dataset based on the randomized household IDs, we see that it becomes impossible to reconstruct the suppressed values based on the values of the neighboring records. Note that in this example the randomization was carried out within the regions and the region number is included in the household ID (first digit).

Table 7.3 Illustration of randomizing order of records in a dataset

Ori ginal		ith randomiz	e house-	Dataset for release ordered by the new randomized					
datas et	hold IDs			household ID					
Hou	Reg	Dis	Ran	Reg	Dis	Ran	Reg	Dis	
sehol d	ion	trict	domiz	ion	trict	domiz	ion	trict	
ID**			ed house			ed house			
			hold			hold			
			ID**			ID**			
101	1	1	108	1	1	101	1	4	
102	1	1	106	1	1	102	1	3	
103	1	2	104	1	2	103	1	5	
104	1	2	112	1	2	104	1	2	
105	1	2	105	1	2	105	1	2	
106	1	3	102	1	3	106	1	1	
107	1	3	109	1	NA	107	1	3	
108	1	3	107	1	3	108	1	1	
109	1	4	101	1	4	109	1	NA	
110	1	5	111	1	5	110	1	NA	
111	1	5	110	1	NA	111	1	5	
112	1	5	103	1	5	112	1	2	
201	2	6	203	2	6	201	2	6	
202	2	6	204	2	6	202	2	6	
203	2	6	201	2	6	203	2	6	
204	2	6	202	2	6	204	2	6	

The randomization is easiest if done before or after the anonymization process with *sdcMicro* and directly on the dataset (data.frame in *R*). To randomize the order, we need an ID, such as an individual ID, household ID or geographical ID. If the dataset does not contain such ID, this should be created first. Example 7.16 shows how to randomize households. "HID" is the household ID and "regionid" is the region ID. First the variable "HID" is replaced by a randomized variable "HIDrandom". Then the file is sorted by region and the randomized ID and the actual order of the records in the dataset is changed. To make the randomization reproducible, it is advisable to set a seed for the random number generator.

Example 7.16: Randomize order of households

n <- length(file\$HID) # number of households

set.seed(123) # set seed

file\$HIDrandom <- sample(1:n, n, replace = FALSE, prob = rep(1/n, n)) # generate random HID

file <- file1[order(file\$regionid, file\$HIDrandom),] # sort file by regionid and random HID

file\$HIDrandom <- 1:n # renumber the households in randomized order to 1-n

# 7.8 Computation time

Some SDC methods can take a very long time to evaluate in terms of computation. For instance, local suppression with the function localSuppression() of the *sdcMicro* package in *R* can take days to execute on large datasets of more than 30,000 individuals that have many categorical quasi-identifiers. Our experiments reveal that computation time is a function of the following factors: the applied SDC method; data size, i.e., number of observations, number of

variables and the number of categories or factor levels of each categorical variable; data complexity (e.g., the number of different combinations of values of key variables in the data); as well as the computer/server specifications.

Table 7.4 gives some indication of computation times for different methods on datasets of different size and complexity based on findings from our experiments. The selected quasi-identifiers and categories for those variables in Table 7.3 are the same in both datasets being compared. Because it is impossible to predict the exact computation time, this table should be used to illustrate how long computations may take. These methods have been executed on a powerful server. Given long computation times for some methods, it is recommended, where possible, to first test the SDC methods on a subset or sample of the microdata, and then choose the appropriate SDC methods. *R* provides functions to select subsets from a dataset. After setting up the code, it can then be run on the entire dataset on a powerful computer or server.

Table 7.4: Computation times of different methods on datasets of different sizes

Dataset with 5,000 obser	vations	Dataset with 45,000 obervations			
Methods	Computation time (hours)	Methods	Computation time (hours)		
Top coding age, local suppression (k=3)	11	Top coding age, local suppression (k=3)	268		
Recoding age, local suppression (k=3)	8	Recoding age, local suppression (k=3)	143		
Recoding age, local suppression (k=5)	10	Recoding age, local suppression (k=5)	156		

The number of categories and the product of the number of categories of all categorical quasi-identifiers give an idea of the number of potential combinations (keys). This is only an indication of the actual number of combinations, which influences the computation time to compute, for example, the frequencies of each key in the dataset. If there are many categories but not so many combinations (e.g., when the variables correlate), the computation time will be shorter.

Table 7.5 shows the number of categories for seven datasets with the same variables but of different complexities that were all processed using the same script on 16 processors, in order of execution time. The table also shows an approximation of the number of unique combinations of quasi-identifiers, as indicated by the percentage of observations violating k-anonymity in each dataset pre-anonymization in relation to processing time. The results in the table clearly indicate that both the number of observations (i.e., sample size) and the complexity of the data play a role in the execution time. Also, using the same script (and hence anonymization methods), the execution time can vary greatly; the longest running time is about 10 times longer than the shortest. Computer specifications also influence the computation time. This includes the processor, RAM and storage media.

Table 7.5: Number of categories (complexity), record uniqueness and computation times

	Nu				Execution	n time in			
Sa	mber of o	cate gori es	per quas	Pe	hours				
mple Size	moor or c	ato gon oc	rcen tage of obse rvat ions viol atin g k anon ymit y befo re	nours					
					anon ymiz				
					atio n				
	Wate r	Toil et	Occu	Reli	Ethn icit	Regi on	k3	k5	
			pati on	gion	у				
20,0 14	10	4	70	5	7	6	74	88	53.7 2
66,2 85	15	6	39	4	0	24	40	49	67.1 9
60,7 47	13	6	70	8	9	4	35	45	74.4 7
26,6 01	19	6	84	10	10	10	77	87	108. 84
38,0 89	17	6	30	5	56	9	70	81	198. 90
35,8 20	19	7	67	6	NA	6	81	90	267. 60
51,9 76	12	6	32	8	50	12	77	87	503. 58

The large-scale experiment executed for this guide utilized 75 microdata files from 52 countries, using surveys on topics including health, labor, income and expenditure. By applying anonymization methods available in the *sdcMicro* package, at least 20 different anonymization scenarios were tested on each dataset. Most of the processing was done using a powerful server and up to 16 - 20 processors (cores) at a time. Other processing platforms included a laptop and desktop computers, each using four processors. Computation times were significantly shorter for datasets processed on the server, compared to those processed on the laptop and desktop.

The use of parallelization can improve performance even on a single computer with one processor with multiple cores. Since *R* does not use multiple cores unless instructed to do so, our anonymization programs allowed for parallelization such that jobs/scenarios in each dataset could be processed simultaneously through efficient allocation of tasks to different processors. Without parallelization, depending on the server/computer, only one core is used when running the jobs sequentially. Running the anonymization program without parallelization leads to significantly longer execution time. Note however, that the parallelization itself also causes overhead. Therefore, a summation of the times it takes to run each task in parallel does not necessarily amount to the time it may take to run them sequentially. The fact that the RAM is shared might, however, slightly reduce the gains of parallelization. If you want to compare the results of different methods on large datasets that require long computation times, using parallel computing can be a solution.[#foot70]\_

Appendix D zooms in on seven selected datasets from a health survey that were processed using the same parallelization program and anonymization methods. Note that the computation times in the appendix are only meant to create awareness for expected computation time, and may vary based on the type of computer used. In our case, although all datasets were anonymized using the parallelization program, computation times were significantly shorter for datasets processed on the server, compared to those processed on the laptop and desktop. Among those datasets processed on the server using the same number of processors (datasets 1, 2 and 6), some variation also exists in the computation times. NOTE: Computation time in the table in Appendix D includes recalculating the risk after applying the anonymization methods, which is automatically done in sdcMicro when using standard methods/functions. Us-

<sup>&</sup>lt;sup>9</sup> Here a scenario refers to a combination of SDC methods and their parameters.

<sup>&</sup>lt;sup>10</sup> The server has 512 GB RAM and four processors each with 16 cores, translating to 64 cores total.

ing the function group Vars(), for instance, is not computationally intensive but can still take a long time if the dataset is large and risk measures have to be recalculated.

### 7.9 Common errors

In this section, we present a few common errors and their causes, which might be encountered when using the *sdcMicro* package in *R* for anonymization of microdata:

- The class of a certain variable is not accepted by the function, e.g., a categorical variable of class numeric should be first recoded to the required class (e.g., factor or data.frame). Section 7.4 shows how to do this.
- After manually making changes to variables the risk did not change, since it is not updated automatically and has to be manually recomputed by using the function calcRisks().

**CHAPTER** 

**EIGHT** 

#### THE SDC PROCESS

This section presents the SDC process in a step-by-step representation and can be used as guidance for the actual SDC process. It should be noted, however, that jumping between steps and returning to previous steps is often required during the actual SDC process, as it is not necessarily a linear step-by-step process. This guidance brings together the different parts of the SDC process as discussed in the previous chapters and links to these chapters. The case studies in the next chapter follow these steps. This presentation is adapted from Hundepool et al. (2012). Figure 8.1 at the end of this chapter presents the entire process in a schematic way.

### 8.1 Step 1: Need for confidentiality protection

Before starting the SDC process for a microdata set, the need for confidentiality protection has to be determined. This is closely linked to the interpretation of laws and regulations on this topic from the country in which the data originates and thus country-specific. A first step is to determine the statistical units in the dataset: if these are individuals, households or legal entities, such as companies, a need for disclosure control is likely. There are also examples of microdata for which there is no need for disclosure control. Examples could be data with climate and weather observations or data with houses as statistical units. Even if the primary statistical units are not natural or legal persons, however, the data can still contain confidential information on natural or legal persons. For example, a dataset with houses as primary statistical units can also contain information on the persons living in these houses and their income or a dataset on hospitalizations can include information about the hospitalized patients. In these cases, there is likely still a need for confidentiality protection. One option to solve this is to remove the information on the natural and legal persons in the datasets for release.

One dataset can also contain more than one type of statistical unit. The standard example here is a dataset containing both information on individuals and households. Another example is data with employees in enterprises. All types of statistical units present in the dataset have to be considered for the need of SDC. This is especially important in case the data has a hierarchical structure, such as individuals in households or employees in enterprises.

In addition, one has to evaluate whether the variables contained in the dataset are confidential or sensitive. Which variables are sensitive or confidential depends again on the applicable legislation and can differ substantially from country to country. In case the dataset includes sensitive or confidential variables, SDC is likely required. The set of sensitive variables and confidential variables together with the statistical units in the dataset determine the need for statistical disclosure control.

## 8.2 Step 2: Data preparation and exploring data characteristics

After assessing the need for statistical disclosure control, we should prepare the data and, if there are multiple, combine and consider all related data files. Then we explore the characteristics and structure in the data, which are important for the users of the data. Compiling an inventory of these characteristics is important for assessing the utility of the data after anonymization and producing an anonymized dataset, which is useful for end users.

The first step in data preparation is classifying the variables as sensitive or non-sensitive, and removing direct identifiers such as full names, passport numbers, addresses, phone numbers and GPS coordinates. In case of survey data, an inspection of the survey questionnaire is useful to classify the variables. Furthermore, it is necessary to select the variables that contain relevant information for end users and should be included in the dataset for release. At this point, it can also be useful to remove variables other than direct identifiers from the microdata set to be released. An example can be a variable with many missing values, e.g., a variable recorded only for a select group of individuals eligible for a particular survey module, and missing values for the rest. Such variables can cause a high level of disclosure risk while adding little information for end users. Examples are variables relating to education (current grade), where a missing value indicates that the individual is not currently in school, or variables relating to childbirth, where a missing value indicates that the individual has not delivered a child in the reference period. Missing values in themselves can be disclosive, especially if they indicate that the variable is not applicable. Often variables with the majority of values missing are deleted at this stage already. Other variables that might be deleted at this stage are those too sensitive to be anonymized and released or those not important to data users and that could increase the risk of disclosure.

Relationships may exist among variables in a dataset for a variety of reasons. For instance, variables can be mutually exclusive in cases where several binary variables are used for each category. An individual not in the labor force will have a missing value for the sector in which this person is employed (or more precisely not applicable). Relationships may also exist if some variables are ratios, sums or other mathematical functions of other variables. Examples are the variable household size (as a count of individuals per household), or aggregate expenditure (as a sum of all expenditure components). A certain value in one variable may also reduce the number of possible or valid values for another variable; for example, the age of an individual attending primary education or the gender of an individual having delivered a child. These relationships are important for two reasons: 1) they can be used by intruders to reconstruct anonymized values. For example, if age is suppressed but another variable indicates that they are in school, then it is still possible to infer a likely age range for that individual. Another example is if an individual is shown to be active in Sector B of the economy. Even if the labor status of this individual is suppressed, it can be inferred that this person is employed. 2) the relationships in the original data should be maintained in the anonymized dataset and inconsistencies should be avoided (e.g., SDC methods should not create 58-year-old school boys, or married 3-year-olds), or the dataset will be invalid for analysis. Another example is the case of expenditures per category, where it is important that the sum of the categories adds up to the total. One way to guarantee this is to anonymize the totals and then recalculate the sub-categories according to the original shares of the anonymized totals.

It is also useful at this stage to consolidate variables that provide the same information where possible, so as reduce the number of variables, reduce the likelihood of inconsistencies and minimize the variables an intruder can use to reconstruct the data. This is especially true if the microdata stems from an elaborate questionnaire and each variable represents one (sub-) question leading to a dataset with hundreds of variables. As an example, we take a survey with several labor force variables indicating whether a person is in the labor force, employed or unemployed, and if employed, in what sector. The data in Table 8.1 illustrates this example. It is possible that each type of sector has its own binary variable. In that case, this set of variables can be summarized in two variables: one variable indicating whether a person is in labor force and another indicating the employment status, as well as the respective sector if a person is employed. These two variables contain all information contained in the previous five variables and make the anonymization process easier. If data users are used to a certain release format where including all five variables has been the norm, then it is possible to transform the variables back after the anonymization process rather than complicating the anonymization process by trying to treat more variables than is necessary. This approach also guarantees that the relationships between the variables are preserved (e.g., no individuals will be employed in several sectors).

Table 8.1: Illustration of merging variables without information loss for SDC process

Before		After				
In labor force	Employed	Sector A	Sector B	Sector C	In labor force	Emplo yed
Yes	Yes	Missing	Yes	Missing	Yes	В
No	No	Missing	Missing	Missing	No	No
Yes	Yes	Yes	Missing	Missing	Yes	A
Yes	Yes	Missing	Yes	Missing	Yes	В
Yes	Yes	Missing	Missing	Yes	Yes	C
Yes	No	Missing	Missing	Missing	Yes	No

Besides relationships between variables, we also gather information about the survey methodology, such as strata, sampling methods, survey design and sample weights. This information is important in later stages, when estimating the disclosure risk and choosing the SDC methods. It is important to distinguish between a full census and a sample. For a full census, it is common practice to publish only a sample, as the risk of disclosure for a full sample is too high, given that we know that everyone in the country or institution is in the data (see also Section 5.6). Strata and sample weights can disclose information about the area or group to which an individual belongs (e.g., the weights can be linked with the geographical area or specific group in case of stratified sampling); this should be taken into account in the SDC process and checked before release of the dataset.

# 8.3 Step 3: Type of release

The type of release is an important factor for determining the required level of anonymization as well as the requirements end users have for the data (e.g., researchers need more detail than students for whom a teaching file might be sufficient) and should be clarified before the start of the anonymization process. Data release or dissemination by statistical agencies and data producers is often guided by the applicable law and dissemination strategies of the statistical agency, which specify the type of data that should be disseminated as well as the fashion.

Generally, there exist three types of data release methods for different target groups (Chapter 3 provides more information on different release types):

- PUF: The data is directly available to anyone interested, e.g., on the website of the statistical agency
- SUF: The data is available to accredited researchers, who have to file their research proposals beforehand and have to sign a contract; this is also known as licensed file or microdata under contract
- Available in a controlled research data center: only on-site access to data on special computers; this is also known as data enclave

There are other data access possibilities besides these, such as teaching files or files for other specific purposes. Obviously, the required level of protection depends on the type of release: a PUF file must be protected to a greater extent than a SUF file, which in turn has to be protected more than a file which is available only in an on-site facility, since the options the intruder can use the data are limited in the latter case.

Besides the applicable legislation, the choice of the type of release depends on the type of the data and the content. **NOTE:** Not every microdata set is suitable for release in any release type, even after SDC. Some data cannot be protected sufficiently – it might always contain information that is too sensitive to be published as SUF or PUF. In such cases, the data can be released in on-site facilities, or the number of variables can be reduced by removing problematic variables.

Generally, the release of two or more anonymized datasets, e.g., tailored for different end users from the same original, is problematic because it can lead to disclosure if the two were later obtained and merged by the same user. The information contained in one dataset that is not contained in the other can lead to unintended disclosure. An exception is the simultaneous release and anonymization of a microdata set as PUF and SUF files. In this case, the PUF file is constructed from the SUF file by further anonymization. In that way, all information in the PUF file is also contained in the SUF file and the PUF file does not provide any additional information for users that have access to the SUF file.

**NOTE:** The anonymization process is an iterative process where steps can be revisited, whereas the publication of an anonymized dataset is a one-shot process. Once the anonymized data is published, it is not possible to revoke and publish another dataset of the same microdata file. This would in fact mean publishing more than one anonymized file from the same microdata set, since some users might have saved the previous file.

# 8.4 Step 4: Intruder scenarios and choice of key variables

After determining the release type of the data, the possibilities of how an individual in the microdata could (realistically) be identified by an intruder under that release type should be examined. For PUF and SUF release the focus is on the use of external datasets from various source. These possibilities are described in disclosure or intruder scenarios, which specify what data an intruder could possibly have access to and how this auxiliary data can be used for identity disclosure. This leads to the specification of quasi-identifiers, which are a set of variables that are available both in the dataset to be released and in auxiliary datasets and need protection. **NOTE: If the number of quasi-identifiers is high, it is recommended to reduce the set of quasi-identifiers by removing some variables from the dataset for release.** This is especially true for PUF releases. Disclosure scenarios can also help define the required level of anonymization.

Drafting disclosure scenarios requires the support of subject matter specialists, assuming the subject specialist is not the same as the person doing the anonymization. Auxiliary datasets may contain information on the identity of the individuals and allow identity disclosure. Examples of such auxiliary data files are population registers and electoral rolls, as well as data collected by specialized firms. **NOTE:** External datasets can come from many sources (e.g., other institutions, private companies) and it is sometimes difficult to make a full list of external data sources. In addition, not all external data sources are in the public domain. Nevertheless, proprietary data can be used by the owner to re-identify individuals and should be taken into account in the SDC process, even if the exact content is not known. Also, the variables or datasets may not coincide perfectly (e.g., different definitions, more or less detailed variables, different survey period). Nevertheless, they should be considered in the SDC process.

Disclosure scenarios include both identity and inferential disclosure. The disclosure depends on the type of release, i.e., different data users have different data available and may use the data in a different way for re-identification. For example, a user in a research data center cannot match with large external datasets as (s)he is not permitted to take these into the data center. A user of a SUF is bound by an agreement specifying the use of the data and consequences if the agreement is breached (see Chapter 3). Furthermore, it should be evaluated whether, in case of a sample, possible intruders have knowledge as to which individuals are in the sample. This can be the case if it is known which schools were visited by the survey team, for example. A few examples of disclosure scenarios are (see Section 4.3 for more information):

- Matching: The intruder uses auxiliary data, e.g., data on region, marital status and age from a population register, and matches them to released microdata. Individuals from the two datasets that match and are unique are successfully identified. This principle is used as an assumption in several disclosure risk measures, such as *k*-anonymity, individual and global risk, as described in Chapter 4. This scenario can apply to both PUFs and SUFs.
- Spontaneous recognition: This scenario should be considered for SUF files, but is especially important for data available in research data centers where outliers are present in the data and data is often not strongly anonymized. The researcher might (unintentionally) recognize some individuals he knows (e.g., his colleagues, neighbors, family members, public figures, famous persons or large companies), while inspecting the data. This is especially true for rare combinations of values, such as outliers or unlikely combinations.

## 8.5 Step 5: Data key uses and selection of utility measures

In this step, we analyze the main uses of the data by the end users of the released microdata file. The data should be useful for the type of statistical analysis for which the data was collected and for which it is mostly used. The uses and requirements of data users will be different for different release types. Contacting data users directly or searching

for scientific studies and papers that use similar data can be useful when collecting this information and making this assessment. Alternatively, this information can be collected from research proposals by researchers when applying for microdata access (SUF) or user groups can be set up. Furthermore, it is important to understand what level of precision the data users need and what types of categories are used. For instance, in the case of global recoding of age in years, one could recode age in groups of 10 years, e.g., 0 - 9, 10 - 19, 20 - 29, ... Many indicators relating to the labor market use categories that span the range 15 - 65, however. Therefore, constructing categories that coincide with the categories used for the indicators keeps the data much more useful while at the same time reducing the risk of disclosure in a similar way. This knowledge is important for the selection of useful utility measures, which in turn are used for selecting appropriate SDC methods.

The uses of the data depend on the release type, too. Researchers using SUF files require a higher level of detail in the data than PUF users. NOTE: Anonymization will always lead to information loss and a PUF file will have reduced utility. If certain users require a high level of detail, release types other than PUF should be considered, such as SUF or release through a research data center. In the case of SUFs, it is easier to find the main uses of the data since access is documented. One way to obtain information on the use of PUF files is to ask for a short description of intended use of the data before supplying the data. This is, however, useful only if microdata has been released previously.

Statistics computed from the anonymized and released microdata file should produce analytical results that agree or almost agree with previously published statistics from the original data. If, for instance, a previously published primary school enrollment rate was computed from these data and published, the released anonymized dataset should produce a very similar result to the officially published result. At the very least, the result should fall within the confidence region of the published result. It might be the case that not all published statistics can be generated from the published data. If this is the case, a choice should be made on which indicators and statistics to focus, and inform the users as to which ones have been selected and why.

As discussed in Chapter 6, it is necessary to compute general utility measures that compare the raw and anonymized data, taking into consideration the end user's need for their analysis. In some cases the utility measures can give contradicting results, for example, a certain SDC method might lead to lower information loss for labor force figures but greater information loss for ratios relating to education. In such cases, the data uses might need to be ranked in order of importance and it should be clearly documented for the user that the prioritization of certain metrics over others means that certain metrics are no longer valid. This may be necessary, as it is not possible to release multiple files for different users. This problem occurs especially in multi-purpose studies. For more details on utility measures, refer to Chapter 6.

### Note on Steps 6 to 10

The following Steps 6 through 10 should be repeated if the data has quasi-identifiers that are on different hierarchical levels, e.g., individual and household. In that case, variables on the higher hierarchical level should be anonymized first, and then merged with the lower-level untreated variables. Subsequently, the merged dataset should be anonymized. This approach guarantees consistency in the treated data. If we neglect this procedure, the values of variables measured on the higher hierarchical level could be treated differently for observations of the same unit. For instance, the variable "region" is the same for all household members. If the value 'rural' would be suppressed for two members but not for the remaining three, this would lead to unintended disclosure; with the household ID the variable region would be easy to reconstruct for the two suppressed values. Sections 4.9 and 7.6 provide more details on how to deal with data with household structure in *R* and *sdcMicro*.

# 8.6 Step 6: Assessing disclosure risk

The next step is to evaluate the disclosure risk of the raw data. Here it is important to distinguish between sample data and census data. In the case of census data, it is possible to directly calculate the risk measures when assuming that the dataset covers the entire population. If working with a sample, or a sample of the census (which is the more common case when releasing sample data), we can use the models discussed in Chapter 4 to estimate the risk in the population. The main inputs for the risk measurement are the set of quasi-identifiers determined from the disclosure scenarios in Step 4 and the thresholds for risk calculations (e.g., the level of k-anonymity or the threshold for which an individual

is considered at risk). If the data has a hierarchical structure (e.g., a household structure), the risk should be measured taking into account this structure as described Section 4.9.

The different risk measures described in Chapter 4 each have advantages and disadvantages. Generally, k-anonymity, individual risk and global risk are used to produce an idea of the disclosure risk. These values can initially be very high but can often very easily be reduced after some simple but appropriate recoding (see Step 8). The thresholds shall be determined according to the release type. Always remember, though, that when using a sample, the risk measures based on the models presented in the literature offer a worst-case risk scenario and might therefore be an exaggeration of the real risks for some cases (see Section 4.5).

# 8.7 Step 7: Assessing utility measures

To quantify the information loss due to the anonymization, we first compute the utility measures selected in Step 6 using the raw data. This creates a base for comparison of results obtained when using the anonymized data - i.e., in Step 10. **NOTE:** If the raw data is a sample, the utility measures are an estimate with a variance and therefore it is useful to construct confidence intervals in addition to the point estimates for the utility measures.

# 8.8 Step 8: Choice and application of SDC methods

The choice of SDC methods depends on the need for data protection (as measured by the disclosure risk), the structure of the data and the type of variables. The influence of different methods on the characteristics of the data important for the users or the data utility should also be taken into account when selecting the SDC methods. In practice, the choice of SDC methods is partially a trial-and-error process: after applying a chosen method, disclosure risk and data utility are measured and compared to other choices of methods and parameters. The choice of methods is bound by legislation on the one hand, and a trade-off between utility and risk on the other.

The classification of methods as presented in Table 5.1 gives a good overview for choosing the appropriate methods. Methods should be chosen according to the type of variable – continuous or categorical – the requirements by the users and the type of release. The anonymization of datasets with both continuous and categorical variables is discussed in Section 4.2.

In general for anonymization of categorical variables, it is useful to restrict the number of suppressions by first applying global recoding and/or removing variables from the microdata set. When the required number of suppressions to achieve the required level of risk is sufficiently low, the few individuals at risk can be treated by suppression. These are generally outliers. It should be noted that possibly not all variables can be released and some must be removed from the dataset (see Step 2). Recoding and minimal use of suppression ensures that already published figures from the raw data can be reproduced sufficiently well from the anonymized data. If suppression is applied without sufficient recoding, the number of suppressions can be very high and the structure of the data can change significantly. This is because suppression mainly affects combinations that are rare in the data.

If the results of recoding and suppression do not achieve the required result, especially in cases where the number of select quasi-identifiers is high, an alternative is using perturbative methods. These can be used without prior recoding of variables. These methods, however, preserve data structure only partially. The preferred method depends on the requirements of the users. We refer to Chapter 5 and especially Section 5.3 for a discussion of perturbative methods implemented in *sdcMicro*.

Finally, the choice of SDC methods depends on the data used since the same methods produce different results on different datasets. Therefore, the comparison of results with respect to risk and utility (Steps 9 and 10) is key to the choice made. Most methods are implemented in the *sdcMicro* package. Nevertheless, it is sometimes useful to use custom-made solutions. A few examples are presented in Chapter 5.

## 8.9 Step 9: Re-measure risk

In this step, we re-evaluate the disclosure risk with the risk measures chosen under Step 6 after applying SDC methods. Besides these risk measures, it is also important to look at individuals with high risk and/or special characteristics, combinations of values or outliers in the data. If the risk is not at an acceptable level, Steps 6 to 10 should be repeated with different methods and/or parameters. **NOTE: Risk measures based on frequency counts (k-anonymity, individual risk, global risk and household risk) cannot be used after applying perturbative methods since their risk estimates are not valid.** These methods are based on introducing uncertainty into the dataset and not on increasing the frequencies of keys in the data and will hence overestimate the risk.

## 8.10 Step 10: Re-measure utility

In this step, we re-measure the utility measures from Step 7 and compare these with the results from the raw data. Also, it is useful here to construct confidence intervals around the point estimates and compare these confidence intervals. The importance of the absolute value of a deviation can only be interpreted knowing the variance of the estimate. Besides these specific utility measures, general utility measures, as discussed in Section 6.1, should be evaluated. This is especially important if perturbative methods have been applied. If the data does not meet the user requirements and deviations are too large, repeat Steps 6 to 10 with different methods and/or different parameters. **NOTE: Anonymization will always lead to at least some information loss.** 

## 8.11 Step 11: Audit and Reporting

After anonymization, it is important to check whether all relationships in the data as identified in Step 2, such as variables that are sums of other variables or ratios, are preserved. Also, any unusual values caused by the anonymization should be detected. Examples of such anomalies are negative income or a pupil in the twentieth grade of school. This can happen after applying perturbative SDC methods. Furthermore, it is necessary to check whether previously published indicators from the raw data are reproducible from the data to be released. If this is not the case, data users might question the credibility of the anonymized dataset.

An important step in the SDC process is reporting, both internal and external. Internal reporting contains the exact description of anonymization methods used, parameters but also the risk measures before and after anonymization. This allows replication of the anonymized dataset and is important for supervisory authorities/bodies to ensure the anonymization process is sufficient to guarantee anonymity according to the applicable legislation.

External reporting informs the user that the data has been anonymized, provides information for valid analysis on the data and explains the limitations to the data as a result of the anonymization. A brief description of the methods used can be included. The release of anonymized microdata should be accompanied by the usual metadata of the survey (survey weight, strata, survey methodology) as well as information on the anonymization methods that allow researchers to do valid analysis (e.g., amount of noise added, transition matrix for PRAM). **NOTE: Care should be taken that this information cannot be used for re-identification (e.g., no release of random seed for PRAM).** The metadata must be updated to comply with the anonymized data. Variable descriptions or value labels might have changed as a result of the anonymization process. In addition, the information loss due to the anonymization process should be explained in detail to the users to make them aware of the limits to the validity of the data and their analyses.

# 8.12 Step 12: Data release

The last step in the SDC process is the actual release of the anonymized data. This step depends on the type of release chosen in Step 3. Changes to the variables made under Step 2, such a merging variables, can be undone to generate a dataset useful for users.

## **Readthedocs Documentation**

## Recommended Reading Material on Risk Measurement

Dupriez, Olivier, and Ernie Boyko. 2010. *Dissemination of Microdata Files; Principles, Procedures and Practices*. IHSN Working Paper No. 005, International Household Survey Network (IHSN). http://www.ihsn.org/HOME/sites/default/files/resources/IHSN-WP005.pdf

# Step 1: Determine the need for confidentiality and protection

- Determine the statistical units (units of observation) in the data set e.g. households, individuals or establishments.
- Evaluate which variables are confidential or sensitive. If none, proceed directly to Step 12.

# Step 2: Data preparation and exploration of raw data

- Classify the sensitive and non-sensitive variables that were identified in Step 1
- Remove direct identifiers e.g. names, addresses, passport numbers.
- Remove any other variables that are not for release.
- Consolidate variables providing same information e.g. age in days/months/years
- Note the key relationships between variables that need to be maintained
- Note the survey methodology, sampling methodology. Identify weights and strata.

### Step 3: Determine the type of release

 Decide on the release type and conditions of release of microdata. For example: Public use file (PUF), scientific use file (SUF) and Data enclave.

NOTE: It is possible to create a version for each release type. See second note after Step 10.

## Step 4: Analyze possible intruder scenarios then select key variables

- Conduct an inventory of available external data sets which might be used by intruders, and based upon this select a set of quasi-identifiers/key variables.
- Combined with the release type, intruder scenarios, will help define the required level of anonymization

### Step 5: Determine the key uses of data and select applicable utility measures

 Identify how end users use the data and identify baseline utility measures (i.e. benchmark).

NOTE: If hierarchical structure e.g. household, extract household level variables and save as separate file. Do steps 6-10 first with the household level then with the combined file.

Figure 8.1: Overview of the SDC process

# Step 6: Compute and assess the raw data disclosure risk

 Compute and evaluate disclosure risk of the raw data using the selected quasi identifiers.
 If acceptable proceed directly to Step 12. If not, proceed to Step 7.

#### Step 7: Assess the raw data utility measures

 Compute and evaluate utility measures for later comparison.

### Step 8: Choose and apply SDC methods

 Choose and apply SDC methods; recoding local suppression, PRAM, shuffling etc., according to the risk and utility measures

#### Step 9: Re-measure risk

 Check whether disclosure risk is now at an acceptable level. If risk is too high return to Step 8, change SDC methods or parameters and re-assess. If satisfactory proceed to Step 10.

#### Step 10: Re-measure utility

 Measure and compare the utility measures against the benchmark values from Step 7. If unsatisfactory, return to Step 8 then re-assess, otherwise if satisfactory proceed to Step 11.

NOTE: If hierarchical data, e.g. household level, merge household level data with individual level data, then repeat Steps 6-10 on merged data before moving to Step 11.

NOTE: If the file generated is a SUF, and a PUF is also required, the PUF must be generated by further anonymizing the SUF. Using the SUF, repeat the process, starting from Step 3.

### Step 11: Audit and reporting

 Check data consistency and quality and prepare internal and external SDC reports

#### Step 12: Data release

 Release anonymized dataset according to the release types determined in Step 3.

**CHAPTER** 

**NINE** 

# **APPENDICES**

# 9.1 Appendix A: Overview of Case Study Variables

	Variable	Description
1	REGION	Region
2	DIST	District
3	URBRUR	Area of residence
4	WGTHH	Individual weighting coefficient (Country-specif ic. Weighting co-efficient to derive individual-leve l i
5	WGTPOP	Population weighting coefficient (Weighting co-efficient to derive population-leve l indicators.)
6	IDH	Household unique identification
7	IDP	Individual identification
8	HHSIZE	Household members
9	GENDER	Sex
10	REL	Relationship to household head
11	MARITAL	Marital status
12	AGEYRS	Age in completed years
13	AGEMTH	Age of child in completed months
14	RELIG	Religion of household head
15	ETHNICITY	Ethnicity
16	LANGUAGE	Language
17	MORBID	Morbidity last x weeks
18	MEASLES	Child immunized against Measles
19	MEDATT	Sought medical attention
20	CHWEIGHTKG	Weight of the child (Kg)
21	CHHEIGHTCM	Height of the child (cms)
22	ATSCHOOL	Current school enrolment
23	EDUCY	Highest level of education completed
24	EDYRS	Years of education
25	EDYRSCURRAT	Years of education for currently enrolled
26	SCHTYP	Type of school attending
27	LITERACY	Literacy status
28	EMPTYP1	Type of employment, Primary job
29	UNEMP1	Unemployed
30	INDUSTRY1	1 digit industry classification, Primary job
31	EMPCAT1	Employment categories, Primary job
32	WHOURSWEEK1	Hours worked last week, Primary job
33	OWNHOUSE	Ownership of dwelling unit
34	ROOF	Main material used for roof
		Contin

Contin

Table 9.1 – continued from previous page

	Variable	Description
35	TOILET	Main toilet facility
36	ELECTCON	Connection of electricity in dwelling
37	FUELCOOK	Main cooking fuel
38	WATER	Main source of water
39	OWNAGLAND	Ownership of agricultural land
40	LANDSIZEHA	Land size owned by household (ha)
41	OWNMOTORCYCLE	Ownership of motorcycle
42	CAR	Ownership of car
43	TV	Ownership of television
44	LIVESTOCK	Number of large-sized livestock owned
45	INCRMT	Total amount of remittances received from remittance sending members
46	INCWAGE	Wage and salaries(annual)
47	INCBONSOCALL	Bonus and social allowance from wage job(annual)
48	INCFARMBSN	Gross income from household farm businesses(annu al)
49	INCNFARMBSN	Gross income from household non-farm businesses(annu al)
50	INCRENT	Rental income(annual)
51	INCFIN	Financial income from savings, loans, tax refunds, maturity payments on insurance
52	INCPENSN	Pension and other social assistance(annu al)
53	INCOTHER	Other income(annual)
54	INCTOTGROSSHH	Total gross household income(annual)
55	FARMEMP	Farm employment
56	THOUSEXP	Total expenditure on housing
57	TFOODEXP	Total food expenditure
58	TALCHEXP	Total alcohol expenditure
59	TCLTHEXP	Total expenditure on clothing and footwear
60	TFURNEXP	Total expenditure on furnishing
61	THLTHEXP	Total expenditure on health
62	TTRANSEXP	Total expenditure on transport
63	TCOMMEXP	Total expenditure on communications
64	TRECEXP	Total expenditure on recreation
65	TEDUEXP	Total expenditure on education
66	TRESTHOTEXP	Total expenditure on restaurants and hotel
67	TMISCEXP	Total miscellaneous expenditure  Total annual nominal household expenditures
68	TANHHEXP	

# 9.2 Appendix B: Example of Blanket Agreement for SUF

Agreement between [providing agency] and [receiving agency] regarding the deposit and use of microdata

1.	
2	
۷.	
3.	
4	

A. This agreement relates to the following microdatasets:

5.

2. Terms of the agreement:

As the owner of the copyright in the materials listed in section A, or as duly authorized by the owner of the copyright in the materials, the representative of [providing agency] grants the [receiving agency] permission for the datasets listed in section A to be used by [receiving agency] employees, subject to the following conditions:

- 1. Microdata (including subsets of the datasets) and copyrighted materials provided by the [providing agency] will not be redistributed or sold to other individuals, institutions or organisations without the [providing agency]'s written agreement. Non-copyrighted materials which do not contain microdata (such as survey questionnaires, manuals, codebooks, or data dictionaries) may be distributed without further authorization. The ownership of all materials provided by the [providing agency] remains with the [providing agency].
- 2. Data will be used for statistical and scientific research purposes only. They will be employed solely for reporting aggregated information, including modeling, and not for investigating specific individuals or organisations.
- 3. No attempt will be made to re-identify respondents, and there will be no use of the identity of any person or establishment discovered inadvertently. Any such discovery will be reported immediately to the [providing agency].
- 4. No attempt will be made to produce links between datasets provided by the [providing agency] or between [providing agency] data and other datasets that could identify individuals or organisations.
- 5. Any books, articles, conference papers, theses, dissertations, reports or other publications employing data obtained from the [providing agency] will cite the source, in line with the citation requirement provided with the dataset.
- 6. An electronic copy of all publications based on the requested data will be sent to the [providing agency].
- 7. The [providing agency] and the relevant funding agencies bear no responsibility the data's use or for interpretation or inferences based upon it.
- 8. An electronic copy of all publications based on the requested data will be sent to the [providing agency].
- 9. Data will be stored in a secure environment, with adequate access restrictions. The [providing agency] may at any time request information on the storage and dissemination facilities in place.
- 10. The [recipient agency] will provide an annual report on uses and users of the listed microdatasets to the [providing agency], with information on the number of researchers having accessed each dataset, and on the output of this research.
- 11. This access is granted for a period of [provide information on this period, or state that the agreement is open ended].
- 3. Communications:

Name of contact person:

The [receiving organisation] will appoint a contact person who

will act as unique focal person for this agreement. Should the focal person be replaced, the [recipient agency] will immediately communicate the name and coordinates of the new contact person to the [providing agency]. Communications for administrative and procedural purposes may be made by email, fax or letter as follows:

Communications made by [providing agency] to [recipient agency] will be directed to:

Title of contact person:
Address of the recipient agency:
Email:
Tel:
Fax:
Communications made by [recipient agency] to [depositor agency]

### **Readthedocs Documentation**

Source: Dupriez and Boyko, 2010

re:
re

# 9.3 Appendix C: Internal and External Reports for Case Studies

This appendix provides example for internal and external reports on the anonymization process for the case studies in Section 9.1. The internal report consists of two parts: the first is on the anonymization of the household-level variables and the second is on the anonymization of the individual-level variables.

Appendix D: Execution Times for Multiple Scenarios Tested using Selected Sample Data

TOTAL ODENTATIONS   40,449   100,429   20,261   34,785   23,022   26,411	Datasets	1	2	3	4	5	6
Quasi-Identifiers							
RIBMAL/MERN   2	Quasi-identifiers	,					
NECON	RURAL/URBAN	2	2			2	2
Description	REGION	8	2	10	5	4	10
REALTON TO HAD OF HOUSEHOLD  REALTON TO HAD OF HOUSEHOLD  REALTON TO HAD OF HOUSEHOLD  SEE OR COMMETTED THANS  999  988  97  98  109  99  988  97  98  109  99  98  99  98  97  98  109  99  99  98  97  98  109  99  99  98  99  98  97  98  109  99  99  99  98  98  97  98  109  99  99  99  98  98  97  98  109  99  99  99  98  98  97  98  109  99  99  99  98  80  97  98  109  99  99  98  80  97  98  109  99  99  99  98  80  97  98  109  99  99  98  80  97  98  109  99  99  98  80  109  99  98  80  97  70  70  60  60  80  80  80  80  80  80  80  8	HOUSEHOLD SIZE	25	26	18	17	20	18
READTON TO READ OF HOUSEHOLD  ARMATEL STATUS  5 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6	OWNHOUSE	2	NA	2	2	3	2
MARTIAL STATUS	SEX	2	2	2	2	2	2
NA	RELATION TO HEAD OF HOUSEHOLD	8	7	8	7	9	7
REMISSION NA NA NA NA NA 10 20 REMISSION NA NA NA NA 4 10 11 11 17 18 14 14 NA REMISSITURED OF EDUCATION - LENGUIDE 16 17 17 18 14 14 NA REMISSITURED OF EDUCATION - LENGUIDE 17 7 7 6 5 5 6 6 7 REMISSION REMISSION 10 464 10 11 10 10 10 REMISSITURED OF EDUCATION - LENGUIDE 10 464 10 11 10 10 10 REMISSITURED OF EDUCATION - LENGUIDE 10 644 10 11 10 11 10 10 REMISSION REMISSION 10 645 10 0 11 10 10 10 REMISSION REMISSION 10 645 10 0 11 10 10 10 REMISSION RE	MARITAL STATUS	5	6	6	6	6	6
NA	AGE IN COMPLETED YEARS	99	98	97	98	109	99
HIGHEST LEVEL OF FORDIC COMMETTED   7	ETHNICITY	NA	NA	NA	NA	10	20
Transfer	RELIGION	NA	NA	NA	4	10	11
IMPLOURINT   TYPE   7	HIGHEST LEVEL OF EDUC COMPLETED	7	7	9	6	8	8
Name	YEARS OF EDUCATION - ENROLLED	16	17	17	18	14	NA
NOUSTRY CODE   10	EMPLOYMENT TYPE	7	7	6	5	6	7
CURRENTITY ENROLLED AT SCHOOL   2   2   2   2   2   2   2   2   2	LITERACY	2	NA	2	3	2	2
Processing Platform   16 Processors   2 Processors   2 Processors   2 Processors   2 Processors   16 Processors   4 Processors   2 Processors   16 Processors   4 Processors   2 Processors   2 Processors   16 Processors   4 Processors   2 Proces	INDUSTRY CODE	10	464	10	11	10	10
Processing Platform   16 Processors   2 Processors   2 Processors   16 Processors   4 Processors   16 Processors   2 Processors   16 Processors   2 Proces	CURRENTLY ENROLLED AT SCHOOL						
Total execution time with parallelization (hours)    10.75   23.70   67.58   77.14   282.58							
Scenarios	Processing Platform	16 Processors	2 Processors	2 Processors	2 Processors	16 Processors	4 Processors
Scenarios  1		6 11	10.75	22.70	67.50	77.14	202 50
Dasse 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0	paranenzation (nours)	0.11	10.75	23./0	07.58	//.14	282.58
Dasse 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0				· · · · · · · · · · · · · · · · · · ·		·	
1 0.35 0.72 0.15 0.31 0.36 0.30 0.30 0.37 0.75 0.16 0.32 0.37 0.31 0.38 0.38 0.38 0.38 0.38 0.38 0.38 0.38		_				_	
2 0.37 0.75 0.16 0.32 0.37 0.31 0.31 0.38 0.38 0.38 0.38 0.38 0.38 0.38 0.39 0.81 0.16 0.33 0.37 0.31 0.31 0.39 0.81 0.16 0.33 0.37 0.31 0.35 0.39 0.81 0.16 0.33 0.38 0.32 0.35 0.31 0.35 0.32 0.39 0.81 0.16 0.33 0.38 0.32 0.32 0.35 0.31 0.31 0.32 0.32 0.35 0.31 0.32 0.32 0.35 0.31 0.32 0.32 0.35 0.35 0.32 0.35 0.35 0.32 0.35 0.35 0.32 0.35 0.35 0.35 0.35 0.35 0.35 0.35 0.35							
0.38	F		l		l	l	
4							
1.14			l		I	l	
1.34	1.		l		I		
7			l		I	l	
8						l	
2.99   2.75   1.77   5.01   43.21   177.47   10	l'						
100							
111			l		I	l	
122		2.02	0.00	1 2120	1.00	27.55	
13			l			l	
144 0.78 0.54 0.96 2.31 15.83 74.24 155 0.90 0.67 0.83 2.13 21.59 123.70 166 2.82 2.84 1.74 5.00 40.50 179.26 177 1.61 0.79 1.21 4.04 19.00 82.80 188 1.48 0.62 1.11 4.31 24.11 118.77 19.00 1.25 1.43 0.96 2.35 21.36 130.83 21.3 21.59 1.36 130.83 21.3 0.78 2.25 1.36 130.83 21.3 0.78 2.25 1.44 0.44 19.00 1.25 1.45 1.45 1.45 1.45 1.45 1.45 1.45 1.4			l				
15							
166							
17			l		I	l	
18       1.48       0.62       1.11       4.31       24.11       118.77         19       4.06       6.51       2.21       6.07       31.79       159.04         20       1.25       1.43       0.96       2.35       21.36       130.83         21       0.78       0.54       0.94       2.36       16.44       69.55         22       0.91       0.68       0.80       2.18       22.74       131.49         23       2.82       2.86       1.73       4.98       42.32       168.45         24       1.49       0.81       1.20       4.04       19.37       84.13         25       1.40       0.64       1.08       4.26       24.54       113.05         26       3.94       6.36       2.19       5.99       31.00       145.07         27       1.36       1.43       0.96       2.31       23.18       107.94         28       0.79       0.54       0.93       2.30       16.96       73.40         29       0.90       0.70       0.81       2.13       22.25       118.66         30       2.86       2.80       1.79       5.05		2.02			0.00	l	
19			l		I		
200	19						
21	20						
222 0.91 0.68 0.80 2.18 22.74 131.49 23 2.82 2.86 1.73 4.98 42.32 168.45 24 1.49 0.81 1.20 4.04 19.37 84.13 25 1.40 0.64 1.08 4.26 24.54 113.05 26 3.94 6.36 2.19 5.99 31.00 145.07 27 1.36 1.43 0.96 2.31 23.18 107.94 28 0.79 0.54 0.93 2.30 16.96 73.40 29 0.90 0.70 0.81 2.13 22.25 118.46 30 2.86 2.80 1.79 5.05 40.63 137.30 31 1.45 0.80 1.24 4.11 19.01 80.12 32 1.60 0.65 1.14 4.39 23.87 95.11 33 3.92 6.40 2.31 6.13 29.99 110.22 34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	21						
23	22		l		I		
24	23		0.00	"""			202110
25	24		l				
26 3.94 6.36 2.19 5.99 31.00 145.07 27 1.36 1.43 0.96 2.31 23.18 107.94 28 0.79 0.54 0.93 2.30 16.96 73.40 29 0.90 0.70 0.81 2.13 22.25 118.46 30 2.86 2.80 1.79 5.05 40.63 137.30 31 1.45 0.80 1.24 4.11 19.01 80.12 32 1.60 0.65 1.14 4.39 23.87 95.11 33 3.92 6.40 2.31 6.13 29.99 110.22 34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	25				I		
27	26						
28 0.79 0.54 0.93 2.30 16.96 73.40 29 0.90 0.70 0.81 2.13 22.25 118.46 31 37.30 30 2.86 2.80 1.79 5.05 40.63 137.30 31 1.45 0.80 1.24 4.11 19.01 80.12 32 1.60 0.65 1.14 4.39 23.87 95.11 33 3.92 6.40 2.31 6.13 29.99 110.22 34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	27						
299 0.90 0.70 0.81 2.13 22.25 118.46 30 2.86 2.80 1.79 5.05 40.63 137.30 31 1.45 0.80 1.24 4.11 19.01 80.12 32 1.60 0.65 1.14 4.39 23.87 95.11 33 3.92 6.40 2.31 6.13 29.99 110.22 34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time If sequentially processed	28		l			l	
30	29		0.70			22.25	
31 1.45 0.80 1.24 4.11 19.01 80.12 32 1.60 0.65 1.14 4.39 23.87 95.11 33 3.92 6.40 2.31 6.13 29.99 110.22 34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46  Approximated execution time if sequentially processed	30		l				
32 1.60 0.65 1.14 4.39 23.87 95.11 33 3.92 6.40 2.31 6.13 29.99 110.22 35 4 6.13 29.99 120.22 55 0.80 0.54 0.99 2.32 16.32 72.85 6 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	31						
33 3.92 6.40 2.31 6.13 29.99 110.22 34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	32						
34 1.27 1.43 1.02 2.35 22.61 97.67 35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	33	3.92		2.31	6.13	29.99	110.22
35 0.80 0.54 0.99 2.32 16.32 72.85 36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	34		1.43			l	
36 0.89 0.70 0.85 2.14 19.68 101.54 37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	35	0.80	0.54	0.99	I	16.32	
37 2.57 2.82 1.92 4.99 37.70 104.72 38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time if sequentially processed	36		0.70	0.85	l	19.68	101.54
38 1.34 0.80 1.35 4.03 18.64 70.92 39 1.42 0.65 1.17 4.17 23.10 70.46 Approximated execution time If sequentially processed	37		l		4.99		
39 1.42 0.65 1.17 4.17 23.10 70.46  Approximated execution time	38		0.80				
Approximated execution time if sequentially processed	39	1.42			4.17		
if sequentially processed	Approximated execution time						
	if sequentially processed						
	(hours)	65.48	69.52	46.46	132.20	880.06	3987.25

\*NA indicates that the variable is not available in that dataset and, hence, could not be selected as quasi-identifier.

Description of anonymization scenarios|image22|

СНАРТЕЯ	
TEN	

# **BIBLIOGRAPHY**

#### **CHAPTER**

## **ELEVEN**

## **LISTS**

May no longer be needed in this format or need to be automated

## 11.1 List of Tables

- Table 4.1: Example dataset showing sample frequencies, population frequencies and individual disclosure risk 20
- Table 4.2: Example dataset to illustrate the effect of missing values on k-anonymity 25
- Table 4.3: 1-diversity illustration 26
- Table 4.4: Sample uniques and special uniques 28
- Table 4.5: Illustrating the calculation of SUDA and DIS-SUDA scores 30
- Table 5.1: SDC methods and corresponding functions in sdcMicro 39
- Table 5.3: Illustration of the effect of recoding on the theoretically possible number of combinations an a dataset 40
- Table 5.3: Illustration of the effect of recoding on the theoretically possible number of combinations an a dataset 41
- Table 5.4: Local suppression illustration sample data before and after suppression 49
- Table 5.5: How importance vectors and k-anonymity thresholds affect running time and total number of suppressions 52.
- Table 5.6 Effect of the all-*m* approach on k-anonymity 53
- Table 5.7: Sample data before and after perturbation 56
- Table 5.8: Tabulation of variable "region" before and after PRAM 57
- Table 5.9: Tabulation of variable "region" before and after (invariant) PRAM 58
- Table 5.10: Cross-tabulation of variable "region" and variable "gender" before and after invariant PRAM 59
- Table 5.11: Illustrating the effect of choosing mean vs. median for microaggregation where outliers are concerned 62
- Table 5.12: Illustration of multivariate microaggregation 63
- Table 5.13: Grouping methods for microaggregation that are implemented in sdcMicro 64
- Table 5.14: Simplified example of the shuffling method 72
- Table 6.1: Description of anonymization methods by scenario 90
- Table 7.1: Packages and functions for reading data in R 95
- Table 7.2: Slot names and slot description of sdcMicro object 100
- Table 7.3 Illustration of randomizing order of records in a dataset 104

### **Readthedocs Documentation**

- Table 7.4: Computation times of different methods on datasets of different sizes 106
- Table 7.5: Number of categories (complexity), record uniqueness and computation times 106
- Table 8.1: Illustration of merging variables without information loss for SDC process 111
- Table 9.1 Overview of variables in dataset 122
- Table 9.2: List of selected quasi-identifiers 126
- Table 9.3: Overview of selected utility measures 127
- Table 9.4: Number and proportion of households violating k-anonymity 129
- Table 9.5: Percentiles 90-100 of the variable LANDSIZE 131
- Table 9.6: 50 largest values of the variable LANDSIZE 131
- Table 9.7: Frequencies of variable HHSIZE (household size) 132
- Table 9.8: Number and proportion of households violating k-anonymity after anonymization 137
- Table 9.9: Univariate frequencies of the PRAMmed variable before and after anonymization 139
- Table 9.10: Multivariate frequencies of the variables WATER with RURURB before and after anonymization 139
- Table 9.11: GINI point estimates and bootstrapped confidence intervals for sum of expenditure components 139
- Table 9.12: Mean monthly expenditure and mean monthly income per capita by rural/urban 140
- Table 9.13 Shares of expenditures components 140
- Table 9.14: k-anonymity violations 142
- Table 9.15: Number of suppressions by variable for different variations of local suppression 145
- Table 9.16: k-anonymity violations 145
- Table 9.17: Net enrollment in primary and secondary education by gender 145
- Table 9.18: Overview of the variables in the dataset 149
- Table 9.19: Overview of selected key variables for PUF file 152
- kTable 9.20: Number and proportion of households violating -anonymity 155
- Table 9.21: Number of suppressions by variable after local suppression with and without importance vector 158
- Table 9.22: Comparison of utility measures 163
- Table 9.23: Number of records violating k-anonimity 164
- Table 9.24: Overview of recodes of categorical variables at individual level 166
- Table 9.25: k-anonymity violations 170

# 11.2 List of Figures

- Figure 2.1: Risk-utility trade-off 7
- Figure 4.1 Classification of variables 17
- Figure 4.2: Visualizations of DIS-SUDA scores 32
- Figure 5.1 Effect of recoding frequency counts before and after recoding 43
- Figure 5.2 Age variable before and after recoding 44

120 Chapter 11. Lists

- Figure 5.3 Age variable before and after recoding 45
- Figure 5.4: Utilizing the frequency distribution of variable age to determine threshold for top coding 47
- Figure 5.5: Changes in labor market indicators after anonymization of I2D2 data 51
- Figure 5.6: Illustration of effect of noise addition to outliers 66
- Figure 5.7: Frequency distribution of a continuous variable before and after noise addition 68
- Figure 5.8: Noise levels and the impact on the value range (percentiles) 69
- Figure 6.1: The trade-off between risk and utility in a hypothetical dataset 77
- Figure 6.2: Effect of anonymization on the point estimates and confidence interval of the gender coefficient in the Mincer equation 86
- Figure 6.3: Histograms of income before and after anonymization 87
- Figure 6.4: Density plots of income before and after anonymization 88
- Figure 6.5: Example of box plots of an expenditure variable before and after anonymization 89
- Figure 6.6: Mosaic plot to illustrate the changes in the WATER variable 90
- Figure 6.7: Comparison of treated vs. untreated gender and relationship status variables with mosaic plots 91
- Figure 6.8: Mosaic plot of the variables ROOF and TOILET before anonymization 92
- Figure 6.9: Mosaic plot of the variables ROOF and TOILET after anonymization 92
- Figure 8.1: Overview of the SDC process 118
- Figure 9.1: Lorenz curve based on positive total expenditures values 140

# 11.3 List of Examples

- $f_k$ Example 4.1: Calculating using sdcMicro 21
- Example 4.2: Calculating the sample and population frequencies using *sdcMicro* 22
- Example 4.3: The individual risk slot in the *sdcMicro* object 23
- Example 4.4: Using the print() function to display observations violating k-anonymity 24
- Example 4.5: Computing k-anonymity violations for other values of k 25
- lExample 4.6: -diversity function in sdcMicro 27
- Example 4.7: Evaluating SUDA scores 31
- Example 4.8: Histogram and density plots of DIS-SUDA scores 31
- Example 4.9 Example with the function dRisk() 33
- Example 4.10: Computing 90 % percentile of variable income 34
- Example 4.11: Computation of the global risk measure 34
- Example 4.12: Computation of expected number of re-identifications 35
- Example 4.13: Number of individuals with individual risk higher than the threshold 0.05 35
- Example 4.14: Computation of household risk and expected number of re-identifications 36
- Example 5.1: Using the sdcMicro function groupVars() to recode a categorical variable 42
- Example 5.2: Using the sdcMicro function globalRecode() to recode a continuous variable (age) 43

### **Readthedocs Documentation**

- Example 5.3: Using globalRecode() to create intervals of unequal width 44
- Example 5.4: Constructing right-open intervals for semi-continuous variables using built-in *sdcMicro* function globalRecode() 46
- Example 5.5: Constructing intervals for semi-continuous and continuous variables using manual recoding in R 46
- Example 5.6: Top coding and bottom coding in sdcMicro using topBotCoding() function 47
- Example 5.7: Application of local suppression with and without importance vector 50
- mExample 5.8 The all-approach in sdcMicro 53
- Example 5.9: Manually suppressing values in linked variables 54
- Example 5.10: Suppressing values in linked variables by specifying ghost variables 54
- Example 5.11: Application of built-in sdcMicro function localSupp() 55
- Example 5.12: Producing reproducible PRAM results by using set.seed() 58
- Example 5.13: Selecting the variable "toilet" to apply PRAM 59
- Example 5.14: Specifying minimum values for diagonal entries in PRAM transition matrix 59
- Example 5.15: Minimizing unlikely combinations by applying PRAM within strata 60
- Example 5.16: Applying univariate microaggregation with sdcMicro function microaggregation() 62
- Example 5.17: Multivariate microaggregation with the Maximum Distance to Average Vector (MDAV) algorithm in *sdcMicro* 64
- Example 5.18: Specifying strata variables for microaggregation 64
- Example 5.19: Uncorrelated noise addition 67
- Example 5.20: Correlated noise addition 69
- Example 5.21: Noise addition for outliers using the 'outdect' method 70
- Example 5.22: Noise addition to aggregates and their components 70
- Example 5.23: Rank swapping using sdcMicro 71
- Example 5.24: Shuffling using a specified regression equation 73
- Example 6.1: Using the print() function to retrieve the total number of suppressions for each categorical key variable 78
- Example 6.2: Displaying the number of missing values for each categorical key variable in an sdcMicro object 78
- Example 6.3: Computing number of records changed per variable 79
- Example 6.4: Comparing contingency tables of categorical variables 79
- Example 6.5: Comparing the means of continuous variables 80
- Example 6.6: Comparing covariance structure and correlation matrices of numeric variables 81
- Example 6.7: Using dUtility() to compute IL1s data utility measure in sdcMicro 82
- Example 6.8: Using dUtility() to compute eigenvalues in sdcMicro 83
- Example 6.9: Computing the GINI coefficient from the income variable to determine income inequality 83
- Example 6.10: Estimating the Mincer equation (regression) to evaluate data utility before and after anonymization 84
- Example 6.11: Plotting histograms and kernel densities 87
- Example 6.12: Creating boxplots for continuous variables 88

122 Chapter 11. Lists

- Example 6.13: Creating univariate mosaic plots 89
- Example 6.14: Creating multivariate mosaic plots 91
- Example 7.1: Loading required packages 94
- Example 7.2: Displaying help for functions 94
- Example 7.3: Reading in a STATA file 95
- Example 7.4: Reading in an Excel file 95
- Example 7.5: Reading in an SPSS file 96
- Example 7.6: Recoding missing values to NA 96
- Example 7.7: Changing the class of an object in R 97
- Example 7.8: Selecting variables and creating an object of class *sdcMicroObj* for the SDC process in *R* 98
- Example 7.9: Displaying slot names and accessing slots of an S4 object 100
- Example 7.10: Saving results of applying SDC methods 101
- Example 7.11: Undo last step in SDC process 101
- Example 7.12: Create a household level file with unique records (remove duplicates) 102
- Example 7.13 Merging anonymized household-level variables with individual-level variables 103
- Example 7.14 Generating the variable household size 103
- Example 7.15 Changing the order of individuals within households 103
- Example 7.16: Randomize order of households 105
- Example 9.1: Loading required packages 120
- Example 9.2: Loading the data 120
- Example 9.3: Number of individuals and variables and variable names 120
- Example 9.4: Tabulation of the variable 'gender' and summary statistics for the variable 'total annual expenditures' in *R* 122
- Example 9.5: Recoding missing value codes 124
- Example 9.6: Dropping variables with only missing values 125
- Example 9.7: Selecting the variables for the household-level anonymization 127
- Example 9.8: Taking a subset with only households 128
- Example 9.9: Creating a *sdcMicro* object for the household variables 128
- Example 9.10: Showing number of households violating k-anonymity for levels 2,3 and 5 129
- Example 9.11: Showing households that violate k-anonymity 130
- Example 9.12: Printing global risk measures 130
- Example 9.13: Observations with individual risk higher than 1% 130
- Example 9.14 Percentiles of LANDSIZE and listing the sizes of the largest 50 plots 131
- Example 9.15: Removing households with large (rare) household sizes 132
- Example 9.16: Local suppression with and without importance vector 134
- Example 9.17: Applying PRAM 135

### **Readthedocs Documentation**

- Example 9.18: Anonymizing the variable LANDSIZEHA 135
- Example 9.19: Anonymizing continuous variables 137
- Example 9.20: Measuring risk of re-identification of continuous variables 138
- Example 9.21: Merging the files with household and individual-level variables and creating an *sdcMicro* object for the anonymization of the individual-level variables 141
- Example 9.22: Global risk of the individual-level variables 142
- Example 9.23: Recoding age in 10-year intervals in the range 15 65 and top code age over 65 years 143
- Example 9.24: Experimenting with different options in local suppression 144
- Example 9.25: Using the report() function for internal and external reports 146
- Example 9.26: Exporting the anonymized dataset 147
- Example 9.27: Loading required packages and datasets 148
- Example 9.28 Number of individuals and variables and variable names 148
- Example 9.29: Selecting the variables for the household-level anonymization 153
- Example 9.30: Taking a subset with only households 154
- Example 9.31: Creating a sdcMicro object for the household variables 154
- kExample 9.32: Showing number of households violating -anonymity for levels 2, 3 and 5 155
- kExample 9.33: Showing records of households that violate -anonymity 156
- Example 9.34: Printing global risk measures 156
- Example 9.35 Determining the highest individual risk 156
- Example 9.36: Local suppression with and without importance vector 158
- Example 9.37: Applying PRAM 160
- Example 9.38: Anonymization of income and expenditure variables 161
- Example 9.39: Measuring risk of re-identification of continuous variables 162
- Example 9.40: Computation of decile dispersion ratio and share of total consumption by the poorest decile 163
- Example 9.41: Merging the files with household and individual-level variables and creating an *sdcMicro* object for the anonymization of the individual-level variables 163
- Example 9.42: Risk measures before anonymization 165
- Example 9.43: Recoding the categorical and continuous variables 166
- Example 9.44: Local suppression to reach 5-anonimity 168
- Example 9.45: Randomizing the order of records within regions 169
- Example 9.46: Exporting the anonymized PUF file 171

124 Chapter 11. Lists

# CASE STUDY (ILLUSTRATING THE SDC PROCESS)

In order to evaluate the use of different SDC methods on different types of survey datasets, we compared the results of the different methods applied to 75 datasets from 52 countries representing six geographic regions: Latin America and the Caribbean (LAC), Sub-Saharan Africa (AFR), South Asia (SA), Europe and Central Asia (ECA), Middle East and North Africa (MENA) as well as East Asia and the Pacific (EAP). The datasets chosen were from a mix of datasets that are already publically available, as well as data made available only to the World Bank for official business. The surveys used included, amongst others, household, demographic, and health surveys. The variables from these surveys used for the experiments were selected based on their relevance for users (e.g. for indicators, MDGs), their sensitivity, and their classification with respect to the SDC process.

The following case studies draw from knowledge gained from these experiments and try to incorporate the lessons learned. The case studies use synthetic data that mimic the structure of the survey types we used in our experiments and present the anonymization of a dataset similar to many surveys designed to measure household income and consumption, labor force participation and general demographic characteristics. The first case study creates a SUF, whereas in the second case study we take this SUF and treat it further to create a PUF.

# 12.1 Case study 1- SUF

This case study shows an example of how the anonymization process might be approached, particularly for a dataset with many continuous variables. We also show how this can be achieved using the open source and free *sdcMicro* package and *R*. A ready-to-run *R* script for this case study and the dataset are also available to reproduce the results and allow the user to adapt the code (see http://ihsn.org/home/projects/sdc-practice). Extracts of this code are presented in this section to illustrate several steps of the anonymization process. **NOTE: The choices of methods and parameters in this case study are based on this particular dataset and the results and choices might be different for other datasets.** The aim is to show the process, not to compare methods per se.

This example uses a dataset with a similar structure to that of a typical social survey with a focus on demographics, labor force participation and income and expenditure patterns. The dataset has been compiled using observations from several datasets from different countries. They are considered synthetic data and as such are used only for illustrative purposes. The source datasets were already treated for disclosure control by their producers. This does not matter, as our concern is to illustrate the process only. The data from which we compiled our case study file was from surveys that contain many variables, but pay particular attention to labor force variables as well as household income and household expenditure variables. The variables in the demo dataset have already been pre-selected from the total set of variables available in the datasets. See Appendix A for the complete overview of all variables.

This case study follows the steps of the SDC process outlined in Chapter 8.

### **Step 1: Need for disclosure control**

The statistical units in this dataset are individuals and households. The household structure provides a hierarchical structure in the data, which should be taken into account when measuring risk and selecting anonymization methods.

The data contains variables with demographic information, income, expenditures, education variables and variables relating to the labor status of the individual. These variables include sensitive and confidential variables. The dataset is an example of a social survey and, due to the nature of the statistical units and the variables, disclosure control is needed before release of the microdata. This is the case regardless of the legal framework, which is not specified here, as this is a hypothetical dataset.

### Step 2: Data preparation and exploring data characteristics

The first step is to explore the data. To analyze the data in R we first have to read the data into R. In our case, the data is saved in a STATA file (.dta). To read STATA files, we need to load the R package foreign (see Section 7.2 on importing other data formats in R). We also load the sdcMicro package and several other packages used later for the computation of the utility measures. If these packages are not yet installed, you should do so before trying to load them. The R code for this case study demonstrates how to do this.

Example 9.1: Loading required packages

```
# Load required packages
library(foreign) # for read/write function for STATA files
library(sdcMicro) # sdcMicro package with functions for the SDC process
```

```
library(laeken) # for GINI
library(reldist) # for GINI
```

```
library(bootstrap) # for bootstrapping
library(ineq) # for Lorenz curves
```

After setting the working directory to the directory where the *STATA* file is stored, we load the data into the object called *file*. All output, unless otherwise specified, is saved in the working directory.

```
Example 9.2: Loading the data
```

setwd("C:/WorldBank/CaseStudy/") # Set working directory

```
# Specify file name
fname <- " case_1_data.dta"
# Read-in file
file <- read.dta(fname, convert.factors = F) # factors as numeric code</pre>
```

We check the number of variables, number of observations and variable names, as shown in Example 9.3.

Example 9.3: Number of individuals and variables and variable names

```
dim(file) # Dimensions of file (observations, variables)
```

colnames(file) # Variable names

## [1] 10574 68

```
## [1] "REGION" "DIST" "URBRUR" "WGTHH"
## [5] "WGTPOP" "IDH" "IDP" "HHSIZE"
```

```
## [9] "GENDER" "REL" "MARITAL" "AGEYRS"
## [13] "AGEMTH" "RELIG" "ETHNICITY" "LANGUAGE"
## [17] "MORBID" "MEASLES" "MEDATT" "CHWEIGHTKG"
## [21] "CHHEIGHTCM" "ATSCHOOL" "EDUCY" "EDYRS"
## [25] "EDYRSCURRAT" "SCHTYP" "LITERACY" "EMPTYP1"
## [29] "UNEMP1" "INDUSTRY1" "EMPCAT1" "WHOURSWEEK1"
## [33] "OWNHOUSE" "ROOF" "TOILET" "ELECTCON"
## [37] "FUELCOOK" "WATER" "OWNAGLAND" "LANDSIZEHA"
## [41] "OWNMOTORCYCLE" "CAR" "TV" "LIVESTOCK"
## [45] "INCRMT" "INCWAGE" "INCBONSOCALL" "INCFARMBSN"
## [49] "INCOTHER" "INCRENT" "INCFIN" "INCPENSN"
## [53] "INCOTHER" "INCTOTGROSSHH" "FARMEMP" "THOUSEXP"
## [57] "TFOODEXP" "TALCHEXP" "TCLTHEXP" "TFURNEXP"
## [61] "THLTHEXP" "TTRANSEXP" "TCOMMEXP" "TRECEXP"
## [65] "TEDUEXP" "TRESTHOTEXP" "TMISCEXP""TANHHEXP"
```

The dataset has 10,574 individuals in 2,000 households and contains 68 variables. The survey corresponds to a population of about 4.3 million individuals, which means that the sample is relatively small and the sample weights are high. This has an impact on the disclosure risk, as we will see in Steps 6a and 6b.

To get an overview of the values of the variables, we use tabulations and cross-tabulations for categorical variables and summary statistics for continuous variables. To include the number of missing values (NA or other), we use the option useNA = "ifany" in the table() function (see

## Example 9.4).

In Table 9.1 the variables in the dataset are listed along with concise descriptions of the variables, the level at which they are collected (individual (IND), household (HH)), the measurement type (continuous, semi-continuous, categorical) and value ranges. Note that the dataset contains a selection of 68 variables (cf. Appendix A) of a total of 112 variables in the survey dataset. The variables have been preselected based on their relevance for data users. This allows to reduce the total numbers of variables to consider in the anonymization process and makes the process easier. The numerical values for many of the categorical variables are codes that refer to values, e.g., in the variable URBRUR, 1 stands for rural and 2 for urban. More information on the meanings of coded values of the categorical variables is available in the *R* code for this case study.

We identified the following sensitive variables in the data: ETHNICITY, RELIGION, variables related to the labor force status of the individual and the variables containing information on income and expenditures of the household. Whether variables can be identified as sensitive may vary across countries and datasets.

The case study dataset does not have any direct identifiers that, if they were present, would need to be removed at this stage. Examples of direct identifiers would be names, telephone numbers, geographical location coordinates, etc.

Example 9.4: Tabulation of the variable 'gender' and summary statistics for the variable 'total annual expenditures' in R

table(file\$GENDER, useNA = "ifany") # tabulation of variable GENDER (sex, categorical)

```
## 0 1
## 5448 5126
```

**summary**(file\$TANHHEXP) # summary statistics for variable TANHHEXP (total annual household expenditures, continuous)

## Min. 1st Qu. Median Mean 3rd Qu. Max. ## 498 15550 17290 28560 29720 353200

Table 9.1 Overview of variables in dataset

No.	Variable name	Description	Level	Measurement	Values
1	IDH	Household ID	НН	•	1-2,000
2	IDP	Individua l ID	IND	•	1-33
3	REGION	Region	HH	categoric al	1-6
4	DISTRICT	District	HH	categoric al	101-1105
5	URBRUR	Area of residence	НН	categoric al	1, 2
6	WGTHH	Individua 1 weighting coefficie nt	НН	weight	31.2-8495 .7
7	WGTPOP	Populatio n weighting coefficie nt	IND	weight	45.8-9345 2.2
8	HHSIZE	Household size	НН	semi-cont inuous	1-33
9	GENDER	Gender	IND	categoric al	0, 1
10	REL	Relations hip to household head	IND	categoric al	1-9
11	MARITAL	Marital status	IND	categoric al	1-6
12	AGEYRS	Age in completed years	IND	semi-cont inuous	0-95 (under 1, 1/12 year increment s)
13	AGEMTH	Age of child in completed years	IND	semi-cont inuous	1-1140
14	RELIG	Religion of household head	НН	categoric al	1, 5-7, 9
15	ETHNICITY	Ethnicity of household head	НН	categoric al	all missing val- ues
16	LANGUAGE	Language of household head	НН	categoric al	all missing val- ues
17	MORBID	Morbidity last x weeks	IND	categoric al	0, 1
18	MEASLES	Child immu- nized against measles	IND	categoric al	0, 1, 9
19	MEDATT	Sought medical attention	IND	categoric al	0, 1
20	CHWEIGHTK G	Weight of the child (Kg)	IND	continuou s	2 – 26.5
21	CHHEIGHTC M	Height of child (cms)	IND	continuou s	7 - 140

Continued on next page

Table 12.1 – continued from previous page

No.	Variable name	Description	Level	Measurement	Values
22	ATSCHOOL	Currently en- rolled in school	IND	categoric al	0, 1
23	EDUCY	Highest level of education attended	IND	categoric al	1-6
24	EDYEARS	Years of educa- tion	IND	semi-cont inuous	0-18
25	EDYRSCURR AT	Years of edu- cation for cur- rently enrolled	IND	semi-cont inuous	1-18
26	SCHTYP	Type of school attending	IND	categoric al	1-3, 9
27	LITERACY	Literacy	IND	categoric al	1-3
28	EMPTYP1	Type of employ- men t	IND	categoric al	1-9
29	UNEMP1	Unemploye d	IND	categoric al	0, 1
30	INDUSTRY1	Industry classific ation (1-digit)	IND	categoric al	1-10
31	EMPCAT1	Employmen t categorie s	IND	categoric al	11, 12, 13, 14, 21, 22
32	WHOURSLAS TWEEK1	Hours worked last week	IND	continuou s	0-154
33	OWNHOUSE	Ownership of dwelling	НН	categoric al	0, 1
34	ROOF	Main material used for roof	IND	categoric al	1-5, 9
35	TOILET	Main toilet fa- cility	НН	categoric al	1-4, 9
36	ELECTCON	Electrici ty	НН	categoric al	0-3
37	FUELCOOK	Main cooking fuel	НН	categoric al	1-5, 9
38	WATER	Main source of water	НН	categoric al	1-9
39	OWNAGLAND	Ownership of agricultu ral land	НН	categoric al	1-3
40	LANDSIZEH A	Land size owned by household (ha) (agric and non agric)	НН	continuou s	0-1214
41	OWNMOTORC YCLE	Ownership of motorcycl e	НН	categoric al	0, 1
42	CAR	Ownership of car	НН	categoric al	0, 1
43	TV	Ownership of televisio n	НН	categoric al	0, 1

Continued on next page

Table 12.1 – continued from previous page

No.	Variable name	Die 12.1 – continue Description	Level	Measurement	Values
44	LIFESTOCK	Number of large-siz ed livestock owned	НН	semi-cont inuous	0-25
45	INCRMT	Income – Remittances	НН	continuou s	•
46	INCWAGE	Income - Wages and salaries	НН	continuou s	•
47	INCBONSOC ALL	Income - Bonuses and social allowance s derived from wage jobs	НН	continuou s	•
48	INCFARMBS N	Income - Gross income from household farm businesse s	НН	continuou s	•
49	INCNFARMB SN	Income -Gross income from household non- farm businesse s	НН	continuou s	•
50	INCRENT	Income - Rent	НН	continuou s	•
51	INCFIN	Income - Finan- cial	НН	continuou s	•
52	INCPENSN	Income - Pen- sions/ social as- sistanc e	НН	continuou s	•
53	INCOTHER	Income - Other	НН	continuou s	•
54	INCTOTGRO SSHH	Income - Total	НН	continuou s	•
55	FARMEMP				•
56	TFOODEXP	Total expenditu re on food	НН	continuou s	•
57	TALCHEXP	Total expenditu re on alcoholic beverages, to- bacco and nar- cotics	НН	continuou s	• ued on next page

Continued on next page

Table 12.1 – continued from previous page

No.	Variable name	Description	Level	Measurement	Values
58	TCLTHEXP	Total expenditu re on clothing	НН	continuou s	•
59	THOUSEXP	Total expenditu re on housing	НН	continuou s	•
60	TFURNEXP	Total expenditu re on furnishin g	НН	continuou s	•
61	THLTHEXP	Total expenditu re on health	НН	continuou s	•
62	TTRANSEXP	Total expenditu re on transport	НН	continuou s	•
63	TCOMMEXP	Total expenditu re on commu- nica tion	НН	continuou s	•
64	TRECEXP	Total expenditu re on recreatio n	НН	continuou s	•
65	TEDUEXP	Total expenditu re on education	НН	continuou s	•
66	TRESHOTEX P	Total expenditu re on restauran ts and hotels	НН	continuou s	•
67	TMISCEXP	Total expenditu re on miscellan eous spending	НН	continuou s	•
68	TANHHEXP	Total annual nominal house- hold expenditu res	НН	continuou s	•

It is always important to ensure that the relationships between variables in the data are preserved during the anonymization process and to explore and take note of these relationships before beginning the anonymization. In the final step in the anonymization process, an audit should be conducted, using these initial results, to check that these relationships are maintained in the anonymized dataset.

In our demo dataset, we identify several relationships between variables that need to be preserved during the anonymization process. The variables TANHHEXP and INCTOTGROSSHH represent the total annual nominal household expenditure and the total gross annual household income, respectively, and these variables are aggregations of existing income and expenditure components in the dataset.

The variables related to education are available only for individuals in the appropriate age groups and missing for other individuals. We make a similar observation for variables relating to children, such as height, weight and age in months. In addition, the household-level variables (cf. fourth column of Table 9.1) have the same values for all members in any particular household. The value of household size corresponds to the actual number of individuals belonging to that household in the dataset. As we proceed, we have to take care that these relationships and structures are preserved in the anonymization process.

When tabulating the variables, we notice that the variables RELIG, EMPTYP1 and LIVESTOCK have missing value

codes different from the *R* standard missing value code NA. Before proceeding, we need to recode these to NA so *R* interprets them correctly. The missing value codes are resp. 99999, 99 and 9999 for these three variables. These are genuine missing value codes and not caused by the variables being not applicable to the individual. Example 9.5 shows how to make these changes. **NOTE:** At the end of the anonymization process, and if desired for users, it is relatively easy to change these values back to their original missing value code.

```
Example 9.5: Recoding missing value codes
# Set different NA codes to R missing value NA
file[,'RELIG'][file[,'RELIG'] == 99999] <- NA
file[,'EMPTYP1'][file[,'EMPTYP1'] == 99] <- NA
file[,'LIVESTOCK'][file[,'LIVESTOCK'] == 9999] <- NA
```

We also take note that the variables LANGUAGE and ETHNICITY have only missing values. Variables that contain only missing values should be removed from the dataset at this stage and excluded from the anonymization process. Removing these variables does not mean loss of data or reduction of the data utility, since these variables did not contain any information. It is, however, necessary to remove them, because keeping them can lead to errors in some of the anonymization methods in *R*. It is always possible to add these variables back into the dataset to be released at the end of the anonymization process. It is useful to reduce the dataset to those variables and records relevant for the anonymization process. This guarantees the best results in *R* and fewer errors. In Example 9.6 we drop the variables that contain all missing values.

Example 9.6: Dropping variables with only missing values

```
# Drop variables containing only missings file <- file[,!names(file) %in% c('LANGUAGE', 'ETHNICITY')]
```

We assume that the data are collected in a survey that uses simple sampling of households. The data contains two weight coefficients: WGTHH and WGTPOP. The relationship between the weights is WGTPOP = WGTHH \* HH-SIZE. WGTPOP is the sampling weight for the households and WGTHH is the sampling weight for the individuals to be used for disclosure risk calculations. WGTHH is used for computing individual-level indicators (such as education) and WGTPOP is used for population level indicators (such as income indicators). There are no strata variables available in the data. We will use WGTPOP for the anonymization of the household variables and WGTHH for the anonymization of the individual-level variables.

### **Step 3: Type of release**

In this case study, we assume that data will be released as a SUF, which will be only available under license to accredited researchers with approved research proposals (see Section 3.2 for more information of the release of SUF). Therefore, the accepted risk level is higher and a broader set of variables can be released than would be the case when releasing a PUF. Since we do not have an overview of the requirements of all users, we restrict the utility measures to a selected number of data uses (see Step 5).

## Step 4: Intruder scenarios and choice of key variables

Next, we analyze possible intruder scenarios and select quasi-identifiers or key variables based on these scenarios. Since the dataset used in this case study is a demo dataset that does not stem from an existing country (and hence we do not have information on external data sources available to possible intruders) and the original data has already been anonymized, it is not possible to define exact disclosure scenarios. Instead, we draft intruder scenarios for this demo dataset based on some hypothetical assumptions about availability of external data sources. We consider two types of disclosure scenarios: 1) matching to other publicly available datasets and 2) spontaneous recognition. The license under which the dataset will be distributed (SUF) prohibits matching to external resources. Still this can

happen. However, the more important scenario is the one of spontaneous recognition. We describe both scenarios in the following two paragraphs.

For the sake of illustration, we assume that population registers are available with the demographic variables gender, age, place of residence (region, urban/rural), religion and other variables such as marital status and variables relating to education and professional status that are also present in our dataset. In addition, we assume that there is a publically available cadastral register on land ownership. Based on this analysis of available data sources, we select the variables REGION, URBRUR, HHSIZE, OWNAGLAND, RELIG, GENDER, REL (relationship to household head), MARITAL (marital status), AGEYRS, INDUSTRY1 and two variables relating to school attendance as categorical quasi-identifiers, the expenditure and income variables as well as LANDSIZEHA as continuous quasi-identifiers. According to our assessment, these variables might enable an intruder to re-identify an individual or household in the dataset by matching with other available datasets.

Table 9.2 gives an overview of the selected quasi-identifiers and their levels of measurement.

The decision to release the dataset as a SUF means the level of anonymization will be relatively low and consequently, the variables are more detailed and a scenario of spontaneous recognition is our main concern. Therefore, we should check for rare combinations or unusual patterns in the variables. Variables that may lead to spontaneous recognition in our sample are amongst others HHSIZE (household size), LANDSIZEHA as well as income and expenditure variables. Large households and large land ownership are easily identifiable. The same holds for extreme outliers in wealth and expenditure variables, especially when combined with other identifying variables such as region. There might be only one or a few households in a certain region with a high income, such as the local doctor. Variables that are easily observable and known by neighbors such as ROOF, TOILET, WATER, ELECTCON, FUELCOOK, OWNMOTORCYCLE, CAR, TV and LIVESTOCK may also need protection depending on what stands out in the community, since a researcher might be able to identify persons (s)he knows. This is called the nosy-neighbor scenario.

Table 9.2: List of selected quasi-identifiers

Name	Measurement
REGION (region)	Household, categorical
URBRUR (area of residence)	Household, categorical
HHSIZE (household size)	Household, categorical
OWNAGLAND (agricultural land ownership)	Household, categorical
RELIG (religion of household head)	Household, categorical
LANDSIZEHA (size of agr. and non-agr. land)	Household, continuous
TANHHEXP (total expenditures)	Household, continuous
T***EXP (expenditures in category ***)	Household, continuous
INCTOTGROSSHH (total income)	Household, continuous
INC*** (income in category ***)	Household, continuous
GENDER (sex)	Individual, categorical
REL (relationship to household head)	Individual, categorical
MARITAL (marital status)	Individual, categorical
AGEYRS (age in completed years)	Individual, semi-continuous
EDYRSCURATT (years of education for currently enrolled)	Individual, semi-continuous
EDUCY (highest level of education completed)	Individual, categorical
ATSCHOOL (currently enrolled in school)	Individual, categorical
INDUSTRY1 (industry classification)	Individual, categorical

Step 5: Data key uses and selection of utility measures

In this case study, our aim is to create a SUF that provides sufficient information for accredited researchers. We know that the primary use of these data will be to evaluate indicators relating to income and inequality. Examples are the GINI coefficient and indicators on what share of income is spent on what type of expenditures. Furthermore, we focus on some education indicators. Table 9.3 gives an overview of the utility measures we selected. Besides these utility measures, which are specific to the data uses, we also do standard checks, such as comparing tabulations,

cross-tabulations and summary statistics before and after anonymization.

Table 9.3: Overview of selected utility measures

Gini point estimates and confidence intervals for total expenditures
Lorenz curves for total expenditures
Mean monthly per capita total expenditures by area of residence
Average share of components for expenditures
Mean monthly per capita total income by area of residence
Average share of components for income
Net enrollment in primary education by gender

There are no published figures and statistics available that are calculated from this dataset because it is a demo. In general, the published figures should be re-computed based on the anonymized dataset and compared to the published figures in Step 11. Large differences would reduce the credibility of the anonymized dataset.

#### Hierarchical (household) structure

Our demo survey collects data on individuals in households. The household structure is important for data users and should be considered in the risk assessment. Since some variables are measured on the household level and thus have identical values for each household member, the values of the household variables should be treated in the same way for each household member (see Section 5.5). Therefore, we first anonymize only the household variables. After this, we merge them with the individual-level variables and then anonymize the individual-level and household-level variables jointly.

Since the data has a hierarchical structure, Steps 6 through 10 are repeated twice: Steps 6a through 10a are for the household-level variables and Steps 6b through 10b for the combined dataset. In this way, we ensure that household-level variable values remain consistent across household members for each household and the household structure cannot be used to re-identify individuals. This is further explained in Sections 4.4 and 7.6.

Before continuing to Step 6a, we select the categorical key variables, continuous key variables and any variables selected for use in PRAM routines, as well as household-level sampling weights. We extract these selected household variables and the households from the dataset and save them as *fileHH*. The choice of PRAM variables is further explained in Step 8a. Example 9.7 illustrates how these steps are done in *R* (see also Section 7.6). **NOTE: In our dataset, some of the categorical variables when imported from the STATA file were not imported as factors. sdcMicro requires that these be converted to factors before proceeding.** Conversion of these variables to factors is also shown in Example 9.7.

Example 9.7: Selecting the variables for the household-level anonymization

```
### Select variables (household level)

# Key variables (household level)

selectedKeyVarsHH = c('URBRUR', 'REGION', 'HHSIZE', 'OWNHOUSE', 'OWNAGLAND', 'RELIG')

# Changing variables to class factor

file$URBRUR <- as.factor(file$URBRUR)

file$REGION <- as.factor(file$REGION)

file$OWNHOUSE <- as.factor(file$OWNHOUSE)

file$OWNAGLAND <- as.factor(file$OWNAGLAND)
```

# Numerical variables

numVarsHH = c('LANDSIZEHA', 'TANHHEXP', 'TFOODEXP', 'TALCHEXP', 'TCLTHEXP', 'THOUSEXP', 'TFURNEXP', 'THLTHEXP', 'TTRANSEXP', 'TCOMMEXP', 'TRECEXP', 'TEDUEXP', 'TRESHOTEXP', 'TMISCEXP', 'INCTOTGROSSHH', 'INCRMT', 'INCWAGE', 'INCFARMBSN', 'INCNFARMBSN', 'INCRENT', 'INCFIN', 'INCPENSN', 'INCOTHER')

# PRAM variables

pramVarsHH = c('ROOF', 'TOILET', 'WATER', 'ELECTCON', 'FUELCOOK', 'OWNMOTORCYCLE', 'CAR', 'TV', 'LIVESTOCK')

# sample weight (WGTPOP) (household)
weightVarHH = c('WGTPOP')
# All household level variables
HHVars <- c('HID', selectedKeyVarsHH, pramVarsHH, numVarsHH, weightVarHH)

We then extract these variables from *file*, the dataframe in *R* that contains all variables. Every household has the same number of entries as it has members (e.g., a household of three will be repeated three times in *fileHH*). Before analyzing the household-level variables, we select only one entry per household, as illustrated in Example 9.8. This is further explained in Section 7.6.

Example 9.8: Taking a subset with only households

# Create subset of file with households and HH variables fileHH <- file[,HHVars]

# Remove duplicated rows based on IDH, select uniques, one row per household in fileHH fileHH <- fileHH[which(!duplicated(fileHH\$IDH)),]

### dim(fileHH)

## [1] 2000 39

The file *fileHH* contains 2,000 households and 39 variables. We are now ready to create our *sdcMicro* object with the corresponding variables we selected in Example 9.7. For our case study, we will create an *sdcMicro* object called *sdcHH* based on the data in *fileHH*, which we will use for steps 6a – 10a (see Example 9.9). **NOTE: When the sdcMicro object is created, the sdcMicro package automatically calculates and stores the risk measures for the data.** This leads us to Step 6a.

Example 9.9: Creating a sdcMicro object for the household variables

# Create initial SDC object for household level variables sdcHH <- createSdcObj(dat = fileHH, keyVars = selectedKeyVarsHH, pramVars = pramVarsHH,

weightVar = weightVarHH, numVars = numVarsHH)

numHH <- length(fileHH[,1]) # number of households</pre>

Step 6a: Assessing disclosure risk (household level)

As a first measure, we evaluate the number of households violating k-anonymity at the levels 2, 3 and 5.

Table 9.4 shows the number of violating households as well as the percentage of the total number of households. Example 9.10 illustrates how to find these values with *sdcMicro*. The print() function in *sdcMicro* shows only the values for thresholds 2 and 3. Values for other thresholds can be calculated manually by summing up the frequencies smaller than the k-anonymity threshold, as shown in Example 9.10.

Table 9.4: Number and proportion of households violating k-anonymity

k-anonymity level	Number of HH violating	Percentage of total number of HH
2	103	5.15 %
3	229	11.45 %
5	489	24.45 %

Example 9.10: Showing number of households violating k-anonymity for levels 2,3 and 5

# Number of observations violating k-anonymity (thresholds 2 and 3) print(sdcHH)

```
## Infos on 2/3-Anonymity:
##
## Number of observations violating
## - 2-anonymity: 103
## - 3-anonymity: 229
##
## Percentage of observations violating
## - 2-anonymity: 5.150 %
## - 3-anonymity: 11.450 %
```

# Calculate sample frequencies and count number of obs. violating k(5) - anonymity kAnon5 <- sum(sdcHH@risk\$individual[,2] < 5)

kAnon5

## [1] 489

# As percentage of total kAnon5 / numHH

## [1] 0.2445

It is often useful to view the values for the household(s) that violate k-anonymity. This might help clarify which variables cause the uniqueness of these households; this can then be used later when choosing appropriate SDC methods. Example 9.11 shows how to assess the values of the households violating 3 and 5-anonymity. It seems that among the categorical key variables, the variable HHSIZE is responsible for many of the unique combinations and the origin of much of the risk. Having determined this, we can flag HHSIZE as a possible first variable to treat to obtain the required risk level. In practice, with a variable like HHSIZE, this will likely involve removing large households from the dataset to be released. As explained in Section 5.5, recoding and local suppression are no valid options for the

variable HHSIZE. The frequencies of household size in Table 9.7 on page 132 show that there are few households with more than 13 household members. This makes these households easily identifiable based on the number of household members and at high risk of re-identification, also in the context of the nosy neighbor scenario.

Example 9.11: Showing households that violate k-anonymity

```
# Show values of key variable of records that violate k-anonymity fileHH[sdcHH@risk$individual[,2] < 3, selectedKeyVarsHH] # for 3-anonymity
```

```
fileHH[sdcHH@risk$individual[,2] < 5, selectedKeyVarsHH] # for 5-anonymity
```

We also assess the disclosure risk of the categorical variables with the individual and global risk measures as described in Sections 4.5 and 4.8. In *fileHH* every entry represents a household. Therefore, we use the individual non-hierarchical risk here, where the individual refers in this case to a household. *fileHH* contains only households and has no hierarchical structure. In Step 6b, we evaluate the hierarchical risk in *file*, the dataset containing both households and individuals. The individual and global risk measures automatically take into consideration the household weights, which we defined in Example 9.7. In our file, the global risk measure calculated using the chosen key variables is 0.05%. This percentage is extremely low and corresponds to 1.03 expected re-identifications. The results are also shown in Example 9.12. This low figure can be explained by the relatively small sample size of 0.25% of the total population. Furthermore, one should keep in mind that this risk measure is based only on the categorical quasi-identifiers at the household level. Example 9.12 illustrates how to print the global risk measure.

Example 9.12: Printing global risk measures

```
print(sdcHH, "risk")
```

```
## Risk measures:
##
## Number of observations with higher risk than the main part of the data: 0
## Expected number of re-identifications: 1.03 (0.05 %)
```

The global risk measure does not provide information about the spread of the individual risk measures. There might be a few households with relatively high risk, while the global (average) risk is low. It is therefore useful as a next step to inspect the observations with relatively high risk. The highest risk is 5.5% and only 14 households have risk larger than 1%. Example 9.13 shows how to display those households with risk over a certain threshold. Here the threshold is 0.01 (1%).

Example 9.13: Observations with individual risk higher than 1%

```
# Observations with risk above certain threshold (0.01) fileHH[sdcHH@risk$individual[, "risk"] > 0.01,]
```

Since the selected key variables at the household level are both categorical and numerical, the individual and global risk measures based on frequency counts do not completely reflect the disclosure risk of the entire dataset. Both categorical and continuous key variables are important for the data users, thus options like recoding the continuous variables (e.g., by creating quantiles of income and expenditure variables) to make all of them categorical will likely not satisfy the data user's needs. We therefore avoid recoding continuous variables and assess the disclosure risk of the categorical and continuous variables separately. This approach can be partly justified by the fact that any potential matching to external data sources for the continuous and categorical variables are available from different external data sources and as such will not be used simultaneously for matching.

### **Continuous variables**

To measure the risk of the continuous variables, we use an interval measure, which measures the number of anonymized values that are too close to their original values. See Section 4.7.2 for more information on interval-based risk measures for continuous variables. This measure is an ex-post measure, meaning that the risk can be evaluated only after anonymization and measures whether the perturbation is sufficiently large. Since it is an ex-post measure, we can evaluate it only in Step 9a after the variables have been treated. Evaluating this measure based on the original data would lead to a risk of 100%; all values would be too close to the original values since they would coincide with the original values, no matter how small the chosen intervals would be.

We also look at the distribution of LANDSIZEHA. In the variable LANDSIZEHA high values are rare and can lead to re-identification. An example is a large landowner in a specific region. To evaluate the distribution of the variable LANDSIZEHA, we look at the percentiles. Every percentile represents approximately 20 households. In addition, we look at the values of the largest 50 plots. Example 9.14 shows how to use R to display the quantiles and the largest landplots. Table 9.5 shows the  $90^{th} - 100^{th}$  percentiles and Table 9.6 displays the largest 50 values for LANDSIZEHA. Based on these values, we conclude that values of LANDSIZEHA over 40 make the household very identifiable. These large households and households with large land plots need extra protection, as discussed in Step 8a.

Example 9.14 Percentiles of LANDSIZE and listing the sizes of the largest 50 plots

# 1st - 100th percentiles of land size quantile(fileHH\$LANDSIZEHA, probs = (1:100)/100, na.rm= TRUE)

# Values of landsize for largest 50 plots tail(sort(fileHH\$LANDSIZEHA), n = 50)

Table 9.5: Percentiles 90-100 of the variable LANDSIZE

Percentile	90	91	92	93	94	95
Value	6.00	8.00	8.09	10.12	10.12	10.12
Percentile	96	97	98	99	100	
Value	12.14	20.23	33.83	121.41	1,214.08	

Table 9.6: 50 largest values of the variable LANDSIZE

12.14	15.00	15.37	15.78	16.19	20.00	20.23	20.23	20.23	20.23
20.23	20.23	20.23	20.23	20.23	20.23	20.23	20.23	20.23	20.23
20.23	20.23	20.50	30.35	32.38	40.47	40.47	40.47	40.47	40.47
40.47	40.47	80.93	80.93	80.93	80.93	121.41	121.41	161.88	161.88
161.88	182.11	246.86	263.05	283.29	404.69	404.69	607.04	809.39	1214.08

### **Step 7a: Assessing utility measures (household level)**

The utility of the data does not only depend on the household level variables, but on the combination of household-level and individual-level variables. Therefore, it is not useful to evaluate all the utility measures selected in Step 5 at this stage, i.e., before anonymizing the individual level variables. We restrict the initial measurement of utility to those measures that are solely based on the household variables. In our dataset, these are the measures related to income and expenditure and their distributions. The results are presented in Step 10a, together with the results after anonymization, which allow direct comparison. If after the next anonymization step it appears that the data utility has been significantly decreased by the suppression of some household level variables, we can return to this step.

Step 8a: Choice and application of SDC methods (household variables)

This step is divided into the anonymization of the variable HHSIZE, as this is a special case, the anonymization of the other selected categorical quasi-identifiers and the anonymization of the selected continuous quasi-identifiers.

#### Variable HHSIZE

The variable HHSIZE poses a problem for the anonymization of the file, since suppressing it will not anonymize this variable: a simple headcount based on the household ID would allow the reconstruction of this variable. Table 9.7 shows the absolute frequencies of HHSIZE. The number of households for each size larger than 13 is 6 or fewer and can be considered outliers with a higher risk of re-identification, as discussed in Step 6a. One way to deal with this is to remove all households of size 14 or larger from the dataset<sup>1</sup>. Removing 29 households of size 14 or larger reduces the number of 2-anonymity violations by 18, of 3-anonymity violations by 26 and of 5-anonymity violations by 29. This means that all removed households violated 5-anonymity due to the value of the variable HHSIZE and many of them 2- or 3-anonymity. In addition, the average individual risk amongst the 29 households is 0.15%, which is almost three times higher than the average individual risk of all households. The impact on the global risk measure of removing these 29 households is, however, limited, due to the relatively small number of removed households in comparison to the total number of 2,000 households. Removing the households is primarily to protect these specific households, not to reduce the global risk. NOTE: When using sdcMicro and manually removing households, the sdcMicro object should be recreated based on the new, manually edited dataset. Changes, such as removing records, cannot be done in the sdcMicro object. Example 9.15 illustrates the way to remove households and recreate the sdcMicro object.

Table 9.7: Frequencies of variable HHSIZE (household size)

HHSIZE	1	2	3	4	5	6	7	8	9	10	11	12
Frequency	152	194	238	295	276	252	214	134	84	66	34	21
HHSIZE	13	14	15	16	17	18	19	20	21	22	33	
Frequency	11	6	6	5	4	2	1	2	1	1	1	

Example 9.15: Removing households with large (rare) household sizes

```
# Tabulation of variable HHSIZE table(sdcHH@manipKeyVars$HHSIZE)
```

# Remove large households (14 or more household members) from file and fileHH file <- file[!file[,'HHSIZE'] >= 14,]

fileHHnew <- fileHH[!fileHH[,'HHSIZE'] >= 14,]

# Create new sdcMicro object based on the file without the removed households sdcHH <- createSdcObj(dat=fileHHnew, keyVars=selectedKeyVarsHH, pramVars=pramVarsHH, weightVar=weightVarHH, numVars = numVarsHH)

## **Categorical variables**

We are now ready to move on to the choice of SDC methods for the categorical variables on the household level in our dataset. As noted in our discussion of the methods, applying perturbative methods and local suppression may lead to large loss of utility. The common approach is to apply recoding to the largest possible extent as a first approach, to reach a prescribed level of risk and reduce the number of suppressions needed. Only after that should methods such as

<sup>&</sup>lt;sup>1</sup> Other methods and guidance on treating datasets where household size is a quasi-identifier are discussed in Section 5.5.

local suppression be applied. If this approach does not already achieve the desired result, we can consider perturbative methods.

Since the file is to be released as a SUF, we can keep a higher level of detail in the data. The selected categorical key variables at the household level are not suitable for recoding at this point. Due to the relatively low risk of reidentification based on the five selected categorical household level variables, it is possible in this case to use an option like local suppression to achieve our desired level of risk. Applying local suppression when initial risk is relatively low will likely only lead to suppression of few observations and thus limit the loss of utility. If, however, the data had been measured to have a relatively high risk, then applying local suppression without previous recoding would likely result in a large number of suppressions and greater information loss. Efforts such a recoding should be taken first before suppressing in cases where risk is initially measured as high. Recoding will reduce risk with little information loss and thus the number of suppressions, if local suppression is applied as an additional step. We apply local suppression to reach 2-anonymity. The choice of the low level of two is based on the overall low re-identification risk due to the high sample weights and the release as SUF. High sample weights mean, ceteris paribus, a low level of re-identification risk. Achieving 2-anonymity is the same as removing sample uniques. This leads to 42 suppressions in the variable HHSIZE and 4 suppressions in the variable REGION. As explained earlier, suppression of the value of the variable HHSIZE does not lead to actual suppression of this information. Therefore, we redo the local suppression, but this time we tell *sdcMicro* to, if possible, not suppress HHSIZE but one of the other variables.

In *sdcMicro* it is possible to tell the algorithm which variables are important and less important for making small changes (see also Section 5.2.2). To prevent HHSIZE being suppressed, we set the importance of HHSIZE in the importance vectors to the highest (i.e., 1). Example 9.16 shows how to apply local suppression and put importance on the variable HHSIZE. The variable REGION is the type of variable that should not have any suppressions either. We also set the importance of REGION to 2 and the importance of RURURB to 3. This leads to an order of the variables to be considered for suppression by the algorithm. Instead of 42 suppressions in the variable HHSIZE, this leads one suppressed value in the variable HHSIZE, and to 6, 1, 48 and 16 suppressions respectively for the variables URBRUR, REGION, OWNAGLAND and RELIG (which we set as less important). The importance is clearly reflected in the number of suppression. The total number of suppressions is higher than without importance vector (71 vs. 46), but 2-anonymity is achieved in the dataset with fewer suppressions in the variables HHSIZE and REGION. We remove the one household with the suppressed value of HHSIZE (13) to protect this household. NOTE: In Example 9.16 we use the undolast() function in sdcMicro to go one step back after we had first applied local suppression with no importance vector. The undolast() function restores the *sdcMicro* object back to the previous state (i.e., before we applied local suppression), which allows us to rerun the same command, but this time with an importance vector set. The undolast() function can only be used to go one step back.

Example 9.16: Local suppression with and without importance vector

```
# Local suppression
sdcHH <- localSuppression(sdcHH, k=2, importance = NULL) # no importance vector

print(sdcHH, "ls")

## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## URBRUR | 0 | 0.000
## REGION | 4 | 0.203
## HHSIZE | 37 | 1.877
## OWNAGLAND | 0 | 0.000
## RELIG | 0 | 0.000</pre>
```

```
sdcHH <- undolast(sdcHH)
sdcHH <- localSuppression(sdcHH, k=2, importance = c(3, 2, 1, 5, 5)) # importance on HHSIZE (1), REGION (2)
and URBRUR (3)
print(sdcHH, "ls")

## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## URBRUR | 6 | 0.304
## REGION | 1 | 0.051
## HHSIZE | 1 | 0.051
## OWNAGLAND | 43 | 2.182
## RELIG | 16 | 0.812</pre>
```

The variables ROOF, TOILET, WATER, ELECTCON, FUELCOOK, OWNMOTORCYCLE, CAR, TV and LIVE-STOCK are not sensitive variables and were not selected as quasi-identifiers because we assumed that there are no external data sources containing this information that could be used for matching. Values can be easily observed or be known to neighbors, however, and therefore are important, together with other variables, for the spontaneous recognition scenario and nosy neighbor scenario. To anonymize these variables, we want to introduce a low level of uncertainty in them. Therefore, we decide to use invariant PRAM for the variables ROOF, TOILET, WATER, ELECTCON, FUELCOOK, OWNMOTORCYCLE, CAR, TV and LIVESTOCK, where we treat LIVESTOCK as a semi-continuous variable due to the low number of different values. Section 5.3.1 provides more information on the PRAM method and its implementation in *sdcMicro*. Example 9.17 illustrates how to apply PRAM. We choose the parameter pd, the lower bound for the probability that a value is not changed, to be relatively high at 0.8. We can choose a high value, because the variables themselves are not sensitive and we only want to introduce a low level of changes to minimize the utility loss. Because the distribution of many of the variables chosen for PRAM depends on the REGION, we decide to use the variable REGION as a strata variable. In this way the transition matrix is computed for each region separately. Because PRAM is a probabilistic method, we set a seed for the random number generator before applying PRAM to ensure reproducibility of the results. NOTE: In practice, it is not advisable to set a seed of 12345, but rather a longer more complex and less easy to guess sequence. The seed should not be released, since it allows for reconstructing the original values if combined with the transition matrix. The transition matrix can be released: this allows for consistent statistical inference by correcting the statistical methods used if the researcher has knowledge about the PRAM method (at this point *sdcMicro* does not allow the retrieval of the transition matrix).

Example 9.17: Applying PRAM

```
## CAR != CAR_pram : 172 (8.73%)
## TV != TV_pram : 137 (6.95%)
## LIVESTOCK != LIVESTOCK_pram : 149 (7.56%)
```

PRAM has changed values within the variables according to the invariant transition matrices. Since we used the invariant PRAM method (see Section 5.3.1), the absolute univariate frequencies remain unchanged. This is not the case for the multivariate frequencies. In Step 10a we compare the changes in the multivariate frequencies for the PRAMmed variables.

#### Continuous variables

We have selected income and expenditures variables and the variable LANDSIZEHA as numerical quasi-identifiers, as discussed in Step 4. In Step 5 we identified variables having high interest for the users of our data: many users use the data for measuring inequality and expenditure patterns.

Based on the risk evaluation in Step 6a, we decide to anonymize the variable LANDSIZEHA by top coding at the value 40 (cf. Table 9.5 and Table 9.6) and round values smaller than 1 to one digit, and values larger than 1 to zero digits. Rounding the values prevents exact matching with the available cadastral register. Furthermore, we group the values between 5 and 40 in the groups 5 - 19 and 20 - 39. After these steps, no household has a unique plot size and the number of households in the sample with the same plot size was increased to at least 7. This is shown by the tabulation of the variable LANDSIZEHA after manipulation in the last line of Example 9.18. In addition, all outliers have been removed by top coding the values. This has reduced the risk of spontaneous recognition as discussed in Step 6. How to recode values in R is introduced in Section 5.2.1 and, for this particular case, shown in Example 9.18.

Example 9.18: Anonymizing the variable LANDSIZEHA

```
# Rounding values of LANDSIZEHA to 1 digit for plots smaller than 1 and to 0 digits for plots larger than 1
sdcHH@manipNumVars$LANDSIZEHA[sdcHH@manipNumVars$LANDSIZEHA <= 1 &
!is.na(sdcHH@manipNumVars$LANDSIZEHA)] <-
round(sdcHH@manipNumVars$LANDSIZEHA[sdcHH@manipNumVars$LANDSIZEHA <= 1 &
!is.na(sdcHH@manipNumVars$LANDSIZEHA)], digits = 1)
sdcHH@manipNumVars\$LANDSIZEHA[sdcHH@manipNumVars\$LANDSIZEHA]\\
!is.na(sdcHH@manipNumVars$LANDSIZEHA)] <- round(sdcHH@manipNumVars$LANDSIZEHA[sdcHH@manipNumVars$LANDSIZEHA]
> 1 & !is.na(sdcHH@manipNumVars$LANDSIZEHA)], digits = 0)
# Grouping values of LANDSIZEHA into intervals 5-19, 20-39
sdcHH@manipNumVars$LANDSIZEHA[sdcHH@manipNumVars$LANDSIZEHA >= 5 &
sdcHH@manipNumVars$LANDSIZEHA < 20 & !is.na(sdcHH@manipNumVars$LANDSIZEHA)] <- 13
sdcHH@manipNumVars$LANDSIZEHA[sdcHH@manipNumVars$LANDSIZEHA
                                                                                      &
                                                                                            sd-
cHH@manipNumVars$LANDSIZEHA < 40 & !is.na(sdcHH@manipNumVars$LANDSIZEHA)] <- 30
# Topcoding values of LANDSIZEHA larger than 40 (also recomputes risk after manual changes)
sdcHH <- topBotCoding(sdcHH, value = 40, replacement = 40, kind = 'top', column = 'LANDSIZEHA')
# Results for LANDSIZEHA
table(sdcHH@manipNumVars$LANDSIZEHA)
```

```
## 0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1 2 3 4 13 30 40 ## 188 109 55 30 24 65 22 7 31 16 154 258 53 60 113 18 25
```

For the expenditure and income variables we compared, **based on the actual case study data**, several methods. As mentioned earlier, the main use of the data is to compute inequality measures, such as the Gini coefficient. Recoding these variables into percentiles creates difficulties computing these measures or changes these measures to a large extent and is hence not a suitable method. Often, income and expenditure variables that are released in a SUF are anonymized by top-coding. This protects the outliers, which are the values that are the most at risk. Top-coding, however, destroys the inequality information in the data, by removing high (and low) incomes. Therefore, we decide to use noise addition. To take into account the higher risk of outliers, we add a higher level of noise to those.

Adding noise can lead to a transformation of the shape of the distribution. Depending on the magnitude of the noise (see Section 5.3.3 for the definition of the magnitude of noise), the values of income can also become negative. One way to solve this would be to cut off the values below zero and set them to zero. This would, however, destroy the properties conserved by noise addition (amongst others the value of the expected mean, see also Section 5.3.3) and we chose to keep the negative values.

As mentioned before, the aggregates of income and expenditures are the sums of the components. Adding noise to each of the components might lead to violation of this condition. Therefore, one solution is to add noise to the aggregates and remove the components. We prefer to keep the components in the data and apply noise addition to each component separately. This allows to apply a lower level of noise than when applying noise only to the aggregates. A noise level of 0.01 seems to be sufficient with extra noise of 0.05 added to the outliers. The outliers are defined by a robust Mahalanobis distance (see Section 5.3.3). After adding noise to the components, we recomputed the aggregates as the sum of the perturbed components. **NOTE: This result is only based on the actual case study dataset and is not necessarily true for other datasets.** The noise addition is shown in Example 9.19. Before applying probabilistic methods such as noise addition, we set a seed for the random number generator. This allows us to reproduce the results.

```
Example 9.19: Anonymizing continuous variables
```

# Add noise to income and expenditure variables by category

```
# Anonymize components
compExp <- c("TFOODEXP", "TALCHEXP", "TCLTHEXP", "THOUSEXP", "TFURNEXP", "THLTHEXP",
"TTRANSEXP", "TCOMMEXP", "TRECEXP", "TEDUEXP", "TRESHOTEXP", "TMISCEXP")
set.seed(123)

# Add noise to expenditure variables
sdcHH <- addNoise(noise = 0.01, obj = sdcHH, variables = compExp, method = "additive")
# Add noise to outliers
sdcHH <- addNoise(noise = 0.05, obj = sdcHH, variables = compExp, method = "outdect")
# Sum over expenditure categories to obtain consistent totals
sdcHH@manipNumVars[, TANHHEXP'] <- rowSums(sdcHH@manipNumVars[, compExp])

compInc <- c('INCRMT', 'INCWAGE', 'INCFARMBSN', 'INCNFARMBSN', 'INCRENT', 'INCRENT', 'INCFIN',
'INCPENSN', 'INCOTHER')
# Add noise to income variables
```

sdcHH <- addNoise(noise = 0.01, obj = sdcHH, variables = compInc, method = "additive")

# Add noise to outliers

sdcHH <- addNoise(noise = 0.05, obj = sdcHH, variables = compInc, method = "outdect")

# Sum over income categories to obtain consistent totals

sdcHH@manipNumVars[,'INCTOTGROSSHH'] <- rowSums(sdcHH@manipNumVars[,compInc])

# recalculate risks after manually changing values in sdcMicro object

calcRisks(sdcHH)

# Step 9a: Re-measure risk

For the categorical variables, we conclude that we have achieved 2-anonymity in the data with local suppression. Only 104 households, or about 5% of the total number, violate 3-anonymity. Table 9.8 gives an overview of these risk measures. The global risk is reduced to 0.02% (expected number of re-identifications 0.36), which is extremely low. Therefore, we conclude that based on the categorical variables, the data has been sufficiently anonymized. No household has a risk of re-identification higher than 0.01 (1%). By removing households with rare values (or outliers) of the variable HHSIZE, we have reduced the risk of spontaneous recognition of these households. This reasoning can also be applied to the result of the risk of recoding the variable LANDSIZEHA and PRAMming the variables identified to be important in the nosy neighbor scenario. An intruder cannot know with certainty whether a household that he recognizes in the data is the correct household, due to the noise.

Table 9.8: Number and proportion of households violating k-anonymity after anonymization

k-anonymity	Number HH violating	Percentage		
2	0	0 %		
3	104	5.28 %		
5	374	18.70 %		

These measures refer only to the categorical variables. To evaluate the risk of the continuous variables we could use an interval measure or closest neighbor algorithm. These risk measures are discussed in Section 4.7. We chose to use an interval measure, since exact value matching is not our largest concern based on the assumed scenarios and external data sources. Instead, datasets with similar values but not the exact same values could be used for matching. Here the main concern is that the values are sufficiently far from the original values, which is measured with an interval measure.

Example 9.20 shows how to evaluate the interval measure for each of the expenditure variables, which are contained in the vector  $compExp^2$ . The different values of the parameter k in the function dRisk() define the size of the interval around the original value, as explained in Section 4.7.2. The larger k, the larger the intervals, the higher the probability that a perturbed value is in the interval around the original value and the higher the risk measure. The result is satisfactory with relatively small intervals (k = 0.01), but not when increasing the size of the intervals. In our case, k = 0.01 is sufficiently large, since we are looking at the components, not the aggregates. We have to pay special attention to the outliers. Here the value 0.01 for k is too small to assume that they are protected when outside this small interval. It would be necessary to check outliers and their perturbed values and there might be a need for a higher level of perturbation for outliers. We conclude that, from a risk perspective and based on the interval measure, the chosen levels of noise are acceptable. In the next step, we will look at the impact on the data utility of the noise addition.

Example 9.20: Measuring risk of re-identification of continuous variables

**dRisk**(sdcHH@origData[,compExp], xm = sdcHH@manipNumVars[,compExp], k = 0.01)

[1] 0.0608828

dRisk(sdcHH@origData[,compExp], xm = sdcHH@manipNumVars[,compExp], k = 0.02)

<sup>&</sup>lt;sup>2</sup> For illustrative purposes, we only show this evaluation for the expenditure variables. It can be easily copied for the income variables. The results are similar.

#### [1] 0.9025875

**dRisk**(sdcHH@origData[,compExp], xm = sdcHH@manipNumVars[,compExp], k = 0.05)

[1] 1

# Step 10a: Re-measure utility

None of the variables has been recoded and the original level of detail in the data is kept, except for the variable LANDSIZEHA. As described in Step 8a, local suppression has only removed a few values in the other variables, which has not greatly reduced the validity of the data.

The univariate frequency distributions of the variables ROOF, TOILET, WATER, ELECTCON, FUELCOOK, OWN-MOTORCYCLE, CAR, TV and LIVESTOCK did not, by definition of the invariant PRAM method (see Section 5.3.1), change to a large extent. The tabulations are presented in Table 9.9 (the values 1 – 9 and NA in the first row are the values of the variables and .m after the variable name refers to the values after anonymization). **NOTE: Although the frequencies are almost the same, this does not mean that the values of particular households did not change.** Values have been swapped between households. This becomes apparent when looking at the multivariate frequencies of the WATER with the variable URBRUR in Table 9.10. The multivariate frequencies of the PRAMmed with the variable URBRUR could be of interest for users, but these are not preserved. Since we applied PRAM within the regions, the multivariate frequencies of the PRAMmed variables with REGION are preserved.

Table 9.9: Univariate frequencies of the PRAMmed variable before and after anonymization

	0	1	2	3	4	5	6	7	8	9	NA
ROOF		27	1	914	307	711				10	1
ROOF.m		25	1	907	319	712				6	1
TOILET		76	594	817	481					3	
TOILET.m		71	597	816	483					4	
WATER		128	323	304	383	562	197	18	21	35	
WATER.m		134	319	308	378	573	188	16	21	34	
ELECTCON	768	216	8	2							977
ELECTCON.m	761	218	8	3							981
FUELCOOK		1289	21	376	55	36				139	55
FUELCOOK.m		1284	22	383	50	39				143	50
OWNMOTORCYCLE	1883	86									2
OWNMOTORCYCLE.m	1882	86									2
CAR	963	31									977
CAR.m	966	25									
TV	1216	264									491
TV.m	1203	272									496

Table 9.10: Multivariate frequencies of the variables WATER with RURURB before and after anonymization

	1	2	3	4	5	6	7	8	9
WATER/URB	11	49	270	306	432	183	12	15	21
WATER/RUR	114	274	32	76	130	14	6	6	14
WATER/URB.m	79	220	203	229	402	125	10	12	19
WATER/RUR.m	54	98	105	147	169	63	6	9	15

For conciseness, we restrict ourselves to the analysis of the expenditure variables. The analysis of the income variables can be done in the same way and leads to similar results.

We look at the effect of anonymization on some indicators as discussed in Step 5. Table 9.11 presents the point

estimates and bootstrapped confidence interval of the GINI coefficient<sup>3</sup> for the sum of the expenditure components. The calculation of the GINI coefficient and the confidence interval are based on the positive expenditure values. We observe very small changes in the Gini coefficient, that are statistically negligible. We use a visualization to illustrate the impact on utility of the anonymization. Visualizations are discussed in Section 6.4 and the specific R code for this case study is available in the R script. The change in the inequality measures is illustrated in Figure 9.1, which shows the Lorenz curves based on the positive expenditure values before and after anonymization.

Table 9.11: GINI point estimates and bootstrapped confidence intervals for sum of expenditure components

	before	after
Point estimate	0.510	0.508
Left bound of CI	0.476	0.476
Right bound of CI	0.539	0.538

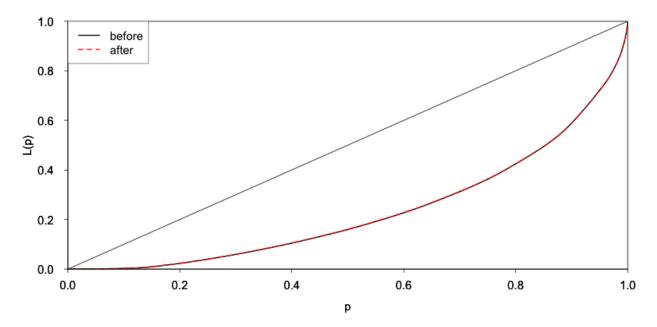


Figure 9.1: Lorenz curve based on positive total expenditures values

We compare the mean monthly expenditures (MME) and mean monthly income (MMI) for rural, urban and total population. The results are shown in Table 9.12. We observe that the chosen levels of noise add only small distortions to the MME and slightly larger changes to the MMI.

Table 9.12: Mean monthly expenditure and mean monthly income per capita by rural/urban

	before	after
MME rural	400.5	398.5
MME urban	457.3	459.9
MME total	412.6	412.6
MMI rural	397.1	402.2
MMI urban	747.6	767.8
MMI total	472.1	478.5

Table 9.13 shows the share of each of the components of the expenditure variables before and after anonymization.

<sup>&</sup>lt;sup>3</sup> To compute the GINI coefficient, bootstrap to construct the confidence intervals and plot the Lorenz curve we used the *R* packages *laeken*, *reldist*, *bootstrap* and *ineq*.

Table 9.13 Shares of expenditures components

	TFOOD EXP	TALCH EXP	TCLTH EXP	THOUS EXP	TFURN EXP	THLTH EXP
be- fore	0.58	0.01	0.03	0.09	0.02	0.03
after	0.59	0.01	0.03	0.09	0.02	0.03
	TTRAN	TCOMM	TRECE XP	TEDUE XP	TRESH OT-	TMISC
	SEXP	EXP			EXP	EXP
be-	0.04	0.02	0.00	0.08	0.03	0.05
fore						
after	0.04	0.02	0.00	0.08	0.03	0.05

Anonymization for the creation of a SUF will inevitably lead to some degree of utility loss. It is important to describe this loss in the external report, so that users are aware of the changes in the data. This is described in Step 11 and presented in Appendix A. Appendix A also shows summary statistics and tabulations of the household level variables before and after anonymization.

# Merging the household- and individual-level variables

The next step is to merge the treated household variables with the untreated individual variables for the anonymization of the individual level variables. Example 9.21 shows the steps to merge these files. This also includes the selection of variables used in the anonymization of the individual-level variables. We create the *sdcMicro* object for the anonymization of the individual variables in the same way as for the household variable in Example 9.9. Subsequently, we repeat Steps 6-10 for the individual-level variables.

Example 9.21: Merging the files with household and individual-level variables and creating an *sdcMicro* object for the anonymization of the individual-level variables

### Select variables (individual level)

```
# Key variables (individual level)
selectedKeyVarsIND = c('GENDER', 'REL', 'MARITAL', 'AGEYRS', 'EDUCY', 'ATSCHOOL', 'INDUSTRY1')
# list of selected key variables
# Sample weight (WGTHH, individual weight)
selectedWeightVarIND = c('WGTHH')
# Household ID
selectedHouseholdID = c('IDH')
# No strata
# Recombining anonymized HH datasets and individual level variables
indVars <- c("IDH", "IDP", selectedKeyVarsIND, "WGTHH") # HID and all non HH variables
fileInd <- file[indVars] # subset of file without HHVars</pre>
```

```
HHmanip <- extractManipData(sdcHH) # manipulated variables HH
```

HHmanip <- HHmanip[HHmanip[,'IDH'] != 1782,]

fileCombined <- merge(HHmanip, fileInd, by.x=c('IDH'))

fileCombined <- fileCombined[order(fileCombined[,'IDH'], fileCombined[,'IDP']),]</pre>

dim(fileCombined)

# SDC objects with all variables and treated HH vars for anonymization of individual level variables sdcCombined <- createSdcObj(dat = fileCombined, keyVars = selectedKeyVarsIND, weightVar = selectedWeightVarIND, hhId = selectedHouseholdID)

#### Step 6b: Assessing disclosure risk (individual level)

All key variables at the individual level are categorical. Therefore, we can use k-anonymity and the individual and global risk measures (see Sections 4.5 and 4.8). The hierarchical risk is now of interest, given the household structure in the dataset *fileCombined*, which includes both household- and individual-level variables. The number of individuals (absolute and relative) that violate k-anonymity at the levels 2, 3 and 5 are shown in Table 9.14. **NOTE: k-anonymity does not consider the household structure and therefore underestimates the risk. Therefore, we are more interested in the individual and global hierarchical risk measures.** 

Table 9.14: k-anonymity violations

k-anonymity Number HH violating		Percentage
2	998	9.91%
3	1,384	13.75%
5	2,194	21.79%

The global risk measures can be found using R as illustrated in Example 9.22. The global risk is 0.24%, which corresponds to 24 expected re-identifications. Accounting for the hierarchical structure, this rises to 1.26%, or 127 expected re-identifications. The global risk measures are low compared to the number of k-anonymity violators due to the low sampling weights. The high number of k-anonymity violators is mainly due to the very detailed age variable. The risk measures are based only on the individual level variables, since we assume that the individual and household level variables are be used simultaneously by an intruder. If we would consider an intruder scenario where these variables are used simultaneously by an intruder to re-identify individuals, the household level variables should also be taken into account here. This would results in a high number of key variables.

Example 9.22: Global risk of the individual-level variables

print(sdcCombined, 'risk')

```
## Risk measures:
##
## Number of observations with higher risk than the main part of the data: 0
## Expected number of re-identifications: 23.98 (0.24 %)
##
## Information on hierarchical risk:
## Expected number of re-identifications: 127.12 (1.26 %)
```

#### Step 7b: Assessing utility (individual level)

We evaluate the utility measures selected in Step 5 besides some general utility measures. The values computed from the raw data are presented in step 10b to allow for direct comparison with the values computed from the anonymized data.

# **Step 8b: Choice and application of SDC methods (individual level)**

We use the same approach for the anonymization of the individual-level categorical key variables as for the household level categorical variables described earlier: first use global recoding to limit the necessary number of suppressions, then apply local suppressions and finally, if necessary, use of perturbative methods.

The variable AGEYRS (i.e., age in years) has many different values (age in months for children 0-1 years and age in years for individuals over 1 year). This level of detail leads to a high level of re-identification risk, given external datasets with exact age as well as knowledge of the exact age of close relatives. We have to reduce the level of detail in the age variables by recoding the age values (see Section 5.2.1 on recoding). First, we recode the values from 15 to 65 in ten-year intervals. Since some indicators related to education are computed from the survey dataset, our first approach is not to recode the age range 0-15 years. For children under the age of 1 year, we reduce the level of detail and recode these to 0 years. These recodes are shown in Example 9.23. We also top-code age at the age of 65 years. This protects individuals with high (rare) age values.

Example 9.23: Recoding age in 10-year intervals in the range 15 – 65 and top code age over 65 years

# Recoding age and top coding age (top code 65), below that 10 year age groups, children aged under 1 are recoded 0 (previously in months)

```
sdcCombined@manipKeyVars$AGEYRS[sdcCombined@manipKeyVars$AGEYRS >= 0 & sdcCom-
bined@manipKeyVars$AGEYRS < 1] <- 0</pre>
```

sdcCombined@manipKeyVars\$AGEYRS[sdcCombined@manipKeyVars\$AGEYRS >= 15 & sdcCombined@manipKeyVars\$AGEYRS < 25] <- 20</pre>

. .

sdcCombined@manipKeyVars\$AGEYRS[sdcCombined@manipKeyVars\$AGEYRS >= 55 & sdcCombined@manipKeyVars\$AGEYRS < 66] <- 60

#topBotCoding also recalculates risk based on manual recoding above sdcCombined <- topBotCoding(obj = sdcCombined, value = 65, replacement = 65, kind = 'top', column = 'AGEYRS')

```
table(sdcCombined@manipKeyVars$AGEYRS) # check results
```

```
## 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14
## 311 367 340 332 260 334 344 297 344 281 336 297 326 299 263
## 20 30 40 50 60 65
## 1847 1220 889 554 314 325
```

These recodes already reduce the risk to 531 individuals violating 3-anonymity. We could recode the values of age in the lower range according to the age categories users require (e.g., 8 - 11 for education). There are many different categories for different indicators, however, including education indicators. This would reduce the utility of the data for some users. Therefore, we decide to look first at the number of suppressions needed in local suppression after this limited recoding. If the number of suppressions is too high, we can go back and recode age in the range 1 - 14 years.

In Example 9.24 we demonstrate how one might experiment with local suppression to find the best option. We use local suppression to achieve 3-anonymity (see Section 5.2.2 on local suppression). On the first attempt, we do not specify any importance vector; this leads to many suppressions in the variable AGEYRS (see Table 9.15 below, first row), however. This is undesirable from a utility point of view. Therefore, we decide to specify an importance vector to prevent suppressions in the variable AGEYRS. Suppressing the variable GENDER is also undesirable from the utility point of view. The variable GENDER is a type of variable that should not have suppressions. We set GENDER as variable with the second highest importance. After specifying the importance vector to prevent suppressions of the age variable, there are no age suppressions (see Table 9.15, second row). The total number of suppressions in the other variables increased, however, from 253 to 323 because of the importance vector. This is to be expected because the algorithm without the importance vector minimizes the total number of suppressions by first suppressing values in variables with many categories – in this case, age and gender. Specifying an importance vector prevents reaching this optimality and hence leads to a higher total number of suppressions. There is a trade-off between which variables are

suppressed and the total number of suppressions. After specifying an importance vector, the variable REL has many suppressions (see Table 9.15, second row). We choose this second option.

Example 9.24: Experimenting with different options in local suppression

```
# Copy of sdcMicro object to later undo steps
sdcCopy <- sdcCombined
# Importance vectors for local suppression (depending on utility measures)
impVec1 <- NULL # for optimal suppression
impVec2 <- rep(length(selectedKeyVarsIND), length(selectedKeyVarsIND))
impVec2[match('AGEYRS', selectedKeyVarsIND)] <- 1 # AGEYRS
impVec2[match('GENDER', selectedKeyVarsIND)] <- 2 # GENDER
# Local suppression without importance vector
sdcCombined <- localSuppression(sdcCombined, k = 2, importance = impVec1)
# Number of suppressions per variable
print(sdcCombined, "ls")
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## GENDER | 0 | 0.000
## REL | 34 | 0.338
## MARITAL | 0 | 0.000
## AGEYRS | 195 | 1.937
## EDUCY | 0 | 0.000
## EDYRSCURRAT | 3 | 0.030
## ATSCHOOL | 0 | 0.000
## INDUSTRY1 | 21 | 0.209
# Number of suppressions per variable for each value of AGEYRS
table(sdcCopy@manipKeyVars$AGEYRS) - table(sdcCombined@manipKeyVars$AGEYRS)
## 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 20 30 40 50 60 65
## 0 0 0 0 0 0 2 0 2 1 0 1 4 1 5 25 53 37 36 15 13
# Undo local suppression
sdcCombined <- undolast(sdcCombined)</pre>
# Local suppression with importance vector on AGEYRS and GENDER
```

sdcCombined <- localSuppression(sdcCombined, k = 2, importance = impVec2)

# Number of suppressions per variable print(sdcCombined, "ls")

```
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## GENDER | 0 | 0.000
## REL | 323 | 3.208
## MARITAL | 0 | 0.000
## AGEYRS | 0 | 0.000
## EDUCY | 0 | 0.000
## EDYRSCURRAT | 0 | 0.000
## ATSCHOOL | 0 | 0.000
## INDUSTRY1 | 0 | 0.000
```

# Number of suppressions for each value of the variable AGEYRS table(sdcCopy@manipKeyVars\$AGEYRS) - table(sdcCombined@manipKeyVars\$AGEYRS)

Table 9.15: Number of suppressions by variable for different variations of local suppression

Loc al suppr essio n optio ns	GENDER	REL	MARITAL	AGEYRS	EDUCY	EDYRS CUF	ATS CHOOL RAT	IND US- TRY1
k = 2, no imp	0	34	0	195	0	3	0	21
k = 2, imp on AGEYRS	0	323	0	0	0	0	0	0

### Step 9b: Re-measure risk (individual level)

We re-evaluate the risk measures selected in Step 6b. Table 9.16 shows that local suppression, not surprisingly, has reduced the number of individuals violating 2-anonymity to 0.

Table 9.16: k-anonymity violations

k-anonymity	Number HH violating	Percentage
2	0	0.00 %
3	197	1.96 %
5	518	5.15 %

The hierarchical global risk was reduced to 0.11%, which corresponds to 11.3 expected re-identifications. The highest individual hierarchical re-identification risk is 1.21%. These risk levels would seem acceptable for a SUF.

**Step 10b: Re-measure utility (individual level)** 

We selected two utility measures for the individual variables: primary and secondary education enrollment, both also by gender. These two measures are sensitive to changes in the variables gender (GENDER), age (AGEYRS) and education (EDUCY and EDYRSATCURR), and therefore give a good overview of the impact of the anonymization. As shown in Table 9.17 the anonymization did not change the results. The results of the tabulations in Appendix A confirm these results.

Table 9.17: Net enrollment in primary and secondary education by gender

	Primary	education	on	Secondary education		
	Total Male Female		le Total Male F		Female	
Before	72.6%	74.2%	70.9%	42.0%	44.8%	39.1%
After	72.6%	74.2%	70.9%	42.0%	44.8%	39.1%

#### Step 11: Audit and reporting

In the audit step, we check whether the data allow for reproduction of published figures from the original dataset and relationships between variables and other data characteristics are preserved in the anonymization process. In short, we check whether the dataset is valid for analytical purposes. There are no figures available that were published from the dataset and need to be reproducible from the anonymized data.

In Step 2, we explored the data characteristics and relationships between variables. These data characteristics and relationships have been mainly preserved, since we took them into account when choosing the appropriate anonymization methods. The variables TANHHEXP and INCTOTGROSSHH are the sums of the individual components, because we added noise to the components and reconstructed the aggregates by summing over the components. Initially, the income variables were all positive. This characteristic has been violated, as a result of noise addition. Since values of the variable AGEYRS were not perturbed, but only recoded and suppressed, we did not introduce unlikely combinations, such as a 60-year-old individual enrolled in primary education. Also, by separating the anonymization process into two parts, one for household-level variables and one for individual-level variables, the values of variables measured at the household level agree for all members of each household.

Furthermore, we drafted two reports, internal and external, on the anonymization of the case study dataset. The internal report includes the methods used, the risk before and after anonymization as well as the reasons for the selected methods and their parameters. The external report focuses on the changes in the data and the loss in utility. Focus here should be on the number of suppressions as well as the perturbative methods (PRAM). This is described in the previous steps. **NOTE:** When creating a SUF, it is inevitable that there will be a loss of information and it is very important for the users to be aware of these changes and release them in a report that accompanies the data. Appendix A provides examples of an internal and external report of the anonymization process of this dataset. Depending on the users and readers of the reports, the content may differ. The code to this case study shows how to obtain the information for the reports. Some measures are also available in the standard reports generated with the report() function. This is shown in Example 9.25. The report() function will only use the data available in the *sdcMicro* object, which does not contain all households for sdcHH.

Example 9.25: Using the report() function for internal and external reports

```
# Create reports with sdcMicro report() function
report(sdcHH, internal = F) # external (brief) report
```

report(sdcHH, internal = T) # internal (extended) report

# Create reports with sdcMicro report() function
report(sdcCombined, internal = F) # external (brief) report

**report**(sdcCombined, internal = T) # internal (extended) report

#### **Step 12: Data release**

The final step is the release of the anonymized dataset together with the external report. Example 9.26 shows how to collect the data from the *sdcMicro* object with the extractManipData() function. Before releasing the file, we add an individual ID to the file (line number in household). We export the anonymized dataset in as *STATA* file. Section 7.2 presents functions for exporting files in other data formats.

Example 9.26: Exporting the anonymized dataset

# Anonymized dataset

# Household variables and individual variables

dataAnon <- **extractManipData**(sdcCombined, ignoreKeyVars = F, ignorePramVars = F, ignoreNumVars = F, ignoreStrataVar = F) # extracts all variables, not just the manipulated ones

# Create STATA file

write.dta(dataframe = dataAnon, file= 'Case1DataAnon.dta', convert.dates=TRUE)

# 12.2 Case study 2 - PUF

This case study is a continuation of case study 1 in Section 9.1. Case study 1 produces a SUF file. In this case study we use this SUF file to produce a PUF file of the same dataset, which can be freely distributed. The structure of the SUF and PUF releases will be the same. However, the PUF will contain fewer variables and less (detailed) information than the SUF. We refer to Section 9.1 for a description of the dataset. **NOTE**: It is also possible to directly produce a PUF from a dataset without first creating a SUF.

As in case study 1, we show how the creation of a PUF can be achieved using the open source and free *sdcMicro* package and *R*. A ready-to-run *R* script for this case study and the dataset are also available to reproduce the results and allow the user to adapt the code (see http://ihsn.org/home/projects/sdc-practice). Extracts of this code are presented in this section to illustrate several steps of the anonymization process. **NOTE: The choices of methods and parameters in this case study are based on this particular dataset and the results and choices might be different for other datasets.** 

This case study follows the steps of the SDC process outlined in Chapter 8.

#### **Step 1: Need for disclosure control**

The same reasoning as in case study 1 applies: the SUF dataset produced in case study 1 contains data on individuals and households and some variables are confidential and/or sensitive. The decisions made in case study 1 are based on the disclosure scenarios for a SUF release. The anonymization applied for the SUF does not provide sufficient protection for the release as PUF and the SUF file cannot be released as PUF without further treatment. Therefore, we have to repeat the SDC process with a different set of disclosure scenarios based on the characteristics of a PUF release (see Step 4). This leads to different risk measures, lower accepted risk levels and different SDC methods.

### Step 2: Data preparation and exploring data characteristics

In order to guarantee consistency between the released PUF and SUF files, which is required to prevent intruders from using the datasets together (SUF users have also access to the PUF file), we have to use the anonymized SUF file to create the PUF file (see also Section 8.3). In this way all information in the PUF file is also contained in the SUF, and the PUF does not provide additional information to an intruder with access to the SUF. We load the required packages to read the data (*foreign* package for *STATA* files) and load the SUF dataset into "file" as illustrated in Example 9.27. We also load the original data file (raw data) as "fileOrig". We need the raw data to undo perturbative methods used in case study 1 (see Step 8) and to compare data utility measures (see Step 5). To evaluate the utility loss in the PUF, we

have to compare the information in the anonymized PUF file with the information in the raw data. For an overview of the data characteristics and a description of the variables in both files, we refer to Step 2 of case study 1 in Section 9.1.

Example 9.27: Loading required packages and datasets

```
# Load required packages

library(foreign) # for read/write function for STATA

library(sdcMicro) # sdcMicro package

# Set working directory - set to the path on your machine

setwd("/Users/CaseStudy2")

# Specify file name of SUF file from case study 1
fname <- "CaseDataAnon.dta"

# Specify file name of original dataset (raw data)
fnameOrig <- "CaseA.dta"

# Read-in files
file <- read.dta(fname, convert.factors = TRUE) # SUF file case study 1 *fileOrig <- **read.dta*(fnameOrig, convert.factors = TRUE) # original data
```

We check the number of variables and number of observations of both files and the variable names of the SUF file, as shown in Example 9.28. The PUF file has fewer records and fewer variables than the original data file, since we removed large households and several variables to generate the SUF file.

Example 9.28 Number of individuals and variables and variable names

# Dimensions of file (observations, variables)

```
dim(file)
## [1] 10068 49
dim(fileOrig)
## [1] 10574 68
colnames(file) # Variable names
```

```
## [1] "IDH" "URBRUR" "REGION" "HHSIZE"
## [5] "OWNAGLAND" "RELIG" "ROOF" "TOILET"
## [9] "WATER" "ELECTCON" "FUELCOOK" "OWNMOTORCYCLE"
## [13] "CAR" "TV" "LIVESTOCK" "LANDSIZEHA"
## [17] "TANHHEXP" "TFOODEXP" "TALCHEXP" "TCLTHEXP"
## [21] "THOUSEXP" "TFURNEXP" "THLTHEXP" "TTRANSEXP"
## [25] "TCOMMEXP" "TRECEXP" "TEDUEXP" "TRESTHOTEXP"
## [29] "TMISCEXP" "INCTOTGROSSHH" "INCRMT" "INCWAGE"
## [33] "INCFARMBSN" "INCNFARMBSN" "INCRENT" "INCFIN"
```

```
## [37] "INCPENSN" "INCOTHER" "WGTPOP" "IDP"
## [41] "GENDER" "REL" "MARITAL" "AGEYRS"
## [45] "EDUCY" "EDYRSCURRAT" "ATSCHOOL" "INDUSTRY1"
## [49] "WGTHH"
```

To get an overview of the values of the variables, we use tabulations and cross-tabulations for categorical variables and summary statistics for continuous variables. To include the number of missing values ('NA' or other), we use the option useNA = "ifany" in the table() function. For some variables, these tabulations differ from the tabulations of the raw data, due to the anonymization of the SUF file.

In Table 9.18 the variables in the dataset "file" are listed along with concise descriptions of the variables, the level at which they are collected (individual level (IND) or household level (HH)), the measurement type (continuous, semi-continuous or categorical) and value ranges. Note that the dataset contains a selection of 49 variables of the 68 variable selected for the SUF release. The variables have been preselected based on their relevance for data users and some variables were removed while creating a SUF file. The numerical values for many of the categorical variables are codes that refer to values, e.g., in the variable "URBRUR", 1 stands for 'rural' and 2 for 'urban'. More information on the meanings of coded values of the categorical variables is available in the *R* code for this case study.

Any data cleaning, such as recoding missing value codes and removing empty variables, was already done in case study 1. The same holds for removing any direct identifiers. Direct identifiers are not released in the SUF file.

We identified the following sensitive variables in the dataset: variables related to schooling and labor force status as well as income and expenditure related variables. These variables need protection. Whether a variable is considered sensitive may depend on the release type, country and the dataset itself.

Table 9.18: Overview of the variables in the dataset

No.	Variable	Description	Level	Measurement	Values
	name				
1	IDH	Household ID	НН	•	1-2,000
2	IDP	Individua l ID	IND	•	1-13
3	REGION	Region	НН	categoric al	1-6
4	URBRUR	Area of residence	НН	categoric al	1, 2
5	WGTHH	Individua I weighting coefficie nt	НН	weight	31.2-8495 .7
6	WGTPOP	Populatio n weighting coefficie nt	IND	weight	45.8-9345 2.2
7	HHSIZE	Household size	НН	semi-cont inuous	1-33
8	GENDER	Gender	IND	categoric al	0, 1
9	REL	Relations hip to household head	IND	categoric al	1-9
10	MARITAL	Marital status	IND	categoric al	1-6

Continued on next page

Table 12.2 – continued from previous page

No.	140	le 12.2 – continue Description	Level	Measurement	Values	
	Variable name	·		Wedsurement		
11	AGEYRS	Age in completed years	IND	semi-cont inuous	0-65	
12	RELIG	Religion of household head	НН	categoric al	1, 5-7, 9	
13	ATSCHOOL	Currently en- rolled in school	IND	categoric al	0, 1	
14	EDUCY	Highest level of education attended	IND	categoric al	1-6	
15	EDYRSCURR AT	Years of edu- cation for cur- rently enrolled	IND	semi-cont inuous	1-18	
16	INDUSTRY1	Industry classific ation (1-digit)	IND	categoric al	1-10	
17	ROOF	Main material used for roof	IND	categoric al	1-5, 9	
18	TOILET	Main toilet fa- cility	НН	categoric al	1-4, 9	
19	ELECTCON	Electrici ty	HH	categoric al	0-3	
20	FUELCOOK	Main cooking fuel	НН	categoric al	1-5, 9	
21	WATER	Main source of water	НН	categoric al	1-9	
22	OWNAGLAND	Ownership of agricultu ral land	НН	categoric al	1-3	
23	LANDSIZEH A	Land size owned by household (ha) (agric and non agric)	НН	continuou s	0-40	
24	OWNMOTORC YCLE	Ownership of motorcycl e	НН	categoric al	0, 1	
25	CAR	Ownership of car	НН	categoric al	0, 1	
26	TV	Ownership of televisio n	НН	categoric al	0, 1	
27	LIVESTOCK	Number of large-siz ed livestock owned	НН	semi-cont inuous	0-25	
28	INCRMT	Income – Remittances	НН	continuou s	•	
29	INCWAGE	Income - Wages and salaries	НН	continuou s	•	

Continued on next page

Table 12.2 – continued from previous page

Table 12.2 – continued from previous page  No. Description Level Measurement Values					
	Variable name	·			values
30	INCFARMBS N	Income - Gross income from household farm businesse s	НН	continuou s	•
31	INCNFARMB SN	Income -Gross income from household non- farm businesse s	НН	continuou s	•
32	INCRENT	Income - Rent	НН	continuou s	•
33	INCFIN	Income - Finan- cial	НН	continuou s	•
34	INCPENSN	Income - Pensions/ social assistance	НН	continuou s	•
35	INCOTHER	Income - Other	НН	continuou s	•
36	INCTOTGRO SSHH	Income - Total	НН	continuou s	•
37	TFOODEXP	Total expenditu re on food	НН	continuou s	•
38	TALCHEXP	Total expenditu re on alcoholic beverages, to- bacco and nar- cotics	НН	continuou s	•
39	TCLTHEXP	Total expenditu re on clothing	НН	continuou s	•
40	THOUSEXP	Total expenditu re on housing	НН	continuou s	•
41	TFURNEXP	Total expenditu re on furnishin g	НН	continuou s	•
42	THLTHEXP	Total expenditu re on health	НН	continuou s	•
43	TTRANSEXP	Total expenditu re on transport	НН	continuou s	•

Continued on next page

Table 12.2 – continued from previous page

No.	Variable	Description	Level	Measurement	Values
	Variable name				
44	TCOMMEXP	Total expenditu re on commu- nica tion	НН	continuou s	•
45	TRECEXP	Total expenditu re on recreatio n	НН	continuou s	•
46	TEDUEXP	Total expenditu re on education	НН	continuou s	•
47	TRESHOTEX P	Total expenditu re on restauran ts and hotels	НН	continuou s	•
48	TMISCEXP	Total expenditu re on miscellan eous spending	НН	continuou s	•
49	ТАМННЕХР	Total annual nominal house- hold expenditu res	НН	continuou s	•

It is always important to ensure that the relationships between variables in the data are preserved during the anonymization process and to explore and take note of these relationships before beginning the anonymization. At the end of the anonymization process before the release of the data, an audit should be conducted, using these initial results, to check that these relationships are maintained in the anonymized dataset (see Step 11).

In our dataset, we identify several relationships between variables that need to be preserved during the anonymization process. The variables "TANHHEXP" and "INCTOTGROSSHH" represent the total annual nominal household expenditure and the total gross annual household income, respectively, and these variables are aggregations of existing income and expenditure components in the dataset.

The variables related to education are available only for individuals in the appropriate age groups and missing for other individuals. In addition, the household-level variables (cf. fourth column of Table 9.18) have the same values for all members in any particular household. The value of household size corresponds to the actual number of individuals belonging to that household in the dataset. As we proceed, we have to take care that these relationships and structures are preserved in the anonymization process.

We assume that the data are collected in a survey that uses simple sampling of households. The data contains two weight coefficients: "WGTHH" and "WGTPOP". The relationship between the weights is WGTPOP = WGTHH \* HHSIZE. "WGTPOP" is the sampling weight for the households and "WGTHH" is the sampling weight for the individuals to be used for disclosure risk calculations. "WGTHH" is used for computing individual-level indicators (such as education) and "WGTPOP" is used for population level indicators (such as income indicators). There are no strata variables available in the data. We will use "WGTPOP" for the anonymization of the household variables and "WGTHH" for the anonymization of the individual-level variables.

#### **Step 3: Type of release**

In this case study, we assume that the file will be released as a PUF, which will be freely available to anyone interested in the data (see Section 3.1 for the conditions and more information on the release of PUFs). The PUF release is intended for users with lower information requirements (e.g., students) and researchers interested in the structure of the data and willing to do preliminary research. The PUF file can give an idea to the researcher whether it is worthwhile for their research to apply for access to the SUF file. Researchers willing to do more in-depth research will most likely

apply for SUF access. Generally, users of a PUF file are not restricted by an agreement that prevents them from using the data to re-identify individuals and hence the accepted risk level is much lower than in the case of the SUF and the set of released variables is limited.

# Step 4: Intruder scenarios and choice of key variables

Next, based on the release type, we reformulate the intruder scenarios for the PUF release. This leads to the selection of a set of quasi-identifiers. Since this case study is based on a demo dataset, we do not have a real context and we cannot define exact disclosure scenarios. Therefore, we make hypothetical assumptions on possible disclosure scenarios. We consider two types of disclosure scenarios: 1) matching with other publically available datasets and 2) spontaneous recognition. Since the dataset will be distributed as PUF, there are de facto no restrictions on the use of the dataset by intruders.

For the sake of illustration, we assume that population registers are available with the demographic variables gender, age, place of residence (region, urban/rural), religion and other variables such as marital status and variables relating to education and professional status that are also present in our dataset. In addition, we assume that there is a publically available cadastral register on land ownership. Based on this analysis of available data sources, we have selected in case study 1 the variables "REGION", "URBRUR", "HHSIZE", "OWNAGLAND", "RELIG", "GENDER", "REL" (relationship to household head), "MARITAL" (marital status), "AGEYRS", "INDUSTRY1" and two variables relating to school attendance as categorical quasi-identifiers, the expenditure and income variables as well as LANDSIZEHA as continuous quasi-identifiers. According to our assessment, these variables might enable an intruder to re-identify an individual or household in the dataset by matching with other available datasets. The key variables for PUF release generally coincide with the key variables for the SUF release. Possibly, more variables could be added, since the user has more possibilities to match the data extensively and is not bound by any contract, as is in the case of the SUF file. Equally, some key variables in the SUF file may not be released in the PUF file and, as a consequence, these variables are removed from the list of key variables.

Upon further consideration, this initial set of identifying variables is too large for a PUF release, as the number of possible combinations (keys) is very high and hence many respondents could be identified based on these variables. Therefore, we decide to limit the set of key variables, by excluding variables from the dataset for PUF release. The choice of variables to be removed is led by the needs of the intended PUF users. Assuming the typical users are mainly interested in aggregate income and expenditure data, we can therefore remove from the initial set of key variables "OWNAGLAND", "RELIG" and "LANDSIZEHA" at the household level and "EDYRSCURRAT" and "ATSCHOOL" at the individual level. **NOTE: These variables will not be released in the PUF file.** We also remove the income and expenditure components from the list of key variables, since we reduce their information content by building proportions (see Step 8a). The list of the remaining key variables is presented in Table 9.19.

Table 9.19:	Overview	of selected	l key var	riables fo	or PUF file
-------------	----------	-------------	-----------	------------	-------------

Variable name	Variable description	Measurement level
REGION	region	Household, categorical
URBRUR	area of residence	Household, categorical
HHSIZE	household size	Household, categorical
TANHHEXP	total expenditure	Household, continuous
INCTOTGROSSHH	total income	Household, continuous
GENDER	gender	Individual, categorical
REL	relationship to household head	Individual, categorical
MARITAL	marital status	Individual, categorical
AGEYRS	age in completed years	Individual, semi-continuous/categ orical
EDUCY	highest level of education completed	Individual, categorical
INDUSTRY1	industry classification	Individual, categorical

The decision to release the dataset as a PUF means the level of anonymization will be relatively high and consequently, the variables are less detailed (e.g., after recoding) and a scenario of spontaneous recognition is less likely. Nevertheless, we should check for rare combinations or unusual patterns in the variables. Variables that may lead to spontaneous recognition in our sample are amongst others "HHSIZE" (household size) as well as "INCTOTGROSSHH"

(aggregate income) and "TANHHEXP" (total expenditure). Large households and households with high income are easily identifiable, especially when combined with other identifying variables such as a geographical identifier ("RE-GION"). There might be only one or a few households/individuals in a certain region with a high income, such as the local doctor. Variables that are easily observable and known by neighbors such as "ROOF", "TOILET", "WATER", "ELECTCON", "FUELCOOK", "OWNMOTORCYCLE", "CAR", "TV" and "LIVESTOCK" may also need protection depending on what stands out in the community, since a user might be able to identify persons (s)he knows. This is called the nosy-neighbor scenario.

#### Step 5: Data key uses and selection of utility measures

A PUF file contains less information and the file is generally used by students as a teaching file, by researchers to get an idea about the data structure, and for simple analyses. The users have generally lower requirements than for a SUF file and the results of analysis may be less precise. The researcher interested in a more detailed dataset, would have to apply for access to the SUF file. Therefore, we select more aggregate utility measures for the PUF file that reflect the intended use of a PUF file. Data intensive measures, such as the Gini coefficient, cannot be computed from the PUF file. Besides the standard utility measures, such as tabulations, we evaluate the decile dispersion ratio and a regression with the income deciles as regressand.

To measure the information loss, we should compare the initial data file before any anonymization (including the anonymization for the SUF) with the file after the anonymization for the PUF. Comparing the files directly before and after the PUF anonymization would underestimate the information loss, as this would omit the information loss due to SUF anonymization. Therefore, in Step 2, we also loaded the raw dataset.

#### Hierarchical (household) structure

As noted in case study 1, the data has a household structure. For the SUF release, we protected large households by removing these from the dataset. Since some variables are measured on the household level and thus have identical values for each household member, the values of the household variables should be treated in the same way for each household member (see Section 5.5). Therefore, we first anonymize only the household variables. After this, we merge them with the individual-level variables and then anonymize the individual-level and household-level variables jointly.

Since the data has a hierarchical structure, Steps 6 through 10 are repeated twice: Steps 6a through 10a are for the household-level variables and Steps 6b through 10b for the combined dataset. In this way, we ensure that household-level variable values remain consistent across household members for each household and the household structure cannot be used to re-identify individuals. This is further explained in Sections 4.4 and 7.6.

Before continuing to Step 6a, we select the categorical key variables, continuous key variables and any variables selected for use in PRAM routines, as well as household-level sampling weights in R. We also collect the variable names of the variables that will not be released. The PRAM variables are variables select for the PRAM routine, which we discuss further in Step 8a. We extract these selected household variables from the SUF dataset and save them as "fileHH". The choice of PRAM variables is further explained in Step 8a. Example 9.29 illustrates how these steps are done in R (see also Section 7.6).

Example 9.29: Selecting the variables for the household-level anonymization

```
# Categorical key variables at household level
selectedKeyVarsHH <- c('URBRUR', 'REGION', 'HHSIZE')

# Continuous key variables
numVarsHH <- c('TANHHEXP', 'INCTOTGROSSHH')

# PRAM variables
pramVarsHH <- c('ROOF', 'TOILET', 'WATER', 'ELECTCON', 'FUELCOOK', 'OWNMOTORCYCLE', 'CAR', 'TV', 'LIVESTOCK')

# Household weight
weightVarHH <- c('WGTPOP')

# Variables not suitable for release in PUF (HH level)
```

```
varsNotToBeReleasedHH <- c("OWNAGLAND", "RELIG", "LANDSIZEHA")

# Vector with names of all HH level variables

HHVars <- c('IDH', selectedKeyVarsHH, pramVarsHH, numVarsHH, weightVarHH)

# Create subset of file with only HH level variables

fileHH <- file[,HHVars]
```

Every household has the same number of entries as it has members (e.g., a household of three will be repeated three times in "fileHH"). Before analyzing the household-level variables, we select only one entry per household, as illustrated in Example 9.30. This is further explained in Section 7.6. In the same way we extract "fileOrigHH" from "fileOrig". "fileOrigHH" contains all variables from the raw data, but contains every household only once. We need "fileOrigHH" in Steps 8a and 10a for undoing some perturbative methods used in the SUF file and computing utility measures from the raw data respectively.

Example 9.30: Taking a subset with only households

```
# Remove duplicated rows based on IDH, one row per household in fileHH
fileHH <- fileHH[which(!**duplicated**(fileHH$IDH)),] # SUF file
fileOrigHH <- fileOrig[which(!**duplicated**(fileOrig$IDH)),] # original dataset
# Dimensions of fileHH
dim(fileHH)

## [1] 1970 16
dim(fileOrigHH)
## [1] 2000 68
```

The file "fileHH" contains 1,970 households and 16 variables. We are now ready to create our *sdcMicro* object with the corresponding variables we selected in Example 9.28. For our case study, we will create an *sdcMicro* object called "sdcHH" based on the data in "fileHH", which we will use for steps 6a – 10a (see Example 9.31). **NOTE: When the sdcMicro object is created, the sdcMicro package automatically calculates and stores the risk measures for the data.** This leads us to Step 6a.

Example 9.31: Creating a *sdcMicro* object for the household variables

```
# Create initial sdcMicro object for household level variables
sdcHH <- createSdcObj(dat = fileHH, keyVars = selectedKeyVarsHH,
pramVars = pramVarsHH, weightVar = weightVarHH, numVars = numVarsHH)
numHH <- length(fileHH[,1]) # number of households
```

#### Step 6a: Assessing disclosure risk (household level)

Based on the key variables selected in the disclosure scenarios, we can evaluate the risk at the household level. The PUF risk measures show a lower risk level than in the SUF file after anonymization in case study 1. The reason is that the set of key variables is smaller, since some variables will not be released in the PUF file. Removing (key) variables reduces the risk, and it is one of the most straightforward SDC methods.

As a first measure, we evaluate the number of households violating k-anonymity at the levels 2, 3 and 5. Table 9.20 shows the number of violating households as well as the percentage of the total number of households. Example

9.32 illustrates how to find these values with sdcMicro. The print() function in sdcMicro shows only the values for thresholds 2 and 3. Values for other thresholds can be calculated manually by summing up the frequencies smaller than the k-anonymity threshold, as shown in Example 9.32. The number of violators is already at a low level, due to the prior anonymization of the SUF file and the reduced set of key variables.

Table 9.20: Number and proportion of households violating k-anonymity

k-anonymity level	Number of HH violating	Percentage of total number of HH
2	0	0.0%
3	18	0.9%
5	92	4.7%

Example 9.32: Showing number of households violating k-anonymity for levels 2, 3 and 5

# Number of observations violating k-anonymity (thresholds 2 and 3) print(sdcHH)

```
## Infos on 2/3-Anonymity:
##
## Number of observations violating
## - 2-anonymity: 0
## - 3-anonymity: 18
##
## Percentage of observations violating
## - 2-anonymity: 0.000 %
## - 3-anonymity: 0.914 %
```

# Calculate sample frequencies and count number of obs. violating k(5) - anonymity kAnon5 <- sum(sdcHH@risk\$individual[,2] < 5)

kAnon5

## [1] 92

# As percentage of total kAnon5 / numHH

```
##[1] 0.04670051
```

It is often useful to view the records of the household(s) that violate k-anonymity. This might help to find which variables cause the uniqueness of these households; this can then be used later when choosing appropriate SDC methods. Example 9.32 shows how to access the values of the households violating 3 and 5-anonymity. Not surprisingly, the variable "HHSIZE" is responsible for many of the unique combinations and the origin of much of the risk. This is even the case after removing large households for the SUF release.

Example 9.33: Showing records of households that violate k-anonymity

# Show values of key variable of records that violate k-anonymity fileHH[sdcHH@risk\$individual[,2] < 3, selectedKeyVarsHH] # for 3-anonymity

fileHH[sdcHH@risk\$individual[,2] < 5, selectedKeyVarsHH] # for 5-anonymity

We also assess the disclosure risk of the categorical variables with the individual and global risk measures as described in Sections 4.5 and 4.8. In "fileHH" every entry represents a household. Therefore, we use the individual non-hierarchical risk here, where the individual refers in this case to a household. "fileHH" is a subset of the complete dataset and contains only households and has, contrary to the complete dataset, no hierarchical structure. In Step 6b, we evaluate the hierarchical risk in the dataset "file", the dataset containing both households and individuals. The individual and global risk measures automatically take into consideration the household weights, which we defined in Example 9.29. In our file, the global risk measure calculated using the chosen key variables is lower than 0.01% (the smallest reported value is 0.01%, in fact the global risk is 0.0000642%). This percentage is extremely low and corresponds to 0.13 expected re-identifications. The results are also shown in Example 9.34. This low figure can be explained by the relatively small sample size of 0.25% of the total population (see case study 1). Furthermore, one should keep in mind that this risk measure is based only on the categorical quasi-identifiers at the household level.

Example 9.34: Printing global risk measures

```
print(sdcHH, "risk")
```

```
## Risk measures:
##
## Number of observations with higher risk than the main part of the data: 0
## Expected number of re-identifications:0.13 (0.01%)
```

The global risk measure does not provide information about the spread of the individual risk measures. There might be a few households with relatively high risk, while the global (average) risk is low. Therefore we check the highest individual risk as shown in Example 9.35. The individual risk of the household with the highest risk is 0.1 %, which is still very low.

Example 9.35 Determining the highest individual risk

```
# Highest individual risk
max(sdcHH@risk$individual[, "risk"])
```

```
## [1] 0.001011633
```

Since the selected key variables at the household level are both categorical and numerical, the individual and global risk measures based on frequency counts do not completely reflect the disclosure risk of the entire dataset. When generating the SUF file, we concluded that recoding of continuous variables to make them all categorical would likely not satisfy the needs of the SUF users. For the PUF file it is acceptable to recode continuous variables, such as income and expenditures since PUF content is typically less detailed. In Step 8a we will recode these variables into deciles and convert them into categorical variables. Therefore, we exclude these variables from the risk calculations now We take these variables into account while remeasuring risk after anonymization.

#### Step 7a: Assessing utility measures (household level)

The utility of the data does not only depend on the household level variables, but on the combination of household-level and individual-level variables. Therefore, it is not useful to evaluate all the utility measures selected in Step 5 at this stage, i.e., before anonymizing the individual level variables. We restrict the initial measurement of utility to those measures that are solely based on the household variables. In our dataset, these are the measures related to

income and expenditure and their distributions. The results are presented in Step 10a, together with the results after anonymization, which allow direct comparison. If after the next anonymization step it appears that the data utility has been significantly decreased by the suppression of some household level variables, we can return to this step. **NOTE:** to analyze the utility loss, the utility measures before anonymization have to be calculated from the raw data and not from the anonymized SUF file. Not all measures from case study 1 can be computed from the PUF file, since the information content is lower. The set of utility measures we use to evaluate the information loss in the PUF file consists of measures that need less detailed variables. This reflects the lower requirements a PUF user has on the dataset.

# Step 8a: Choice and application of SDC methods (household level)

This step is divided into the anonymization of the categorical key variables and the continuous key variables, since different methods are used for both sets of variables. As already discussed in Step 4, we do not release all variables in the PUF file. At the household level "RELIG" (religion of household head), "OWNAGLAND" (land ownership) and "LANDSIZEHA" (plot size in ha) are not released in addition to the variables removed for the SUF release. For the sake of illustration, we assume that the variable "RELIG" is too sensitive and the variables "OWNAGLAND" and "LANDSIZEHA" are too identifying.

#### **Categorical variables**

We are now ready to move on to the choice of SDC methods for the categorical variables on the household level in our dataset. In the SUF file we already recoded some of the key variables and used local suppression. We only have three categorical key variables at the household level; "URBRUR", "REGION" and "HHSIZE". The selected categorical key variables at the household level are not suitable for recoding at this point, since the values cannot be grouped further. "URBRUR" has only two distinct categories and "REGION" has only six non-combinable regions. As noted before, the variable "HHSIZE" can be reconstructed by a headcount per household. Therefore, recoding of this variable alone does not lead to disclosure control.

Due to the relatively low risk of re-identification based on the three selected categorical household level variables, it is possible in this case to use an option like local suppression to achieve our desired level of risk. Applying local suppression when initial risk is relatively low will likely only lead to suppression of few observations and thus limit the loss of utility. If, however, the data had been measured to have a relatively high risk, then applying local suppression without previous recoding would likely result in a large number of suppressions and greater information loss. Efforts such a recoding should be taken first before suppressing values in cases where risk is initially measured as high. Recoding will reduce risk with little information loss and thus the number of suppressions, if local suppression is applied as an additional step.

We apply local suppression to reach 5-anonymity. The chosen level of five is higher than in the SUF release and is based on the release type as PUF. This leads to a total of 39 suppressions, all in the variable "HHSIZE". As explained earlier, suppression of the value of the variable "HHSIZE" does not lead to actual suppression of this information. Therefore, we redo the local suppression, but this time we tell *sdcMicro* to, if possible, not suppress "HHSIZE" but one of the other variables. Alternatively, we could remove households with suppressed values of the variable "HHSIZE", remove large households or split households.

In *sdcMicro* it is possible to tell the algorithm which variables are important and less important for making small changes (see also Section 5.2.2). To prevent values of the variable "HHSIZE" being suppressed, we set the importance of "HHSIZE" in the importance vectors to the highest (i.e., 1). We try two different importance vectors: the first where "REGION" is more important than "URBRUR" and the second with the importance of "REGION" and "URBRUR" swapped. Example 9.36 shows how to apply local suppression and put importance on the variable "HHSIZE". **NOTE:** In Example 9.36 we use the undolast() function in sdcMicro to go one step back after we had first applied local suppression with no importance vector. The undolast() function restores the *sdcMicro* object back to the previous state (i.e., before we applied local suppression), which allows us to rerun the same command, but this time with an importance vector set. The undolast() function can only be used to go one step back.

The suppression patterns of the three different options are shown in Table 9.21. The importance is clearly reflected in the number of suppressions per variable. The total number of suppressions is with an importance vector higher than without an importance vector (44/73 vs. 39), but 5-anonymity is achieved in the dataset with no suppressions in the variable "HHSIZE". This means that we do not have to remove or split households. The variable "REGION" is the

type of variable that should not have any suppressions either. From that perspective we chose the third option. This leads to more suppressions, but no suppressions in "HHSIZE" and as few as possible in "REGION".

Table 9.21: Number of suppressions by variable after local suppression with and without importance vector

Key vari- able	Number of suppressions and proportion of total			
	No importance	Importance HHSIZE, URBRUR, RE-	Importance HHSIZE, REGION, UR-	
	vector	GION	BRUR	
URBRUR	0 (0.0 %)	2 (0.1 %)	61 (3.1 %)	
REGION	0 (0.0 %)	42 (2.1 %)	12 (0.6 %)	
HHSIZE	39 (2.0 %)	0 (0.0 %)	0 (0.0 %)	

Example 9.36: Local suppression with and without importance vector

```
# Local suppression to achieve 5-anonimity
sdcHH <- localSuppression(sdcHH, k = 5, importance = NULL) # no importance vector
print(sdcHH, "ls")
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## URBRUR | 0 | 0.000
## REGION | 0 | 0.000
## HHSIZE | 39 | 1.980
## -----
sdcHH <- undolast(sdcHH) # undo suppressions to see the effect of an importance vector
# Redo local suppression minimizing the number of suppressions in HHSIZE
sdcHH \leftarrow localSuppression(sdcHH, k = 5, importance = c(2, 3, 1))
print(sdcHH, "ls")
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## URBRUR | 2 | 0.102
## REGION | 42 | 2.132
## HHSIZE | 0 | 0.000
sdcHH <- undolast(sdcHH) # undo suppressions to see the effect of a different importance vector
# Redo local suppression minimizing the number of suppressions in HHSIZE
sdcHH \leftarrow localSuppression(sdcHH, k = 5, importance = c(3, 2, 1))
print(sdcHH, "ls")
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
```

```
## URBRUR |61|3.096
## REGION |12 |0.609
## HHSIZE | 0 | 0.000
## -----
```

In case study 1 we applied invariant PRAM to the variables "ROOF", "TOILET", "WATER", "ELECTCON", "FU-ELCOOK", "OWNMOTORCYCLE", "CAR", "TV" and "LIVESTOCK", since these variables are not sensitive and were not selected as quasi-identifiers because we assumed that there are no external data sources containing this information that could be used for matching. Values can be easily observed or be known to neighbors, however, and therefore are important, together with other variables, for the nosy neighbor scenario. For the PUF release we would like to level of uncertainty by increasing the number of changes. Therefore, we redo PRAM with a different transition matrix. As discussed in Section 5.3.1, the invariant PRAM method has the property that the univariate distributions do not change. To maintain this property, we reapply PRAM to the raw data, rather than to the already PRAMmed variables in the SUF file.

Example 9.37 illustrates how to apply PRAM. We use the original values to apply PRAM and replace the values in the *sdcMicro* object with these values. We choose the parameter 'pd', the lower bound for the probability that a value is not changed, to be relatively low at 0.6. This is a lower value than the 0.8 used in the SUF file and will lead to a higher number of changes (cf. Example 9.17 on page 135). This is acceptable for a PUF file and introduces more uncertainty as required for a PUF release. Example 9.37 also shows the number of changed records per variables. Because PRAM is a probabilistic method, we set a seed for the random number generator before applying PRAM to ensure reproducibility of the results. **Note: In some cases the choice of the seed matters. The choice of seed changes the results.** The seed should not be released, since it allows for reconstructing the original values if combined with the transition matrix. The transition matrix can be released: this allows for consistent statistical inference by correcting the statistical methods used if the researcher has knowledge about the PRAM method (at this point *sdcMicro* does not allow to retrieve the transition matrix).

Example 9.37: Applying PRAM

## TV != TV\_pram : 275 (13.96%)

## LIVESTOCK != LIVESTOCK\_pram : 109 (5.53%)

PRAM has changed values within the variables according to the invariant transition matrices. Since we used the invariant PRAM method (see Section 5.3.1), the absolute univariate frequencies remain approximately unchanged. This is not the case for the multivariate frequencies. In Step 10a we compare the multivariate frequencies before and after anonymization for the PRAMmed variables.

#### **Continuous variables**

We have selected the variables "INCTOTGROSSHH" (total income) and "TANHHEXP" (total expenditure) as numerical quasi-identifiers, as discussed in Step 4. In Step 5 we identified variables having high interest for the users of our data: many users use the data for measuring inequality and expenditure patterns. The noise addition in the SUF file does not protect these variables sufficiently, especially, because outliers are not protected. Therefore, we decide to recode total income and total expenditure into deciles.

As with PRAM, we want to compute the deciles from the raw data rather than from the perturbed values in the SUF file. We compute the deciles directly from the raw data and overwrite these values in the *sdcMicro* object. Subsequently, we compute the mean of each decile from the raw data and replace the values for total income and total expenditures with the mean of the respective decile. In this way the mean of both variables does not change. This approach can be interpreted as univariate microaggregation with very large groups (group size n/10) with the mean as replacement value (see Section 5.3.2).

The information in the income and expenditure variables by component is too sensitive to release as PUF, and, summing those variables would allow an intruder to reconstruct the totals. PUF users might however be interested in the shares. Therefore, we decide to keep the income and expenditure components as proportions of the raw totals, rounded to two digits. The anonymization of the income and expenditure variables is shown in Example 9.38.

Example 9.38: Anonymization of income and expenditure variables

```
# Create bands (deciles) for income and expenditure variables (aggregates) based on the original data
decExp <- as.numeric(cut(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "TANHHEXP"],
quantile(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "TANHHEXP"], (0:10)/10, na.rm = T), include.lowest
= TRUE, labels = \mathbf{c}(1, 2, 3, 4, 5, 6, 7, 8, 9, 10))
decInc <- as.numeric(cut(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "INCTOTGROSSHH"],
quantile(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "INCTOTGROSSHH"], (0:10)/10, na.rm = T),
include.lowest = TRUE, labels = c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10))
# Mean values of deciles
decExpMean <- round(sapply(split(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "TANHHEXP"], decExp),
decIncMean <- round(sapply(split(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "INCTOTGROSSHH"],
decInc), mean))
# Replace with mean value of decile
sdcHH@manipNumVars$TANHHEXP <- decExpMean[match(decExp, names(decExpMean))]
sdcHH@manipNumVars$INCTOTGROSSHH <- decIncMean[match(decInc, names(decIncMean))]
# Recalculate risks after manually changing values in sdcMicro object
calcRisks(sdcHH)
```

# Extract data from sdcHH

HHmanip <- extractManipData(sdcHH) # manipulated variables HH

```
# Keep components of expenditure and income as share of total, use original data since previous data was perturbed compExp <- c('TFOODEXP', 'TALCHEXP', 'TCLTHEXP', 'THOUSEXP', 'TFURNEXP', 'THLTHEXP', 'TTRANSEXP', 'TCOMMEXP', 'TRECEXP', 'TEDUEXP', 'TRESTHOTEXP', 'TMISCEXP') compInc <- c('INCRMT', 'INCWAGE', 'INCFARMBSN', 'INCNFARMBSN', 'INCRENT', 'INCFIN', 'INCPENSN', 'INCOTHER')
```

```
HHmanip <- cbind(HHmanip, round(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), compExp] / fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "TANHHEXP"], 2))

HHmanip <- cbind(HHmanip, round(fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), compInc] / fileOrigHH[match(fileHH$IDH, fileOrigHH$IDH), "INCTOTGROSSHH"], 2))
```

#### Step 9a: Re-measure risk (household level)

For the categorical variables, we conclude that we have achieved 5-anonymity in the data with local suppression. 5-anonymity also implies 2- and 3-anonymity. The global risk stayed close to zero (as the expected number of reidentifications), which is very low. Therefore, we conclude that based on the categorical variables, the data has been sufficiently anonymized. One should keep in mind that the anonymization methods applied are complementing the ones used for the SUF. **NOTE: The methods selected methods in this case study alone would not be sufficient to protect the data set for a PUF release.** 

We have reduced the risk of spontaneous recognition of households, by removing the variable "LANDSIZEHA" and PRAMming the variables identified to be important in the nosy neighbor scenario. An intruder cannot know with certainty whether a household that (s)he recognizes in the data is the correct household, due to the noise in these variables.

These measures refer only to the categorical variables. To evaluate the risk of the continuous variables we could use an interval measure or closest neighbor algorithm. These risk measures are discussed in Section 4.7. We chose to use an interval measure, since exact value matching is not our largest concern based on the assumed scenarios and external data sources. Instead, datasets with similar values but not the exact same values could be used for matching. Here the main concern is that the values are sufficiently far from the original values, which is measured with an interval measure.

Example 9.39 shows how to evaluate the interval measure for the variables "INCTOTGROSSHH" and "TANHHEXP" (total income and expenditure). The different values of the parameter k in the function dRisk() define the size of the interval around the original value as a function of the standard deviation, as explained in Section 4.7.2. The larger k, the larger the intervals, the higher the probability that a perturbed value is in the interval around the original value and the higher the risk measure. The results are satisfactory, especially when keeping in mind that there are only 10 distinct values in the dataset (the means of each of the deciles). All outliers have been recoded. Looking at the proportions of the components, we do not detect any outliers (households with an unusual high or low spending pattern in one component).

Example 9.39: Measuring risk of re-identification of continuous variables

Step 10a Re-measure utility (household level)

```
# Risk evaluation continuous variables
dRisk(sdcHH@origData[,**c**("TANHHEXP", "INCTOTGROSSHH")], xm =
sdcHH@manipNumVars[,c("TANHHEXP", "INCTOTGROSSHH")], k = 0.01)
## [1] 0.4619289
dRisk(sdcHH@origData[,**c**("TANHHEXP",
                                              "INCTOTGROSSHH")],
                                                                                          sd-
                                                                          xm
cHH@manipNumVars[,c("TANHHEXP", "INCTOTGROSSHH")], k = 0.02)
## [1] 0.642132
dRisk(sdcHH@origData[,**c**("TANHHEXP",
                                              "INCTOTGROSSHH")],
                                                                          xm
                                                                                          sd-
cHH@manipNumVars[,c("TANHHEXP", "INCTOTGROSSHH")], k = 0.05)
## [1] 0.8258883
```

The utility in the data has decreased compared to the raw data, mainly because variables were completely removed. Many of the utility measures used in case study 1 are not applicable to the PUF file. However, by replacing the deciles with their means, we can still use the income and expenditure variables for arithmetic operations. Also the shares of the income and expenditure components can still be used, since they are based on the raw data.

We select two additional utility measures: the decile dispersion ratio and the share of total consumption by the poorest decile. The decile dispersion ratio is the ratio of the average income of the top decile and the average income of the bottom decile. Example 9.40 shows how to compute these from the raw data and the household variables after anonymization. Table 9.22 presents the estimated values. The differences are small and mainly due to the removed households.

Table 9.22: Comparison of utility measures

	Raw data	PUF file
Decile dispersion ratio	24.12	23.54
Share of consumption by the poorest decile	0.0034	0.0035

Example 9.40: Computation of decile dispersion ratio and share of total consumption by the poorest decile

# Decile dispersion ratio

# raw data

mean(tail(sort(fileOrigHH\$INCTOTGROSSHH), n = 200)) / mean(head(sort(fileOrigHH\$INCTOTGROSSHH), n = 200))

```
## [1] 24.12152
```

**mean(tail(sort**(HHmanip\$INCTOTGROSSHH), n = 197)) / **mean(head(sort**(HHmanip\$INCTOTGROSSHH), n = 197))

```
## [1] 23.54179
```

# Share of total consumption by the poorest decile households sum(head(sort(fileOrigHH\$TANHHEXP), n = 200)) / sum(fileOrigHH\$TANHHEXP)

```
## [1] 0.003411664
```

sum(head(sort(HHmanip\$TANHHEXP), n = 197)) / sum(HHmanip\$TANHHEXP)

```
## [1] 0.003530457
```

#### Merging the household- and individual-level variables

The next step is to merge the treated household variables with the untreated individual variables for the anonymization of the individual level variables. Example 9.41 shows the steps to merge these files. This also includes the selection of variables used in the anonymization of the individual-level variables. We create the *sdcMicro* object for the anonymization of the individual variables in the same way as for the household variable in Example 9.31. Generally, at this stage, the household level and individual level variables should be combined and quasi-identifiers at both levels be used (see Section 4.4). Unfortunately, in our dataset, this leads to long computation times. Therefore, we create two *sdcMicro* objects, one with all key variables ("sdcCombinedAll") and one with only the individual level key variables ("sdcCombined"). In Step 6b we compare the risk measures for both cases and in Step 8b we discuss alternative approaches to keeping the complete set of variables. We now repeat Steps 6-10 for the individual-level variables.

Example 9.41: Merging the files with household and individual-level variables and creating an *sdcMicro* object for the anonymization of the individual-level variables

```
### Select variables (individual level)
selectedKeyVarsIND = c('GENDER', 'REL', 'MARITAL', 'AGEYRS', 'EDUCY', 'INDUSTRY1') # list of selected
key variables
# sample weight (WGTHH, individual weight)
selectedWeightVarIND = c('WGTHH')
# Household ID
selectedHouseholdID = c('IDH')
# Variables not suitable for release in PUF (IND level)
varsNotToBeReleasedIND <- c("ATSCHOOL", "EDYRSCURRAT")
# All individual level variables
INDVars <- c(selectedKeyVarsIND)
# Recombining anonymized HH data sets and individual level variables
indVars <- c("IDH", "IDP", selectedKeyVarsIND, "WGTHH") # HID and all non HH vars
fileInd <- file[indVars] # subset of file without HHVars
fileCombined \leftarrow merge(HHmanip, fileInd, by.x = \mathbf{c}('\text{IDH'}))
fileCombined <- fileCombined[order(fileCombined[,'IDH'], fileCombined[,'IDP']),]
dim(fileCombined)
## [1] 10068 44
# SDC objects with only IND level variables
sdcCombined <- createSdcObj(dat = fileCombined, keyVars = c(selectedKeyVarsIND), weightVar =
selectedWeightVarIND, hhId = selectedHouseholdID)
# SDC objects with both HH and IND level variables
sdcCombinedAll < -createSdcObj(dat = fileCombined, keyVars = c(selectedKeyVarsIND, selectedKeyVarsHH),
weightVar = selectedWeightVarIND, hhId = selectedHouseholdID)
```

# Step 6b: Assessing disclosure risk (individual level)

As first measure, we evaluate the number of records violating k-anonymity at the levels 2, 3 and 5. Table 9.23 shows the number of violating individuals as well as the percentage of the total number of households. The second and third column refer to "sdcCombined" and the fourth and fifth column to "sdcCombinedAll". We see that combining the individual level and household level variables leads to a large increase in the number of k-anonymity violators. The choice not to include the household level variables is pragmatically driven by the computation time and can be justified by the different type of variables on the household level and individual level. One could assume that these variables are not available in the same dataset and can therefore not simultaneously be used by an intruder to re-identify individuals.

Table 9.23: Number of records violating k-anonimity

	sdcCombined		sdcCombinedAll	
k-	Number of records vio-	Percentage	Number of records vio-	Percentage
anonymity	lating	of total records	lating	of total records
2	0	0.0 %	4,048	40.2 %
3	167	1.7 %	6,107	60.7 %
5	463	4.6 %	8,292	82.4 %

The global hierarchical risk measure is 0.095%, which corresponds to approximately 10 expected re-identifications. We use here the hierarchical risk measure, since the re-identification of a single household member would lead to the re-identification of the other members of the same household too. This number is low compared to the number of k-anonymity violations, due to the high sample weights, which protect the data already to a large extent. Only 24 observations have an individual hierarchical risk higher than 1%, with a maximum of 1.17%. This is mainly because of the lower sample weights of these records. Example 9.42 shows how to retrieve these measures in R.

Example 9.42: Risk measures before anonymization

numIND <- length(fileCombined[,1]) # number of households</pre>

# Number of observations violating k-anonymity print(sdcCombined)

# Calculate sample frequencies and count number of obs. violating k (3,5) - anonymity kAnon5 <- sum(sdcCombined@risk\$individual[,2] <5) kAnon5

```
## [1] 463
```

# As percentage of total kAnon5 / numIND

```
## [1] 0.04598729
```

# Global risk on individual level print(sdcCombined, 'risk')

#### Step 7b: Assessing utility measures (individual level)

We evaluate the utility measures as discussed in Step 5 based on the raw data (before applying any anonymization measures). The results are presented in Step 10b together with the values after anonymization to allow for direct comparison.

# Step 8b: Choice and application of SDC methods (individual level)

In this step, we discuss four different techniques used for anonymization: 1) removing variables from the dataset to be released, 2) recoding of categorical variables to reduce the level of detail, 3) local suppression to achieve the required level of k-anonymity, 4) randomization of the order of the records in the file. Finally, we discuss some alternative options for treating the household structure in the dataset.

#### Removing variables

Additional to the variables removed from the dataset for the SUF release (see case study 1), we further reduce the number of variables in the dataset to be released. This is normal practice for PUF releases. Sensitive or identifying variables are removed, which allows to release other variables at a more detailed level. In a PUF release, the set of key variables should be limited.

In our case, we decide to remove at the individual level the variables "EDYRSCURRAT", as this variable is too identifying (identifies whether there are school-going children in the household). We keep the variable "EDUCY" (highest level of education attended) for information on education. NOTE: As an alternative to removing the variables from the dataset, one could also set all values to missing. This would allow the user to see the structure and variables contained in the SUF file.

# Recoding

As noted before, PUF users require a lower level of information and therefore we can recode the key variables even further to reduce the disclosure risk. The recoding of variables in case study 1 is not sufficient for a PUF release. Therefore, we recode most of the categorical key variables from Table 9.19 to reduce the risk and number of necessary suppressions by local suppression. Table 9.24 gives an overview of the recodes made. All new categories are formed with the needs of the data user in mind. Example 9.43 shows how to do this in *R* and also shows value labels and the univariate tabulations of these variables before and after recoding.

Table 9.24: Overview of recodes of categorical variables at individual level

Variable	Recoding
REL (relation	recode 'Father/Mother', 'Grandchild', 'Son/Daughter in law', 'Other relative' to 'Other relative'
to household	and recode 'Domestic help' and 'Non-relative' to 'Other'
head)	
MARITAL	recode 'Married monogamous', 'Married polygamous', 'Common law, union coutumiere, union
(marital status)	libre, living together' to 'Married/living together' and 'Divorced/Separated' and 'Widowed' to
	'Divorced/Separated/Widowed'
AGEYRS (age	recode values under 15 to 7 (other values have been recoded for SUF)
in completed	
years)	
EDUCY (high-	recode 'Completed lower secondary (or post-primary vocational education) but less than com-
est level of	pleted upper secondary', 'Completed upper secondary (or extended vocational/technical educa-
education	tion)', 'Post secondary technical' and 'University and higher' to 'Completed lower secondary or
completed)	higher'
INDUSTRY1	recode to 'primary', 'secondary' and 'tertiary'

Example 9.43: Recoding the categorical and continuous variables

```
# Recode REL (relation to household head)
table(sdcCombined@manipKeyVars$REL, useNA = "ifany")
##
## 1 2 3 4 5 6 7 8 9 <NA>
## 1698 1319 4933 52 765 54 817 40 63 327
#1 - Head, 2 - Spouse, 3 - Child, 4 - Father/Mother, 5 - Grandchild, 6 - Son/Daughter in law
#7 - Other relative, 8 - Domestic help, 9 - Non-relative
sdcCombined <- groupVars(sdcCombined, var = "REL", before = c("4", "5", "6", "7"), after = c("7", "7", "7", "7"))
# other relative
sdcCombined <- groupVars(sdcCombined, var = "REL", before = c("8", "9"), after = c("9", "9")) # other
table(sdcCombined@manipKeyVars$REL, useNA = "ifany")
##
## 1 2 3 7 9 <NA>
## 1698 1319 4933 1688 103 327
# Recode MARITAL (marital status)
table(sdcCombined@manipKeyVars$MARITAL, useNA = "ifany")
##
## 1 2 3 4 5 6 <NA>
## 3542 2141 415 295 330 329 3016
```

```
#1 - Never married, 2 - Married monogamous, 3 - Married polygamous,
#4 - Common law, union coutumiere, union libre, living together, 5 - Divorced/Separated, 6 - Widowed
sdcCombined \leftarrow c("2", "3", "4"), after = c("2", "2", "2")) #
married/living together
sdcCombined <- groupVars(sdcCombined, var = "MARITAL", before = c("5", "6"), after = c("9", "9")) #
divorced/seperated/widowed
table(sdcCombined@manipKeyVars$MARITAL, useNA = "ifany")
##
## 1 2 9 <NA>
## 3542 2851 659 3016
# Recode AGEYRS (0-15 years)
table(sdcCombined@manipKeyVars$AGEYRS, useNA = "ifany")
##
## 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14
## 311 367 340 332 260 334 344 297 344 281 336 297 326 299 263
## 20 30 40 50 60 65 <NA>
## 1847 1220 889 554 314 325 188
sdcCombined <- groupVars(sdcCombined, var = "AGEYRS", before = c("0", "1", "2", "3", "4", "5", "6", "7", "8",
"9", "10", "11", "12", "13", "14")
, after = rep("7", 15))
table(sdcCombined@manipKeyVars$AGEYRS, useNA = "ifany")
##
## 7 20 30 40 50 60 65 <NA>
## 4731 1847 1220 889 554 314 325 188
sdcCombined <- calcRisks(sdcCombined)</pre>
# Recode EDUCY (highest level of educ compl)
table(sdcCombined@manipKeyVars$EDUCY, useNA = "ifany")
##
## 0 1 2 3 4 5 6 <NA>
## 1582 4755 1062 330 139 46 104 2050
#0 - No education, 1 - Pre-school/ Primary not completed, 2 - Completed primary, but less than completed lower
secondary
#3 - Completed lower secondary (or post-primary vocational education) but less than completed upper secondary
```

#4 - Completed upper secondary (or extended vocational/technical education), 5 - Post secondary technical

#6 - University and higher

```
"3")) # completed lower secondary or higher
table(sdcCombined@manipKeyVars$EDUCY, useNA = "ifany")
##
## 0 1 2 3 <NA>
## 1582 4755 1062 619 2050
# Recode INDUSTRY1 ()
table(sdcCombined@manipKeyVars$INDUSTRY1, useNA = "ifany")
##
## 1 2 3 4 5 6 7 8 9 10 <NA>
## 5300 16 153 2 93 484 95 17 70 292 3546
#1 - Agriculture and Fishing, 2 - Mining, 3 - Manufacturing, 4 - Electricity and Utilities
#5 - Construction, 6 - Commerce, 7 - Transportation, Storage and Communication, 8 - Financial, Insurance and
Real Estate
#9 - Services: Public Administration, 10 - Other Services, 11 - Unspecified
sdcCombined <- groupVars(sdcCombined, var = "INDUSTRY1", before = c("1", "2"), after = c("1", "1")) # primary
sdcCombined <- groupVars(sdcCombined, var = "INDUSTRY1", before = c("3", "4", "5"), after = c("2", "2", "2"))
# secondary
sdcCombined <- groupVars(sdcCombined, var = "INDUSTRY1", before = c("6", "7", "8", "9", "10"), after = c("3",
"3", "3", "3", "3")) # tertiary
table(sdcCombined@manipKeyVars$INDUSTRY1, useNA = "ifany")
##
## 1 2 3 <NA>
## 5316 248 958 3546
```

# Local suppression

The recoding has reduced the risk already considerably. We use local suppression to achieve the required level of k-anonymity. Generally, the required level of k-anonymity for PUF files is 3 or 5. In this case study, we require 5-anonimity. Example 9.44 shows the suppression pattern without specifying an importance vector. All suppressions are made in the variable "AGEYRS". This is the variable with the highest number of different values, and hence considered first by the algorithm. We try different suppression patterns by specifying importance vectors, but we decide that the pattern without importance vector yields the best result. This is also the result with the lowest total number of suppressions. Less than 1 percent suppression in the age variable is acceptable. We could reduce this number by further recoding the variable "AGEYRS".

Example 9.44: Local suppression to reach 5-anonimity

```
# Local suppression without importance vector
sdcCombined <- localSuppression(sdcCombined, k = 5, importance = NULL)
```

# Number of suppressions per variable print(sdcCombined, "ls")

```
## Local Suppression:
## KeyVar | Suppressions (#) | Suppressions (%)
## GENDER | 0 | 0.000
## REL | 0 | 0.000
## MARITAL | 0 | 0.000
## AGEYRS | 91 | 0.904
## EDUCY | 0 | 0.000
## INDUSTRY1 | 0 | 0.000
```

#### Randomization of order of records

The records in the dataset are ordered by region and household ID. There is a certain geographical order of the households within the regions, due to the way the households IDs were assigned. Intruders could reconstruct suppressed values by using this structure. To prevent this, we randomly reorder the records within the regions. Example 9.45 shows how to do this in R. We first count the number of records per region (NOTE: Some records have their region value suppressed, so we include the count of NAs). Subsequently, we draw randomly household IDs, in such way that the regional division is respected. Finally, we sort the file by the new, randomized, individual ID ("IDP"). Households with suppressed values for "REGION" will be last in the reordered file. Before randomizing the order, we extract the data from the sdcMicro object "sdcCombined" as shown in Example 9.45.

Example 9.45: Randomizing the order of records within regions

```
# Randomize order of households dataAnon and recode IDH to random number (sort file by region)

set.seed(97254)

# Sort by region

dataAnon <- dataAnon[order(dataAnon$REGION),]

# Number of households per region

hhperregion <- table(dataAnon[match(unique(dataAnon$IDH), dataAnon$IDH), "REGION"], useNA = "ifany")

# Randomized IDH (household ID)

randomHHid <- c(sample(1:hhperregion[1], hhperregion[1]), unlist(lapply(1:(length(hhperregion)-1), function(i){sample((sum(hhperregion[1:i]) + 1): sum(hhperregion[1:(i+1)]), hhperregion[(i+1)])})))

dataAnon$IDH <- rep(randomHHid, table(dataAnon$IDH)[match(unique(dataAnon$IDH), as.numeric(names(table(dataAnon$IDH))))])

# Sort by IDH (and region)

dataAnon <- dataAnon[order(dataAnon$IDH),]
```

# Alternative options for dealing with household structure

In Step 6b we compared the disclosure risk for two cases: one with only individual level key variables and another with individual level and household level key variables combined. We decided to use only the individual level key variables to reduce the computation time and justified this choice by arguing that intruders cannot use household and individual level variables simultaneously. This might not always be the case. Therefore we explore other options to

reduce the risk when taking both individual level and household level variables into account. We present two options: removing the household structure; and using options in the local suppression algorithm.

#### Removing household structure

We consider the risk emanating from the household structure in the dataset to be very high. We can remove the hierarchical household structure completely and treat all variables at the individual level. This entails, besides removing the household id ("IDH"), also treating variables that could be used for reconstructing households. These are, for instance, "REL" (relation to household head), "HHSIZE" (household size), and any of the household level variables, such as income and expenditure. However, not all household level variables need to be treated. For example, "REGION" is a household level variable, but the probability that this variable leads to the reconstruction of a household is low. Also, we need to reorder the records in the file, as they are sorted by households. Note that by removing the household structure, we interpret all variables as individual level variables for measuring disclosure risk. This leads to a lower need for recoding and suppression, since the hierarchical risk disappears. The reason why we did not opt for this approach is the loss of utility for the user. The household structure is an important feature of the data, and should be kept in the PUF file.

# Using different options for local suppression

The long running time is mainly due to the local suppression algorithm. In Section 5.2.2 we discuss options to reduce the running time of the local suppression in case of many key variables. The all-m approach reduces the running time by first considering subsets of the complete set of key variables. This reduces the complexity of the problem and leads to lower computation times. However, the total number of suppressions made is likely to be higher. Also, if not explicitly specified, it is not guaranteed that the required level for k-anonymity is automatically achieved on the complete set of key variables. It is therefore important to check the results.

#### Step 9b: Re-measure risk

We re-evaluate the risk measures selected in Step 6. Table 9.25 shows that local suppression, not surprisingly, has reduced the number of individuals violating 5-anonymity to 0. The global hierarchical risk was reduced to 0.02%, which corresponds to approximately 2 correct re-identifications. The highest individual hierarchical re-identification risk is 0.2%. These risk levels are acceptable for a PUF release. Furthermore, the recoding has removed any unusual combinations in the data. **NOTE: The risk may be underestimated by excluding the household level variables.** 

Table 9.25: k-anonymity violations

k-anonymity	Number of records violating	Percentage
2	0	0.0 %
3	0	0.0 %
5	0	0.0 %

### Step 10b: Re-measure utility

We compare (cross-)tabulations before and after anonymization, which are illustrated in the *R* code to this case study. We note that due to the recoding in Step 8b, the detail in the variables is reduced. This reduces the number of necessary suppressions and is acceptable for a public use file.

#### Step 11: Audit and reporting

In the audit step, we check whether the data allow for reproduction of published figures from the original dataset and relationships between variables and other data characteristics are preserved in the anonymization process. In short, we check whether the dataset is valid for analytical purposes. There are no figures available that were published from the dataset and need to be reproducible from the anonymized data.

In Step 2, we explored the data characteristics and relationships between variables. These data characteristics and relationships have been mainly preserved, since we took them into account when choosing the appropriate anonymization methods. Since values of the variable "AGEYRS" were not perturbed, but only recoded and suppressed, we did not introduce unlikely combinations, such as a 60-year-old individual enrolled in primary education. Also, by separating

the anonymization process into two parts, one for household-level variables and one for individual-level variables, the values of variables measured at the household level agree for all members of each household.

Furthermore, we drafted two reports, internal and external, on the anonymization of the case study dataset. The internal report includes the methods used, the risk before and after anonymization as well as the reasons for the selected methods and their parameters. The external report focuses on the changes in the data and the loss in utility. Focus here should be on the number of suppressions as well as the perturbative methods (PRAM). This is described in the previous steps. **NOTE: When creating a PUF, it is inevitable that there will be a loss of information and it is very important for the users to be aware of these changes and release them in a report that accompanies the data.** Appendix C provides examples of an internal and external report of the anonymization process of this dataset. Depending on the users and readers of the reports, the content may differ. **NOTE: The report() function in sdcMicro is at this point not useful, since this will only report on the SDC measures in the second case study.** However, the report should contain the entire process, including the measures applied in case study 1.

#### Step 12: Data release

The final step is the release of the anonymized dataset together with the external report. Example 9.46 shows how to export the anonymized dataset as *STATA* file. Section 7.2 presents functions for exporting files in other data formats.

Example 9.46: Exporting the anonymized PUF file

# Create STATA file

write.dta(dataframe = dataAnon, file= 'Case2DataAnon.dta', convert.dates=TRUE)

# **BIBLIOGRAPHY**

- [Bran02] Brand, R. (2002). **Microdata Protection through Noise Addition.** In J. Domingo-Ferrer (Ed.), Inference Control in Statistical Databases From Theory to Practice (Vol. Lecture Notes in Computer Science Series Volume 2316, pp. 97-116). Berlin Heidelberg, Germany: Springer.
- [BCRZ13] Burgert, C. R., Colston, J., Roy, T., & Zachary, B. (2013). Geographic Displacement Procedure and Georeferenced Data Release Policy for the Demographic and Health Surveys. DHS Spatial Analysis Report No. 7.
- [DaRe78] Dalenius, T., & Reiss, S. (1978). **Data-swapping: A Technique for Disclosure Control (extended abstract).** American Statistical Association Proceedings of the Section on Survey Research Methods, (pp. 191-194).
- [WaWi99] de Waal, A., & Willenborg, L. (1999). **Information Loss Through Global Recoding and Local Suppression.** Netherlands Official Statistics, 14 (Special issue on SDC), 17-20.
- [WGKW98] de Wolf, P., Gouweleew, J. M., Kooiman, P., & Willenborg, L. (1998). **Reflections on PRAM.** Proceedings of the conference "Statistical Data Protection". Lisbon.
- [Wolf15] de Wolf, P.-P. (2015). Public Use Files of EU-SILC and EU-LFS data.
- [DoMa02] Domingo-Ferrer, J., & Mateo-Sanz, J. M. (2002). **Practical data-oriented microaggregation for statistical disclosure control.** IEEE Transactions on Knowledge and Data Engineering, 14, 189-201.
- [DoTo01a] Domingo-Ferrer, J., & Torra, V. (2001). A Quantitative Comparison of Disclosure Control Methods for Microdata. In P. Doyle, J. Lane, J. Theeuwes, & L. Zayatz (Eds.), Confidentiality, Disclosure and Data Access: Theory and Practical Applications for Statistical Agencies (pp. 111-133). Amsterdam, North-Holland: Elsevier Science.
- [DoTo01b] Domingo-Ferrer, J., & Torra, V. (2001). **Disclosure Protection Methods and Information Loss for Microdata.** In P. Doyle, J. Lane, J. Theeuwes, & Z. L., Theory and Practical Applications for Statistical Agencies (pp. 91-110). Amsterdam.
- [DoMT01] Domingo-Ferrer, J., Mateo-Sanz, J., & Torra, V. (2001). Comparing SDC Methods for Microdata on the basis of Information Loss and Disclosure Risk. Pre-proceedings of ETK-NTTS 2001, (pp. 807-826). Crete, Greece.
- [DuBo10] Dupriez, O., & Boyko, E. (2010). **Dissemination of Microdata Files; Principles, Procedures and Practices.** International Household Survey Network (IHSN).
- [ElMa03] Elliot, M., & Manning, A. M. (2003). **Using DIS to Modify the Classification of Special Uniques.** Invited Paper. Joint ECE/Eurostat Work Session on Statistical Data Confidentiality. Luxemboug 2-9 April 2003.
- [ElMF02] Elliot, M. J., Manning, A. M., & Ford, R. W. (2002). A Computational Algorithm for Handling the Special Uniques Problem. International Journal of Uncertainty, Fuzziness and Knowledge Based System, 10 (5), 493-509.

- [EMMG05] Elliot, M. J., Manning, A., Mayes, K., Gurd, J., & Bane, M. (2005). SUDA: A Program for Detecting Special Uniques. Joint UNECE/Eurostat Work Session on Statistical Data Confidentiality. Geneva.
- [GKWW98a] Gouweleeuw, J. M., Kooiman, P., Willenborg, L., & de Wolf, P. (1998a). **Post Randomization for Statistical Disclosure Control:Theory and Implementation.** Journal of Official Statistics, 14 (4), 463-478.
- [GKWW98b] Gouweleeuw, J. M., Kooiman, P., Willenborg, L., & de Wolf, P. P. (1998b). **The Post Randomization Method for Protecting Microdata.** Qüestiió, Quaderns d'Estadística i Investigació Operativa, 22 (1), 145-156.
- [HaMu03] Hansen, S. L., & Mukherjee, S. (2003). A polynomial algorithm for univariate optimal. IEEE Transactions on Knowledge and Data Engineering, 15, 1043-1044.
- [HuDr15] Hu, J., & Drechsler, J. (2015). **Generating synthetic geocoding infromation for public release.** NTTS Conferences on New Techniques and Technologies for Statistics. Brussels.
- [HDFG06] Hundepool, A., Domingo-Ferrer, J., Franconi, L., Giessing, S., Lenz, R., Naylor, J., et al. (2006). **Handbook on Statistical Disclosure Control.** ESSNet SDC.
- [HDFG12] Hundepool, A., Domingo-Ferrer, J., Franconi, L., Giessing, S., Nordholt, E. S., Spicer, K., et al. (2012). **Statistical Disclosure Control.** Chichester, UK: John Wiley & Sons Ltd.
- [HWRF08] Hundepool, A., Van de Wetering, A., Ramaswamy, R., Franconi, L., Polettini, S., Capobianchi, A., et al. (2008). μ-Argus User's Manual. The Hague: Statistics Netherlands.
- [KiWi03] Kim, J. J., & Winkler, W. W. (2003, April 17). Multiplicative Noise for Masking Continuous Data. Research Report Series.
- [KTMF13] Kowarik, A., Templ, M., Meindl, B., & Fonteneau, F. (2013). **Graphical User Interface for Package sdcMicro.** Retrieved from http://cran.r-project.org/web/packages/sdcMicroGUI/index.html
- [Lamb93] Lambert, D. (1993). **Measures of Disclosure Risk and Harm.** Journal of Official Statistics , 9 (2), 313-331.
- [MKGV07] Machanavajjhala, A., Kifer, D., Gehrke, J., & Venkitasubramaniam, M. (2007). L-diversity: Privacy Beyond K-anonymity. ACM Trans. Knowl. Discov. Data, 1 (Article 3) (1556-4681).
- [MaHK08] Manning, A. M., Haglin, D. J., & Keane, J. A. (2008). A Recursive Search Algorithm for Statistical Disclosure Assessment. Data Mining and Knowledge Discovery, 16 (2), 165-196.
- [MaTo10] Marés, J., & Torra, V. (2010). **PRAM Optimization Using an Evolutionary Algorithm.** In J. Domingo-Ferrer, & E. Magkos, Privacy in Statistical Databases (pp. 97-106). Corfú, Greece: Springer.
- [Mivu12] Mivule, K. (July 16-19, 2012). **Utilizing Noise Addition for Data Privacy, An Overview.** International Conference on Information and Knowledge Engineering (IKE 2012), (pp. 65-71). Las Vegas, Nevada.
- [MuSa06] Muralidhar, K., & Sarathy, R. (2006). **Data Shuffling- A New Masking Approach for Numerical Data.** Management Science, 658-670.
- [TeMK14] Templ, M., Meindl, B., & Kowarik, A. (2014, August). **Tutorial for SDCMicroGUI.** Retrieved from International Household Survey Network (IHSN): http://www.ihsn.org/home/software/disclosure-control-toolbox
- [TMKC14] Templ, M., Meindl, B., Kowarik, A., & Chen, S. (2014, August 1). Introduction to Statistical Disclosure Control (SDC). Retrieved November 13, 2014, from http://www.ihsn.org/home/software/disclosure-control-toolbox.
- [ToCa11] Torra, V., & Cano, I. (2011). **Edit Constraints on Microaggregation and Additive Noise.** In A. G.-D. Christos Dimitrakakis (Ed.), Privacy and Security Issues in Data Mining and Machine Learning (Vol. Lecture Notes in Computer Science Volume 6549, pp. 1-14). Berlin Heidelberg: Springer.
- [TuLS15] Tudor, C., Lowthian, P., & Spicer, K. (2015). **Opening up governemnt data while maintaining data privacy.** EDBT/ICDT 2015 Joint Conference.
- [UKIn] UK Information Commissioner's Office. Anonymisation: Managing Data Protection Risk Code of Practice. UK Information Commissioner's Office.

180 Bibliography

- [Unit07] United Nations. (2007). Managing Statistical Confidentiality & Microdata Access: Principles and Guidelines of Good Practice. New York and Geneva: United Nations.
- [Warn65] Warner, S. (1965). Randomized Response: A Survey Technique for Eliminating Evasive Answer Bias. Journal of American Statistical Association, 57, 622-627.
- [YaWC02] Yancey, W. W., Winkler, W. E., & Creecy, R. H. (2002). **Disclosure Risk Assessment in Perturbative Microdata Protection.** Research Report Series, Statistics 2002-01.
- [YaSB08] Yang, A., Shan, Y., & Bui, L. T. (2008). Success in Evolutionary Computation. Springer Science & Business Media.

Bibliography 181